

Best of LessWrong: April 2016

1. [Turning the Technical Crank](#)

Best of LessWrong: April 2016

1. [Turning the Technical Crank](#)

Turning the Technical Crank

A few months ago, Vaniver wrote [a really long post](#) speculating about potential futures for Less Wrong, with a focus on the idea that the spread of the Less Wrong diaspora has left the site weak and fragmented. I wasn't here for our high water mark, so I don't really have an informed opinion on what has *socially* changed since then. But a number of complaints are technical, and as an IT person, I thought I had some useful things to say.

I [argued at the time](#) that many of the technical challenges of the diaspora were solved problems, and that the solution was NNTP -- an ancient, yet still extant, discussion protocol. I am something of a crank on the subject and didn't expect much of a reception. I was pleasantly surprised by the 18 karma it generated, and tried to write up a full post arguing the point.

I failed. I was trying to write a manifesto, didn't really know how to do it right, and kept running into a vast inferential distance I couldn't seem to cross. I'm a product of a prior age of the Internet, from before the http prefix assumed its imperial crown; I kept wanting to say things that I knew would make no sense to anyone who came of age this millennium. I got bogged down in irrelevant technical minutia about how to implement features X, Y, and Z. Eventually I decided I was attacking the wrong problem; I was thinking about 'how do I promote NNTP', when really I should have been going after 'what would an ideal discussion platform look like and how does NNTP get us there, if it does?'

So I'm going to go after that first, and work on the inferential distance problem, and *then* I'm going to talk about NNTP, and see where that goes and what could be done better. I still believe it's the closest thing to a *good, available* technological schelling point, but it's going to take a lot of words to get there from here, and I might change my mind under persuasive argument. We'll see.

Fortunately, this is Less Wrong, and sequences are a thing here. This is the first post in an intended sequence on *mechanisms of discussion*. I know it's a bit off the beaten track of Less Wrong subject matter. I posit that it's both relevant to our difficulties and probably more useful and/or interesting than most of what comes through these days. I just took the 2016 survey and it has a couple of sections on the effects of the diaspora, so I'm guessing it's on topic for meta purposes if not for site-subject purposes.

Less Than Ideal Discussion

To solve a problem you must first define it. Looking at the LessWrong 2.0 post, I see the following *technical* problems, at a minimum; I'll edit this with suggestions from comments.

1. Aggregation of posts. Our best authors have formed their own fiefdoms and their work is not terribly visible here. We currently have limited support for this via the sidebar, but that's it.
2. Aggregation of comments. You can see diaspora authors in the sidebar, but you can't comment from here.

3. Aggregation of *community*. This sounds like a social problem but it isn't. You can start a new blog, but unless you plan on *also* going out of your way to market it then your chances of starting a discussion boil down to "hope it catches the attention of Yvain or someone else similarly prominent in the community." Non-prominent individuals can theoretically post here; yet this is the place we are decrying as moribund.
4. Incomplete and poor curation. We currently do this via Promoted, badly, and via the diaspora sidebar, also badly.
5. Pitiful interface feature set. This is not so much a Less Wrong-specific problem as a 2010s-internet problem; people who inhabit SSC have probably seen me respond to feature complaints with "they had something that did that in the 90s, but nobody uses it." (my own bugbear is searching for comments by author-plus-content).
6. Changes are hamstrung by the existing architecture, which gets you volunteer reactions like [this one](#).

I see these *meta-technical* problems:

1. Expertise is scarce. Few people are in a position to technically improve the site, and those that are, have other demands on their time.
2. The [Trivial Inconvenience Problem](#) limits the *scope* of proposed changes to those that are not inconvenient to commenters or authors.
3. Getting cooperation from diaspora authors is a coordination problem. Are we better than average at handling those? I don't know.

Slightly Less Horrible Discussion

"Solving" community maintenance is a hard problem, but to the extent that pieces of it can be solved technologically, the solution might include these ultra-high-level elements:

1. Centralized *from the user perspective*. A reader should be able to interact with the entire community in one place, and it should be recognizable as a community.
2. Decentralized *from the author perspective*. Diaspora authors seem to like having their own fiefdoms, and the social problem of "all the best posters went elsewhere" can't be solved without their cooperation. Therefore any technical solution must allow for it.
3. Proper division of labor. Scott Alexander probably should not have to concern himself with user feature requests; that's not his comparative advantage and I'd rather he spend his time [inventing moral cosmologies](#). I suspect he would prefer the same. The same goes for Eliezer Yudkowsky or any of our still-writing-elsewhere folks.
4. Really good moderation tools.
5. Easy entrance. New users should be able to join the discussion without a lot of hassle. Old authors that want to return should be able to do so and, preferably, bring their existing content with them.
6. Easy exit. Authors who don't like the way the community is heading should be able to jump ship -- and, crucially, bring their content with them to their new ship. Conveniently. This is essentially what has happened, except old content is hostage here.
7. [Separate policy and mechanism](#) within the site architecture. Let this one pass for now if you don't know what it means; it's the first big inferential hurdle I need to

cross and I'll be starting soon enough.

As with the previous, I'll update this from the comments if necessary.

Getting There From Here

As I said at the start, I feel on firmer ground talking about technical issues than social ones. But I have to acknowledge one strong social opinion: I believe the greatest factor in Less Wrong's decline is *the departure of our best authors for personal blogs*. Any plan for revitalization has to provide an *improved substitute* for a personal blog, because that's where everyone seems to end up going. You need something that looks and behaves like a blog to the author or casual readers, but integrates seamlessly into a community discussion gateway.

I argue that this can be achieved. I argue that the technical challenges are solvable and the inherent coordination problem is also solvable, provided the people involved still have an interest in solving it.

And I argue that it can be done -- and done *better* than what we have now -- using technology that has existed since the '90s.

I *don't* argue that this actually will be achieved in anything like the way I think it ought to be. As mentioned up top, I am a crank, and I have no access whatsoever to anybody with any community pull. My odds of pushing through this agenda are basically nil. But we're all about crazy thought experiments, right?

This topic is something I've wanted to write about for a long time. Since it's not typical Less Wrong fare, I'll take the karma on this post as a referendum on whether the community would like to see it here.

Assuming there's interest, the sequence will look something like this (subject to reorganization as I go along, since I'm pulling this from some lengthy but horribly disorganized notes; in particular I might swap subsequences 2 and 3):

1. Technical Architecture
 1. Your Web Browser Is Not Your Client
 2. Specialized Protocols: or, NNTP and its Bastard Children
 3. Moderation, Personal Gardens, and Public Parks
 4. Content, Presentation, and the Division of Labor
 5. The Proper Placement of User Features
 6. Hard Things that are Suddenly Easy: or, what does client control gain us?
 7. Your Web Browser Is Still Not Your Client (but you don't need to know that)
2. Meta-Technical Conflicts (or, obstacles to adoption)
 1. Never Bet Against Convenience
 2. Conflicting Commenter, Author, and Admin Preferences
 3. Lipstick on the Configuration Pig
 4. Incremental Implementation and the Coordination Problem.
 5. Lowering Barriers to Entry and Exit
3. Technical and Social Interoperability
 1. Benefits and Drawbacks of Standards
 2. Input Formats and Quoting Conventions
 3. Faking Functionality
 4. Why Reddit Makes Me Cry

5. What NNTP Can't Do
4. Implementation of Nonstandard Features
 1. Some desirable feature #1
 2. Some desirable feature #2
 3. ...etc. This subsequence is only necessary if someone actually wants to try and do what I'm arguing for, which I think unlikely.

(Meta-meta: This post was written in Markdown, converted to HTML for posting using Pandoc, and took around four hours to write. I can often be found lurking on #lesswrong or #slatestarcodex on workday afternoons if anyone wants to discuss it, but I don't promise to answer quickly because, well, workday)

[Edited to add: At +10/92% karma I figure continuing is probably worth it. After reading comments I'm going to try to slim it down a lot from the outline above, though. I still want to hit all those points but they probably don't all need a full post's space. Note that I'm not Scott or Eliezer, I write like I bleed, so what I do post will likely be spaced out]