

Categorisation and Concepts

1. [Diseased thinking: dissolving questions about disease](#)
2. [The Categories Were Made For Man, Not Man For The Categories](#)
3. [The noncentral fallacy - the worst argument in the world?](#)
4. [Ethnic Tension And Meaningless Arguments](#)
5. [The Moral Of The Story](#)

Diseased thinking: dissolving questions about disease

Related to: [Disguised Queries](#), [Words as Hidden Inferences](#), [Dissolving the Question](#), [Eight Short Studies on Excuses](#)

Today's therapeutic ethos, which celebrates curing and disparages judging, expresses the liberal disposition to assume that crime and other problematic behaviors reflect social or biological causation. While this absolves the individual of responsibility, it also strips the individual of personhood, and moral dignity

-- George Will, townhall.com

Sandy is a morbidly obese woman looking for advice.

Her husband has no sympathy for her, and tells her she obviously needs to stop eating like a pig, and would it kill her to go to the gym once in a while?

Her doctor tells her that obesity is primarily genetic, and recommends the diet pill orlistat and a consultation with a surgeon about gastric bypass.

Her sister tells her that obesity is a perfectly valid lifestyle choice, and that fat-ism, equivalent to racism, is society's way of keeping her down.

When she tells each of her friends about the opinions of the others, things really start to heat up.

Her husband accuses her doctor and sister of absolving her of personal responsibility with feel-good platitudes that in the end will only prevent her from getting the willpower she needs to start a real diet.

Her doctor accuses her husband of ignorance of the real causes of obesity and of the most effective treatments, and accuses her sister of legitimizing a dangerous health risk that could end with Sandy in hospital or even dead.

Her sister accuses her husband of being a jerk, and her doctor of trying to medicalize her behavior in order to turn it into a "condition" that will keep her on pills for life and make lots of money for Big Pharma.

Sandy is fictional, but similar conversations happen every day, not only about obesity but about a host of other marginal conditions that some consider character flaws, others diseases, and still others normal variation in the human condition. Attention deficit disorder, internet addiction, social anxiety disorder (as one skeptic said, didn't we used to call this "shyness"?), alcoholism, chronic fatigue, oppositional defiant disorder ("didn't we used to call this being a teenager?"), compulsive gambling, homosexuality, Aspergers' syndrome, antisocial personality, even depression have all been placed in two or more of these categories by different people.

Sandy's sister may have a point, but this post will concentrate on the debate between her husband and her doctor, with the understanding that the same techniques will apply to evaluating her sister's opinion. The disagreement between Sandy's husband and doctor centers around the idea of "disease". If obesity, depression, alcoholism,

and the like are diseases, most people default to the doctor's point of view; if they are not diseases, they tend to agree with the husband.

The debate over such marginal conditions is in many ways a debate over whether or not they are "real" diseases. The usual surface level arguments trotted out in favor of or against the proposition are generally inconclusive, but this post will apply a host of techniques previously discussed on Less Wrong to illuminate the issue.

What is Disease?

In [Disguised Queries](#), Eliezer demonstrates how a word refers to a cluster of objects related upon multiple axes. For example, in a company that sorts red smooth translucent cubes full of vanadium from blue furry opaque eggs full of palladium, you might invent the word "rube" to designate the red cubes, and another "blegg", to designate the blue eggs. Both words are useful because they "carve reality at the joints" - they refer to two completely separate classes of things which it's practically useful to keep in separate categories. Calling something a "blegg" is a quick and easy way to describe its color, shape, opacity, texture, and chemical composition. It may be that the odd blegg might be purple rather than blue, but in general the characteristics of a blegg remain sufficiently correlated that "blegg" is a useful word. If they weren't so correlated - if blue objects were equally likely to be palladium-containing-cubes as vanadium-containing-eggs, then the word "blegg" would be a waste of breath; the characteristics of the object would remain just as mysterious to your partner after you said "blegg" as they were before.

"Disease", like "blegg", suggests that certain characteristics always come together. A rough sketch of some of the characteristics we expect in a disease might include:

1. Something caused by the sorts of thing you study in biology: proteins, bacteria, ions, viruses, genes.
2. Something involuntary and completely immune to the operations of free will
3. Something rare; the vast majority of people don't have it
4. Something unpleasant; when you have it, you want to get rid of it
5. Something discrete; a graph would show two widely separate populations, one with the disease and one without, and not a normal distribution.
6. Something commonly treated with science-y interventions like chemicals and radiation.

Cancer satisfies every one of these criteria, and so we have no qualms whatsoever about classifying it as a disease. It's a type specimen, [the sparrow as opposed to the ostrich](#). The same is true of heart attack, the flu, diabetes, and many more.

Some conditions satisfy a few of the criteria, but not others. Dwarfism seems to fail (5), and it might get its status as a disease only after studies show that the supposed dwarf falls way out of normal human height variation. Despite the best efforts of transhumanists, it's hard to convince people that aging is a disease, partly because it fails (3). Calling homosexuality a disease is a poor choice for many reasons, but one of them is certainly (4): it's not necessarily unpleasant.

The marginal conditions mentioned above are also in this category. Obesity arguably sort-of-satisfies criteria (1), (4), and (6), but it would be pretty hard to make a case for (2), (3), and (5).

So, is obesity really a disease? Well, is Pluto really a planet? Once we state that obesity satisfies some of the criteria but not others, it is meaningless to talk about an additional fact of whether it "really deserves to be a disease" or not.


If it weren't for those pesky [hidden inferences](#)...

Hidden Inferences From Disease Concept

The state of the disease node, meaningless in itself, is used to predict several other nodes with non-empirical content. In English: we make value decisions based on whether we call something a "disease" or not.

If something is a real disease, the patient deserves our sympathy and support; for example, [cancer sufferers must universally be described as "brave"](#). If it is not a real disease, people are more likely to get our condemnation; for example Sandy's husband who calls her a "pig" for her inability to control her eating habits. The difference between "shyness" and "social anxiety disorder" is that people with the first get called "weird" and told to man up, and people with the second get special privileges and the sympathy of those around them.

And if something is a real disease, it is socially acceptable (maybe even mandated) to seek medical treatment for it. If it's not a disease, medical treatment gets derided as a "quick fix" or an "abdication of personal responsibility". I have talked to several doctors who are uncomfortable suggesting gastric bypass surgery, even in people for whom it is medically indicated, because they believe it is morally wrong to turn to medicine to solve a character issue.



While a condition's status as a "real disease" ought to be meaningless as a "hanging node" after the status of all other nodes have been determined, it has acquired

political and philosophical implications because of its role in determining whether patients receive sympathy and whether they are permitted to seek medical treatment.

If we can determine whether a person should get sympathy, and whether they should be allowed to seek medical treatment, independently of the central node "disease" or of the criteria that feed into it, we will have successfully unasked the question "are these marginal conditions real diseases" and cleared up the confusion.

Sympathy or Condemnation?

Our attitudes toward people with marginal conditions mainly reflect a deontologist libertarian (libertarian as in "free will", not as in "against government") model of blame. In this concept, people make decisions using their free will, a spiritual entity operating free from biology or circumstance. People who make good decisions are intrinsically good people and deserve good treatment; people who make bad decisions are intrinsically bad people and deserve bad treatment. But people who make bad decisions for reasons that are outside of their free will may not be intrinsically bad people, and may therefore be absolved from deserving bad treatment. For example, if a normally peaceful person has a brain tumor that affects areas involved in fear and aggression, they go on a crazy killing spree, and then they have their brain tumor removed and become a peaceful person again, many people would be willing to accept that the killing spree does not reflect negatively on them or open them up to deserving bad treatment, since it had biological and not spiritual causes.

Under this model, deciding whether a condition is biological or spiritual becomes very important, and the rationale for worrying over whether something "is a real disease" or not is plain to see. Without figuring out this extremely difficult question, we are at risk of either blaming people for things they don't deserve, or else letting them off the hook when they commit a sin, both of which, to libertarian deontologists, would be terrible things. But determining whether marginal conditions like depression have a spiritual or biological cause is difficult, and no one knows how to do it reliably.

Determinist consequentialists can do better. We believe it's biology all the way down. Separating spiritual from biological illnesses is impossible and unnecessary. Every condition, from brain tumors to poor taste in music, is "biological" insofar as it is encoded in things like cells and proteins and follows laws based on their structure.

But determinists don't just ignore the very important differences between brain tumors and poor taste in music. Some biological phenomena, like poor taste in music, are encoded in such a way that they are extremely vulnerable to what we can call social influences: praise, condemnation, introspection, and the like. Other biological phenomena, like brain tumors, are completely immune to such influences. This allows us to develop a more useful model of blame.

The consequentialist model of blame is very different from the deontological model. Because all actions are biologically determined, none are more or less metaphysically blameworthy than others, and none can mark anyone with the metaphysical status of "bad person" and make them "deserve" bad treatment. Consequentialists don't on a primary level want anyone to be treated badly, full stop; thus [is it written](#): "Saddam Hussein doesn't deserve so much as a stubbed toe." But if consequentialists don't believe in punishment for its own sake, they do believe in punishment for the sake of, well, consequences. Hurting bank robbers may not be a good in and of itself, but it will prevent banks from being robbed in the future. And, one might infer, although

alcoholics may not deserve condemnation, societal condemnation of alcoholics makes alcoholism a less attractive option.

So here, at last, is a rule for which diseases we offer sympathy, and which we offer condemnation: if giving condemnation instead of sympathy decreases the incidence of the disease enough to be worth the hurt feelings, condemn; otherwise, sympathize. Though the rule is based on philosophy that the majority of the human race would disavow, it leads to intuitively correct consequences. Yelling at a cancer patient, shouting "How dare you allow your cells to divide in an uncontrolled manner like this; is that the way your mother raised you??" will probably make the patient feel pretty awful, but it's not going to cure the cancer. Telling a lazy person "Get up and do some work, you worthless bum," very well might cure the laziness. The cancer is a biological condition immune to social influences; the laziness is a biological condition susceptible to social influences, so we try to socially influence the laziness and not the cancer.

The question "Do the obese deserve our sympathy or our condemnation," then, is asking whether condemnation is such a useful treatment for obesity that its utility outweighs the disutility of hurting obese people's feelings. This question may have different answers depending on the particular obese person involved, the particular person doing the condemning, and the availability of other methods for treating the obesity, which brings us to...

The Ethics of Treating Marginal Conditions

If a condition is susceptible to social intervention, but an effective biological therapy for it also exists, is it okay for people to use the biological therapy instead of figuring out a social solution? My gut answer is "Of course, why wouldn't it be?", but apparently lots of people find this controversial for some reason.

In a libertarian deontological system, throwing biological solutions at spiritual problems might be disrespectful or dehumanizing, or a band-aid that doesn't affect the deeper problem. To someone who believes it's biology all the way down, this is much less of a concern.

Others complain that the existence of an easy medical solution prevents people from learning personal responsibility. But here [we see the status-quo bias at work, and so can apply a preference reversal test](#). If people really believe learning personal responsibility is more important than being not addicted to heroin, we would expect these people to support deliberately addicting schoolchildren to heroin so they can develop personal responsibility by coming off of it. Anyone who disagrees with this somewhat shocking proposal must believe, on some level, that having people who are not addicted to heroin is more important than having people develop whatever measure of personal responsibility comes from kicking their heroin habit the old-fashioned way.

But the most convincing explanation I have read for why so many people are opposed to medical solutions for social conditions is a signaling explanation by Robin Hans...wait! no!...by Katja Grace. On [her blog](#), she says:

...the situation reminds me of a pattern in similar cases I have noticed before. It goes like this. Some people make personal sacrifices, supposedly toward solving problems that don't threaten them personally. They sort recycling, buy free range eggs, buy fair trade, campaign for wealth redistribution etc. Their actions are seen as virtuous. They see those who don't join them as uncaring and immoral. A more

efficient solution to the problem is suggested. It does not require personal sacrifice. People who have not previously sacrificed support it. Those who have previously sacrificed object on grounds that it is an excuse for people to get out of making the sacrifice. The supposed instrumental action, as the visible sign of caring, has become virtuous in its own right. Solving the problem effectively is an attack on the moral people.

A case in which some people eat less enjoyable foods and exercise hard to avoid becoming obese, and then campaign against a pill that makes avoiding obesity easy demonstrates some of the same principles.

There are several very reasonable objections to treating any condition with drugs, whether it be a classical disease like cancer or a marginal condition like alcoholism. The drugs can have side effects. They can be expensive. They can build dependence. They may later be found to be placebos whose efficacy was overhyped by dishonest pharmaceutical advertising.. They may raise ethical issues with children, the mentally incapacitated, and other people who cannot decide for themselves whether or not to take them. But these issues do not magically become more dangerous in conditions typically regarded as "character flaws" rather than "diseases", and the same good-enough solutions that work for cancer or heart disease will work for alcoholism and other such conditions (but see [here](#)).

I see no reason why people who want effective treatment for a condition should be denied it or stigmatized for seeking it, whether it is traditionally considered "medical" or not.

Summary

People commonly debate whether social and mental conditions are real diseases. This masquerades as a medical question, but its implications are mainly social and ethical. We use the concept of disease to decide who gets sympathy, who gets blame, and who gets treatment.

Instead of continuing the fruitless "disease" argument, we should address these questions directly. Taking a determinist consequentialist position allows us to do so more effectively. We should blame and stigmatize people for conditions where blame and stigma are the most useful methods for curing or preventing the condition, and we should allow patients to seek treatment whenever it is available and effective.

The Categories Were Made For Man, Not Man For The Categories

I.

“Silliest internet atheist argument” is a hotly contested title, but I have a special place in my heart for the people who occasionally try to prove Biblical fallibility by pointing out whales are not a type of fish.

(this is going to end up being a metaphor for something, so bear with me)

The argument goes like this. Jonah got swallowed by a whale. But the Bible says Jonah got swallowed by a big fish. So the Bible seems to think whales are just big fish. Therefore the Bible is fallible. Therefore, the Bible was not written by God.

The first problem here is that “whale” is just our own modern interpretation of the Bible. For all we know, Jonah was swallowed by a really really really big herring.

The second problem is that if the ancient Hebrews want to call whales a kind of fish, let them call whales a kind of fish.

I’m not making the weak and boring claim that since they’d never discovered genetics they don’t know better. I am making the much stronger claim that, even if the ancient Hebrews had taken enough of a break from murdering Philistines and building tabernacles to sequence the genomes of all known species of aquatic animals, there’s nothing whatsoever wrong, false, or incorrect with them calling a whale a fish.

Now, there’s something wrong with saying “whales are phylogenetically just as closely related to bass, herring, and salmon as these three are related to each other.” What’s wrong with the statement is that it’s false. But saying “whales are a kind of fish” isn’t.

Suppose you travel back in time to ancient Israel and try to explain to King Solomon that whales are a kind of mammal and not a kind of fish.

Your translator isn’t very good, so you pause to explain “fish” and “mammal” to Solomon. You tell him that fish is “the sort of thing herring, bass, and salmon are” and mammal is “the sort of thing cows, sheep, and pigs are”. Solomon tells you that your word “fish” is Hebrew *dag* and your word “mammal” is Hebrew *behemah*.

So you try again and say that a whale is a *behemah*, not a *dag*. Solomon laughs at you and says you’re an idiot.

You explain that you’re not an idiot, that in fact all kinds of animals have things called genes, and the genes of a whale are much closer to those of the other *behemah* than those of the *dag*.

Solomon says he’s never heard of these gene things before, and that maybe genetics is involved in your weird foreign words “fish” and “mammal”, but *dag* are just finned creatures that swim in the sea, and *behemah* are just legged creatures that walk on the Earth.

(like the *kelev* and the *parah* and the *gavagai*)

You try to explain that no, Solomon is wrong, *dag* are actually defined not by their swimming-in-sea-with-fins-ness, but by their genes.

Solomon says you didn't even *know* the word *dag* ten minutes ago, and now suddenly you think you know what it means better than he does, who has been using it his entire life? Who died and made *you* an expert on Biblical Hebrew?

You try to explain that whales actually have tiny little hairs, too small to even see, just as cows and sheep and pigs have hair.

Solomon says oh God, you are so annoying, who the hell cares whether whales have tiny little hairs or not. In fact, the only thing Solomon cares about is whether responsibilities for his kingdom's production of blubber and whale oil should go under his Ministry of Dag or Ministry of Behemah. The Ministry of Dag is based on the coast and has a lot of people who work on ships. The Ministry of Behemah has a strong presence inland and lots of people who hunt on horseback. So please (he continues) keep going about how whales have little tiny hairs.

It's easy to see that Solomon has a point, and that if he wants to define *behemah* as four-legged-land-dwellers that's his right, and no better or worse than your definition of "creatures in a certain part of the phylogenetic tree". Indeed, it might even be that if you spent ten years teaching Solomon all about the theory of genetics and evolution (which would be hilarious - think how annoyed the creationists would get) he might still say "That's very interesting, and I can see why we need a word to describe creatures closely related along the phylogenetic tree, but make up your own word, because *behemah* already means 'four-legged-land-dweller'."

Now imagine that instead of talking to King Solomon, you're talking to that guy from Duck Dynasty with the really crazy beard (I realize that may describe more than one person), who stands in for all uneducated rednecks in the same way King Solomon stands in for all Biblical Hebrews.

"Ah course a whale is a feesh, ya moron" he says in his heavy Southern accent.

"No it isn't," you say. "A fish is a creature phylogenetically related to various other fish, and with certain defining anatomical features. It says so right here in this biology textbook."

"Well," Crazy Beard Guy tells you, "Ah reckon that might be what a fish is, but a *feesh* is some'in that swims in the orshun."

With a sinking feeling in your stomach, you spend ten years turning Crazy Beard Guy into a world expert on phylogenetics and evolutionary theory. Although the Duck Dynasty show becomes *much* more interesting, you fail to budge him a bit on the meaning of "feesh".

It's easy to see here that "fish" and "feesh" can be different just as "fish" and "*dag*" can be different.

You can point out how many important professors of ichthyology in fancy suits use your definition, and how only a couple of people with really weird facial hair use his. But now you're making a status argument, not a factual argument. Your argument is "conform to the way all the cool people use the word 'fish'", not "a whale is really and truly not a fish".

There are facts of the matter on each individual point – whether a whale has fins, whether a whale lives in the ocean, whether a whale has tiny hairs, et cetera. But there is no fact of the matter on whether a whale is a fish. The argument is entirely semantic.

So this is the second reason why this particular objection to the Bible is silly. If God wants to call a whale a big fish, stop telling God what to do.

(also, [bats](#))

II.

When terms are *not* defined directly by God, we need our own methods of dividing them into categories.

The essay [“How An Algorithm Feels From The Inside”](#) is a gift that keeps on giving. You can get a reputation as a daring and original thinker just by copy-pasting it at different arguments with a couple of appropriate words substituted for one another, mad-libs like. It is the solution to something like 25% of extant philosophical problems.

It starts with a discussion of whether or not Pluto is a planet. Planets tend to share many characteristics in common. For example, they are large, round, have normal shaped orbits lined up with the plane of the ecliptic, have cleared out a certain area of space, and are at least kind of close to the Sun as opposed to way out in the Oort Cloud.

One could imagine a brain that thought about these characteristics like Network 1 here:



Obligatory [Less Wrong](#) picture

One could imagine this model telling you everything you need to know. If an object is larger, it's more likely to be round and in cis-Neptunian space. If an object has failed to clear its orbit of debris, it's more likely to have a skewed orbit relative to the plane of the ecliptic. We could give each of these relationships Bayesian weights and say things like large objects have a 32% chance of being in cis-Neptunian space and small objects an 86% chance. Or whatever.

But Network 1 has some big problems. For one thing, if you inscribe it in blood, you might accidentally summon the Devil. But for another, it's computationally very complicated. Each attribute affects each other attribute which affects it in turn and so on in an infinite cycle, so that its behavior tends to be chaotic and unpredictable.

What people actually seem to do is more like Network 2: sweep all common correlations into one big category in the middle, thus dividing possibility-space into large round normal-orbit solitary inner objects, and small irregular skewed-orbit crowded outer objects. It calls the first category “planets” and the second category “planetoids”.

You can then sweep minor irregularities under the rug. Neptune is pretty far from the sun, but since it's large, round, normal-orbit, and solitary, we know which way the evidence is leaning.

When an object satisfies about half the criteria for planet and half the criteria for planetoid, *then* it's awkward. Pluto is the classic example. It's relatively large, round, skewed orbit, solitary...ish? and outer-ish. What do you do?

The *practical* answer is you convene some very expensive meeting of prestigious astronomers and come to some official decision which everyone agrees to follow so they're all on the same page.

But the *ideal* answer is you say "Huh, the assumption encoded in the word 'planet' that the five red criteria always went together and the five blue criteria always went together doesn't hold. Whatever."

Then you divide the solar system into three types of objects: planets, planetoids, and dammit-our-categorization-scheme-wasn't-as-good-as-we-thought.

(psychiatry, whose philosophy of categorization is light years ahead of a lot of the rest of the world, conveniently abbreviates this latter category as "NOS")

The situation with whales and fish is properly understood in the same context. Fish and mammals differ on a lot of axes. Fish generally live in the water, breathe through gills, have tails and fins, possess a certain hydrodynamic shape, lay eggs, and are in a certain part of the phylogenetic tree. Mammals generally live on land, breathe through lungs, have legs, give live birth, and are in another part of the phylogenetic tree. Most fish conform to all of the fish desiderata, and most mammals conform to all of the mammal desiderata, so there's no question of how to categorize them. Occasionally you get something weird (a platypus, a lungfish, or a whale) and it's a judgment call which you have to decide by fiat. In our case, that fiat is "use genetics and ignore all other characteristics" but some other language, culture, or scientific community might make a different fiat, and then the borders between their categories would look a little bit different.

III.

Since I shifted to a borders metaphor, let's follow that and see where it goes.

Imagine that Israel and Palestine agree to a two-state solution with the final boundary to be drawn by the United Nations. You're the head of the United Nations committee involved, so you get out a map and a pencil. Both sides have sworn by their respective gods to follow whatever you determine.

Your job is not to draw "the correct border". There is no one correct border between Israel and Palestine. There are a couple of very strong candidates (for example, the pre-1967 line of control), but both countries have suggested deviations from that (most people think an actual solution would involve Palestine giving up some territory that has since been thoroughly settled by Israel in exchange for some territory within Israel proper, or perhaps for a continuous "land bridge" between the West Bank and Gaza). Even if you wanted to use the pre-1967 line as a starting point, there would still be a lot of work to do deciding what land swaps should and shouldn't be made.

Instead you'd be making a series of trade-offs. Giving all of Jerusalem to the Israelis would make them very happy but anger Palestine. Creating a contiguous corridor between Gaza and the West Bank makes some sense, but then you'd be cutting off Eilat from the rest of Israel. Giving all of the Israeli settlements in the West Bank back to Palestine would satisfy a certain conception of property rights, but also leave a lot of Jews homeless.

There are also much stupider decisions you could make. You could give Tel Aviv to Palestine. You could make the Palestinian state a perfect circle five miles in radius centered on Rishon LeZion. You could just split the territory in half with a straight line, and give Israel the north and Palestine the south. All of these things would be really dumb.

But, crucially, they would not be *false*. They would not be *factually incorrect*. They would just be failing to achieve pretty much any of the goals that we would expect a person solving land disputes in the Middle East to have. You can think of alternative arrangements in which these wouldn't be dumb. For example, if you're a despot, and you want to make it very clear to both the Israelis and Palestinians that their opinions don't matter and they should stop bothering you with annoying requests for arbitration, maybe splitting the country in half north-south is the way to go.

This is now unexpectedly a geography blog again.

The border between Turkey and Syria follows a mostly straight-ish line near-ish the 36th parallel, except that about twenty miles south of the border Turkey controls a couple of square meters in the middle of a Syrian village. This is the tomb of the ancestor of the Ottoman Turks, and Turkey's border agreement with Syria stipulates that it will remain part of Turkey forever. And the Turks take this very seriously; they maintain a platoon of special forces there and have recently been [threatening war against Syria](#) if their "territory" gets "invaded" in the current conflict.



Pictured: Turkey (inside fence), Syria (outside)

The border between Bangladesh and India is complicated at the best of times, but it becomes absolutely ridiculous in a place called Cooch-Bihar, which I guess is as good a name as any for a place full of ridiculous things. In at least one spot there is an 'island' of Indian territory within a larger island of Bangladeshi territory within a larger island of Indian territory within Bangladesh. According to [mentalfloss.com](#):

So why'd the border get drawn like that? It can all be traced back to power struggles between local kings hundreds of years ago, who would try to claim pockets of land inside each other's territories as a way to leverage political power. When Bangladesh became independent from India in 1947 (as East Pakistan until 1971), all those separate pockets of land were divvied up. Hence the polka-dotted mess.



Namibia is a very weird-looking country with a very thin three-hundred-mile-long panhandle (eg about twice as long as Oklahoma's). Apparently during the Scramble For Africa, the Germans who colonized Namibia really wanted access to the Zambezi River so they could reach the Indian Ocean and trade their colonial resources. They kept pestering the British who colonized Botswana until the Brits finally agreed to give up a tiny but very long strip of territory ending at the riverbank. This turned out to be not so useful, as *just* after Namibia's Zambezi access sits Victoria Falls, the largest waterfall in the world – meaning that any Germans who tried to traverse the Zambezi to reach the Indian Ocean would last a matter of minutes before suddenly encountering a four hundred foot drop and falling to pretty much certain death. The moral of the story is not to pester the British Empire too much, especially if they've explored Africa and you haven't.



But the other moral of the story is that borders are weird. Although we think of borders as nice straight lines that separate people of different cultures, they can form giant panhandles, distant islands, and enclaves-within-enclaves-within-enclaves. They can depart from their usual course to pay honor to national founders, to preserve records of ancient conquests, or to connect to trade routes.

Hume's ethics restrict "bad" to an instrumental criticism – you can condemn something as a bad way to achieve a certain goal, but not as morally bad independent of what the goal is. In the same way, borders can be bad at fulfilling your goals in drawing them, but not bad in an absolute sense or factually incorrect. Namibia's border is bad from the perspective of Germans who want access to the Indian Ocean. But it's *excellent* from the perspective of Englishmen who want to watch Germans plummet into the Lower Zambezi and get eaten by hippos.

Breaking out of the metaphor, the same is true of conceptual boundaries. You *may* draw the boundaries of the category "fish" any way you want. A category "fish" containing herring, dragonflies, and asteroids is going to be stupid, but only in the same sense that a Palestinian state centered around Tel Aviv would be stupid – it fails to fulfill any conceivable goals of the person designing it. Categories "fish" that do or don't include whales may be appropriate for different people's purposes, the same way Palestinians might argue about whether the borders of their state should be optimized for military defensibility or for religious/cultural significance.

Statements like "the Zambezi River is full of angry hippos" are brute facts. Statements like "the Zambezi River is the territory of Namibia" are negotiable.

In the same way, statements like "whales have little hairs" are brute facts. Statements like "whales are not a kind of fish" are negotiable.

So it's important to keep these two sorts of statements separate, and remember that in no case can an agreed-upon set of borders or a category boundary be factually incorrect.

IV.

I usually avoid arguing LGBT issues on here, not because I don't have strong opinions about them but because I assume so many of my readers already agree with me that it would be a waste of time. I'm pretty sure I'm right about this – on the recent survey, readers of this blog who were asked to rate their opinion of gay marriage from 1 (strongly against) to 5 (strongly in favor) gave an average rating of 4.32.

Nevertheless, I've seen enough anti-transgender comments recently that the issue might be worth a look.

In particular, I've seen one anti-transgender argument around that I take very seriously. The argument goes: we are rationalists. Our *entire shtick* is trying to believe what's actually true, not on what we wish were true, or what our culture tells us is true, or what it's popular to say is true. If a man thinks he's a woman, then we might (empathetically) wish he were a woman, other people might demand we call him a woman, and we might be much more popular if we say he's a woman. But if we're going to be rationalists who focus on believing what's actually true, then we've got to call him a man and take the consequences.

Thus Abraham Lincoln's famous riddle: "If you call a tail a leg, how many legs does a dog have?" And the answer: "Four - because a tail isn't a leg regardless of what you call it."

(if John Wilkes Booth had to suffer through that riddle, then I don't blame him)

I take this argument very seriously, because sticking to the truth really is important. But having taken it seriously, I think it's seriously wrong.

An alternative categorization system is not an error, and borders are not objectively true or false.

Just as we can come up with criteria for a definition of "planet", we can come up with a definition of "man". Absolutely typical men have Y chromosomes, have male genitalia, appreciate manly things like sports and lumberjackery, are romantically attracted to women, personally identify as male, wear male clothing like blue jeans, sing baritone in the opera, et cetera.

Some people satisfy some criteria of manhood and not others, in much the same way that Pluto satisfies only some criteria of planethood and whales satisfy only some criteria of mammalhood. For example, gay men might date other men and behave in effeminate ways. People with congenital androgen insensitivity syndrome might have female bodies, female external genitalia, and have been raised female their entire life, but when you look into their cells they have Y chromosomes.

Biologists defined by fiat that in cases of ambiguous animal grouping like whales, phylogenetics will be the tiebreaker. This was useful to resolve ambiguity, and it's worth sticking to as a Schelling point so everyone's using their words the same way, but it's kind of arbitrary and mostly based on biologists caring a lot about phylogenetics. If we let King Solomon make the decision, he might decide by fiat that whether animals lived in land or water would be the tiebreaker, since he's most interested in whether the animal is hunted on horseback or by boat.

Likewise, astronomers decided by fiat that something would be a planet if and only if it meets the three criteria of orbiting, round, and orbit-clearing. But here we have a pretty neat window into how these kinds of decisions take place - you can [read the history](#) of the International Astronomical Union meeting where they settled on the definition and learn about all the alternative proposals that were floated and rejected and which particular politics resulted in the present criteria being selected among all the different possibilities. Here it is *obvious* that the decision was by fiat.

Without the input of any prestigious astronomers at all, most people seem to assume that the ultimate tiebreaker in man vs. woman questions is presence of a Y chromosome. I'm not sure this is a very principled decision, because I expect most people would classify congenital androgen insensitivity patients (XY people whose bodies are insensitive to the hormone that makes them look male, and so end up looking 100% female their entire lives and often not even knowing they have the condition) as women.

The project of the transgender movement is to propose a switch from using chromosomes as a tiebreaker to using self-identification as a tiebreaker.

(This isn't actually the whole story - some of the more sophisticated people want to split "sex" and "gender", so that people who want to talk about what chromosomes they've got have a categorization system to do that with, and a few people even want

to split “chromosomal sex” and “anatomical sex” and “gender” and goodness knows what else – and I support all of these as very important examples of the virtue of precision – but to a first approximation, they want to define gender as self-identification)

This is not something that can be “true” or “false”. It’s a boundary-redrawing project. It can make for some boundaries that look a little bit weird – like a small percent of men being able to get pregnant – but as far as weird boundaries go that’s probably not as bad as having a tiny exclave of Turkish territory in the middle of a Syrian village.

(Ozy tells me this is sort of what queer theory is getting at, but in a horrible unreadable postmodernist way. They assure me you’re better off just reading the darned Sequences.)

You draw category boundaries in specific ways to capture tradeoffs you care about. If you care about the sanctity of the tomb of your country’s founder, sometimes it’s worth having a slightly weird-looking boundary in order to protect and honor it. And if you care about...

I’ve lived with a transgender person for six months, so I probably should have written this earlier. But I’m writing it now because I just finished accepting a transgender man to the mental hospital. He alternates between trying to kill himself and trying to cut off various parts of his body because he’s so distressed that he is biologically female. We’ve connected him with some endocrinologists who can hopefully get him started on male hormones, after which maybe he’ll stop doing that and [hopefully](#) be able to lead a normal life.

If I’m willing to accept an unexpected chunk of Turkey deep inside Syrian territory to honor some random dead guy – and I better, or else a platoon of Turkish special forces will want to have a word with me – then I ought to accept an unexpected man or two deep inside the conceptual boundaries of what would normally be considered female if it’ll save someone’s life. There’s no rule of rationality saying that I shouldn’t, and there are plenty of rules of human decency saying that I should.

V.

I’ve made this argument before and gotten a reply something like this:

“Transgender is a psychiatric disorder. When people have psychiatric disorders, certainly it’s right to sympathize and feel sorry for them and want to help them. But the way we try to help them is by treating their disorder, not by indulging them in their delusion.”

I think these people expect me to argue that transgender “isn’t really a psychiatric disorder” or something. But “psychiatric disorder” is just another category boundary dispute, and one that I’ve already [written enough about elsewhere](#). At this point, I don’t care enough to say much more than “If it’s a psychiatric disorder, then attempts to help transgender people get covered by health insurance, and most of the ones I know seem to want that, so sure, gender dysphoria is a psychiatric disorder.”

And then I think of the Hair Dryer Incident.

The Hair Dryer Incident was probably the biggest dispute I’ve seen in the mental hospital where I work. Most of the time all the psychiatrists get along and have pretty

much the same opinion about important things, but people were at each other's *throats* about the Hair Dryer Incident.

Basically, this one obsessive compulsive woman would drive to work every morning and worry she had left the hair dryer on and it was going to burn down her house. So she'd drive back home to check that the hair dryer was off, then drive back to work, then worry that maybe she hadn't *really* checked well enough, then drive back, and so on ten or twenty times a day.

It's a pretty typical case of obsessive-compulsive disorder, but it was really interfering with her life. She worked some high-powered job – I think a lawyer – and she was *constantly* late to everything because of this driving back and forth, to the point where her career was in a downspin and she thought she would have to quit and go on disability. She wasn't able to go out with friends, she wasn't even able to go to restaurants because she would keep fretting she left the hair dryer on at home and have to rush back. She'd seen countless psychiatrists, psychologists, and counselors, she'd done all sorts of therapy, she'd taken every medication in the book, and none of them had helped.

So she came to my hospital and was seen by a colleague of mine, who told her "Hey, have you thought about just bringing the hair dryer with you?"

And it *worked*.

She would be driving to work in the morning, and she'd start worrying she'd left the hair dryer on and it was going to burn down her house, and so she'd look at the seat next to her, and there would be the hair dryer, right there. And she only had the one hair dryer, which was now accounted for. So she would let out a sigh of relief and keep driving to work.

And approximately half the psychiatrists at my hospital thought this was *absolutely scandalous*, and This Is Not How One Treats Obsessive Compulsive Disorder, and what if it got out to the broader psychiatric community that instead of giving all of these high-tech medications and sophisticated therapies we were just telling people to *put their hair dryers on the front seat of their car*?

I, on the other hand, thought it was the best fucking story I had ever heard and the guy deserved a medal. Here's someone who was totally untreatable by the normal methods, with a debilitating condition, and a drop-dead simple intervention that nobody else had thought of gave her her life back. If one day I open up my own psychiatric practice, I am half-seriously considering using a picture of a hair dryer as the logo, just to let everyone know where I stand on this issue.

Miyamoto Musashi is quoted as saying:

The primary thing when you take a sword in your hands is your intention to cut the enemy, whatever the means. Whenever you parry, hit, spring, strike or touch the enemy's cutting sword, you must cut the enemy in the same movement. It is essential to attain this. If you think only of hitting, springing, striking or touching the enemy, you will not be able actually to cut him.

Likewise, the primary thing in psychiatry is to help the patient, whatever the means. Someone can concern-troll that the hair dryer technique leaves something to be desired in that it might have prevented the patient from seeking a more thorough cure

that would prevent her from having to bring the hair dryer with her. But compared to the alternative of “nothing else works” it seems clearly superior.

And that’s the position from which I think a psychiatrist should approach gender dysphoria, too.

Imagine if we could give depressed people a much higher quality of life merely by giving them cheap natural hormones. I don’t think there’s a psychiatrist in the world who wouldn’t celebrate that as one of the biggest mental health advances in a generation. Imagine if we could ameliorate schizophrenia with one safe simple surgery, just snip snip you’re not schizophrenic anymore. Pretty sure that would win *all* of the Nobel prizes. Imagine that we could make a serious dent in bipolar disorder just by calling people different pronouns. I’m pretty sure the entire mental health field would join together in bludgeoning anybody who refused to do that. We would bludgeon them over the head with big books about the side effects of lithium.

Really, are you *sure* you want your opposition to accepting transgender people to be “I think it’s a mental disorder”?

VI.

Some people can’t leave well enough alone, and continue to push the mental disorder angle. For example:



There are a lot of things I could say here.

I could point out that trans-Napoleonism seem to be mysteriously less common than transgender.

I could relate this mysterious difference to the various heavily researched [apparent biological correlates of transgender](#), including unusual variants of the androgen receptor, birth-sex-discordant sizes of various brain regions, birth-sex-discordant responses to various pheromones, high rates of something [seemingly like body integrity identity disorder](#), and of course our old friend altered digit ratios. If our hypothetical trans-Napoleon came out of the womb wearing a French military uniform and clutching a list of 19th century Grand Armee positions in his cute little baby hands, I think I’d take him more seriously.

I could argue that questions about gender are questions about category boundaries, whereas questions about Napoleon – absent some kind of philosophical legwork that I would very much like to read – are questions of fact.

I could point out that if the extent of somebody’s trans-Napoleonness was wanting to wear a bicorne hat, and he was going to be suicidal his entire life if he couldn’t but pretty happy if I could, let him wear the damn hat.

I could just link people to [other sites’](#) pretty good objections to the same argument.

But I think what I actually want to say is that there was once a time somebody tried pretty much exactly this, silly hat and all. Society shrugged and played along, he led a rich and fulfilling life, his grateful Imperial subjects came to love him, and it’s one of the most heartwarming episodes in the history of one of my favorite places in the world.

Sometimes when you make a little effort to be nice to people, even people you might think are weird, [really good things happen](#).

The noncentral fallacy - the worst argument in the world?

Related to: [Leaky Generalizations](#), [Replace the Symbol With The Substance](#), [Sneaking In Connotations](#)

David Stove once [ran a contest](#) to find the Worst Argument In The World, but he awarded the prize to his own entry, and one that shored up his politics to boot. It hardly seems like an objective process.

If he can unilaterally declare a Worst Argument, then so can I. I declare the Worst Argument In The World to be this: "X is in a category whose archetypal member gives us a certain emotional reaction. Therefore, we should apply that emotional reaction to X, even though it is not a central category member."

Call it the Noncentral Fallacy. It sounds dumb when you put it like that. Who even does that, anyway?

It sounds dumb only because we are talking soberly of categories and features. As soon as the argument gets framed in terms of *words*, it becomes so powerful that somewhere between many and most of the bad arguments in politics, philosophy and culture take some form of the noncentral fallacy. Before we get to those, let's look at a simpler example.

Suppose someone wants to build a statue honoring Martin Luther King Jr. for his nonviolent resistance to racism. An opponent of the statue objects: "But Martin Luther King was a *criminal*!"

Any historian can confirm this is correct. A criminal is technically someone who breaks the law, and King knowingly broke a law against peaceful anti-segregation protest - hence his famous Letter from Birmingham Jail.

But in this case calling Martin Luther King a criminal is the noncentral. The archetypal criminal is a mugger or bank robber. He is driven only by greed, preys on the innocent, and weakens the fabric of society. Since we don't like these things, calling someone a "criminal" naturally lowers our opinion of them.

The opponent is saying "Because you don't like criminals, and Martin Luther King is a criminal, you should stop liking Martin Luther King." But King doesn't share the important criminal features of being driven by greed, preying on the innocent, or weakening the fabric of society that made us dislike criminals in the first place. Therefore, even though he is a criminal, there is no reason to dislike King.

This all seems so nice and logical when it's presented in this format. Unfortunately, it's also one hundred percent contrary to instinct: the urge is to respond "Martin Luther King? A criminal? No he wasn't! You take that back!" This is why the noncentral is so successful. As soon as you do that you've fallen into their trap. Your argument is no longer about whether you should build a statue, it's about whether King was a criminal. Since he was, you have now lost the argument.

Ideally, you should just be able to say "Well, King was the good kind of criminal." But

that seems pretty tough as a debating maneuver, and it may be even harder in some of the cases where the noncentral Fallacy is commonly used.

Now I want to list some of these cases. Many will be political¹, [for which I apologize](#), but it's hard to separate out a bad argument from its specific instantiations. None of these examples are meant to imply that the position they support is wrong (and in fact I myself hold some of them). They only show that certain particular arguments for the position are flawed, such as:

"Abortion is murder!" The archetypal murder is Charles Manson breaking into your house and shooting you. This sort of murder is bad for a number of reasons: you prefer not to die, you have various thoughts and hopes and dreams that would be snuffed out, your family and friends would be heartbroken, and the rest of society has to live in fear until Manson gets caught. If you define murder as "killing another human being", then abortion is technically murder. But it has none of the downsides of murder Charles Manson style. Although you can criticize abortion for many reasons, insofar as "abortion is murder" is an invitation to apply one's feelings in the Manson case directly to the abortion case, it [ignores](#) the latter's lack of the features that generated those intuitions in the first place².

"Genetic engineering to cure diseases is eugenics!" Okay, you've got me there: since eugenics means "trying to improve the gene pool" that's clearly right. But what's wrong with eugenics? "What's wrong with eugenics? Hitler did eugenics! Those unethical scientists in the 1950s who sterilized black women without their consent did eugenics!" "And what was wrong with what Hitler and those unethical scientists did?" "What do you mean, what was wrong with them? Hitler killed millions of people! Those unethical scientists ruined people's lives." "And does using genetic engineering to cure diseases kill millions of people, or ruin anyone's life?" "Well...not really." "Then what's wrong with it?" "It's *eugenics*!"

"Evolutionary psychology is sexist!" If you define "sexist" as "believing in some kind of difference between the sexes", this is true of at least some evo psych. For example, [Bateman's Principle](#) states that in species where females invest more energy in producing offspring, mating behavior will involve males pursuing females; this posits a natural psychological difference between the sexes. "Right, so you admit it's sexist!" "And why exactly is sexism bad?" "Because sexism claims that men are better than women and that women should have fewer rights!" "Does Bateman's principle claim that men are better than women, or that women should have fewer rights?" "Well...not really." "Then what's wrong with it?" "It's *sexist*!"

A second, subtler use of the noncentral fallacy goes like this: "X is in a category whose archetypal member gives us an emotional reaction. Therefore, we should apply that same emotional reaction to X even if X gives some benefit that outweighs the harm."

"Capital punishment is murder!" Charles Manson-style murder is solely harmful. This kind of murder produces really strong negative feelings. The proponents of capital punishment believe that it might decrease crime, or have some other attending benefits. In other words, they believe it's "the good kind of murder"³, just like the introductory example concluded that Martin Luther King was "the good kind of criminal". But since normal murder is so taboo, it's really hard to take the phrase "the good kind of murder" seriously, and just mentioning the word "murder" can call up exactly the same amount of negative feelings we get from the textbook example.

"Affirmative action is racist!" True if you define racism as "favoring certain people based on their race", but once again, our immediate negative reaction to the archetypal example of racism (the Ku Klux Klan) cannot be generalized to an immediate negative reaction to affirmative action. Before we generalize it, we have to check first that the problems that make us hate the Ku Klux Klan (violence, humiliation, divisiveness, lack of a meritocratic society) are still there. Then, even if we do find that some of the problems persist (like disruption of meritocracy, for example) we have to prove that it doesn't produce benefits that outweigh these harms.

"Taxation is theft!" True if you define theft as "taking someone else's money regardless of their consent", but though the archetypal case of theft (breaking into someone's house and stealing their jewels) has nothing to recommend it, taxation (arguably) does. In the archetypal case, theft is both unjust and socially detrimental. Taxation keeps the first disadvantage, but arguably subverts the second disadvantage if you believe being able to fund a government has greater social value than leaving money in the hands of those who earned it. The question then hinges on the relative importance of these disadvantages. Therefore, you can't dismiss taxation without a second thought just because you have a natural disgust reaction to theft in general. You would also have to prove that the supposed benefits of this form of theft don't outweigh the costs.

Now, because most arguments are rapid-fire debate-club style, sometimes it's still useful to say "Taxation isn't theft!" At least it beats saying "Taxation is theft but nevertheless good", then having the other side say "Apparently my worthy opponent thinks that theft can be good; we here on this side would like to bravely take a stance *against* theft", and then having the moderator call time before you can explain yourself. If you're in a debate club, do what you have to do. But if you have the luxury of philosophical clarity, you would do better to forswear the [Dark Arts](#) and look a little deeper into what's going on.

Are there ever cases in which this argument pattern can be useful? Yes. For example, it may be a groping attempt to suggest a [Schelling fence](#); for example, a principle that one must never commit theft even when it would be beneficial because that would make it harder to distinguish and oppose the really bad kinds of theft. Or it can be an attempt to spark conversation by pointing out a potential contradiction: for example "Have you noticed that taxation really does contain some of the features you dislike about more typical instances of theft? Maybe you never even thought about that before? Why do your moral intuitions differ in these two cases? Aren't you being kind of hypocritical?" But this usage seems pretty limited - once your interlocutor says "Yes, I considered that, but the two situations are different for reasons X, Y, and Z" the conversation needs to move on; there's not much point in continuing to insist "But it's theft!"

But in most cases, I think this is more of an *emotional* argument, or even an argument from "You would look silly saying that". You really *can't* say "Oh, he's the good kind of criminal", and so if you have a potentially judgmental audience and not much time to explain yourself, you're pretty trapped. You have been forced to round to the archetypal example of that word and subtract exactly the information that's most relevant.

But in all other cases, the proper response to being asked to subtract relevant information is "No, why should I?" - and that's why this is the worst argument in the world.

Footnotes

1: On advice from the community, I have deliberately included three mostly-liberal examples and three-mostly conservative examples, so save yourself the trouble of counting them up and trying to speculate on this article's biases.

2: This should be distinguished from deontology, the belief that there is some provable moral principle about how you can never murder. I don't think this is *too* important a point to make, because only a tiny fraction of the people who debate these issues have thought that far ahead, and also because my personal and admittedly controversial opinion is that much of deontology is just an attempt to formalize and justify this fallacy.

3: Some people "solve" this problem by saying that "murder" only refers to "non-lawful killing", which is exactly as creative a solution as redefining "criminal" to mean "person who breaks the law and is not Martin Luther King." Identifying the noncentral fallacy is a more complete solution: for example, it covers the related (mostly sarcastic) objection that "imprisonment is kidnapping".

4: EDIT 8/2013: I've edited this article a bit after getting some feedback and complaints. In particular I tried to remove some LW jargon which turned off some people who were being linked to this article but were unfamiliar with the rest of the site.

5: EDIT 8/2013: The other complaint I kept getting is that this is an uninteresting restatement of some other fallacy (no one can agree which, but [poisoning the well](#) comes up particularly often). The question doesn't seem too interesting to me - I never claimed particular originality, a lot of fallacies blend into each other, and the which-fallacy-is-which game isn't too exciting anyway - but for the record I don't think it is. Poisoning the well is a presentation of two different facts, such as "Martin Luther King was a plagiarist...oh, by the way, what do you think of Martin Luther King's civil rights policies?" It may have no relationship to categories, and it's usually something someone else does to you as a conscious rhetorical trick. Noncentral fallacy is presenting a single fact, but using category information to frame it in a misleading way - and it's often something people do to themselves. The above plagiarism example of poisoning the well is *not* noncentral fallacy. If you think this essay is about bog-standard poisoning the well, then either there is an alternative meaning to poisoning the well I'm not familiar with, or you are missing the point.

Ethnic Tension And Meaningless Arguments

I.

Part of what bothers me – and apparently several others – about [yesterday's motte-and-bailey discussion](#) is that here's a fallacy – a pretty successful fallacy – that depends entirely on people not being entirely clear on what they're arguing about. Somebody says God doesn't exist. Another person objects that God is just a name for the order and beauty in the universe. Then this somehow helps defend the position that God is a supernatural creator being. How does that even happen?

"Sir, you've been accused of murdering your wife. We have three witnesses who said you did it. What do you have to say for yourself?"

"Well, your honor, I think it's quite clear I didn't murder the President. For one thing, he's surrounded by Secret Service agents. For another, check the news. The President's still alive."

"Huh. For some reason I vaguely remember thinking you didn't have a case. Yet now that I hear you talk, everything you say is incredibly persuasive. You're free to go."

While motte-and-bailey is less subtle, it seems to require a similar sort of misdirection. I'm not saying it's impossible. I'm just saying it's a fact that needs to be explained.

When everything works the way it's supposed to in philosophy textbooks, arguments are supposed to go one of a couple of ways:

1. Questions of empirical fact, like "Is the Earth getting warmer?" or "Did aliens build the pyramids?". You debate these by presenting factual evidence, like "An average of global weather station measurements show 2014 is the hottest year on record" or "One of the bricks at Giza says 'Made In Tau Ceti V' on the bottom." Then people try to refute these facts or present facts of their own.

2. Questions of morality, like "Is it wrong to abort children?" or "Should you refrain from downloading music you have not paid for?" You can only debate these *well* if you've already agreed upon a moral framework, like a particular version of natural law or consequentialism. But you can *sort of* debate them by comparing to examples of agreed-upon moral questions and trying to maintain consistency. For example, "You wouldn't kill a one day old baby, so how is a nine month old fetus different?" or "You wouldn't download a *car*."

If you are very lucky, your philosophy textbook will also admit the existence of:

3. Questions of policy, like "We should raise the minimum wage" or "We should bomb Foreignistan". These are combinations of competing factual claims and competing values. For example, the minimum wage might hinge on factual claims like "Raising the minimum wage would increase unemployment" or "It is very difficult to live on the minimum wage nowadays, and many poor families cannot afford food." But it might also hinge on value claims like "Corporations owe it to their workers to pay a living wage," or "It is more important that the poorest be protected than that the economy be strong." Bombing Foreignistan might depend on factual claims like "The

Foreignistanis are harboring terrorists”, and on value claims like “The safety of our people is worth the risk of collateral damage.” If you can resolve all of these factual and value claims, you should be able to agree on questions of policy.

None of these seem to allow the sort of vagueness of topic mentioned above.

II.

A question: are you pro-Israel or pro-Palestine? Take a second, actually think about it.

Some people probably answered pro-Israel. Other people probably answered pro-Palestine. Other people probably said they were neutral because it’s a complicated issue with good points on both sides.

Probably very few people answered: *Huh? What?*

This question doesn’t fall into any of the three Philosophy 101 forms of argument. It’s not a question of fact. It’s not a question of particular moral truths. It’s not even a question of policy. There are closely related policies, like whether Palestine should be granted independence. But if I support a very specific two-state solution where the border is drawn upon the somethingth parallel, does that make me pro-Israel or pro-Palestine? At exactly which parallel of border does the solution under consideration switch from pro-Israeli to pro-Palestinian? Do you think the crowd of people shouting and waving signs saying “SOLIDARITY WITH PALESTINE” have an answer to that question?

But it’s even worse, because this question covers much more than just the borders of an independent Palestinian state. Was Israel justified by responding to Hamas’ rocket fire by bombing Gaza, even with the near-certainty of collateral damage? Was Israel justified in building a wall across the Palestinian territories to protect itself from potential terrorists, even though it severely curtails Palestinian freedom of movement? Do Palestinians have a “right of return” to territories taken in the 1948 war? Who should control the Temple Mount?

These are four very different questions which one would think each deserve independent consideration. But in reality, what percent of the variance in people’s responses do you think is explained by a general “pro-Palestine vs. pro-Israel” factor? 50%? 75%? More?

In a way, when we round people off to the Philosophy 101 kind of arguments, we are failing to respect their self-description. People aren’t out on the streets saying “By my cost-benefit analysis, Israel was in the right to invade Gaza, although it may be in the wrong on many of its other actions.” They’re waving little Israeli flags and holding up signs saying “ISRAEL: OUR STAUNCHEST ALLY”. Maybe we should take them at face value.

This is starting to look related to the original question in (I). Why is it okay to suddenly switch points in the middle of an argument? In the case of Israel and Palestine, it might be because people’s support for any particular Israeli policy is better explained by a General Factor Of Pro-Israeliness than by the policy itself. As long as I’m arguing in favor of Israel in *some way*, it’s still considered by everyone to be on topic.

III.

Some moral philosophers got fed up with nobody being able to explain what the heck a moral truth was and invented emotivism. Emotivism says there *are* no moral truths, just expressions of little personal bursts of emotion. When you say “Donating to charity is good,” you don’t mean “Donating to charity increases the sum total of utility in the world,” or “Donating to charity is in keeping with the Platonic moral law” or “Donating to charity was commanded by God” or even “I like donating to charity”. You’re just saying “Yay charity!” and waving a little flag.

Seems a lot like how people handle the Israel question. “I’m pro-Israel” doesn’t necessarily imply that you believe any empirical truths about Israel, or believe any moral principles about Israel, or even support any Israeli policies. It means you’re waving a little flag with a Star of David on it and cheering.

So here is Ethnic Tension: A Game For Two Players.

Pick a vague concept. “Israel” will do nicely for now.

Player 1 tries to associate the concept “Israel” with as much good karma as she possibly can. Concepts get good karma by doing good moral things, by being associated with good people, by being linked to the beloved in-group, and by being oppressed underdogs [in bravery debates](#).

“Israel is the freest and most democratic country in the Middle East. It is one of America’s strongest allies and shares our Judeo-Christian values.

Player 2 tries to associate the concept “Israel” with as much bad karma as she possibly can. Concepts get bad karma by committing atrocities, being associated with bad people, being linked to the hated out-group, and by being oppressive big-shots in bravery debates. Also, she obviously needs to neutralize Player 1’s actions by disproving all of her arguments.

“Israel may have some level of freedom for its most privileged citizens, but what about the millions of people in the Occupied Territories that have no say? Israel is involved in various atrocities and has often killed innocent protesters. They are essentially a neocolonialist state and have allied with other neocolonialist states like South Africa.”

The prize for winning this game is the ability to win the other three types of arguments. If Player 1 wins, the audience ends up with a strongly positive General Factor Of Pro-Israeliness, and vice versa.

Remember, people’s capacity for [motivated reasoning](#) is pretty much infinite. Remember, a [motivated skeptic](#) asks if the evidence *compels* them to accept the conclusion; a motivated credulist asks if the evidence *allows* them to accept the conclusion. Remember, Jonathan Haidt and his team [hypnotized](#) people to have strong disgust reactions to the word “often”, and then tried to hold in their laughter when people in the lab came up with convoluted yet plausible-sounding arguments against any policy they proposed that included the word “often” in the description.

I’ve never heard of the experiment being done the opposite way, but it sounds like the sort of thing that might work. Hypnotize someone to have a very positive reaction to the word “often” (for most hilarious results, have it give people an orgasm). “Do you think governments should raise taxes more often?” “Yes. Yes yes YES YES OH GOD YES!”

Once you finish the Ethnic Tension Game, you're replicating Haidt's experiment with the word "Israel" instead of the word "often". Win the game, and any pro-Israel policy you propose will get a burst of positive feelings and tempt people to try to find some explanation, any explanation, that will justify it, whether it's invading Gaza or building a wall or controlling the Temple Mount.

So this is the fourth type of argument, the kind that doesn't make it into Philosophy 101 books. The [trope namer](#) is Ethnic Tension, but it applies to anything that can be identified as a Vague Concept, or paired opposing Vague Concepts, which you can use emotivist thinking to load with good or bad karma.

IV.

Now motte-and-bailey stands revealed:

Somebody says God doesn't exist. Another person objects that God is just a name for the order and beauty in the universe. Then this somehow helps defend the position that God is a supernatural creator being. How does that even happen?

The two-step works like this. First, load "religion" up with good karma by pitching it as persuasively as possible. "Religion is just the belief that there's beauty and order in the universe."

Wait, / think there's beauty and order in the universe!

"Then you're religious too. We're all religious, in the end, because religion is about the common values of humanity and meaning and compassion sacrifice beauty of a sunrise Gandhi Buddha Sufis St. Francis awe complexity humility wonder Tibet the Golden Rule love."

Then, once somebody has a strongly positive General Factor Of Religion, it doesn't really matter whether someone believes in a creator God or not. If they have any predisposition whatsoever to do so, they'll find a reason to let themselves. If they can't manage it, they'll say it's true "metaphorically" and continue to act upon every corollary of it being true.

("God is just another name for the beauty and order in the universe. But Israel definitely belongs to the Jews, because the beauty and order of the universe promised it to them.")

If you're an atheist, you probably have a lot of important issues on which you want people to consider non-religious answers and policies. And if somebody can maintain good karma around the "religion" concept by believing God is the order and beauty in the universe, then that can still be a victory for religion even if it is done by jettisoning many traditionally "religious" beliefs. In this case, it is useful to think of the "order and beauty" formulation as a "motte" for the "supernatural creator" formulation, since it's allowing the *entire concept* to be defended.

But even this is giving people too much credit, because the existence of God is a (sort of) factual question. From yesterday's post:

Suppose we're debating feminism, and I defend it by saying it really is important that women are people, and you attack it by saying that it's not true that all men are terrible. What is the real feminism we should be debating? Why would you even ask that question? What is this, some kind of dumb high school debate club?

Who the heck thinks it would be a good idea to say 'Here's a vague poorly-defined concept that mind-kills everyone who touches it - quick, should you associate it with positive affect or negative affect?!'

Who the heck thinks that? Everybody, all the time.

Once again, if I can load the concept of "feminism" with good karma by making it so obvious nobody can disagree with it, then I have a massive "home field advantage" when I'm trying to convince anyone of any particular policy that can go under the name "feminism", even if it's unrelated to the arguments that gave feminism good karma in the first place.

Or if I'm against feminism, I just post quotes from the ten worst feminists on Tumblr again and again until the entire movement seems ridiculous and evil, and then you'll have trouble convincing anyone of *anything* feminist. "That seems reasonable...but wait, isn't that a feminist position? Aren't those the people I hate?"

(compare: [most Americans](#) oppose Obamacare, but most Americans support each individual component of Obamacare when it is explained without using the word "Obamacare")

V.

Little flow diagram things make everything better. Let's make a little flow diagram thing.

We have our node "Israel", which has either good or bad karma. Then there's another node close by marked "Palestine". We would expect these two nodes to be pretty anti-correlated. When Israel has strong good karma, Palestine has strong bad karma, and vice versa.

Now suppose you listen to Noam Chomsky talk about how strongly he supports the Palestinian cause and how much he dislikes Israel. One of two things can happen:

"Wow, a great man such as Noam Chomsky supports the Palestinians! They must be very deserving of support indeed!"

or

"That idiot Chomsky supports Palestine? Well, screw him. And screw them!"

So now there is a third node, Noam Chomsky, that connects to both Israel and Palestine, and we have discovered it is positively correlated with Palestine and negatively correlated with Israel. It probably has a pretty low weight, because there are a lot of reasons to care about Israel and Palestine other than Chomsky, and a lot of reasons to care about Chomsky other than Israel and Palestine, but the connection is there.

I don't know anything about neural nets, so maybe this system isn't actually a neural net, but whatever it is I'm thinking of, it's a structure where eventually the three nodes reach some kind of equilibrium. If we start with someone liking Israel and Chomsky, but not Palestine, then either that's going to shift a little bit towards liking Palestine, or shift a little bit towards disliking Chomsky.

Now we add more nodes. Cuba seems to really support Palestine, so they get a positive connection with a little bit of weight there. And I think Noam Chomsky supports Cuba, so we'll add a connection there as well. Cuba is socialist, and that's one of the most salient facts about it, so there's a heavily weighted positive connection between Cuba and socialism. Palestine kind of makes noises about socialism but I don't think they have any particular economic policy, so let's say very weak direct connection. And Che is heavily associated with Cuba, so you get a pretty big Che - Cuba connection, plus a strong direct Che - socialism one. And those pro-Palestinian students who threw rotten fruit at an Israeli speaker also get a little path connecting them to "Palestine" - hey, why not - so that if you support Palestine you might be willing to excuse what they did and if you oppose them you might be a little less likely to support Palestine.

Back up. This model produces crazy results, like that people who like Che are more likely to oppose Israel bombing Gaza. That's such a weird, implausible connection that it casts doubt upon the entire...

Oh. Wait. Yeah. Okay.

I think this kind of model, in its efforts to sort itself out into a ground state, might settle on some kind of General Factor Of Politics, which would probably correspond pretty well to the left-right axis.

In [Five Case Studies On Politicization](#), I noted how fresh new unpoliticized issues, like the Ebola epidemic, were gradually politicized by connecting them to other ideas that were already part of a political narrative. For example, a quarantine against Ebola would require closing the borders. So now there's a weak negative link between "Ebola quarantine" and "open borders". If your "open borders" node has good karma, now you're a little less likely to support an Ebola quarantine. If "open borders" has bad karma, a little more likely.

I also tried to point out how you could make different groups support different things by changing your narrative a little:

Global warming has gotten inextricably tied up in the Blue Tribe narrative: Global warming proves that unrestrained capitalism is destroying the planet. Global warming disproportionately affects poor countries and minorities. Global warming could have been prevented with multilateral action, but we were too dumb to participate because of stupid American cowboy diplomacy. Global warming is an important cause that activists and NGOs should be lauded for highlighting. Global warming shows that Republicans are science denialists and probably all creationists. Two lousy sentences on "patriotism" aren't going to break through that.

If I were in charge of convincing the Red Tribe to line up behind fighting global warming, here's what I'd say:

In the 1950s, brave American scientists shunned by the climate establishment of the day discovered that the Earth was warming as a result of greenhouse gas emissions, leading to potentially devastating natural disasters that could destroy American agriculture and flood American cities. As a result, the country mobilized against the threat. Strong government action by the Bush administration outlawed the worst of these gases, and brilliant entrepreneurs were able to discover and manufacture new cleaner energy sources. As a result of these brave decisions, our emissions stabilized and are currently declining.

Unfortunately, even as we do our part, the authoritarian governments of Russia and China continue to industrialize and militarize rapidly as part of their bid to challenge American supremacy. As a result, Communist China is now by far the world's largest greenhouse gas producer, with the Russians close behind. Many analysts believe Putin secretly welcomes global warming as a way to gain access to frozen Siberian resources and weaken the more temperate United States at the same time. These countries blow off huge disgusting globs of toxic gas, which effortlessly cross American borders and disrupt the climate of the United States. Although we have asked them to stop several times, they refuse, perhaps egged on by major oil producers like Iran and Venezuela who have the most to gain by keeping the world dependent on the fossil fuels they produce and sell to prop up their dictatorships.

We need to take immediate action. While we cannot rule out the threat of military force, we should start by using our diplomatic muscle to push for firm action at top-level summits like the Kyoto Protocol. Second, we should fight back against the liberals who are trying to hold up this important work, from big government bureaucrats trying to regulate clean energy to celebrities accusing people who believe in global warming of being 'racist'. Third, we need to continue working with American industries to set an example for the world by decreasing our own emissions in order to protect ourselves and our allies. Finally, we need to punish people and institutions who, instead of cleaning up their own carbon, try to parasitize off the rest of us and expect the federal government to do it for them.

In the first paragraph, "global warming" gets positively connected to concepts like "poor people and minorities" and "activists and NGOs", and gets negatively connected to concepts like "capitalism", "American cowboy diplomacy", and "creationists". That gives global warming really strong good karma if (and only if) you like the first two concepts and hate the last three.

In the next three paragraphs, "global warming" gets positively connected to "America", "the Bush administration" and "entrepreneurs", and negatively connected to "Russia", "China", "oil producing dictatorships like Iran and Venezuela", "big government bureaucrats", and "welfare parasites". This is going to appeal to, well, a different group.

Notice two things here. First, the exact connection isn't that important, as long as we can hammer in the existence of a connection. I could probably just say GLOBAL WARMING! COMMUNISM! GLOBAL WARMING! COMMUNISM! GLOBAL WARMING! COMMUNISM! several hundred times and have the same effect if I could get away with it (this is the principle behind attack ads which link a politician's face to scary music and a very concerned voice).

Second, there is no attempt whatsoever to challenge the idea that the issue at hand is the positive or negative valence of a concept called "global warming". At no point is it debated what the solution is, which countries the burden is going to fall on, or whether any particular level of emission cuts would do more harm than good. It's just accepted as obvious by both sides that we debate "for" or "against" global warming, and if the "for" side wins then they get to choose some solution or other or whatever oh god that's so boring can we get back to Israel vs. Palestine.

Some of the scientists working on IQ have started talking about "hierarchical factors", meaning that there's a general factor of geometry intelligence partially correlated

with other things into a general factor of mathematical intelligence partially correlated with other things into a general factor of total intelligence.

I would expect these sorts of things to work the same way. There's a General Factor Of Global Warming that affects attitudes toward pretty much all proposed global warming solutions, which is very highly correlated with a lot of other things to make a General Factor Of Environmentalism, which itself is moderately highly correlated with other things into the General Factor Of Politics.

VI.

Speaking of politics, a fruitful digression: what the heck was up with the Ashley Todd mugging hoax in 2008?

Back in the 2008 election, a McCain campaigner [claimed](#) (falsely, it would later turn out) to have been assaulted by an Obama supporter. She said he slashed a "B" (for "Barack") on her face with a knife. This got a lot of coverage, and according to Wikipedia:

John Moody, executive vice president at Fox News, commented in a blog on the network's website that "this incident could become a watershed event in the 11 days before the election," but also warned that "if the incident turns out to be a hoax, Senator McCain's quest for the presidency is over, forever linked to race-baiting."

Wait. One Democrat, presumably not acting on Obama's direct orders, attacks a Republican woman. And this is supposed to *alter the outcome of the entire election*? In what universe does one crime by a deranged psychopath change whether Obama's tax policy or job policy or bombing-scary-foreigners policy is better or worse than McCain's?

Even *if* we're willing to make the irresponsible leap from "Obama is supported by psychopaths, therefore he's probably a bad guy," there are like a hundred million people on each side. Psychopaths are usually estimated at about 1% of the population, so any movement with a million people will already have 10,000 psychopaths. Proving the existence of a single one changes *nothing*.

I think insofar as this affected the election – and everyone seems to have agreed that it might have – it hit President Obama with a burst of bad karma. Obama something something psychopath with a knife. Regardless of the exact content of those something somethings, *is that the kind of guy you want to vote for?*

Then when it was discovered to be a hoax, it was McCain something something race-baiting hoaxer. Now *he's* got the bad karma!

This sort of conflation between a cause and its supporters really only makes sense in the emotivist model of arguing. I mean, this shouldn't even get dignified with the name *ad hominem* fallacy. Ad hominem fallacy is "McCain had sex with a goat, therefore whatever he says about taxes is invalid." At least it's still the same *guy*. This is something the philosophy textbooks can't bring themselves to believe really exists, even as a fallacy.

But if there's a General Factor Of McCain, then anything bad remotely connected to the guy – goat sex, lying campaigners, whatever – reflects on everything else about him.

This is the same pattern we see in Israel and Palestine. How many times have you seen a news story like this one: "Israeli speaker hounded off college campus by pro-Palestinian partisans throwing fruit. Look at the intellectual bankruptcy of the pro-Palestinian cause!" It's clearly intended as an argument for *something* other than just not throwing fruit at people. The causation seems to go something like "These particular partisans are violating the usual norms of civil discussion, therefore they are bad, therefore something associated with Palestine is bad, therefore your General Factor of Pro-Israeliness should become more strongly positive, therefore it's okay for Israel to bomb Gaza." Not usually said in those *exact words*, but the thread can be traced.

VII.

Here is a prediction of this model: we will be obsessed with what concepts we can connect to other concepts, even when the connection is totally meaningless.

Suppose I say: "Opposing Israel is anti-Semitic". Why? Well, the Israelis are mostly Jews, so in a sense by definition being anti- them is "anti-Semitic", broadly defined. Also, $p(\text{opposes Israel} | \text{is anti-Semitic})$ is probably pretty high, which sort of lends some naive plausibility to the idea that $p(\text{is anti-Semitic} | \text{opposes Israel})$ is at least higher than it otherwise *could* be.

Maybe we do our research and we find exactly what percent of opponents of Israel endorse various anti-Semitic statements like "I hate all Jews" or "Hitler had some bright ideas". We've [replaced the symbol with the substance](#). Problem solved, right?

Maybe not. In the same sense that people can agree on all of the characteristics of Pluto – its diameter, the eccentricity of its orbit, its number of moons – and still disagree on the question "Is Pluto a planet", one can agree on every characteristic of every Israel opponent and still disagree on the definitional question "Is opposing Israel anti-Semitic?"

(fact: it wasn't until proofreading this essay that I realized I had originally written "Is Israel a planet?" and "Is opposing Pluto anti-Semitic?" I would like to see Jonathan Haidt hypnotize people until they can come up with positive arguments for those propositions.)

What's the point of this useless squabble [over definitions](#)?

I think it's about drawing a line between the concept "anti-Semitism" and "oppose Israel". If your head is screwed on right, you assign anti-Semitism some very bad karma. So if we can stick a thick line between "anti-Semitism" and "oppose Israel", then you're going to have very bad feelings about opposition to Israel and your General Factor Of Pro-Israeliness will go up.

Notice that this model *is transitive, but shouldn't be*.

That is, let's say we're arguing over the definition of anti-Semitism, and I say "anti-Semitism just means anything that hurts Jews". This is a dumb definition, but let's roll with it.

First, I load "anti-Semitism" with lots of negative affect. Hitler was anti-Semitic. The pogroms in Russia were anti-Semitic. The Spanish Inquisition was anti-Semitic. Okay, negative affect achieved.

Then I connect “wants to end the Israeli occupation of Palestine” to “anti-Semitism”. Now wanting to end the Israeli occupation of Palestine has lots of negative affect attached to it.

It sounds dumb when you put it like that, but when you put it like “You’re anti-Semitic for wanting to end the occupation” it’s a pretty damaging argument.

This is *trying* to be transitive. It’s trying to say “anti-occupation = anti-Semitism, anti-Semitism = evil, therefore anti-occupation = evil”. If this were arithmetic, it would work. But there’s no Transitive Property Of Concepts. If anything, concepts are more like sets. The logic is “anti-occupation is a member of the set anti-Semitic, the set anti-Semitic contains members that are evil, therefore anti-occupation is evil”, which obviously doesn’t check out.

(compare: “I am a member of the set ‘humans’, the set ‘humans’ contains the Pope, therefore I am the Pope”.)

Anti-Semitism is generally considered evil because a lot of anti-Semitic things involve killing or dehumanizing Jews. Opposing the Israel occupation of Palestine doesn’t kill or dehumanize Jews, so even if we call it “anti-Semitic” by definition, there’s no reason for our usual bad karma around anti-Semitism to transfer over. But by an unfortunate rhetorical trick, it does – you can gather up bad karma into “anti-Semitic” and then shoot it at the “occupation of Palestine” issue just by clever use of definitions.

This means that if you can come up with sufficiently clever definitions and convince your opponent to accept them, you can win any argument by default just by having a complex system of mirrors in place to reflect bad karma from genuinely evil things to the things you want to tar as evil. This is essentially the point I make in [Words, Words, Words](#).

If we kinda tweak the definition of “anti-Semitism” to be “anything that inconveniences Jews”, we can pull a trick where we leverage people’s dislike of Hitler to make them support the Israeli occupation of Palestine – but in order to do that, we need to get everyone on board with our *slightly* non-standard definition. Likewise, the social justice movement insists on their own novel definitions of words like “racism” that don’t match common usage, any dictionary, or etymological history – but which do perfectly describe a mirror that reflects bad karma toward opponents of social justice while making it impossible to reflect any bad karma back. Overreliance on this mechanism explains why so many social justice debates end up being about whether a particular mirror can be deployed to transfer bad karma in a specific case (“are trans people privileged?!”) rather than any feature of the real world.

But they are hardly alone. Compare: “Is such an such an organization a *cult*?”, “Is such and such a policy *socialist*?”, “Is abortion or capital punishment or war *murder*?” All entirely about whether we’re allowed to reflect bad karma from known sources of evil to other topics under discussion.

Look around you. Just look around you. Have you worked out what we’re looking for? Correct. The answer is [The Worst Argument In The World](#). Only now, we can explain why it works.

VIII.

From the self-esteem literature, I gather that the self is also a concept that can have good or bad karma. From the cognitive dissonance literature, I gather that the self is actively involved in maintaining good karma around itself through as many biases as it can manage to deploy.

I've mentioned [this study](#) before. Researchers make ~~victims~~ participants fill out a questionnaire about their romantic relationships. Then they pretend to "grade" the questionnaire, actually assigning scores at random. Half the participants are told their answers indicate they have the tendency to be very faithful to their partner. The other half are told they have very low faithfulness and their brains just aren't built for fidelity. Then they ask the ~~participants~~ victims their opinion on staying faithful in a relationship – very important, moderately important, or not so important?

There is a strong signal of people who are told they are bad at fidelity to state fidelity is unimportant, and another strong signal of people who are told they are especially faithful stating that fidelity is a great and noble virtue that must be protected.

The researchers conclude that people want to have high self-esteem. If I am terrible at fidelity, and fidelity is the most important virtue, that makes me a terrible person. If I am terrible at fidelity and fidelity doesn't matter, I'm fine. If I am great at fidelity, and fidelity is the most important virtue, I can feel pretty good about myself.

This doesn't seem too surprising. It's just the more subtle version of the effect where white people are a lot more likely to be white supremacists than members of any other race. Everyone likes to hear that they're great. The question is whether they can defend it and fit it in with their other ideas. The answer is "usually yes, because people are capable of pretty much any contortion of logic you can imagine and a lot that you can't".

I had a bad experience when I was younger where a bunch of feminists attacked and threatened me because of something I wrote. It left me kind of scarred. More importantly, the shape of that scar was a big anticorrelated line between self-esteem and the "feminism" concept. If feminism has lots of good karma, then I have lots of bad karma, because I am a person feminists hate. If feminists have lots of bad karma, then I look good by comparison, the same way it's pretty much a badge of honor to be disliked by Nazis. The result was a permanent haze of bad karma around "feminism" unconnected to any specific feminist idea, which I have to be constantly on the watch for if I want to be able to evaluate anything related to feminism fairly or rationally.

Good or bad karma, when applied to yourself, looks like high or low self-esteem; when applied to groups, it looks like high or low status. In the giant muddle of a war for status that we politely call "society", this makes beliefs into weapons and the karma loading of concepts into the difference between lionization and dehumanization.

The Trope Namer for emotivist arguments is "ethnic tension", and although it's most obvious in the case of literal ethnicities like the Israelis and the Palestinians, the ease with which concepts become attached to different groups creates a whole lot of "proxy ethnicities". I've [written before](#) about how American liberals and conservatives are seeming less and less like people who happen to have different policy prescriptions, and more like two different tribes engaged in an ethnic conflict quickly approaching Middle East level hostility. More recently, a friend on Facebook described the-thing-whose-name-we-do-not-speak-lest-it-appear-and-destroy-us-all, the one involving reproductively viable worker ants, as looking more like an ethnic conflict about who is oppressing whom than any real difference in opinions.


Once a concept has joined up with an ethnic group, either a real one or a makeshift one, it's impossible to oppose the concept without simultaneously lowering the status of the ethnic group, which is going to start at least a *little* bit of a war. Worse, once a concept has joined up with an ethnic group, one of the best ways to argue against the concept is to dehumanize the ethnic group it's working with. Dehumanizing an ethnic group has always been easy – just associate them with a disgust reaction, [portray](#) them as conventionally unattractive and unlovable and full of all the worst human traits – and now it is profitable as well, since it's one of the fastest ways to load bad karma into an idea you dislike.

IX.

According to [The Virtues Of Rationality](#):

The tenth virtue is precision. One comes and says: The quantity is between 1 and 100. Another says: the quantity is between 40 and 50. If the quantity is 42 they are both correct, but the second prediction was more useful and exposed itself to a stricter test. What is true of one apple may not be true of another apple; thus more can be said about a single apple than about all the apples in the world. The narrowest statements slice deepest, the cutting edge of the blade. As with the map, so too with the art of mapmaking: The Way is a precise Art. Do not walk to the truth, but dance. On each and every step of that dance your foot comes down in exactly the right spot. Each piece of evidence shifts your beliefs by exactly the right amount, neither more nor less. What is exactly the right amount? To calculate this you must study probability theory. Even if you cannot do the math, knowing that the math exists tells you that the dance step is precise and has no room in it for your whims.

The official description is of literal precision, as specific numerical precision in probability updates. But is there a secret interpretation of this virtue?



Precision as separation. Once you're debating "religion", you've already lost. Precision as sticking to a precise question, like "Is the first chapter of Genesis literally true?" or "Does Buddhist meditation help treat anxiety disorders?" and trying to keep these issues as separate from any General Factor Of Religiousness as humanly possible. Precision such that "God the supernatural Creator exists" and "God the order and

beauty in the Universe exists” are as carefully sequestered from one another as “Did the defendant kill his wife?” and “Did the defendant kill the President?”

I want to end by addressing a point a commenter made in my last post on motte-and-bailey:

In the real world, the particular abstract questions aren’t what matter – the groups and people are what matter. People get things done, and they aren’t particularly married to particular abstract concepts, they are married to their values and their compatriots. In order to deal with reality, we must attack and defend groups and individuals. That does not mean forsaking logic. It requires dealing with obfuscating tactics like those you outline above, but that’s not even a real downside, because if you flee into the narrow, particular questions all you’re doing is covering your eyes to avoid perceiving the the monsters that will still make mincemeat of your attempts to change things.

I don’t entirely disagree with this. But I think we’ve [been over this territory before](#).

The world is a scary place, full of bad people who want to hurt you, and in the state of nature you’re pretty much [obligated](#) to engage in whatever it takes to survive.

But instead of sticking with the state of nature, we have the ability to form communities built on mutual disarmament and mutual cooperation. Despite artificially limiting themselves, these communities become stronger than the less-scrupulous people outside them, because they can work together effectively and because they can boast a better quality of life that attracts their would-be enemies to join them. At least in the short term, these communities can resist races to the bottom and prevent the use of personally effective but negative-sum strategies.

One such community is the kind where members try to stick to rational discussion as much as possible. These communities are definitely better able to work together, because they have a [powerful method](#) of resolving empirical disputes. They’re definitely better quality of life, because you don’t have to deal with constant [insult wars and personal attacks](#). And the existence of such communities provides positive externalities to the outside world, since they are better able to resolve difficult issues and find truth.

But forming a rationalist community isn’t just about having the *will* to discuss things well. It’s also about having the *ability*. Overcoming bias is really hard, and so the members of such a community need to be constantly trying to advance the art and figure out how to improve their discussion tactics.

As such, it’s acceptable to try to determine and discuss negative patterns of argument, even if those patterns of argument are useful and necessary weapons in a state of nature. If anything, understanding them makes them *easier* to use if you’ve got to use them, and makes them easier to recognize and counter from others, giving a slight advantage in battle if that’s the kind of thing you like. But moving them from unconscious to conscious also gives you the crucial *choice* of when to deploy them and allows people to try to root out ethnic tension in particular communities.

The Moral Of The Story

[content warning: puns. This is mostly self-plagiarism from my Tumblr and Twitter]

Once upon a time there was a small desert village with a single well outside town. One day a young woman went to the well to fetch water, and the well heard her crying, and asked "What's wrong?"

She stopped her sobbing and asked the well "You can talk?"

"Yes," said the well. "Long ago, the witch who lives in this town gave me life so I could serve as a guardian to the townspeople."

"Alas," said the young woman. "I am the daughter of that witch. She lived in peace with the townsfolk for many years. But the new mayor, who is a violent and hateful man, riled the people up against her, and they burned her at the stake. I am young and still do not know very much magic. I tried to curse them, but my curses fizzled. Now I worry I will never avenge my mother's death."

"Do not be afraid," said the well. "I will take care of this."

The next morning, when the Mayor came to fetch water from the well, he heard an odd noise coming from the bottom. He peered over as far as he could to see what was happening. Then an impossibly long arm shot up from the bottom of the well, grabbed the mayor, and pulled him into the well shaft. There was a horrible crunching sound, and nobody ever saw the Mayor again. The townsfolk apologized to the witch's daughter, and they all lived happily ever after.

Moral of the story: Living well is the best revenge.

Pixar's movie *Up* won the Academy Award for "Best Picture" and was widely hailed as one of the best children's films of the decade. In fact, some people argued it was *too* good, and that kids were ignoring school, chores, and other responsibilities to watch it again and again. They said that along with the cute plot, the short, catchy name gave it an almost drug-like addictive quality. This made a lot of people very angry, and Pixar agreed to give its addictive must-watch movies longer names in the future.

Moral of the story: Do not call *Up* what you cannot put down.

There's a new report out of CERN that a team of scientists has unraveled the structure of the photon. Apparently this started years ago when some equations showed that photons acted like tiny "hands" - structures with a "palm" and radiating "fingers" - which "crawl" across time/space and "grab" the solid particles they interact with. This explained most of the properties of light but wasn't an exact match for the data. The latest result is that single photons are actually made up of hundreds of these shapes, all joined together into a single particle, and this is how they're able to travel so quickly.

Moral of the story: Many hands make light work.

Once upon a time there was an ugly duckling. All of the other ducklings had grown their beautiful white soft downy feathers, but this duckling had no down feathers at all and was bald and ugly and all the other ducklings teased him.

He went to the mysterious crow who lived in the woods and asked for help. The crow said to repeat the magic words “HOCUS POCUS” at midnight with a full moon, and then he would grow his down feathers. The duckling tried that, but the moon just laughed at him and said the magic had no power here.

So he went to the creepy raven who lived in the swamp and asked for help. The raven said to repeat the magic words “ABRA CADABRA” at high noon on a sunny day, and then he would grow his down feathers. The duckling tried that, but the sun just laughed at him and said he wasn’t bound by the magic.

So he went to the wise old owl who lived in the tallest tree and asked for help. The owl explained that the duckling should just ignore the mockery of the other birds and accept that he was okay just the way he was, because there were no magic spells to make ducklings grow feathers.

Moral of the story: You are beautiful, no matter what they say. Words can’t bring you down.

Once upon a time a young lady died and went to Hell. At the check-in desk, Satan asked her age. She was in her twenties, but looked much younger; she thought quick and realized that even in Hell, they probably wouldn’t be mean to children. So she told Satan that she was twelve, and sure enough he said she wasn’t old enough to be held accountable for her sins, and ushered her off to a more peaceful part of Hell reserved for ages eleven through thirteen. She met the other sinners there and realized that many of them, like her, were older people who had lied to get out of their punishment.

Satan began to suspect something like this was going on, so he set up hidden cameras in the 11-13 wing of Hell, trying to catch people acting like adults or admitting to one another that they had lied about their age. But there were hundreds of millions of sinners and Satan couldn’t monitor all the cameras himself. So he went up to the mortal world and asked for the best supercomputer they had. The mortals recommended a newer model of Deep Blue, the supercomputer that had first beaten a human world champion at chess. Satan picked one up from IBM and went back to Hell, where he programmed the Deep Blue to monitor all of the hidden camera feeds at once and report any suspicious activity.

Sure enough, after a few days, he got thousands of reports of people acting older than thirteen. He hunted them down and removed them to Hell proper, where there was much wailing and gnashing of teeth. And it all could have been avoided if they had just stuck to their charade and acted as young as they said they were.

Moral of the story: Don’t get caught, be tween – the Devil and the Deep Blue see.

By 2050, screens have shrunk and become more flexible until the dream of “programmable paper” becomes a reality. Citizens of the future read newspapers like the ones in Harry Potter that include moving images and even videos of important events. This new technology even makes it as far as the US Post Office, where they decide to include programmable stamps. Instead of a static picture of eg George

Washington's head, it will have a moving image of Washington speaking and giving his famous Farewell Address.

Unfortunately, the technology isn't ready for the kind of abuse that envelopes undergo on their travel throughout the country and the world. Most of the computerized stamps become corrupted and "crash"; in a particularly common bug, they try to reload but just end up displaying "GENERATING IMAGE..." permanently. The government has no money to fix the problem, so people just get used to stamps on their letters that say "GENERATING IMAGE..." instead of having interesting pictures on them.

Moral of the story: If you want a vision of the future, imagine a human face booting on a stamp forever.