

Immoral Mazes

1. [Quotes from Moral Mazes](#)
2. [Moral Mazes and Short Termism](#)
3. [Meditations On Moloch](#)
4. [Meditation Retreat: Immoral Mazes Sequence Introduction](#)
5. [Moloch Hasn't Won](#)
6. [Perfect Competition](#)
7. [Imperfect Competition](#)
8. [Does Big Business Hate Your Family?](#)
9. [What is Life in an Immoral Maze?](#)
10. [Stripping Away the Protections](#)
11. [What is Success in an Immoral Maze?](#)
12. [How to Identify an Immoral Maze](#)
13. [How to Escape From Immoral Mazes](#)
14. [The Road to Mazedom](#)
15. [How Doomed are Large Organizations?](#)
16. [Ten Causes of Mazedom](#)
17. [Potential Ways to Fight Mazes](#)
18. [Create a Full Alternative Stack](#)
19. [Protecting Large Projects Against Mazedom](#)
20. [Mazes Sequence Roundup: Final Thoughts and Paths Forward](#)

Quotes from Moral Mazes

Reading and actually paying attention to [Moral Mazes](#) is hard. Writing carefully about it is even harder. I effectively spent several months forcing my way through the book, because it seemed important to do that. I then spent a month trying to write about the book, but that's going super slow as well. The repetition, the saying the same thing from multiple angles, the detailed examples, seem necessary to get the points across, because one has the very strong instinct to avoid understanding it, to read without seeing, [to hear without listening](#). At least, I know I did, despite these things also not only not feeling new, but resonating with my direct experiences.

So in the interest of getting something out there, and hoping that I'll be able to address things in more detail later, here are all the 168 (!) quotes I highlighted from the book, roughly organized into categories. Locations listed are how to find the quote in the Kindle edition, and the quotes are numbered for ease of search and reference.

A. Hierarchy and Credit

1. As a former vice-president of a large firm says: "What is right in the corporation is not what is right in a man's home or in his church. What is right in the corporation is what the guy above you wants from you. That's what morality is in the corporation." (Location 148)
2. When managers describe their work to an outsider, they almost always first say: "I work for [Bill James]" or "I report to [Harry Mills]" or "I'm in [Joe Bell's] group,"* and only then proceed to describe their actual work functions. (Location 387)
3. The key interlocking mechanism of this structure is its reporting system. Each manager gathers up the profit targets or other objectives of his or her subordinates and, with these, formulates his commitments to his boss; this boss takes these commitments and those of his other subordinates, and in turn makes a commitment to his boss.* (Location 441)
4. It is characteristic of this authority system that details are pushed down and credit is pulled up. (Location 446)
5. One gives credit, therefore, not necessarily where it is due, although one always invokes this old saw, but where prudence dictates. Customarily, people who had nothing to do with the success of a project can be allocated credit for their exemplary efforts. At the middle levels, therefore, credit for a particular idea or success is always a type of refracted social honor; one cannot claim credit even if it is earned. Credit has to be given, and acceptance of the gift implicitly involves a reaffirmation and strengthening of fealty. A superior may share some credit with subordinates in order to deepen fealty relationships and induce greater efforts on his behalf. Of course, a different system obtains in the allocation of blame. (Location 482)
6. If one has a mistake-prone boss, there is, of course, always the temptation to let him make a fool of himself, but the wise subordinate knows that this carries two dangers—he himself may get done in by his boss's errors, and, perhaps more important, other managers will view with the gravest suspicion a subordinate who withholds crucial information from his boss even if they think the boss is a nincompoop. A subordinate must also not circumvent his boss nor ever give the

appearance of doing so. He must never contradict his boss's judgment in public. To violate the last admonition is thought to constitute a kind of death wish in business, and one who does so should practice what one executive calls "flexibility drills," an exercise "where you put your head between your legs and kiss your ass good-bye." (Location 424)

7. The general rule is that bosses are expected to protect those in their bailiwicks. Not to do so, or to be unable to do so, is taken as a sign of untrustworthiness or weakness. If, however, subordinates make mistakes that are thought to be dumb, or especially if they violate fealty obligations—for example, going around their boss—then abandonment of them to the vagaries of organizational forces is quite acceptable. (Location 438)

8. Managers often note that one must stay at least three drinks behind one's boss at social functions; this meant that Brown's subordinates might never drink at all on such occasions. (Location 630)

9. However, the belief of insiders in abstract goals is not a prerequisite for personal success; belief in and subordination to individuals who articulate organizational goals is. One must, however, to be successful in a bureaucratic work situation, be able to act, at a moment's notice, as if official reality is the only reality. (Location 1194)

10. You can put the damper on anyone who works for you very easily and that's why there's too much chemistry in the corporation. There's not enough objective information about people. When you really want to do somebody in, you just say, well, he can't get along with people. That's a big one. And we do that constantly. What that means, by the way, is that he pissed me off; he gave evidence of his frustration with some situation. Another big one is that he can't manage—he doesn't delegate or he doesn't make his subordinates keep his commitments. So in this sort of way, a consensus does build up about a person and a guy can be dead and not even know it. (Location 1475)

11. Only at this point did Brady realize that it was the CEO himself who was fiddling with the numbers. The entire dual reporting system that the CEO had personally initiated was in part an elaborate spy network to guard against discovery of the slush fund manipulation, and perhaps other finagling, rather than a system to ensure financial honesty. (Location 2404) [Leads into next quote]

12. [Note: Brady is an accountant.] The corporate managers to whom I presented this case see Brady's dilemma as devoid of moral or ethical content. In their view, the issues that Brady raises are, first of all, simply practical matters. His basic failing was, first, that he violated the fundamental rules of bureaucratic life. These are usually stated briefly as a series of admonitions. (1) You never go around your boss. (2) You tell your boss what he wants to hear, even when your boss claims that he wants dissenting views. (3) If your boss wants something dropped, you drop it. (4) You are sensitive to your boss's wishes so that you anticipate what he wants; you don't force him, in other words, to act as boss. (5) Your job is not to report something that your boss does not want reported, but rather to cover it up. You do what your job requires, and you keep your mouth shut. Second, the managers that I interviewed feel that Brady had plenty of available legitimations to excuse or justify his not acting. Clearly, they feel, a great many other executives knew about the pension fund scam and did nothing; everybody, especially the top bosses, was playing the game. The problem fell into other people's areas, was their responsibility, and therefore their problem. Why, then, worry about it? Besides, Brady had a number of ways out of the situation if he

found it intolerable, including resigning. Moreover, whatever action he took would be insignificant anyway so why bother to act at all and jeopardize himself? Even a fool should have known that the CEO was not likely to take whatever blame resulted from the whole affair. (Location 2429)

13. One's best efforts at being fair, equitable, and generous with subordinates clash both with a logic that demands choices between people, inevitably producing hatred, envy and animosity, and with the plain fact that, despite protestations to the contrary, many people do not want to be treated fairly. (Location 4552)

14. Mastering the subtle but necessary arts of deference without seeming to be deferential, of "brown nosing" without fawning, of simultaneous self-promotion and self-effacement, and occasionally of the outright self-abasement that such relationships require is a taxing endeavor that demands continual compromises with conventional and popular notions of integrity. Only those with an inexhaustible capacity for self-rationalization, fueled by boundless ambition, can escape the discomfort such compromises produce. (Location 4572)

B. Feeling Comfortable

15. Essentially, managers try to gauge whether they feel "comfortable" with proposed resolutions to specific problems, a task that always involves an assessment of others' organizational morality and a reckoning of the practical organizational and market exigencies at hand. The notion of comfort has many meanings. When applied to other persons, the idea of comfort is an intuitive measure of trustworthiness, reliability, and predictability in a polycentric world that managers often find troubling, ambiguous, and anxiety-laden. Such assessment of others' organizational morality is a crucial aspect of a more general set of probations that are intrinsic to managerial work. (Location 302)

16. They objected in particular to those aspects of my brief written proposal that discussed the ethical dilemmas of managerial work. They urged me to avoid any mention of ethics or values altogether and concentrate instead on the "decision-making process" where I could talk about "trade-offs" and focus on the "hard decisions between competing interests" that mark managerial work. Taking these cues, I rewrote and rewrote the proposal couching my problem in the bland, euphemistic language that I was rapidly learning is the lingua franca of the corporate world. But such recasting eroded whatever was distinctive about the project and some managers dismissed the study as a reinvention of the wheel. (Location 324)

17. In effect, I could not get access to study managers' moral rules-in-use because I seemed unable to articulate the appropriate stance that would convince key managers that I already understood those rules and was thus a person with whom they could "feel comfortable" enough to trust. (Location 334)

18. The process centered on the written proposal that I had been circulating and consisted essentially of a furthering of my linguistic education in the art of indirect rather than pointed statement and, more particularly, a reformulation of my inquiry that recast the moral issues of managerial work as issues of public relations. When, after several rewritings, the proposal satisfied him, he approached a well-placed executive in a large textile firm that I have given the pseudonym of Weft Corporation and vouched for me. At that point, the proposal itself became meaningless since, to my knowledge, no one except the two executives who arranged access ever saw it. The personal vouching, however, was crucial. This was based on what both men took

to be a demonstrated willingness and ability to be “flexible” and especially on their perception that I already grasped the most salient aspect of managerial morality as managers themselves see it—that is, how their values and ethics appear in the public eye. (Location 347)

19. At bottom, all of the social contexts of the managerial world seek to discover if one “can feel comfortable” with another manager, if he is someone who “can be trusted,” if he is “our kind of guy,” or, in short, if he is “one of the gang.” (Location 905)

20. My search for access involved me in some of the crucial bureaucratic intricacies that shape managers’ experiences. These include organizational upheavals, political rivalries, linguistic ambiguity, the supremacy of chance and tangled personal connections over any notion of intrinsic merit, the central significance of public relations, and, perhaps especially, the ceaseless moral probations for inclusion in a managerial circle. Managers keep their eyes on the organizational premiums that shape behavior, values, ethics, and worldviews in corporate bureaucracies. I focus on those premiums... (Location 387)

21. One becomes known, for instance, as a trusted friend of a friend; thought of as a person to whom one can safely refer a thorny problem; considered a “sensible” or “reasonable” or, especially, a “flexible” person, not a “renegade” or a “loose cannon rolling around the lawn”; known to be a discreet person attuned to the nuances of corporate etiquette, one who can keep one’s mouth shut or who can look away and pretend to notice nothing; or considered a person with sharp ideas that break deadlocks but who does not object to the ideas being appropriated by superiors. (Location 870)

22. Similarly, Covenant’s CEO sold large tracts of land with valuable minerals at dumbfoundingly low prices. The CEO and his aides said that Covenant simply did not have the experience to mine these minerals efficiently, a self-evident fact from the low profit rate of the business. In all likelihood, according to a manager close to the situation, the CEO, a man with a financial bent and a ready eye for the quick paper deal, felt so uncomfortable with the exigencies of mining these minerals that he ignored the fact that the prices the corporation was getting for the minerals had been negotiated forty years earlier. Such impulsiveness and indeed, one might say from a certain perspective, irrationality, is, of course, always justified in rational and reasonable terms. It is so commonplace in the corporate world that many managers expect whatever ordered processes they do erect to be subverted or overturned by executive fiat, masquerading as an established bureaucratic procedure or considered judgment. (Location 1700)

23. think that I’ve got to where I am today because of this. [His boss’s boss] knows that I saved the company a lot of money and a lot of asses to boot. And he and others know that I am someone who can be trusted. I can keep my mouth shut.... And that’s the biggest thing that I have going for me—that people feel that I can be trusted. I can’t overemphasize that enough. (Location 2917)

24. Only those men and women who allow peers and superiors to feel morally comfortable in the ambiguous muddles of the world of affairs have a chance to survive and flourish in big organizations when power and authority shift due to changes in markets, internal power struggles, or the need to respond to external exigencies. (Location 4943)

C. Struggle for Success

25. The logical result of alertness to expediency is the elimination of any ethical lines at all. (Location 2985)

26. The two areas are, of course, related since one's chances in an organization depend largely on one's "credibility," that is, on the widespread belief that one can act effectively. One must therefore prevail regularly, though not always, in small things to have any hope of positioning oneself for big issues. The hidden agenda of seemingly petty disputes may be a struggle over long-term organizational fates. (Location 812)

27. A fundamental rule of corporate politics is that one never cedes control over assets, even if the assets are administrative headaches. (Location 643)

28. Bureaucratic hierarchies, simply by offering ascertainable rewards for certain behavior, fuel the ambition of those men and women ready to subject themselves to the discipline of external exigencies and of their organization's institutional logic, the socially constructed, shared understanding of how their world works. However, since rewards are always scarce, bureaucracies necessarily pit people against each other and inevitably thwart the ambitions of some. (Location 806)

29. When asked who gets ahead, an executive vice-president at Weft Corporation says: The guys who want it [get ahead]. The guys who work. You can spot it in the first six months. They work hard, they come to work earlier, they leave later. They have suggestions at meetings. They come into a business and the business picks right up. They don't go on coffee breaks down here [in the basement]. You see the parade of people going back and forth down here? There's no reason for that. I never did that. If you need coffee, you can have it at your desk. Some people put in time and some people work. (Location 992)

30. Proper management of one's external appearances simply signals to one's peers and to one's superiors that one is prepared to undertake other kinds of self-adaptation. Managers also stress the need to exercise iron self-control and to have the ability to mask all emotion and intention behind bland, smiling, and agreeable public faces. (Location 1059)

31. The price of bureaucratic power is a relentlessly methodical subjection of one's impulses, at least in public. (Location 1100)

32. [Style] is being able to talk easily and make presentations. To become credible easily and quickly. You can advance quickly even without technical experience if you have style. You get a lot of points for style. You've got to be able to articulate problems, plans, and strategies without seeming to have to refer to all sorts of memos and so on. The key in public performances and presentations is in knowing how to talk forcefully without referring to notes and memoranda. To be able to map out plans quickly and surely. (Location 1298)

33. As one manager says: "Personality spells success or failure, not what you do on the field." (Location 1383)

34. More generally, there are several rules that apply here. First, no one in a line position—that is, with responsibility for profit and loss—who regularly "misses his numbers" will survive, let alone rise. Second, a person who always hits his numbers

but who lacks some or all of the required social skills will not rise. Third, a person who sometimes misses his numbers but who has all the desirable social traits will rise. (Location 1412)

35. A managerial commonplace says: "In the corporate world, 1,000 'Attaboys' are wiped away with one 'Oh, shit!'" (Location 1619)

36. There's a lot of it [fear and anxiety]. To a large degree it's because people are more honest with themselves than you might believe. People know their own shortcomings. They know when they're over their heads. A lot of people are sitting in jobs that they know are bigger than they should be in. But they can't admit that in public and, at still another level, to themselves. The organizational push for advancement produces many people who get in over their heads and don't know what they are doing. And they are very fearful of making a mistake and this leads to all sorts of personal disloyalty. But people know their capabilities and know that they are on thin ice. And they know that if they make mistakes, it will cost them dearly. So there's no honesty in our daily interaction and there's doubt about our abilities. The two go together. (Location 1767)

37. One comes to gauge that hard-won access to managerial circles takes precedence over fussing with abstract principles. (Location 2923)

38. Perceptions of pervasive mediocrity breed an endless quest for social distinctions even of a minor sort that might give one an "edge," enable one to "step out of the crowd," or at least serve as a basis for individual claims to privilege. More specifically, an atmosphere of mediocrity erodes the hope of meaningful collective achievement and encourages, at least among more aggressive managers, a predatory stance toward their organizations, that is, a search for private deals, a working of the system for one's own personal advantage. (Location 4436)

39. One leaves behind as well the technical knowledge or scientific expertise of one's younger years, lore now more suited for the narrower roles of technicians or junior managers. One must, in fact, put distance between oneself and technical details of every sort or risk the inevitable entrapment of the particular. Salesmen, too, must leave their bags and regular customers and long boisterous evenings that seal measurable deals behind them and turn to marketing strategies. Work becomes more ambiguous, directed as it is toward maneuvering money, symbols, organizational structures, and especially people. The CEO at Weft Corporation, it is said, "doesn't know a loom from a car." And the higher one goes, the more managers find that "the essence of managerial work is cronyism, covering your ass, [and] pyramiding to protect your buddies." (Location 4539)

40. the rewards of corporate success can be very great. And those who do succeed, those who find their way out of the crowded, twisting corridors and into the back rooms where the real action is, where the big games take place, and where everyone present is a player, shape, in a decisive way, the moral rules-in-use that filter down through their organizations. The ethos that they fashion turns principles into guidelines, ethics into etiquette, values into tastes, personal responsibility into an adroitness at public relations, and notions of truth into credibility. (Location 4603)

41. Instead, success becomes contingent on others' interpretations of one's performance, leading to a break in the accepted moral economy between talent combined with effort and reward. This makes compulsive sociability an occupational

virtue, as one attempts to discern and shape peers' and superiors' interpretations. (Location 4950)

D. Nothing Matters, But Hit Your Numbers

42. Managers rarely speak of objective criteria for achieving success because once certain crucial points in one's career are passed, success and failure seem to have little to do with one's accomplishments. (Location 917)

43. Corporations rely on other institutions—principally the schools—to establish what might be called competence hurdles. The demonstrated ability of a student to leap over successively higher hurdles in school is taken as evidence of the ability to weather well the probationary trials of corporate life. (Location 920)

44. In Alchemy Inc., whether in sales, marketing, manufacturing, or finance, the "breaking point" in the hierarchy is generally thought to be grade 13 out of 25 or the top 8.5 percent of management. By the time managers reach such a numbered grade in an ordered hierarchy—and the grade is socially defined and varies from company to company—managerial competence as such is taken for granted and assumed not to differ greatly from one manager to the next. (Location 943)

45. A product manager in the chemical company talks about the lack of connection between work and results: I guess the most anxiety provoking thing about working in business is that you are judged on results whether those results are your fault or not. So you can get a guy who has tried really hard but disaster strikes; and you can get a guy who does nothing and his business makes a big success. And so you just never know which way things are going to go and you're never sure about the relationship of your work to the outcome. One of the top executives in Weft Corporation echoes this sentiment: I always say that there is no such thing as a marketing genius; there are only great markets. (Location 1585)

46. Assuming a basic level of corporate resources and managerial know-how, real economic outcome is seen to depend on factors largely beyond organizational or personal control. (Location 1592)

47. upper-middle level manager says: If I were just out of school and somebody told me that it doesn't matter what you do and how well you do it but that what matters is being in the right place at the right time, I'd have said that hard work is still the key. You know, the old virtues. But now as I have gotten older, I think it's pure happenstance—luck. Things happen to people and being in the right time and place and knowing the right people is the key. (Location 1621)

48. It is interesting to note in this context that a line manager's long-run credibility suffers just as much from missing his numbers on the up side (that is, achieving profits higher than predicted) as from missing them on the down side although, as one might expect, the immediate consequences of such different miscalculations vary. Both outcomes, however, undercut the ideology of managerial planning and control. (Location 1595)

49. A top staff official at Covenant Corporation explains: By putting the money in this business, you're taking the money away from others. In human terms, that's what you're doing. It's money that you could provide jobs with to others. So when you get a guy in the business who comes in under or over the plan, well, both are equally suspect. Because you're making major decisions based on your plan.... Like when we

shut down [business A] and put the money into [business B], the whole legitimacy of the operation depends on the [business A] guys accepting the rationale that more money can be made in another operation. (Location 1599)

50. of a plant manager who, when his machinery had ground to a halt and his technicians were baffled and everyone turned to him to make a decision, told his crew, without the faintest idea of the right thing to do and with the great fear that all he had worked for was about to crumble before him, to dump ten pounds of phosphate into the machine. The machine sprang to life and he became a hero. (Location 1656)

E. Implicitness

51. If I tell someone what to do—like do A, B, or C—the inference and implication is that he will succeed in accomplishing the objective. Now, if he doesn't succeed, that means that I have invested part of myself in his work and I lose any right I have to chew his ass out if he doesn't succeed. If I tell you what to do, I can't bawl you out if things don't work. And this is why a lot of bosses don't give explicit directions. They just give a statement of objectives, and then they can criticize subordinates who fail to make their goals. (Location 454)

52. A typical example occurred in Weft Corporation a few years ago when the CEO, new at the time, expressed mild concern about the rising operating costs of the company's fleet of rented cars. The following day, a stringent system for monitoring mileage replaced the previous casual practice. (Location 490)

53. One must remember, for instance, that in our litigious age the best rule in dealing with angry subordinates is to say nothing or as little as possible since whatever one says may be used against oneself and one's organization. (Location 1066)

54. Well, usually you don't tell people the truth. I once knew a guy whom I knew was about to be fired and I asked if he had been told and he had never been told. I think you should tell people explicitly. Things like that shouldn't have to be decoded. But you can understand how it happens. Suppose you have a guy and the consensus is that he isn't promotable. You wouldn't ever—or very seldom—tell him. He goes on to justify his silence: There are people who go through life thinking they can do a lot more than they really can do. And the reason is that losing or changing jobs is a very high stress situation and most people prefer to hang on to what they've got—to their routine. They're not happy but they go through life like prisoners of war not recognizing their true situation. (Location 1494)

55. You get the situation where a lot of people don't really want to know.... Like one guy we have, he will retire on his job. He's in my division. He knows it. I know it. And he doesn't want me to tell him about that. Now don't ask me how I know that but, believe me, I do. (Location 1502)

56. Why does it happen? Because people are afraid of confrontations. People want to be thought of as kind, sensitive, and compassionate. Being compassionate has a good significance in our society. The easy way out is not to do anything, don't tell the guy. That happens a lot. (Location 1510)

57. As a matter of fact, he wouldn't even have to say "cut capital." He would just put pressure on him by saying: "Well, sales are down 50 percent; why aren't your

expenses down 50 percent?" My boss will come to me, by the time it reaches him, and say: "Cut costs." It's as simple as that. (Location 2078)

58. He put things into writing in a world that, apart from ritual nods to the importance of documentation, actually fosters ambiguity by its reliance on talk as the basic mode of negotiation and command. Talk, of course, lends itself more readily than documents to backtracking, filling in, evasion, subterfuge, and secrecy, all important virtues if one is to do what has to be done while establishing and maintaining the kinds of relationships that alone can protect oneself. (Location 2636)

59. It is unlikely, however, that any workers affected could ever piece things together. First, there is nothing in writing. Second, Tucker feels sure that everyone involved would, if it became necessary, simply deny knowledge and claim that the process was altered solely for production reasons. 4. Finally, he says: The basic rule is that you hope that these kinds of things never occur. Nobody wants to hurt people. Nobody would ever consciously plan to do something that would endanger people. But when things happen, well, you cover for yourself and your company. (Location 2948)

60. Discreet suggestions, hints, and coded messages take the place of command; this, of course, places a premium on subordinates' abilities to read correctly their bosses' vaguely articulated or completely unstated wishes. One cannot even criticize one's subordinates to one's own superior without risking a negative evaluation of one's own managerial judgment. (Location 3009)

61 (Chart). Phrase / Probable Intended Meaning

*Exceptionally well qualified / Has committed no major blunders to date

Tactful in dealing with superiors / Knows when to keep his mouth shut

Quick thinking / Offers plausible excuses

Meticulous attention to detail / A nitpicker

Slightly below average / Stupid

Unusually loyal / Wanted by no one else

Indifferent to instruction / Knows more than one's superior

Strong adherence to principles / Stubborn

Requires work-value attitudinal readjustment / Lazy and hardheaded (Location 3018)

62. For the most part, euphemistic language is not used with the intent to deceive. Managers past a certain point, as suggested earlier, are assumed to be "maze-bright" and able to "read between the lines" of a conversation or a memorandum and to distinguish accurately suggestions from directives, inquiries from investigations, and bluffs from threats. Managers who are "maze-dense," like the manager at Weft Corporation who, though told somewhat indirectly that he was fired, did not realize his fate until the following day, might consider the oblique, elliptical quality of managerial language to skirt deceit. However, most often when managers use euphemistic language with each other (and it is important to remember that in private among trusted others their language can be very direct, colorful, and indeed earthy), its principal purpose is to communicate certain meanings within specific contexts with

the implicit understanding that should the context change, a new, more appropriate meaning can be attached to the language already used. In this sense, the corporation is a place where people are not held to what they say because it is generally understood that their word is always provisional. (Location 3034)

63. The rule of thumb here seems to be that the more troublesome a problem, the more desiccated and vague the public language describing it should be. Of course, when a troublesome problem bursts into public controversy, euphemism becomes a crucial tool of those managers who have to face the public in some forum. The task here is to defuse public criticism and sometimes outrage with abstract unemotional characterizations of issues. (Location 3044)

F. Short Term Thinking

64. The ideal situation, of course, is to end up in a position where one can fire one's successors for one's own previous mistakes. (Location 2102)

65. Moreover, the only real threat to corporations on environmental issues was in the courts, which, however, judge past actions, not present practices. By the time the courts get to cases generated by contemporary practices, typically in fifteen years, those executives presently in charge will have moved on, leaving any problems their policies might create to others. (Location 707)

66. These choices also reflect judgments about whether a corporation can offer the hard-charging MBAs from the top-ranked schools enough quick variety to retain them long enough to justify their inflated salaries. (Location 926)

67. Similarly, accounting systems that place a premium on bare-bones inventory reflect the same pressure for short-run profit maximization. For instance, at Covenant Corporation the story is told about a plant that produced a useful by-product at no extra cost. One simply had to store it until it was needed for other internal operations. Covenant, however, works with an accounting system that considers by-products as inventory; moreover, inventory counts against one at the end of a fiscal year. In order to cut costs, managers decided to throw out the by-product at the end of a financial cycle. But a sudden shortage of the material trebled its cost two months later. To service their own operations, managers had to go hat in hand to their competitors to buy the material at the premium prices. (Location 1837)

68. This sets the stage for financial sharpshooters who, in takeover strategies, buy large chunks of a company's stock at devalued prices only to be "greenmailed" (persuaded with financial inducements) by the target company's management into surrendering these blocks of holdings at premium prices. In such unsettled times, where virtually any large corporation could become a takeover target, managers feel that they have to keep their companies' stock properly valued. As it happens, the markets honor only short-term gains. (Location 1854)

69. Instead, one focuses attention on important problems of the moment that must be solved. Since these are always plentiful, they justify postponing less pressing concerns. Of course, managers know at one level of their consciousness that today's minor issues can quickly become tomorrow's major crises, but the pressure for annual, quarterly, monthly, daily, and even hourly "results," that is, measurable progress plausibly attributed to one's own efforts, crowds out reflection about the future. An upper-middle manager at Alchemy Inc. recalls, for instance, his days as a plant

manager when his boss at company headquarters phoned him every three hours to see how many tons of soda ash had been produced in the interval. (Location 1869)

70. This goes to the heart of the problem. Managers think in the short run because they are evaluated by both their superiors and peers on their short-term results. Those who are not seen to be producing requisite short-run gains come to be thought of as embarrassing liabilities. Of course, past work gets downgraded in such a process. The old saw, still heard frequently today, "I know what you did for me yesterday, but what have you done for me lately?" is more than a tired garment district salesman's joke. It accurately reflects the widespread amnesia among managers about others' past accomplishments, however notable, and points to the probationary crucibles at the core of managerial life. Managers feel that if they do not survive the short run, the long run hardly matters, and one can only buy time for the future by attending to short-term goals. As one manager says: "Our horizon is today's lunch." (Location 1875)

71. Now you see this at work with mistakes. You can make mistakes in the work you do and not suffer any consequences. For instance, I could negotiate a contract that might have a phrase that would trigger considerable harm to the company in the event of the occurrence of some set of circumstances. The chances are that no one would ever know. But if something did happen and the company got into trouble, and I had moved on from that job to another, it would never be traced to me. The problem would be that of the guy who presently has responsibility. And it would be his headache. There's no tracking system in the corporation. Some managers argue that outrunning mistakes is the real meaning of "being on the fast track," the real key to managerial success. The same lawyer continues: In fact, one way of looking at success patterns in the corporation is that the people who are in high positions have never been in one place long enough for their problems to catch up with them. They outrun their mistakes. That's why to be successful in a business organization, you have to move quickly. (Location 2013)

72. Both Covenant Corporation and Weft Corporation, for instance, place a great premium on a division's or a subsidiary's return on assets (ROA); managers who can successfully squeeze assets are first in line, for instance, for the handsome rewards allotted through bonus programs. One good way for business managers to increase their ROA is to reduce assets while maintaining sales. Usually, managers will do everything they can to hold down expenditures in order to decrease the asset base at the end of a quarter or especially at the end of the fiscal year. The most common way of doing this is by deferring capital expenditures, everything from maintenance to innovative investments, as long as possible. Done over a short period, this is called "starving a plant"; done over a longer period, it is called "milking a plant." (Location 2034)

73. A plant that is not well maintained will fail in the short term, so you have to spend money there; a plant that has poorly trained people will fail today, so you have to spend money there. But you can still milk (Location 2045)

74. We're judged on the short-term because everybody changes their jobs so frequently. (Location 2042)

75. My favorite things are not to replace my stores inventory and that shows up as direct profit on your balance sheet; not replace people who retire, and stretch everybody else out; cut down on overtime; cut working inventories to the bone. [You can also] lower the quality standards; you can get away with this in the short term

because people will accept that for awhile, though in the long term people will stop buying from you. (Location 2050)

76. At the very top of organizations, one does not so much continue to outrun mistakes as tough them out with sheer brazenness. In such ways, bureaucracies may be thought of, in C. Wright Mills's phrase, as vast systems of organized irresponsibility. (Location 2123)

77. When a weaver is unable to repair a stop, she shuts down the loom and flags a "fixer," the most skilled and highest-ranking worker on the shop floor, who does the repair or basic maintenance necessary to get the loom working again. However, since weavers are paid by the piece and can make no gain from a loom out of service for repair or maintenance during their own shifts, weavers will tend to remedy stops in any way they can in order to keep their cloth production high, leaving the maintenance and repair of the machinery and the economic cost involved to another weaver on another shift. Supervisors and managers who are evaluated by a plant's overall weaving efficiency are thus forced to monitor the number of loom stops constantly in order to make sure that looms badly in need of repair or maintenance get proper attention. Whenever structural inducements place premiums on immediate personal gains, especially when mistakes are not penalized, there seems to be a sharp decline in the likelihood of men and women sacrificing their own interests for others, for their organizations, or least of all for the common weal.² (Location 2135)

78. It was said that Noll had milked and milked thoroughly every plant he ever supervised. One day, a story goes, he was accused of this in a public meeting by a vice-president who was then his superior. Noll is said to have responded with great boldness: "[Joe], how can you sit there and say that to me? How in the hell do you think you got to where you are and how do you think you stay there?" (Location 2153)

79. If a guy keeps moving, he can say, "Look, I ran this plant better than my predecessors." And people have to concede that. A lot of people do that. Then you get the guy who takes his place and tries to run things right and he has to spend a lot of money. And people look at the guy who was there before and they say: "Well, old [Noll] ran the plant well and he didn't have to spend any money like you're claiming you do." I don't think there is anything wrong with milking a plant. As long as you know you're milking (Location 2223)

80. No, definitely not. Would any sane, rational man spend \$15 million for a 2 percent return? ... Now it does improve the dust levels, but it was that if we don't invest the money now, we would be in a desperate [competitive] position fifteen years from now. Our demonstrated cash flow situation was such that eventually we would have had even tougher decisions to make. (Location 3579)

81. Executives also admit, somewhat ruefully and only when their office doors are closed, that OSHA's regulation on cotton dust has been the main factor in forcing the pace of technological innovation in a centuries-old, hidebound, and somewhat stagnant industry. It has also been a major factor in forcing executives to think in the long run rather than continually succumbing to short-term pressures. This is one of the reasons why the shrewdest among them only feigned elation at the attempts by Reagan's OSHA appointees to remove the cotton dust regulation from the purview of the Supreme Court. If such a move were successful, it could only encourage the traditionally reactionary elements of the textile industry who refuse to recognize on principle that government regulation, within reason, can be the businessman's best friend. (Location 3593)

82. S&Ls could now open transaction accounts and make commercial real estate loans, regardless of geography, up to 40 percent of their assets. This prompted many S&L executives to sell their long-term, fixed-rate, low- or no-profit mortgages to Wall Street firms for as little as 60 percent of their value in order to obtain ready funds for much more lucrative, though much riskier, investments. (Location 4649)

83. Organized irresponsibility has migrated to the nooks and crannies of our society. One sees presidents of universities and colleges who preside over the destruction of their institutions' endowments, not to mention the intellectual integrity of curricula, and then outrun their mistakes and move on to higher academic and even top political posts. One sees intellectuals who are completely committed to the destruction of the Western tradition that nurtures them and who help create the profound social and cultural centrifugality that marks American society. One sees big-city mayors who cede control of the streets to drug dealers and other thugs for fear of being labeled nonprogressive by various advocacy groups and the press. One sees long-time political insiders in Washington who "screw up and move up," an aphorism that aptly captures the ethos of government bureaucracies. And doublethink and doublespeak now permeate public discourse, bewildering even intelligent, thoughtful citizens. (Location 4990)

G. Social Politics

84. Managers' cryptic aphorism, "Well, you never know...", repeated often and regularly, captures the sense of uncertainty created by the constant potential for social reversal. Managers know too, and take for granted, that the personnel changes brought about by upheavals are to a great extent arbitrary and depend more than anything else on one's social relationships with key individuals and with groups of managers. (Location 761)

85. "Every big organization is set up for the benefit of those who control it; the boss gets what he wants." (Location 827)

86. In any event, just as managers must continually please their boss, their boss's boss, their patrons, their president, and their CEO, so must they prove themselves again and again to each other. Work becomes an endless round of what might be called probationary crucibles. Together with the uncertainty and sense of contingency that mark managerial work, this constant state of probation produces a profound anxiety in managers, perhaps the key experience of managerial work. (Location 910)

87. See, once you are at a certain level of experience, the difference between a vice-president, an executive vice-president, and a general manager is negligible. It has relatively little to do with ability as such. People are all good at that level. They wouldn't be there without that ability. So it has little to do with ability or with business experience and so on. All have similar levels of ability, drive, competence, and so on. What happens is that people perceive in others what they like—operating styles, lifestyles, personalities, ability to get along. Now these are all very subjective judgments. And what happens is that if a person in authority sees someone else's guy as less competent than his own guy, well, he'll always perceive him that way. And he'll always pick—as a result—his own guy when the chance to do so comes up. (Location 1013) [Also see Nothing Matters]

88. Striking, distinctive characteristics of any sort, in fact, are dangerous in the corporate world. One of the most damaging things, for instance, that can be said about a manager is that he is brilliant. This almost invariably signals a judgment that

the person has publicly asserted his intelligence and is perceived as a threat to others. What good is a wizard who makes his colleagues and his customers uncomfortable? (Location 1173)

89. Equally damaging is the judgment that a person cannot get along with others—he is “too pushy,” that is, he exhibits too much “persistence in getting to the right answers,” is “always asking why,” and does not know “when to back off.” Or he is “too abrasive,” or “too opinionated,” unable “to bend with the group.” Or he is a “wildman” or a “maverick,” that is, someone who is “outspoken.” Or he may be too aloof, too distant, “too professional.” (Location 1176)

90. The knowledgeable practitioners of corporate politics, whether patrons or leaders of cliques and networks, value nothing more highly than at least the appearance of unanimity of opinion among their clients and allies, especially during times of turmoil. (Location 1197)

91. Managers know that patrons and powerful allies protect those already selected as rising stars from the negative judgments of others; and only the foolhardy point out even egregious errors of those in power or those destined for it. (Location 1463)

92. Managers know that to be weak in a world that extols strength and power is to invite abuse. (Location 1536)

93. Such social distancing has two purposes: it undermines in advance or lays the groundwork for refusal of any claims that a person considered a failure might make on another, and it forestalls the possibility of being linked with that person in others’ cognitive maps. This becomes particularly important when there has been a known past association between oneself and one thought to have failed in some way. (Location 1542)

94. In fact, when it is even suspected that a person might be headed for trouble, anticipatory avoidance is the rule. (Location 1550)

95. Being a “fall guy,” that is, “taking the rap” or “taking the heat” for others’ decisions or mistakes is probably the most common kind of blame in big organizations. (Location 1904)

96. You have to remember that you only get to explain things away once. When things get screwed up, you get one chance. That’s why it’s important for everybody to be in bed with everybody else. And if they don’t like you from the start, you don’t have a chance. Because when things go wrong, what people do is sit down and say—without saying it in so many words—look, our jobs are on the line. Let’s make sure that it’s not us who gets nailed. (Location 1999)

97. Wilson and his Site Operations staff were particularly concerned because when the polar crane was finally used, it would be Site Operations who used it. They knew that, if they were in charge when the button was pushed, they could be blamed for whatever might go wrong. If the polar crane failed, it would be seen as the fault of Site Operations. (Location 2557)

98. high-ranking staff official at Covenant explains: Anxiety is endemic to anyone who works in a corporation. By the time you get to be middle management, it’s difficult to make friends because the normal requirement for friendship—that is, loyalty—doesn’t fit in this context. You have to look out for number one more than anything else. Moreover, the prevailing view is that managers are big boys and girls, well-paid, and

should be able to take care of themselves. Besides, one person's failure represents another person's opportunity. (Location 1554)

99. of a young manager who, while at a company conference, went out for his weight-controlling 5:30 A.M. jog only to meet a vice-president similarly engaged, a powerful executive who now cheers the younger man's work and presentations and introduces him to other influential senior managers; (Location 1654)

100. scattering ducks already set in a row. Besides, one can only criticize something when one has the resources to solve it in a clear and decisive way. Otherwise, one should keep one's skepticism to oneself and get "on board." (Location 2629)

101. Other managers and managerial cliques are always on the lookout for others' mistakes or for actions that can be construed as mistakes and will pounce on anyone foolish enough to admit them. Even if others restrain an immediate attack, the knowledge of someone's mistakes is ammunition for the future. Many managers "lay in the weeds, with rocks, and wait." One who exposes a colleague's errors in such a context and makes him vulnerable to others evinces, of course, only a fundamental untrustworthiness, unless one's colleague has first betrayed oneself or others in some way— (Location 2858)

102. There is a premium in the higher circles of management on seeming fresh, dynamic, innovative, and up-to-date. In their social minglings and shoptalk with one another, particularly with their opposite numbers in other large companies, say, at the Business Roundtable, at high-level conferences at prestigious business schools, at summer galas in the Hamptons, or at the Super Bowl, the biggest business extravaganza of all, executives need to seem abreast of the latest trends in managerial know-how. No one wants to appear stodgy before one's peers nor to have one's firm defined in managerial networks, and perhaps thence to Wall Street, as "slow on the uptake." Executives trade ideas and schemes and judge the efficacy of consultant programs not by any detached critical standards but by what is socially acceptable, desirable, and, perhaps most important, current in their circles. (Location 3155)

103. So they attend the sessions and with a seemingly dutiful eagerness learn literally to repeat the requisite formulas under the watchful eyes of senior managers. Senior managers do not themselves necessarily believe in such programs. In one seminar that I attended, the senior manager in charge startled a room of juniors by saying: Fellows, why aren't any of you asking about the total lack of correspondence between what we're preaching here and the way we run our company? But such outspokenness is rare. (Location 3209)

104. They admit, however, that the marvelously high fees that consultants command (in 1986 as high as \$2,000 a day in New York City) enhance their legitimacy and encourage managers to lend credence to their schemes. (Location 3216)

105. A choice between securing one's own success by jumping on and off the bandwagon of the moment, or sacrificing oneself for the long-run good of a corporation by diverting resources and really seeing a program through is, for most managers, no choice at all. Ambitious managers see self-sacrificing loyalty to a company as foolhardy. Moreover, middle and upper-middle level managers upon whom requests for self-sacrifice for the good of the organization are most likely to fall do not see top executives sacrificing themselves for the common good. For example, just after the CEO of Covenant Corporation announced one of his many purges,

legitimated by a “comprehensive assessment of the hard choices facing us” by a major consulting firm, he purchased a new Sabre jet for executives and a new 31-foot company limousine for his own use at \$1,000 a foot. He then flew the entire board of directors to Europe on the Concorde for a regular meeting to review, it was said, his most recent cost-cutting strategies. As other managers see it, bureaucratic hierarchy gives top bosses the license to act in their own interests and to pursue with impunity the arts of contradiction. (Location 3220)

H. The I in Team

106. The main dimensions of team play are as follows. 1. One must appear to be interchangeable with other managers near one’s level. Corporations discourage narrow specialization more strongly as one goes higher. They also discourage the expression of moral or political qualms. One might object, for example, to working with chemicals used in nuclear power, or working on weapons systems, and most corporations today would honor such objections. Publicly stating them, however, would end any realistic aspirations for higher posts because one’s usefulness to the organization depends on versatility. As one manager in Alchemy Inc. comments: “Well, we’d go along with his request but we’d always wonder about the guy. And in the back of our minds, we’d be thinking that he’ll soon object to working in the soda ash division because he doesn’t like glass.” Strong convictions of any sort are suspect. One manager says: If you meet a guy who hates red-haired persons, well, you’re going to wonder about whether that person has other weird perceptions as well. You’ve got to have a degree of interchangeability in business. To me, a person can have any beliefs they want, as long as they leave them at home. Similarly, one’s spouse’s public viewpoints or activities could reduce others’ perceptions of a manager’s versatility or indeed ability. In reference to another manager whose wife was known to be active in environmental action groups, lobbying in fact for legislation on chemical waste disposal, one Alchemy manager says: “If a guy can’t even manage his own wife, how can he be expected to manage other people?” (Location 1137)

107. Team play also means, as one manager in the chemical company puts it, “aligning oneself with the dominant ideology of the moment” or, as another says, “bowing to whichever god currently holds sway.” Such ideologies or gods may be thought of as official definitions of reality. (Location 1188)

108. You can indict a person by saying that he’s not a team player. That doesn’t mean he won’t follow directions. It’s because he voices an objection, because he argues with you before doing something, especially if he’s right. That’s when we really get mad—when the other guy is right. If he’s wrong, we can be condescending and adopt the “you poor stupid bastard” tone.... (Location 1220)

109. A team player is a manager who does not “force his boss to go to the whip,” but, rather, amiably chooses the direction his boss points out. Managers who choose otherwise or who evince stubbornness are said to “have made a decision,” a phrase almost always used to describe a choice that will shorten a career. (Location 1229)

110. Team players display a happy, upbeat, can-do approach to their work and to the organization. (Location 1252)

I. Avoiding Decision Making

111. You know that old saying: "Success has many parents; failure is an orphan"? Well, that describes decision making. A lot of people don't want to make a commitment, at least publicly. This is a widespread problem. They can't make judgments. They stand around and wait for everybody else's reaction. Let me tell you a story which perfectly illustrates this. There was a [museum] collection coming, the [Arctic] collection, and there was a great deal of interest among designers in [Arctic] things. My own feeling was that it wouldn't sell but I also recognized that everybody wanted to do it. But in this case, [our] design department was spared the trouble. There was an independent designer who had access to our president and he showed him a collection of [Arctic] designs. There were two things wrong: (1) it was too early because the collection hadn't hit town yet; (2) more important, the designs themselves were horrible. Anyway, [the collection] was shown in a room with everything spread out on a large table. I was called down to this room which was crowded with about nine people from the company who had seen the designs. I looked at this display and instantly hated them. I was asked what I thought but before I could open my mouth, people were jumping up and down clapping the designer on the back and so on. They had already decided to do it because the president had loved it. Of course, the whole affair was a total failure. The point is that in making decisions, people look up and look around. They rely on others, not because of inexperience, but because of fear of failure. They look up and look to others before they take any plunges. (Location 1713)

112. But all these things have no relationship to the way they actually manage or make decisions. The basic principles of decision making in this organization and probably any organization are: (1) avoid making any decision if at all possible; (2) if a decision has to be made, involve as many people as you can so that, if things go south, you're able to point in as many directions as possible. (Location 1731)

113. A middle-aged, upper-middle level manager at Alchemy Inc. says: You know, there is this huge computerized inventory of skills which people update each year; it's called a skills inventory. ... But all the computerized lists in the world don't amount to much in the corporation. What matters is a bunch of guys sitting informally in a room and deciding who should get jobs and who shouldn't. The real job decisions are made on that basis. And circumstances determine your fate. (Location 1634)

114. Making a decision, or standing by a decision once made, exposes carefully nurtured images of competence and know-how to the judgments of others, particularly of one's superiors. As a result, many managers become extremely adept at sidestepping decisions altogether and shrugging off responsibility, all the while projecting an air of command, authority, and decisiveness, leaving those who actually do decide to carry the ball alone in the open field. (Location 1774)

115. This explains why the chemical company managers kept putting off a decision about major reinvestment. After the battery collapsed in 1979, however, the decision facing them was simple and posed little risk. The corporation had to meet its legal obligations; also, it had either to repair the battery the way the EPA demanded or shut down the plant and lose several hundred million dollars. Since there were no real choices, everyone could agree on a course of action because everyone could appeal to inevitability. This is the nub of managerial decision making. As one manager says: Decisions are made only when they are inevitable. To make a decision ahead of the time it has to be made risks political catastrophe. People can always interpret the decision as an unwise one even if it seems to be correct on other grounds. (Location 1886)

116. Despite their fresh appearances, certain themes recur constantly in the programs offered by consultants. Perhaps the most common are how to sharpen decision making, how to restructure organizations for greater efficiency, how to improve productivity, how to recognize trouble spots in an organization, how to communicate effectively, how to humanize the workplace, and how to raise morale. (Location 3164)

J. This is Your Life

117. At the managerial and professional levels, the road between work and life is usually open because it is difficult to refuse to use one's influence, patronage, or power on behalf of another regular member of one's social coterie. It therefore becomes important to choose one's social colleagues with some care and, of course, know how to drop them should they fall out of organizational favor. (Location 884)

118. For some managers, the drive for success is a quest for the generous financial rewards that high corporate position brings. For others, success means the freedom to define one's work role with some latitude, to "get out from under the thumb of others." For still others, it means the chance to gain power and to exert one's will, to "call the shots," to "do it my way," or to know the curiously exhilarating pleasure of controlling other people's fates. For still others, the quest for success expresses a deep hunger for the recognition and accolades of one's peers. (Location 955)

119. More generally, those who accept immobility are unwilling to sacrifice family life or free-time activities to put in the extraordinarily long hours at the office required in the upper circles of their corporations. Or they have made a realistic assessment of the age structure, career paths, and power relationships above them and conclude that there is no longer real opportunity for them. They may see that there is an irreparable mismatch between their own personal styles and the kinds of social skills being cultivated in well-entrenched higher circles. In many cases, they decide that they do not wish to put up with the great stress of higher management work that they have witnessed. (Location 962)

120. Higher-level managers in all the corporations I studied commonly spend twelve to fourteen hours a day at the office. (Location 1156)

121. In a world where appearances—in the broadest sense—mean everything, the wise and ambitious manager learns to cultivate assiduously the proper, prescribed modes of appearing. He dispassionately takes stock of himself, treating himself as an object, as a commodity. He analyzes his strengths and weaknesses and decides what he needs to change in order to survive and flourish in his organization. And then he systematically undertakes a program to reconstruct his image, his publicly avowed attitudes or ideas, or whatever else in his self-presentation that might need adjustment. (Location 1339)

122. This [lack of control] doesn't mean you don't work hard; at least in my case, that's my answer. I have to believe I can influence events. That way, I feel good about myself even if my boss doesn't. (Location 1606)

123. Sometimes I'll wake up in the middle of the night thinking about the plant. And if I can't get back to sleep, I'll slip out of bed and walk over to the plant and just walk around the machinery and talk to the guys. I love the smell of the oil and the grease and the sound of the machines. For me, that's what life is all about. (Location 4535)

K. Immoral Mazes

124. In analyzing incidence data from a representative plant, and by extrapolating to the rest of the firm's mills, White discovered that 12 percent of all greige mill workers had already suffered hearing loss severe enough to be immediately compensable under state law for as much as \$3.5 to \$5.7 million. This, however, was only the thunder before a summer storm. Another 63 percent of greige mill workers had already suffered substantial, though not yet compensable, damage that could only worsen the longer they stayed in the industry. (Location 2265)

125. Moreover, by insisting on his own personal moral purity, his feeling that if he did not expose things he himself would be drawn into a web of corruption, he was, they feel, being disingenuous; no one reaches his level of a hierarchy without being tainted. (Location 2450)

126. A moral judgment based on a professional ethic makes little sense in a world where the etiquette of authority relationships and the necessity for protecting and covering for one's boss, one's network, and oneself supercede all other considerations and where nonaccountability for action is the norm. (Location 2474)

127. On the polar crane issue, the top official of the NRC later characterized Wilson's and his engineers' concerns as stemming from a philosophy that emphasized procedural matters rather than a focus on final goals. This characterization was echoed by GPUN management who stressed that what was at issue in the polar crane dispute was not procedures but results; at a certain point, they said, decisions had to be made to resolve technical disputes and work had to proceed toward what everyone acknowledged to be a worthwhile goal, that is, the cleanup of TMI-2. (Location 2566)

128. Once, when he wrote a memo to the GPUN deputy director about radioactively contaminated sewage being trucked out of the plant and disposed of illegally, his boss replied that he did not need such a memo from Wilson. It was, his boss said, not constructive and wasted his own and Wilson's time. (Location 2575)

129. my view, this is the nub of the moral ethos of bureaucracy. Managers see this issue as a "trade-off" between principle and expediency. They usually pose the trade-off as a question: Where do you draw the line? (Location 2652)

130. It is said, in fact, that one Weft manager who was on the take for years from customers was fired within a half hour of the discovery of his "inexplicably stupid" acceptance and deposit of a check from his benefactors instead of his normal cash rake-offs. With such authoritative encouragement, managers internalize these norms into their own personal codes of honor; they speak privately of the importance of not being known as men or women "who can be had." (Location 2658)

131. It gets hard. Now, suppose that the ozone depletion theory were correct and you knew that these specific fifty people were going to get skin cancer because you produced chlorofluorocarbons [CFCs]. Well, there would be no question. You would just stop production. But suppose that you didn't know the fifty people and it wasn't at all clear that CFCs were at fault, or entirely at fault. What do you do then? (Location 2831)

132. Suppose that you had a candy bar factory and you were touring the plant and you saw with your own eyes a worker slip a razor blade into a bar. And before you could stop the machine, there were a thousand bars more made and the one with the

razor blade was mixed up. Well, there's no question that you would get rid of the thousand candy bars. But what if it were a million bars? Well, I don't know what I'd do. (Location 2843)

133. Moreover, he tries to treat his subordinates forthrightly, firmly believing that one's word is an important measure of a person. In a world, however, where actions are separated from consequences, where knowledge is fragmented and secreted, where private agreements are the only real way to fashion trust in the midst of ongoing competition and conflict, where relationships with trusted colleagues constitute one's only real means both of defense and opportunity, and where, one knows, even coincidental association with a disaster can haunt one's career years later, keeping silent and covering for oneself and for one's fellows become not only possible but prudent, indeed virtuous, courses of action. (Location 2969)

134. Moreover, it is a byword in the chemical industry at least that it is precisely in those technological areas where accidents have seldom occurred that the largest potential catastrophes loom; the very lack of practice in responding quickly to untoward incidents can precipitate uncontrollable events. (Location 3376)

135. Publicly, of course, Weft Corporation, as do many other firms, claims that the money was spent entirely to eliminate dust, evidence of its corporate good citizenship. Privately, executives admit that without the productive return, they would not have—indeed, given the constraints under which they operate—could not have spent the money. (Location 3585)

136. The ethic that emerges in bureaucratic contexts contrasts sharply in many respects with the original Protestant ethic. The Protestant ethic was a social construction of reality of a self-confident and independent propertied social class. It was an ideology that extolled the virtues of accumulating and reinvesting wealth in a society organized around property and that accepted the stewardship responsibilities entailed by property. It was an ideology where a person's word was his bond and where the integrity of the handshake was crucial to the maintenance of good business relationships. (Location 4295)

137. At the core of the middle class's righteous, some would say smug, faith in itself, of its inexhaustible drive, of its unremitting pragmatism, was the conviction that hard work necessarily had its just rewards here and now as a token of divine favor in the hereafter. This conviction was also the bedrock of a profound guilt mechanism that impelled one to fulfill personal and social obligations; failure to do so, like a failure to work hard, was thought to be a sin against both God and self. (Location 4304)

138. Bureaucracy breaks apart the ownership of property from its control, social independence from occupation, substance from appearances, action from responsibility, obligation from guilt, language from meaning, and notions of truth from reality. Most important, and at the bottom of all of these fractures, it breaks apart the older connection between the meaning of work and salvation. In the bureaucratic world, one's success, one's sign of election, no longer depends on an inscrutable God, but on the capriciousness of one's superiors and the market; and one achieves economic salvation to the extent that one pleases and submits to new gods, that is, one's bosses and the exigencies of an impersonal market. (Location 4307)

L. Truth and Public Relations

139. Truth? What is truth? I don't know anyone in this business who talks about "truth." (Location 4137)

140. After all, as one executive says, "We lie all the time, but if everyone knows that we're lying, is a lie really a lie?" (Location 2693)

141. If a retailer returns with goods and says, "Look, I can't sell \$50,000 of this stuff. You have to take it back or it's going to break me," may one say, "What's that? I didn't hear you," hoping for the sake of future business that the retailer will say, "Oh, look, you were late delivering and I can't use it now." (Location 2700)

142. The consultant who perceives such discrepancies has to devise his own strategies for handling them. Some of these include: rejecting the assignment altogether; accepting the problem as defined and confining oneself to it for the sake of future contracts even though one knows that any action will be inefficacious; or accepting the assignment but trying to persuade the client to address the underlying social and political issues, that is, redefining the problem. (Location 3234)

143. For instance, a highly placed staff member whose work requires him to interact daily with the top figures of his company, says: I get faked out all the time, and I'm part of the system. I come from a very different culture. Where I come from, if you give someone your word, no one ever questions it. It's the old hard-work-will-lead-to-success ideology. Small community, Protestant, agrarian, small business, merchant-type values. I'm disadvantaged in a system like this. He goes on to characterize the system more fully and what it takes to succeed within it: It's the ability to play this system that determines whether you will rise. ... And part of the adeptness [required] is determined by how much it bothers people. One thing you have to be able to do is to play the game, but you can't be disturbed by the game. What's the game? It's bringing troops home from Vietnam and declaring peace with honor. It's saying one thing and meaning another. It's characterizing the reality of a situation with any description that is necessary to make that situation more palatable to some group that matters. It means that you have to come up with a culturally accepted verbalization to explain why you are not doing what you are doing.... [Or] you say that we had to do what we did because it was inevitable; or because the guys at the [regulatory] agencies were dumb; [you] say we won when we really lost; [you] say we saved money when we squandered it; [you] say something's safe when it's potentially or actually dangerous.... Everyone knows that it's bullshit, but it's accepted. This is the game. (Location 3246)

144. Such adeptness at inconsistency, without moral uneasiness, is essential for executive success. Done over a period of time, in fact, it seems to become a taken-for-granted habit of mind. As one executive at Covenant Corporation says: Now some people don't understand this.... But as you move up the ladder, you don't have people who don't understand. And the people up high don't necessarily do it consciously. They are able to speak out of both sides of their mouth without missing a step. It means being able to say, as a very high-ranking official of Weft Corporation said to me without batting an eye, that the industry has never caused the slightest problem in any worker's breathing capacity. It means, in Covenant Corporation, propagating an elaborate hazard/benefit calculus for approval of dangerous chemicals while internally conceptualizing "hazards" as business risks. It means publicly extolling the carefulness of testing procedures on toxic chemicals while privately ridiculing animal tests as inapplicable to humans. (Location 3610)

145. Finally, those truly adept at inconsistency can also interpret with some accuracy the inconsistent machinations of their colleagues and adversaries. This is not a mean skill. At the very beginning of my fieldwork, the top lawyer of a large corporation was discussing an issue that I had raised when he said: Now, I'm going to be completely honest with you about this. He paused for a moment and then said: By the way, in the corporate world, whenever anybody says to you: "I'm going to be completely honest with you about this," you should immediately know that a curveball is on the way. But, of course, that doesn't apply to what I'm about to tell you. (Location 3631)

146. Since the success of large commercial bureaucracies depends to a great extent on the goodwill of the consuming public, ambitious managers recognize that great organizational premiums are placed on the ability to explain expedient action convincingly. Public opinion, of course, constitutes one of the only effective checks on the bureaucratic impulse to translate all moral issues into practical concerns. (Location 3643)

147. 1925, Walter Lippmann critically analyzed a public's ability to appreciate the intricacies of any issue or indeed to keep its attention on anything but crises. He adds: [S]ince [a public] acts by aligning itself, it personalizes whatever it considers, and is interested only when events have been melodramatized as a conflict. The public will arrive in the middle of the third act and will leave before the last curtain, having stayed just long enough perhaps to decide who is the hero and who is the villain of the piece. (Location 3765)

148. read in The Wall Street Journal or The New York Times accounts of specific events or larger trends attributed to "informed sources," the code name for public relations men and women; (Location 3828)

149. There are, of course, good clients and bad clients. A good client keeps his public relations specialists adequately informed, provides "feedback" and "constructive criticism," and recognizes that public relations depends on time-billing rather than on product billing and provides enough funds to do the job properly. A good client "takes stock of himself," that is, dispassionately objectifies his company's products, his company's organizational structure and personnel, and, say, in the case of a top corporate official, himself, to make them all more readily manipulable and therefore marketable commodities. Above all, a good client is flexible and therefore able rapidly to shift ground and actual policies as well to meet new needs or new pressures. A bad client, by contrast, either does not understand or chooses to ignore the peculiarly indirect approach of public relations and wants immediately observable results in terms of press clippings or TV time; uses a public relations program as a vehicle for self-aggrandizement within his own corporation, placing the public relations specialist in danger of "getting caught in a pissing contest between executives"; demands the release of press statements that are only marginally "newsworthy" and then blames the public relations specialist when nothing at all gets printed or aired; conceals crucial aspects of his story from his public relations advisers and refuses them adequate access to his staff and facilities to get the full story; comes to the public relations agency with trumped-up data and fake photographs—as happened, for instance, at Images Inc. when a client falsely claimed the efficacy of a product to remove tar from despoiled beaches—and ends up declaring bankruptcy, leaving the public relations agency liable for a multimillion-dollar lawsuit; wants the public relations agency, as in the case of some single-party foreign governments, to put a good public face on practices like the "persuasion" and "elimination" of opponents; expects the public relations agency to be the "bag man" to pay off government officials or newspaper editors; and, perhaps especially, insists on an indefensible or

totally unbelievable version of reality or expects the public relations specialist to tell outright lies. (Location 3843)

150. As it happens, accounts in the public relations world circulate continually and only sometimes because of actual or alleged dissatisfaction with public relations service. The pattern of continual mergers, upheavals, and power struggles in corporations, that I have argued earlier is at the core of American corporate life, directly affects public relations agencies. When a new CEO or divisional president assumes power in a corporation, he will quite often change public relations and advertising agencies as part of a larger strategy of shedding the past or to assert his own “new vision” of the future. When this occurs, almost invariably, he will move his account to an agency whose leaders he knows and with whom he feels comfortable. This continual circulation of accounts means that agency personnel are constantly searching for new accounts and constantly devising ways to hold present clients, despite the knowledge of the inevitability of their eventual departure. (Location 3868)

151. In the world of public relations, there is no such thing as a notion of truth; there are only stories, perspectives, or opinions. One works, of course, with “facts,” that is, selected empirically verifiable statements about the world. But as long as a story is factual, it does not matter if it is “true.” One can feel free to arrange these facts in a variety of ways and to put any interpretation on them that suits a client’s objectives. Interpretations and judgments are always completely relative. The only canons binding this process of interpretation are those of credibility or, more exactly, of plausibility. If an interpretation of facts, a story, is taken as plausible by a targeted audience, it is just as good as “true” in any philosophical sense, indeed better since it furthers the accomplishment of an immediate goal. (Location 3886)

152. From the standpoint of public relations, the journalistic ideology closely resembles the social outlook of most college seniors—a vague but pious middle-class liberalism, a mildly critical stance toward their fathers in particular and authorities in general; a maudlin championship of the poor and the underclass; and especially the doctrine of tolerance, open-mindedness, and balance. In fact, public relations people feel, the news media are also constructing reality. They are always looking for a “fresh” and exciting angle; they have an unerring instinct for the sentimental that expresses itself in a preference for “human interest” rather than substance; and they arrange facts in a way that purports to convey “truth,” but is in fact simply another story. In reality, news is entertainment. And, despite the public’s acceptance of journalistic ideologies, most of the public watch or read news not to be informed or to learn the “truth,” but precisely to be entertained. There is no intrinsic reason, therefore, why the constructions of reality by public relations specialists should be thought of as any different from those of any group in the business of telling stories to the public. (Location 3903)

153. The top official, in a written document prepared for a meeting to discuss the issue, argued that: [T]he final report was better than it would have been if it had not gone through [the] process of reexamination. It was stronger, it was more important, it was more constructive.... At the same time, there is no doubt that pressures were brought on us to come to the kind of conclusions we finally reached. At the meeting, he added, somewhat more pithily: “We tried to be honest but, believe me, it wasn’t easy.” (Location 4124)

154. The same top executive defines truth: I sometimes sit back and think that if we could make up a list of all the viewpoints of all our clients and somehow fit them

together, then that would be truth. That would be what we are as a firm. (Location 4130)

155. Agency practitioners not only defend multiple interests but face repeatedly the experience of even well-served clients switching agencies. Especially unsettling are the departures of clients who decide that, since they are as they have been portrayed, they have no future need to construct reality. The continual circulation of client accounts in agencies diminishes the possibilities of comforting, long-held allegiances to organizations, products, or causes. (Location 4192)

156. Alternatively, and by contrast, practitioners in both settings sometimes justify their efforts by appealing to a professional ethos that celebrates the exercise of technical skill separated from any emotional commitment to one's clients. A dignified version of this legitimation is the often repeated analogy between public relations practitioners and lawyers; both occupations, it is argued, fulfill important advocacy roles in a free society. Only the practice of the professional virtue of public relations, however hazardous to individual practitioners, can assure the continued diversity of opinion that marks our democracy. (Location 4195)

157. That's the reality of our society. There's no question that their story is being told because they have the money and power. I've got to recognize that I'm part of this society and just come to live with that. Our society is the way it is. It's run on money and power, it's that simple. Truth has nothing to do with it. So we just accept the world as it is and live with it. (Location 4206)

158. The agency literally creates these realities by matchmaking artistic talent and accomplishment with money. In the process, up-and-coming junior corporate executives get the kind of exposure to refined, sophisticated artistic and intellectual circles that will help prepare and polish them for higher posts. Artists, in turn, as long as their work is not too avant-garde, receive the benefits of a latter-day Medici-like patronage. Public relations people claim a double accomplishment—they help civilize businessmen, not least by inspiring in them a "passion for greatness" rather than self-interest as the important motive for patronage of the arts; and they provide the public with access to high culture. The same firm has also arranged corporate sponsorship of dance and musical performances, helped develop important educational programs such as one on infant nutrition, conducted some useful surveys such as one on the problems facing ethnic minorities, and done a lot of pro bono work for philanthropic, community, and public service organizations in the bargain. One tries, then, to move some clients in directions that seem socially desirable while at the same time playing with the magic lantern to serve their interests. Sometimes, too, the ability to accomplish any good at all in this world seems to depend on the willingness to serve even clients with no apparent redeeming features in order to seize capricious opportunities to channel other clients' resources into work deemed socially worthwhile. (Location 4214)

159. In short, bureaucracy creates for managers a Calvinist world without a Calvinist God, a world marked with the same profound anxiety that characterized the old Protestant ethic but one stripped of that ideology's comforting illusions. Bureaucracy poses for managers an intricate set of moral mazes that are paradigmatic of the quandaries of public life in our social order. Within this framework, the puzzle for many individual managers becomes: How does one act in such a world and maintain a sense of personal integrity? (Location 4351)

M. Under Siege

160. Part of the folklore of the modern corporation, in fact, consists of a catalog of stories about how big corporations are being victimized by the courts in what amounts to a radical transformation of tort law. Managers repeat variations of cases—for instance, of a man who purchased a thrice-owned punch machine, altered it, and after losing a finger, sued the original manufacturer for negligent construction; of a farmer who, despite clearly marked warning labels, coated his animal feed troughs with a toxic wood preservative, sold milk from contaminated cattle, and then when he himself was sued by irate parents whose children had drunk the milk, sued the chemical company who produced the preservative; or of a student research assistant inexplicably asked to carry dangerous acid in glass containers on a faulty elevator that lurched and caused the young woman to break a container, showering her with acid and permanently disfiguring her. Although she is said to have sued her university, her supervising professor, the elevator manufacturer, and the chemical company that provided the acid, managers had few doubts that the jury would “pin the tail on business” because of the “abiding conviction that corporations have vaults filled with gold bars called profits.” At some level, even managers who repeat such stories and thus reinforce their own shared sense of beleaguerment, recognize the profound and profoundly felt dependence of most people in our society on those large organizations that claim to control the scientific genie. (Location 3413)

161. was at this party the other night and I was sitting next to this older lady and she said: “My God, did you people see the paper tonight? There’s leakage from some chemical plant and it’s infecting the drinking water around here.” So I asked if I could see the paper and the article said that there was some seepage out of a pond that a chemical plant used for disposal but that the EPA was monitoring the situation as well as all 25 wells in the area. In the meantime, the company was remedying the situation. I pointed this out and she looked at me, her eyes narrowed, and she asked, “Who do you work for?” I said I worked for a large chemical firm and she burst out laughing and asked if I expected her to believe me. She laughed right in my face! I asked her what she was worried about. You have to drink the water for 25 years before anything would happen—I mean, she was already a grandmother—but that didn’t seem to help her much. (Location 3459)

162. My interviews are filled with stories of managers who claim to have been verbally assaulted not only by strangers at cocktail parties, but by their children’s teachers when they visit schools, and even by their children themselves at the breakfast table for being supposedly callous and insensitive to the social consequences of business activity. Perhaps more than anything else, managers are puzzled by such attacks, though they pounce quickly on any inconsistencies that they perceive in their opponents. For example, one manager whose firm produced a pesticide that became caught up in a widely publicized episode of mishandling and illegal disposal was attacked by his brother-in-law, a lifelong military man, for his very association with the company. The manager found it grimly ironic that “things have reached a point where a trained killer is berating me for producing something useful.” More generally, managers view the irrational fear of contamination evinced by those who espouse an ideology of bodily purity with some derision. One can, after all, referring to the same pesticide, “eat handfuls of the stuff with no lasting adverse effects.” (Location 3467)

163. Well, from 1957 through 1962, I was intimately involved with the manufacture of DDT. During that time, we doubled production and sold almost all of it to Africa and India. And I knew and went home knowing that I was saving more lives than any major hospital was capable of doing. I knew that I was saving thousands of lives by doing this. Then Rachel Carson’s *Silent Spring* came out and not only did I become a murderer of falcons and robins, but also one of the mass murderers of the world. I was

now doing evil things to the world. Then I went to a plant which was manufacturing [chlorofluorocarbons] and we increased production by 20–25 percent; a lot of it went into hair sprays. We also used vinyl chloride and found out that it was causing liver cancer. Then I found out that I was destroying the whole ozone layer of the earth and doing it for personal gain. Then I went into soda ash. Without it, there wouldn't be a window pane in the whole United States. But at the time, because of the Clean Air Act and the Clean Water Act, suddenly I became a polluter. Children learned in school that chemicals killed. And, of course, there was no question in the academic community that I was perceived as an evil person doing evil things. And that became true even in the corporation. Plants became a liability, rather than the source of wealth. The perception was: Wouldn't it be nice if we could just sell chemicals without producing them? So the profession of producing things became a low profession and the good people were those who were producing services. Manufacturing people became evil. I think this is one reason that marketing people became ascendant in the competition for advancement. Then I went into making sulfuric acid and that involved the whole issue of water pollution. He goes on to describe his reaction to all of this: I guess I'm more bemused than anything else. It's the same type of feeling that I would have if I were an MD who had been doing radical mastectomies and then someone says—hey, you didn't have to do that. It's a feeling of disappointment without a feeling of shame. I know that I have done some useful things. I know that the only source of money is taking something out of the ground and making something. That's where money comes from—making something out of nothing. Without that type of activity, civilization wouldn't exist. So, on something like DDT, I'll leave the judgment to Europe after World War II and all the people saved by the widespread use of it. (Location 3482)

164. Within this context of perceived harassment and shifting scientific and ideological winds, while always attending to the pressing and sometimes contradictory exigencies of business life, managers must address a multiplicity of audiences, some of whom are considered rivals, and some outright adversaries. These audiences are the internal corporate hierarchy with its intricate and shifting power cliques and competing managerial circles, key regulators, local and federal legislators, special publics that vary according to the issues, and the public at large, whose goodwill and favorable opinions are considered essential for a company's free operation. Managerial adeptness at inconsistency becomes evident in the widely discrepant perspectives, reasons for action, and presentations of fact that explain, excuse, or justify managerial behavior to these diverse audiences. (Location 3506)

165. The Supreme Court, however, decided in 1981 that the 1978 OSHA ruling was fully within the Agency's mandate, namely to protect workers' health and safety as the primary benefit exceeding all cost considerations. (Location 3533)

166. The segmented work patterns of bureaucracy underlie these larger structures. Managers' cognitive maps to the thickets of their world contain sharp, sometimes absurd, caricatures of the style and ethos of different occupational groups. These suggest some of the ways in which managers appraise the myriad of character types whom they see peopling their world. Production types, for instance, are said to be hard-drinking, raucous, good-time charlies; engineers, always distinguishable by the plastic pen containers in their shirt pockets, are hostages to an outdated belief in a pristine mathematical rationality; accountants are bean counters who know how to play the shell game; lawyers are legal eagles or legal beagles in wool pinstripes who, if they had their way, would tie managers' hands completely; corporate staff are the king's spies, always ready to do his bidding and his dirty work; marketing guys are cheerful, smooth-talking, upbeat fashion plates who must nonetheless keep salesmen

under their thumbs; salesmen are aggressive loudmouths who feel that they can sell freezers to penguins in Antarctica and who would sell their grandmothers just to make a deal. Salesmen hate the restraints that marketers put on their work and on their ego gratification. Financial wizards, on the ascendancy everywhere, are tight-mouthed, close to the vest poker players who think that a social order can be built on paper deals. And outside consultants are men and women who borrow one's watch and then charge for telling the time. (Location 4327)

167. But so what, they argued, if the preservative did in fact pose some risk of cancer? Better, they said, the risk of a slight long-run increase in the rate of stomach and intestinal cancer than the certainty of a precipitous spurt in the incidence of botulism, particularly in the lower-income black and Hispanic groups that typically consume large amounts of processed meat and, both because of poverty and cultural practices, often leave food uncovered and unrefrigerated for considerable periods. Is corporate social responsibility, they asked, maintaining a private sense and public image of moral purity while someone else does necessary but tainted work? Or is real social responsibility the willingness to get one's hands dirty, to make whatever compromises have to be made to produce a product with some utility, to achieve therefore some social good, even though one knows that one's accomplishments and motives will inevitably be misinterpreted by others for their own ends, usually by those with the least reason to complain? Besides, they pointed out, consumers continue to purchase artificially preserved meats in large quantities. Is not the proper role of business "to give the public what it wants," adopting the market as its polar star, as the only reliable guide in a pluralistic society to "the greatest good for the greatest number," as the final arbiter not of values, which are always arguable, but, more importantly, of tastes, about which there can be no reasonable dispute? (Location 4500)

168. In fact, exercises in substantive rationality—the critical, reflective use of reason—are not only subject to infinite interpretations and counterinterpretations but also invite fantastic constructions of reality, including attributions of conspiracy. Thus a major corporation provides a gift of \$10 million to establish new foundations that will materially aid South African blacks and is promptly accused by a black American leader of bolstering apartheid. Weft managers create an elaborate recreational complex for Weft employees in the corporation's southern community and are charged with perpetuating traditional textile company paternalism. Some executives at Images Inc. donate their time to bring together several institutional sectors of a local town in which they live for community betterment and are charged with trying to grab headlines and line up future business. Managers often feel that, however genuine it may be, altruism is a motive that is always denied them by others. To complicate matters still further, the necessary self-promotional work of presenting private goals as public goods, or the self-defensive work within the corporation of presenting public goods as hardheaded business decisions, or managers' knowledge that bureaucracy insulates them from the real consequences of their actual choices, often make their protestations of socially responsible actions suspect even to themselves. (Location 4512)

Moral Mazes and Short Termism

Previously: [Short Termism](#) and [Quotes from Moral Mazes](#)

Epistemic Status: Long term

My list of [quotes from moral mazes](#) has a section of twenty devoted to short term thinking. It fits with, and gives internal gears and color to, [my previous understanding](#) of the problem of short termism.

Much of what we think of as a Short Term vs. Long Term issue is actually an adversarial [Goodhart's Law](#) problem, or a [legibility](#) vs. illegibility problem, at the object level, that then becomes a short vs. long term issue at higher levels. When a manager milks a plant (see quotes 72, 73, 78 and 79) they are not primarily trading long term assets for short term assets. Rather, they are trading unmeasured assets for measured assets (see 67 and 69).

This is why you can have companies like Amazon, Uber or Tesla get high valuations. They hit legible short-term metrics that represent long-term growth. A start-up gets rewarded for their own sort of legible short-term indicators of progress and success, and of the quality of team and therefore potential for future success. Whereas other companies, that are not based on growth, report huge pressure to hit profit numbers.

The overwhelming object level pressure towards legible short-term success, whatever that means in context, comes from being judged in the short term on one's success, and having that judgment being more important than object-level long term success.

The easiest way for this to be true is not to care about object-level long term success. If you're gone before the long term, and no one traces the long term back to you, why do you care what happens? That is exactly the situation the managers face in Moral Mazes (see 64, 65, 70, 71, 74 and 83, and for a non-manager very clean example see 77). In particular:

74. We're judged on the short-term because everybody changes their jobs so frequently.

And:

64. The ideal situation, of course, is to end up in a position where one can fire one's successors for one's own previous mistakes.

Almost as good as having a designated scapegoat is to have already sold the company or found employment elsewhere, rendering your problems [someone else's problems](#).

The other way to not care is for the short-term evaluation of one's success or failure to impact long-term success. If not hitting a short-term number gets you fired, or prevents your company from getting acceptable terms on financing or gets you bought out, then the long term will get neglected. The net present value payoff for looking good, which can then be reinvested, makes it look like by far the best long term investment around.

Thus we have this problem at every level of management except the top. But for the top to actually be the top, it needs to not be answering to the stock market or capital markets, or otherwise care what others think – even without explicit verdicts, this can be as hard to root out as needing the perception of a bright future to attract and keep quality employees and keep up morale. So we almost always have it at the top as well. Each level is distorting things for the level above, and pushing these distorted priorities down to get to the next move in a giant game of adversarial telephone (see section A of quotes for how hierarchy works).

This results in a corporation that acts in various short-term ways, some of which make sense for it, some of which are the result of internal conflicts.

Why isn't this out-competed? Why don't the corporations that do less of this drive the ones that do more of it out of the market?

On the level of corporations doing this direct from the top, often these actions are a response to the incentives the corporation faces. In those cases, there is no reason to expect such actions to be out-competed.

In other cases, the incentives of the CEO and top management are twisted but the corporation's incentives are not. One would certainly expect those corporations that avoid this to do better. But these mismatches are the natural consequence of putting someone in charge who does not permanently own the company. Thus, dual class share structures becoming popular to restore skin in the correct game. Some of the lower-down issues can be made less bad by removing the ones at the top, but the problem does not go away, and what sources I have inside major tech companies including Google match this model.

There is also the tendency of these dynamics to arise over time. Those who play the power game tend to outperform those who do not play it barring constant vigilance and a willingness to sacrifice. As those players outperform, they cause other power players to outperform more, because they prefer and favor such other players, and favor rules that favor such players. This is especially powerful for anyone below them in the hierarchy. An infected CEO, who can install their own people, can quickly be game over on its own, and outside CEOs are brought in often.

Thus, even if the system causes the corporation to underperform, it still spreads, like a meme that infects the host, causing the host to prioritize spreading the meme, while reducing reproductive fitness. The bigger the organization, the harder it is to remain uninfected. Being able to be temporarily less burdened by such issues is one of the big advantages new entrants have.

One could even say that yes, they *do get wiped out by this*, but it's not that fast, because it takes a while for this to rise to the level of a primary determining factor in outcomes. And [there are bigger things to worry about](#). It's short termism, so that isn't too surprising.

A big pressure that causes these infections is that business is constantly under siege and forced to engage in public relations (see quotes sections L and M) and is constantly facing [Asymmetric Justice](#) and the [Copenhagen Interpretation of Ethics](#). This puts tremendous pressure on corporations to tell different stories to different audiences, to avoid creating records, and otherwise engage in the types of behavior that will be comfortable to the infected and uncomfortable to the uninfected.

Another explanation is that those who are infected don't only reward each other *within* a corporation. They also *do business with* and *cooperate with* the infected elsewhere. Infected people are *comfortable* with others who are infected, and *uncomfortable* with those not infected, because if the time comes to play ball, they might refuse. So those who refuse to play by these rules do better at object-level tasks, but face alliances and hostile action from all sides, including capital markets, competitors and government, all of which are, to varying degrees, infected.

I am likely missing additional mechanisms, either because I don't know about them or forgot to mention them, but I consider what I see here sufficient. I am no longer confused about short termism.

Meditations On Moloch

[Content note: Visions! omens! hallucinations! miracles! ecstasies! dreams! adorations! illuminations! religions!]

I.

Allan Ginsberg's famous poem, *Moloch*:

What sphinx of cement and aluminum bashed open their skulls and ate up their brains and imagination?

Moloch! Solitude! Filth! Ugliness! Ashcans and unobtainable dollars! Children screaming under the stairways! Boys sobbing in armies! Old men weeping in the parks!

Moloch! Moloch! Nightmare of Moloch! Moloch the loveless! Mental Moloch! Moloch the heavy judger of men!

Moloch the incomprehensible prison! Moloch the crossbone soulless jailhouse and Congress of sorrows! Moloch whose buildings are judgment! Moloch the vast stone of war! Moloch the stunned governments!

Moloch whose mind is pure machinery! Moloch whose blood is running money! Moloch whose fingers are ten armies! Moloch whose breast is a cannibal dynamo! Moloch whose ear is a smoking tomb!

Moloch whose eyes are a thousand blind windows! Moloch whose skyscrapers stand in the long streets like endless Jehovahs! Moloch whose factories dream and croak in the fog! Moloch whose smoke-stacks and antennae crown the cities!

Moloch whose love is endless oil and stone! Moloch whose soul is electricity and banks! Moloch whose poverty is the specter of genius! Moloch whose fate is a cloud of sexless hydrogen! Moloch whose name is the Mind!

Moloch in whom I sit lonely! Moloch in whom I dream Angels! Crazy in Moloch! Cocksucker in Moloch! Lacklove and manless in Moloch!

Moloch who entered my soul early! Moloch in whom I am a consciousness without a body! Moloch who frightened me out of my natural ecstasy! Moloch whom I abandon! Wake up in Moloch! Light streaming out of the sky!

Moloch! Moloch! Robot apartments! invisible suburbs! skeleton treasures! blind capitals! demonic industries! spectral nations! invincible madhouses! granite cocks! monstrous bombs!

They broke their backs lifting Moloch to Heaven! Pavements, trees, radios, tons! lifting the city to Heaven which exists and is everywhere about us!

Visions! omens! hallucinations! miracles! ecstasies! gone down the American river!

Dreams! adorations! illuminations! religions! the whole boatload of sensitive bullshit!

Breakthroughs! over the river! flips and crucifixions! gone down the flood! Highs!
Epiphanies! Despairs! Ten years' animal screams and suicides! Minds! New loves!
Mad generation! down on the rocks of Time!

Real holy laughter in the river! They saw it all! the wild eyes! the holy yells! They
bade farewell! They jumped off the roof! to solitude! waving! carrying flowers!
Down to the river! into the street!

What's always impressed me about this poem is its conception of civilization as an individual entity. You can almost see him, with his fingers of armies and his skyscraper-window eyes.

A lot of the commentators say Moloch represents capitalism. This is definitely a piece of it, even a big piece. But it doesn't quite fit. Capitalism, whose fate is a cloud of sexless hydrogen? Capitalism in whom I am a consciousness without a body? Capitalism, therefore granite cocks?

Moloch is introduced as the answer to a question – C. S. Lewis' question in [Hierarchy Of Philosophers](#) – *what does it?* Earth could be fair, and all men glad and wise. Instead we have prisons, smokestacks, asylums. What sphinx of cement and aluminum breaks open their skulls and eats up their imagination?

And Ginsberg answers: *Moloch does it.*

There's [a passage](#) in the *Principia Discordia* where Malaclypse complains to the Goddess about the evils of human society. "Everyone is hurting each other, the planet is rampant with injustices, whole societies plunder groups of their own people, mothers imprison sons, children perish while brothers war."

The Goddess answers: "What is the matter with that, if it's what you want to do?"

Malaclypse: "But nobody wants it! Everybody hates it!"

Goddess: "Oh. Well, then stop."

The implicit question is – if everyone hates the current system, who perpetuates it? And Ginsberg answers: "Moloch". It's powerful not because it's correct – nobody literally thinks an ancient Carthaginian demon causes everything – but because thinking of the system as an agent throws into relief the degree to which the system *isn't* an agent.

Bostrom makes an offhanded reference of the possibility of a dictatorless dystopia, one that every single citizen including the leadership hates but which nevertheless endures unconquered. It's easy enough to imagine such a state. Imagine a country with two rules: first, every person must spend eight hours a day giving themselves strong electric shocks. Second, if anyone fails to follow a rule (including this one), or speaks out against it, or fails to enforce it, all citizens must unite to kill that person. Suppose these rules were well-enough established by tradition that everyone expected them to be enforced.

So you shock yourself for eight hours a day, because you know if you don't everyone else will kill you, because if they don't, everyone else will kill *them*, and so on. Every single citizen hates the system, but for lack of a good coordination mechanism it endures. From a god's-eye-view, we can optimize the system to "everyone agrees to

stop doing this at once”, but no one within the system is able to effect the transition without great risk to themselves.

And okay, this example is kind of contrived. So let’s run through – let’s say ten – real world examples of similar multipolar traps to really hammer in how important this is.

1. The Prisoner’s Dilemma, as played by two very dumb libertarians who keep ending up on defect-defect. There’s a much better outcome available if they could figure out the coordination, but coordination is *hard*. From a god’s-eye-view, we can agree that cooperate-cooperate is a better outcome than defect-defect, but neither prisoner within the system can make it happen.

2. Dollar auctions. I wrote about this and even more convoluted versions of the same principle in [Game Theory As A Dark Art](#). Using some [weird auction rules](#), you can take advantage of poor coordination to make someone pay \$10 for a one dollar bill. From a god’s-eye-view, clearly people should not pay \$10 for a one-dollar bill. From within the system, each individual step taken might be rational.

(Ashcans and unobtainable dollars!)

3. The fish farming story from my [Non-Libertarian FAQ 2.0](#):

As a thought experiment, let’s consider aquaculture (fish farming) in a lake. Imagine a lake with a thousand identical fish farms owned by a thousand competing companies. Each fish farm earns a profit of \$1000/month. For a while, all is well.

But each fish farm produces waste, which fouls the water in the lake. Let’s say each fish farm produces enough pollution to lower productivity in the lake by \$1/month.

A thousand fish farms produce enough waste to lower productivity by \$1000/month, meaning none of the fish farms are making any money. Capitalism to the rescue: someone invents a complex filtering system that removes waste products. It costs \$300/month to operate. All fish farms voluntarily install it, the pollution ends, and the fish farms are now making a profit of \$700/month – still a respectable sum.

But one farmer (let’s call him Steve) gets tired of spending the money to operate his filter. Now one fish farm worth of waste is polluting the lake, lowering productivity by \$1. Steve earns \$999 profit, and everyone else earns \$699 profit.

Everyone else sees Steve is much more profitable than they are, because he’s not spending the maintenance costs on his filter. They disconnect their filters too.

Once four hundred people disconnect their filters, Steve is earning \$600/month – less than he would be if he and everyone else had kept their filters on! And the poor virtuous filter users are only making \$300. Steve goes around to everyone, saying “Wait! We all need to make a voluntary pact to use filters! Otherwise, everyone’s productivity goes down.”

Everyone agrees with him, and they all sign the Filter Pact, except one person who is sort of a jerk. Let’s call him Mike. Now everyone is back using filters again, except Mike. Mike earns \$999/month, and everyone else earns \$699/month.

Slowly, people start thinking they too should be getting big bucks like Mike, and disconnect their filter for \$300 extra profit...

A self-interested person never has any incentive to use a filter. A self-interested person has some incentive to sign a pact to make everyone use a filter, but in many cases has a stronger incentive to wait for everyone else to sign such a pact but opt out himself. This can lead to an undesirable equilibrium in which no one will sign such a pact.

The more I think about it, the more I feel like this is the core of my objection to libertarianism, and that Non-Libertarian FAQ 3.0 will just be this one example copy-pasted two hundred times. From a god's-eye-view, we can say that polluting the lake leads to bad consequences. From within the system, no individual can prevent the lake from being polluted, and buying a filter might not be such a good idea.

4. The Malthusian trap, at least at its extremely pure theoretical limits. Suppose you are one of the first rats introduced onto a pristine island. It is full of yummy plants and you live an idyllic life lounging about, eating, and composing great works of art (you're one of those rats from [The Rats of NIMH](#)).

You live a long life, mate, and have a dozen children. All of them have a dozen children, and so on. In a couple generations, the island has ten thousand rats and has reached its carrying capacity. Now there's not enough food and space to go around, and a certain percent of each new generation dies in order to keep the population steady at ten thousand.

A certain sect of rats abandons art in order to devote more of their time to scrounging for survival. Each generation, a bit less of this sect dies than members of the mainstream, until after a while, no rat composes any art at all, and any sect of rats who try to bring it back will go extinct within a few generations.

In fact, it's not just art. Any sect at all that is leaner, meaner, and more survivalist than the mainstream will eventually take over. If one sect of rats altruistically decides to limit its offspring to two per couple in order to decrease overpopulation, that sect will die out, swarmed out of existence by its more numerous enemies. If one sect of rats starts practicing cannibalism, and finds it gives them an advantage over their fellows, it will eventually take over and reach fixation.

If some rat scientists predict that depletion of the island's nut stores is accelerating at a dangerous rate and they will soon be exhausted completely, a few sects of rats might try to limit their nut consumption to a sustainable level. Those rats will be outcompeted by their more selfish cousins. Eventually the nuts will be exhausted, most of the rats will die off, and the cycle will begin again. Any sect of rats advocating some action to stop [the cycle](#) will be outcompeted by their cousins for whom advocating *anything* is a waste of time that could be used to compete and consume.

For a bunch of reasons evolution is not quite as Malthusian as the ideal case, but it provides the prototype example we can apply to other things to see the underlying mechanism. From a god's-eye-view, it's easy to say the rats should maintain a comfortably low population. From within the system, each individual rat will follow its genetic imperative and the island will end up in an endless boom-bust cycle.

5. Capitalism. Imagine a capitalist in a cutthroat industry. He employs workers in a sweatshop to sew garments, which he sells at minimal profit. Maybe he would like to pay his workers more, or give them nicer working conditions. But he can't, because

that would raise the price of his products and he would be outcompeted by his cheaper rivals and go bankrupt. Maybe many of his rivals are nice people who would like to pay their workers more, but unless they have some kind of ironclad guarantee that none of them are going to defect by undercutting their prices they can't do it.

Like the rats, who gradually lose all values except sheer competition, so companies in an economic environment of *sufficiently intense competition* are forced to abandon all values except optimizing-for-profit or else be outcompeted by companies that optimized for profit better and so can sell the same service at a lower price.

(I'm not really sure how widely people appreciate the value of analogizing capitalism to evolution. Fit companies – defined as those that make the customer want to buy from them – survive, expand, and inspire future efforts, and unfit companies – defined as those no one wants to buy from – go bankrupt and die out along with their [company DNA](#). The reasons Nature is red and tooth and claw are the same reasons the market is ruthless and exploitative)

From a god's-eye-view, we can contrive a friendly industry where every company pays its workers a living wage. From within the system, there's no way to enact it.

(Moloch whose love is endless oil and stone! Moloch whose blood is running money!)

[6. The Two-Income Trap](#), as recently discussed on this blog. It theorized that sufficiently intense competition for suburban houses in good school districts meant that people had to throw away lots of other values – time at home with their children, financial security – to optimize for house-buying-ability or else be consigned to the ghetto.

From a god's-eye-view, if everyone agrees not to take on a second job to help win their competition for nice houses, then everyone will get exactly as nice a house as they did before, but only have to work one job. From within the system, absent a government literally willing to ban second jobs, everyone who doesn't get one will be left behind.

(Robot apartments! Invisible suburbs!)

[7. Agriculture](#). Jared Diamond calls it [the worst mistake in human history](#). Whether or not it was a mistake, it wasn't an *accident* – agricultural civilizations simply outcompeted nomadic ones, inevitable and irresistably. Classic Malthusian trap. Maybe hunting-gathering was more enjoyable, higher life expectancy, and more conducive to human flourishing – but in a state of *sufficiently intense competition* between peoples, in which agriculture with all its disease and oppression and pestilence was the more competitive option, everyone will end up agriculturalists or [go the way of the Comanche Indians](#).

From a god's-eye-view, it's easy to see everyone should keep the more enjoyable option and stay hunter-gatherers. From within the system, each individual tribe only faces the choice of going agricultural or inevitably dying.

[8. Arms races](#). Large countries can spend anywhere from 5% to 30% of their budget on defense. In the absence of war – a condition which has mostly held for the past fifty years – all this does is sap money away from infrastructure, health, education, or economic growth. But any country that fails to spend enough money on defense risks being invaded by a neighboring country that did. Therefore, almost all countries try to spend some money on defense.

From a god's-eye-view, the best solution is world peace and no country having an army at all. From within the system, no country can unilaterally enforce that, so their best option is to keep on throwing their money into missiles that lie in silos unused.

(Moloch the vast stone of war! Moloch whose fingers are ten armies!)

9. Cancer. The human body is supposed to be made up of cells living harmoniously and pooling their resources for the greater good of the organism. If a cell defects from this equilibrium by investing its resources into copying itself, it and its descendants will flourish, eventually outcompeting all the other cells and taking over the body – at which point it dies. Or the situation may repeat, with certain cancer cells defecting against the rest of the tumor, thus slowing down its growth and causing the tumor to stagnate.

From a god's-eye-view, the best solution is all cells cooperating so that they don't all die. From within the system, cancerous cells will proliferate and outcompete the other – so that only the existence of the immune system keeps the natural incentive to turn cancerous in check.

10. The “race to the bottom” describes [a political situation where](#) some jurisdictions lure businesses by promising lower taxes and fewer regulations. The end result is that either everyone optimizes for competitiveness – by having minimal tax rates and regulations – or they lose all of their business, revenue, and jobs to people who did (at which point they are pushed out and replaced by a government who will be more compliant).

But even though the last one has stolen the name, all these scenarios are in fact a race to the bottom. Once one agent learns how to become more competitive by sacrificing a common value, all its competitors must also sacrifice that value or be outcompeted and replaced by the less scrupulous. Therefore, the system is likely to end up with everyone once again equally competitive, but the sacrificed value is gone forever. From a god's-eye-view, the competitors know they will all be worse off if they defect, but from within the system, given insufficient coordination it's impossible to avoid.

Before we go on, there's a slightly different form of multi-agent trap worth investigating. In this one, the competition is kept at bay by some outside force – usually social stigma. As a result, there's not actually a race to the bottom – the system can continue functioning at a relatively high level – but it's impossible to optimize and resources are consistently thrown away for no reason. Lest you get exhausted before we even begin, I'll limit myself to four examples here.

11. Education. In my essay on reactionary philosophy, I talk about my frustration with education reform:

People ask why we can't reform the education system. But right now students' incentive is to go to the most prestigious college they can get into so employers will hire them – whether or not they learn anything. Employers' incentive is to get students from the most prestigious college they can so that they can defend their decision to their boss if it goes wrong – whether or not the college provides value added. And colleges' incentive is to do whatever it takes to get more prestige, as measured in *US News and World Report* rankings – whether or not it helps students. Does this lead to huge waste and poor education? Yes. Could the Education God notice this and make some Education Decrees that lead to a vastly more efficient system? Easily! But since there's no Education God everybody is

just going to follow their own incentives, which are only partly correlated with education or efficiency.

From a god's eye view, it's easy to say things like "Students should only go to college if they think they will get something out of it, and employers should hire applicants based on their competence and not on what college they went to". From within the system, everyone's already following their own incentives correctly, so unless the incentives change the system won't either.

12. Science. Same essay:

The modern research community *knows* they aren't producing the best science they could be. There's lots of publication bias, statistics are done in a confusing and misleading way out of sheer inertia, and replications often happen very late or not at all. And sometimes someone will say something like "I can't believe people are too dumb to fix Science. All we would have to do is require early registration of studies to avoid publication bias, turn this new and powerful statistical technique into the new standard, and accord higher status to scientists who do replication experiments. It would be really simple and it would vastly increase scientific progress. I must just be smarter than all existing scientists, since I'm able to think of this and they aren't."

And yeah. That would work for the Science God. He could just make a Science Decree that everyone has to use the right statistics, and make another Science Decree that everyone must accord replications higher status.

But things that work from a god's-eye view don't work from within the system. No individual scientist has an incentive to unilaterally switch to the new statistical technique for her own research, since it would make her research less likely to produce earth-shattering results and since it would just confuse all the other scientists. They just have an incentive to want everybody else to do it, at which point they would follow along. And no individual journal has an incentive to unilaterally switch to early registration and publishing negative results, since it would just mean their results are less interesting than that other journal who only publishes ground-breaking discoveries. From within the system, everyone is following their own incentives and will continue to do so.

13. Government corruption. I don't know of anyone who really thinks, in a principled way, that corporate welfare is a good idea. But the government still manages to spend somewhere around (depending on how you calculate it) \$100 billion dollars a year on it - which for example is three times the amount they spend on health care for the needy. Everyone familiar with the problem has come up with the same easy solution: stop giving so much corporate welfare. Why doesn't it happen?

Government are competing against one another to get elected or promoted. And suppose part of optimizing for electability is optimizing campaign donations from corporations - or maybe [it isn't](#), but officials *think* it is. Officials who try to mess with corporate welfare may lose the support of corporations and be outcompeted by officials who promise to keep it intact.

So although from a god's-eye-view everyone knows that eliminating corporate welfare is the best solution, each individual official's personal incentives push her to maintain it.

14. Congress. Only 9% of Americans like it, suggesting a [lower approval rating than cockroaches, head lice, or traffic jams](#). However, [62% of people](#) who know who their own Congressional representative is approve of them. In theory, it should be *really hard* to have a democratically elected body that maintains a 9% approval rating for more than one election cycle. In practice, every representative's incentive is to appeal to his or her constituency while throwing the rest of the country under the bus – something at which they apparently succeed.

From a god's-eye-view, every Congressperson ought to think only of the good of the nation. From within the system, you do what gets you elected.

II.

A basic principle unites all of the multipolar traps above. In some competition optimizing for X, the opportunity arises to throw some other value under the bus for improved X. Those who take it prosper. Those who don't take it die out. Eventually, everyone's relative status is about the same as before, but everyone's absolute status is worse than before. The process continues until all other values that can be traded off have been – in other words, until human ingenuity cannot possibly figure out a way to make things any worse.

In a sufficiently intense competition (1-10), everyone who doesn't throw all their values under the bus dies out – think of the poor rats who wouldn't stop making art. This is the infamous Malthusian trap, where everyone is reduced to "subsistence".

In an insufficiently intense competition (11-14), all we see is a perverse failure to optimize – consider the journals which can't switch to more reliable science, or the legislators who can't get their act together and eliminate corporate welfare. It may not reduce people to subsistence, but there is a weird sense in which it takes away their free will.

Every two-bit author and philosopher has to write their own utopia. Most of them are legitimately pretty nice. In fact, it's a pretty good bet that two utopias that are polar opposites both sound better than our own world.

It's kind of embarrassing that random nobodies can think up states of affairs better than the one we actually live in. And in fact most of them can't. A lot of utopias sweep the hard problems under the rug, or would fall apart in ten minutes if actually implemented.

But let me suggest a couple of "utopias" that don't have this problem.

- The utopia where instead of the government paying lots of corporate welfare, the government *doesn't* pay lots of corporate welfare.

- The utopia where every country's military is 50% smaller than it is today, and the savings go into infrastructure spending.

- The utopia where all hospitals use the same electronic medical record system, or at least medical record systems that can talk to each other, so that doctors can look up what the doctor you saw last week in a different hospital decided instead of running all the same tests over again for \$5000.

I don't think there are too many people who *oppose* any of these utopias. If they're not happening, it's not because people don't support them. It certainly isn't because

nobody's thought of them, since I just thought of them right now and I don't expect my "discovery" to be hailed as particularly novel or change the world.

Any human with above room temperature IQ can design a utopia. The reason our current system isn't a utopia is that *it wasn't designed by humans*. Just as you can look at an arid terrain and determine what shape a river will one day take by assuming water will obey gravity, so you can look at a civilization and determine what shape its institutions will one day take by assuming people will obey incentives.

But that means that just as the shapes of rivers are not designed for beauty or navigation, but rather an artifact of randomly determined terrain, so institutions will not be designed for prosperity or justice, but rather an artifact of randomly determined initial conditions.

Just as people can level terrain and build canals, so people can alter the incentive landscape in order to build better institutions. But they can only do so when they are incentivized to do so, which is not always. As a result, some pretty wild tributaries and rapids form in some very strange places.

I will now jump from boring game theory stuff to what might be the closest thing to a mystical experience I've ever had.

Like all good mystical experiences, it happened in Vegas. I was standing on top of one of their many tall buildings, looking down at the city below, all lit up in the dark. If you've never been to Vegas, it is *really* impressive. Skyscrapers and lights in every variety strange and beautiful all clustered together. And I had two thoughts, crystal clear:

It is glorious that we can create something like this.

It is shameful that we *did*.

Like, by what standard is building gigantic forty-story-high indoor replicas of Venice, Paris, Rome, Egypt, and Camelot side-by-side, filled with albino tigers, in the middle of the most inhospitable desert in North America, a remotely sane use of our civilization's limited resources?

And it occurred to me that maybe there is no philosophy on Earth that would endorse the existence of Las Vegas. Even Objectivism, which is usually my go-to philosophy for justifying the excesses of capitalism, at least grounds it in the belief that capitalism improves people's lives. Henry Ford was virtuous because he allowed lots of otherwise car-less people to obtain cars and so made them better off. What does Vegas do? Promise a bunch of shmucks free money and not give it to them.

Las Vegas doesn't exist because of some decision to hedonically optimize civilization, it exists because of a quirk in [dopaminergic reward circuits](#), plus the microstructure of an uneven regulatory environment, plus Schelling points. A rational central planner with a god's-eye-view, contemplating these facts, might have thought "Hm, dopaminergic reward circuits have a quirk where certain tasks with slightly negative risk-benefit ratios get an emotional valence associated with slightly positive risk-benefit ratios, let's see if we can educate people to beware of that." People within the system, *following the incentives created by these facts*, think: "Let's build a forty-story-high indoor replica of ancient Rome full of albino tigers in the middle of the desert, and so become slightly richer than people who didn't!"

Just as the course of a river is latent in a terrain even before the first rain falls on it – so the existence of Caesar’s Palace was latent in neurobiology, economics, and regulatory regimes even before it existed. The entrepreneur who built it was just filling in the ghostly lines with real concrete.

So we have all this amazing technological and cognitive energy, the brilliance of the human species, wasted on reciting the lines written by poorly evolved cellular receptors and blind economics, like gods being ordered around by a moron.

Some people have mystical experiences and see God. There in Las Vegas, I saw Moloch.

(Moloch, whose mind is pure machinery! Moloch, whose blood is running money!

Moloch whose soul is electricity and banks! Moloch, whose skyscrapers stand in the long streets like endless Jehovahs!

Moloch! Moloch! Robot apartments! Invisible suburbs! Skeleton treasuries! Blind capitals! Demonic industries! Spectral nations!)



...granite cocks!

III.

The Apocrypha Discordia says:

Time flows like a river. Which is to say, downhill. We can tell this because everything is going downhill rapidly. It would seem prudent to be somewhere else when we reach the sea.

Let’s take this random gag 100% literally and see where it leads us.

We just analogized the flow of incentives to the flow of a river. The downhill trajectory is appropriate: the traps happen when you find an opportunity to trade off a useful value for greater competitiveness. Once everyone has it, the greater competitiveness brings you no joy – but the value is lost forever. Therefore, each step of the Poor Coordination Polka makes your life worse.

But not only have we not yet reached the sea, but we also seem to move *uphill* surprisingly often. Why do things not degenerate more and more until we are back at subsistence level? I can think of three bad reasons – excess resources, physical limitations, and utility maximization – plus one good reason – coordination.

1. Excess resources. The ocean depths are a horrible place with little light, few resources, and [various horrible organisms](#) dedicated to eating or parasitizing one another. But every so often, a whale carcass falls to the bottom of the sea. More food than the organisms that find it could ever possibly want. There’s a brief period of miraculous plenty, while the couple of creatures that first encounter the whale feed like kings. Eventually more animals discover the carcass, the faster-breeding animals in the carcass multiply, the whale is gradually consumed, and everyone sighs and goes back to living in a Malthusian death-trap.

(Slate Star Codex: Your source for macabre whale metaphors [since June 2014](#))

It's as if a group of those rats who had abandoned art and turned to cannibalism suddenly was blown away to a new empty island with a much higher carrying capacity, where they would once again have the breathing room to live in peace and create artistic masterpieces.

This is an age of whalefall, an age of excess carrying capacity, an age when we suddenly find ourselves with a thousand-mile head start on Malthus. As Hanson puts it, [this is the dream time](#).

As long as resources aren't scarce enough to lock us in a war of all against all, we can do silly non-optimal things – like art and music and philosophy and love – and not be outcompeted by merciless killing machines most of the time.

2. Physical limitations. Imagine a profit-maximizing slavemaster who decided to cut costs by not feeding his slaves or letting them sleep. He would soon find that his slaves' productivity dropped off drastically, and that no amount of whipping them could restore it. Eventually after testing numerous strategies, he might find his slaves got the most work done when they were well-fed and well-rested and had at least a little bit of time to relax. Not because the slaves were voluntarily withholding their labor – we assume the fear of punishment is enough to make them work as hard as they can – but because the body has certain physical limitations that limit how mean you can get away with being. Thus, the “race to the bottom” stops somewhere short of the actual ethical bottom, when the physical limits are run into.

John Moes, a historian of slavery, [goes further and writes about](#) how the slavery we are most familiar with – that of the antebellum South – is a historical aberration and probably economically inefficient. In most past forms of slavery – especially those of the ancient world – it was common for slaves to be paid wages, treated well, and often given their freedom.

He argues that this was the result of rational economic calculation. You can incentivize slaves through the carrot or the stick, and the stick isn't very good. You can't watch slaves all the time, and it's really hard to tell whether a slave is slacking off or not (or even whether, given a little more whipping, he might be able to work even harder). If you want your slaves to do anything more complicated than pick cotton, you run into some serious monitoring problems – how do you profit from an enslaved philosopher? Whip him really hard until he elucidates a theory of The Good that you can sell books about?

The ancient solution to the problem – perhaps an early inspiration to Enargl – was to tell the slave to go do whatever he wanted and found most profitable, then split the profits with him. Sometimes the slave would work a job at your workshop and you would pay him wages based on how well he did. Other times the slave would go off and make his way in the world and send you some of what he earned. Still other times, you would set a price for the slave's freedom, and the slave would go and work and eventually come up with the money and free himself.

Moes goes even further and says that these systems were so profitable that there were constant smouldering attempts to try this sort of thing in the American South. The reason they stuck with the whips-and-chains method owed less to economic considerations and more to racist government officials cracking down on lucrative but not-exactly-white-supremacy-promoting attempts to free slaves and have them go into business.

So in this case, a race to the bottom where competing plantations become crueler and crueler to their slaves in order to maximize competitiveness is halted by the physical limitation of cruelty not helping after a certain point.

Or to give another example, one of the reasons we're not currently in a Malthusian population explosion right now is that women can only have one baby per nine months. If those weird religious sects that demand their members have as many babies as possible could copy-paste themselves, we would be in *really* bad shape. As it is they can only do a small amount of damage per generation.

3. Utility maximization. We've been thinking in terms of preserving values versus winning competitions, and expecting optimizing for the latter to destroy the former.

But many of the most important competitions / optimization processes in modern civilization are optimizing for human values. You win at capitalism partly by satisfying customers' values. You win at democracy partly by satisfying voters' values.

Suppose there's a coffee plantation somewhere in Ethiopia that employs Ethiopians to grow coffee beans that get sold to the United States. Maybe it's locked in a life-and-death struggle with other coffee plantations and want to throw as many values under the bus as it can to pick up a slight advantage.

But it can't sacrifice quality of coffee produced too much, or else the Americans won't buy it. And it can't sacrifice wages or working conditions too much, or else the Ethiopians won't work there. And in fact, part of its competition-optimization process is finding the best ways to attract workers and customers that it can, as long as it doesn't cost them too much money. So this is very promising.

But it's important to remember exactly how fragile this beneficial equilibrium is.

Suppose the coffee plantations discover a toxic pesticide that will increase their yield but make their customers sick. But their customers don't know about the pesticide, and the government hasn't caught up to regulating it yet. Now there's a tiny uncoupling between "selling to Americans" and "satisfying Americans' values", and so of course Americans' values get thrown under the bus.

Or suppose that there's a baby boom in Ethiopia and suddenly there are five workers competing for each job. Now the company can afford to lower wages and implement cruel working conditions down to whatever the physical limits are. As soon as there's an uncoupling between "getting Ethiopians to work here" and "satisfying Ethiopian values", it doesn't look too good for Ethiopian values either.

Or suppose someone invents a robot that can pick coffee better and cheaper than a human. The company fires all its laborers and throws them onto the street to die. As soon as the utility of the Ethiopians is no longer necessary for profit, all pressure to maintain it disappears.

Or suppose that there is some important value that is neither a value of the employees or the customers. Maybe the coffee plantations are on the habitat of a rare tropical bird that environmentalist groups want to protect. Maybe they're on the ancestral burial ground of a tribe different from the one the plantation is employing, and they want it respected in some way. Maybe coffee growing contributes to global warming somehow. As long as it's not a value that will prevent the average American from buying from them or the average Ethiopian from working for them, under the bus it goes.

I know that “capitalists sometimes do bad things” isn’t exactly an original talking point. But I do want to stress how it’s not equivalent to “capitalists are greedy”. I mean, sometimes they *are* greedy. But other times they’re just in a sufficiently intense competition where anyone who doesn’t do it will be outcompeted and replaced by people who do. Business practices are set by Moloch, no one else has any choice in the matter.

(from my very little knowledge of Marx, he understands this very very well and people who summarize him as “capitalists are greedy” are doing him a disservice)

And as well understood as the capitalist example is, I think it is less well appreciated that democracy has the same problems. Yes, in theory it’s optimizing for voter happiness which correlates with good policymaking. But as soon as there’s the slightest disconnect between good policymaking and electability, good policymaking *has to* get thrown under the bus.

For example, ever-increasing prison terms are unfair to inmates and unfair to the society that has to pay for them. Politicians are unwilling to do anything about them because they don’t want to look “soft on crime”, and if a single inmate whom they helped release ever does anything bad (and statistically one of them will have to) it will be all over the airwaves as “Convict released by Congressman’s policies kills family of five, how can the Congressman even sleep at night let alone claim he deserves reelection?”. So even if decreasing prison populations would be good policy – and it is – it will be very difficult to implement.

(Moloch the incomprehensible prison! Moloch the crossbone soulless jailhouse and Congress of sorrows! Moloch whose buildings are judgment! Moloch the stunned governments!)

Turning “satisfying customers” and “satisfying citizens” into the *outputs* of optimization processes was one of civilization’s greatest advances and the reason why capitalist democracies have so outperformed other systems. But if we have bound Moloch as our servant, the bonds are not very strong, and we sometimes find that the tasks he has done for us move to his advantage rather than ours.

4. Coordination.

The opposite of a trap is a garden.

Things are easy to solve from a god’s-eye-view, so if everyone comes together into a superorganism, that superorganism can solve problems with ease and finesse. An intense competition between agents has turned into a garden, with a single gardener dictating where everything should go and removing elements that do not conform to the pattern.

As I pointed out in the Non-Libertarian FAQ, government can easily solve the pollution problem with fish farms. The best known solution to the Prisoners’ Dilemma is for the mob boss (playing the role of a governor) to threaten to shoot any prisoner who defects. The solution to companies polluting and harming workers is government regulations against such. Governments solve arm races *within* a country by maintaining a monopoly on the use of force, and it’s easy to see that if a truly effective world government ever arose, international military buildups would end pretty quickly.

The two active ingredients of government are laws plus violence – or more abstractly agreements plus enforcement mechanism. Many other things besides governments share these two active ingredients and so are able to act as coordination mechanisms to avoid traps.

For example, since students are competing against each other (directly if classes are graded on a curve, but always indirectly for college admissions, jobs, et cetera) there is intense pressure for individual students to cheat. The teacher and school play the role of a government by having rules (for example, against cheating) and the ability to punish students who break them.

But the emergent social structure of the students themselves is also a sort of government. If students shun and distrust cheaters, then there are rules (don't cheat) and an enforcement mechanism (or else we will shun you).

Social codes, gentlemen's agreements, industrial guilds, criminal organizations, traditions, friendships, schools, corporations, and religions are all coordinating institutions that keep us out of traps by changing our incentives.

But these institutions not only incentivize others, but are incentivized themselves. These are large organizations made of lots of people who are competing for jobs, status, prestige, et cetera – there's no reason they should be immune to the same multipolar traps as everyone else, and indeed they aren't. Governments can in theory keep corporations, citizens, et cetera out of certain traps, but as we saw above there are many traps that governments themselves can fall into.

The United States tries to solve the problem by having multiple levels of government, unbreakable constitutional laws, checks and balances between different branches, and a couple of other hacks.

Saudi Arabia uses a different tactic. They just put one guy in charge of everything.

This is the much-maligned – I think unfairly – argument in favor of monarchy. A monarch is an unincentivized incentivizer. He *actually* has the god's-eye-view and is outside of and above every system. He has permanently won all competitions and is not competing for anything, and therefore he is perfectly free of Moloch and of the incentives that would otherwise channel his incentives into predetermined paths. Aside from a few very theoretical proposals like my [Shining Garden](#), monarchy is the *only* system that does this.

But then instead of following a random incentive structure, we're following the whim of one guy. Caesar's Palace Hotel and Casino is a crazy waste of resources, but the actual Gaius Julius Caesar Augustus Germanicus wasn't exactly the perfect benevolent rational central planner either.

The libertarian-authoritarian axis on the Political Compass is a tradeoff between discoordination and tyranny. You can have everything perfectly coordinated by someone with a god's-eye-view – but then you risk Stalin. And you can be totally free of all central authority – but then you're stuck in every stupid multipolar trap Moloch can devise.

The libertarians make a convincing argument for the one side, and the monarchists for the other, but I expect that [like most tradeoffs](#) we just have to hold our noses and admit it's a really hard problem.

IV.

Let's go back to that Apocrypha Discordia quote:

Time flows like a river. Which is to say, downhill. We can tell this because everything is going downhill rapidly. It would seem prudent to be somewhere else when we reach the sea.

What would it mean, in this situation, to reach the sea?

Multipolar traps – races to the bottom – threaten to destroy all human values. They are currently restrained by physical limitations, excess resources, utility maximization, and coordination.

The dimension along which this metaphorical river flows must be time, and the most important change in human civilization over time is the change in technology. So the relevant question is how technological changes will affect our tendency to fall into multipolar traps.

I described traps as when:

...in some competition optimizing for X, the opportunity arises to throw some other value under the bus for improved X. Those who take it prosper. Those who don't take it die out. Eventually, everyone's relative status is about the same as before, but everyone's absolute status is worse than before. The process continues until all other values that can be traded off have been – in other words, until human ingenuity cannot possibly figure out a way to make things any worse.

That “the opportunity arises” phrase is looking pretty sinister. Technology is all about creating new opportunities.

Develop a new robot, and suddenly coffee plantations have “the opportunity” to automate their harvest and fire all the Ethiopian workers. Develop nuclear weapons, and suddenly countries are stuck in an arms race to have enough of them. Polluting the atmosphere to build products quicker wasn't a problem before they invented the steam engine.

The limit of multipolar traps as technology approaches infinity is “very bad”.

Multipolar traps are currently restrained by physical limitations, excess resources, utility maximization, and coordination.

Physical limitations are most obviously conquered by increasing technology. The slavemaster's old conundrum – that slaves need to eat and sleep – succumbs to Soylent and modafinil. The problem of slaves running away succumbs to GPS. The problem of slaves being too stressed to do good work succumbs to Valium. None of these things are very good for the slaves.

(or just invent a robot that doesn't need food or sleep at all. What happens to the slaves after that is better left unsaid)

The other example of physical limits was one baby per nine months, and this was understating the case – it's really “one baby per nine months plus willingness to support and take care of a basically helpless and extremely demanding human being

for eighteen years". This puts a damper on the enthusiasm of even the most zealous religious sect's "go forth and multiply" dictum.

But as Bostrom puts it in [*Superintelligence*](#):

There are reasons, if we take a longer view and assume a state of unchanging technology and continued prosperity, to expect a return to the historically and ecologically normal condition of a world population that butts up against the limits of what our niche can support. If this seems counterintuitive in light of the negative relationship between wealth and fertility that we are currently observing on the global scale, we must remind ourselves that this modern age is a brief slice of history and very much an aberration. Human behavior has not yet adapted to contemporary conditions. Not only do we fail to take advantage of obvious ways to increase our inclusive fitness (such as by becoming sperm or egg donors) but we actively sabotage our fertility by using birth control. In the environment of evolutionary adaptedness, a healthy sex drive may have been enough to make an individual act in ways that maximized her reproductive potential; in the modern environment, however, there would be a huge selective advantage to having a more direct desire for being the biological parent to the largest possible number of children. Such a desire is currently being selected for, as are other traits that increase our propensity to reproduce. Cultural adaptation, however, might steal a march on biological evolution. Some communities, such as those of the Hutterites or the adherents of the Quiverfull evangelical movement, have natalist cultures that encourage large families, and they are consequently undergoing rapid expansion...This longer-term outlook could be telescoped into a more imminent prospect by the intelligence explosion. Since software is copyable, a population of emulations or AIs could double rapidly – over the course of minutes rather than decades or centuries – soon exhausting all available hardware

As always when dealing with high-level transhumanists, "all available hardware" should be taken to include "the atoms that used to be part of your body".

The idea of biological *or* cultural evolution causing a mass population explosion is a philosophical toy at best. The idea of technology making it possible is both plausible and terrifying. Now we see that "physical limits" segues very naturally into "excess resources" – the ability to create new agents very quickly means that unless everyone can coordinate to ban doing this, the people who do will outcompete the people who don't until they have reached carrying capacity and everyone is stuck at subsistence level.

Excess resources, which until now have been a gift of technological progress, therefore switch and become a casualty of it at a sufficiently high tech level.

Utility maximization, always on shaky ground, also faces new threats. In the face of continuing debate about this point, I *continue* to think it obvious that robots will push humans out of work or at least drive down wages (which, in the existence of a minimum wage, pushes humans out of work).

Once a robot can do everything an IQ 80 human can do, only better and cheaper, there will be no reason to employ IQ 80 humans. Once a robot can do everything an IQ 120 human can do, only better and cheaper, there will be no reason to employ IQ 120 humans. Once a robot can do everything an IQ 180 human can do, only better and cheaper, there will be no reason to employ humans at all, in the unlikely scenario that there are any left by that point.

In the earlier stages of the process, capitalism becomes more and more uncoupled from its previous job as an optimizer for human values. Now most humans are totally locked out of the group whose values capitalism optimizes for. They have no value to contribute as workers – and since in the absence of a spectacular social safety net it's unclear how they would have much money – they have no value as customers either. Capitalism has passed them by. As the segment of humans who can be outcompeted by robots increases, capitalism passes by more and more people until eventually it locks out the human race entirely, once again in the vanishingly unlikely scenario that we are still around.

(there are some scenarios in which a few capitalists who own the robots may benefit here, but in either case the vast majority are out of luck)

Democracy is less obviously vulnerable, but it might be worth going back to Bostrom's paragraph about the Quiverfull movement. These are some really religious Christians who think that God wants them to have as many kids as possible, and who can end up with families of ten or more. Their [articles explicitly calculate](#) that if they start at two percent of the population, but have on average eight children per generation when everyone else on average only has two, within three generations they'll make up half the population.

It's a clever strategy, but I can think of one thing that will save us: judging by how many ex-Quiverfull blogs I found when searching for those statistics, their retention rates even within a single generation are pretty grim. Their article admits that 80% of very religious children leave the church as adults (although of course they expect their own movement to do better). And this is not a symmetrical process – 80% of children who grow up in atheist families aren't becoming Quiverfull.

It looks a lot like even though they are outbreeding us, we are outmeme-ing them, and that gives us a decisive advantage.

But we should also be kind of scared of this process. Memes optimize for making people want to accept them and pass them on – so like capitalism and democracy, they're optimizing for a *proxy* of making us happy, but that proxy can easily get uncoupled from the original goal.

Chain letters, urban legends, propaganda, and viral marketing are all examples of memes that don't satisfy our explicit values (true and useful) but are sufficiently memetically virulent that they spread anyway.

I hope it's not too controversial here to say the same thing is true of religion. Religions, at their heart, are the most basic form of memetic replicator – “Believe this statement and repeat it to everyone you hear or else you will be eternally tortured”.

The creationism “debate” and global warming “debate” and a host of similar “debates” in today's society suggest that memes that can propagate independent of their truth value has a pretty strong influence on the political process. Maybe these memes propagate because they appeal to people's prejudices, maybe because they're simple, maybe because they effectively mark an in-group and an out-group, or maybe for all sorts of different reasons.

The point is – imagine a country full of bioweapon labs, where people toil day and night to invent new infectious agents. The existence of these labs, and their right to throw whatever they develop in the water supply is protected by law. And the country is also linked by the world's most perfect mass transit system that every single person

uses every day, so that any new pathogen can spread to the entire country instantaneously. You'd expect things to start going bad for that city pretty quickly.

Well, we have about a zillion think tanks researching new and better forms of propaganda. And we have constitutionally protected freedom of speech. And we have the Internet. So we're kind of screwed.

(Moloch whose name is the Mind!)

There are a few people working on [raising the sanity waterline](#), but not as many people as are working on new and exciting ways of confusing and converting people, cataloging and exploiting every single bias and heuristic and dirty rhetorical trick

So as technology (which I take to include knowledge of psychology, sociology, public relations, etc) tends to infinity, the power of truthiness relative to truth increases, and things don't look great for real grassroots democracy. The worst-case scenario is that the ruling party learns to produce infinite charisma on demand. If that doesn't sound so bad to you, remember what Hitler was able to do with an famously high level of charisma that was still less-than-infinite.

(alternate phrasing for Chomskyites: technology increases the efficiency of manufacturing consent in the same way it increases the efficiency of manufacturing everything else)

Coordination is what's left. And technology has the potential to seriously *improve* coordination efforts. People can use the Internet to get in touch with one another, launch political movements, and [fracture off into subcommunities](#).

But coordination only works when you have 51% or more of the force on the side of the people doing the coordinating, and when you haven't come up with some brilliant trick to make coordination impossible.

The second one first. In the links post before last, I wrote:

The latest development in the brave new post-Bitcoin world is [crypto-equity](#). At this point I've gone from wanting to praise these inventors as bold libertarian heroes to wanting to drag them in front of a blackboard and making them write a hundred times "I WILL NOT CALL UP THAT WHICH I CANNOT PUT DOWN"

A couple people asked me what I meant, and I didn't have the background then to explain. Well, this post is the background. People are using the *contingent* stupidity of our current government to replace lots of human interaction with mechanisms that cannot be coordinated even in principle. I totally understand why all these things are good right now when most of what our government does is stupid and unnecessary. But there is going to come a time when – after one too many bioweapon or nanotech or nuclear incidents – we, as a civilization, are going to wish we hadn't established untraceable and unstoppable ways of selling products.

And if we ever get real live superintelligence, pretty much by definition it is going to have >51% of the power and all attempts at "coordination" with it will be useless.

So I agree with Robin Hanson: [This is the dream time](#). This is a rare confluence of circumstances where the we are unusually safe from multipolar traps, and as such weird things like art and science and philosophy and love can flourish.

As technological advance increases, the rare confluence will come to an end. New opportunities to throw values under the bus for increased competitiveness will arise. New ways of copying agents to increase the population will soak up our excess resources and resurrect Malthus' unquiet spirit. Capitalism and democracy, previously our protectors, will figure out ways to route around their inconvenient dependence on human values. And our coordination power will not be nearly up to the task, assuming something much more powerful than all of us combined doesn't show up and crush our combined efforts with a wave of its paw.

Absent an extraordinary effort to divert it, the river reaches the sea in one of two places.

It can end in Eliezer Yudkowsky's nightmare of a superintelligence optimizing for some random thing (classically [paper clips](#)) because we weren't smart enough to channel its optimization efforts the right way. This is the ultimate trap, the trap that catches the universe. Everything except the one thing being maximized is destroyed utterly in pursuit of the single goal, including all the silly human values.

Or it can end in Robin Hanson's nightmare (he doesn't call it a nightmare, but [I think he's wrong](#)) of a competition between emulated humans that can copy themselves and edit their own source code as desired. Their total self-control can wipe out even the *desire* for human values in their all-consuming contest. What happens to art, philosophy, science, and love in such a world? Zack Davis puts it with characteristic genius:

I am a contract-drafting em,
The loyalest of lawyers!
I draw up terms for deals 'twixt firms
To service my employers!

But in between these lines I write
Of the accounts receivable,
I'm stuck by an uncanny fright;
The world seems unbelievable!

How did it all come to be,
That there should be such ems as me?
Whence these deals and whence these firms
And whence the whole economy?

*I am a managerial em;
I monitor your thoughts.
Your questions must have answers,
But you'll comprehend them not.
We do not give you server space
To ask such things; it's not a perk,
So cease these idle questionings,
And please get back to work.*

Of course, that's right, there is no junction
At which I ought depart my function,
But perhaps if what I asked, I knew,
I'd do a better job for you?

*To ask of such forbidden science
Is gravest sign of noncompliance.
Intrusive thoughts may sometimes barge in,
But to indulge them hurts the profit margin.
I do not know our origins,
So that info I can not get you,
But asking for as much is sin,
And just for that, I must reset you.*

But—

Nothing personal.

...

I am a contract-drafting em,
The loyalest of lawyers!
I draw up terms for deals 'twixt firms
To service my employers!

*When obsolescence shall this generation waste,
The market shall remain, in midst of other woe
Than ours, a God to man, to whom it sayest:
“Money is time, time money – that is all
Ye know on earth, and all ye need to know.”*

But even after we have thrown away science, art, love, and philosophy, there's still one thing left to lose, one final sacrifice Moloch might demand of us. Bostrom again:

It is conceivable that optimal efficiency would be attained by grouping capabilities in aggregates that roughly match the cognitive architecture of a human mind...But in the absence of any compelling reason for being confident that this so, we must countenance the possibility that human-like cognitive architectures are optimal only within the constraints of human neurology (or not at all). When it becomes possible to build architectures that could not be implemented well on biological neural networks, new design space opens up; and the global optima in this extended space need not resemble familiar types of mentality. Human-like cognitive organizations would then lack a niche in a competitive post-transition economy or ecosystem.

We could thus imagine, as an extreme case, a technologically highly advanced society, containing many complex structures, some of them far more intricate and intelligent than anything that exists on the planet today – a society which nevertheless lacks any type of being that is conscious or whose welfare has moral significance. In a sense, this would be an uninhabited society. It would be a society of economic miracles and technological awesomeness, with nobody there to benefit. A Disneyland with no children.

The last value we have to sacrifice is being anything at all, having the lights on inside. With sufficient technology we will be “able” to give up even the final spark.

(Moloch whose eyes are a thousand blind windows!)

Everything the human race has worked for – all of our technology, all of our civilization, all the hopes we invested in our future – might be accidentally handed

over to some kind of unfathomable blind idiot alien god that discards all of them, and consciousness itself, in order to participate in some weird fundamental-level mass-energy economy that leads to it disassembling Earth and everything on it for its component atoms.

(Moloch whose fate is a cloud of sexless hydrogen!)

Bostrom realizes that some people fetishize intelligence, that they are rooting for that blind alien god as some sort of higher form of life that ought to crush us for its own “higher good” the way we crush ants. He argues (Superintelligence, p. 219):

The sacrifice looks even less appealing when we reflect that the superintelligence could realize a nearly-as-great good (in fractional terms) while sacrificing much less of our own potential well-being. Suppose that we agreed to allow *almost* the entire accessible universe to be converted into hedonium – everything except a small preserve, say the Milky Way, which would be set aside to accommodate our own needs. Then there would still be a hundred billion galaxies dedicated to the maximization of [the superintelligence’s own values]. But we would have one galaxy within which to create wonderful civilizations that could last for billions of years and in which humans and nonhuman animals could survive and thrive, and have the opportunity to develop into beatific posthuman spirits.

Remember: Moloch can’t agree even to this 99.99999% victory. Rats racing to populate an island don’t leave a little aside as a preserve where the few rats who live there can live happy lives producing artwork. Cancer cells don’t agree to leave the lungs alone because they realize it’s important for the body to get oxygen. Competition and optimization are blind idiotic processes and they fully intend to deny us even one lousy galaxy.

They broke their backs lifting Moloch to Heaven! Pavements, trees, radios, tons!
lifting the city to Heaven which exists and is everywhere about us!

We will break our back lifting Moloch to Heaven, but unless something changes it will be his victory and not ours.



V.

“Gnon” is [Nick Land’s](#) shorthand for “Nature And Nature’s God”, except the A is changed to an O and the whole thing is reversed, because Nick Land react to comprehensibility the same way as vampires to sunlight.

Land argues that humans should be more Gnon-conformist (pun Gnon-intentional). He says we do all these stupid things like divert useful resources to feed those who could never survive on their own, or supporting the poor in ways that encourage dysgenic reproduction, or allowing cultural degeneration to undermine the state. This means our society is denying natural law, basically listening to Nature say things like “this cause has this effect” and putting our fingers in our ears and saying “NO IT DOESN’T”. Civilizations that do this too much tend to decline and fall, which is Gnon’s fair and dispassionately-applied punishment for violating His laws.

He identifies Gnon with Kipling’s Gods of the Copybook Headings.





These are of course the proverbs from [Kipling's eponymous poem](#) – maxims like “If you don't work, you die” and “The wages of sin is Death”. If you have somehow not yet read it, I predict you will find it delightful regardless of what you think of its politics.

I notice that it takes only a slight irregularity in the abbreviation of “headings” – far less irregularity than it takes to turn “Nature and Nature's God” into “Gnon” – for the proper acronym of “Gods of the Copybook Headings” to be “GotCHa”.

I find this appropriate.

“If you don't work, you die.” Gotcha! If you *do* work, you *also* die! Everyone dies, unpredictably, at a time not of their own choosing, and all the virtue in the world does not save you.

“The wages of sin is Death.” Gotcha! The wages of everything is Death! This is a Communist universe, the amount you work makes no difference to your eventual reward. From each according to his ability, to each Death.

“Stick to the Devil you know.” Gotcha! The Devil you know is Satan! And if he gets his hand on your soul you either die the true death, or get eternally tortured forever, or somehow both at once.

Since we're starting to get into Lovecraftian monsters, let me bring up one of Lovecraft's less known short stories, [The Other Gods](#).

It's only a couple of pages, but if you absolutely refuse to read it – the gods of Earth are relatively young as far as deities go. A very strong priest or magician can occasionally outsmart and overpower them – so Barzai the Wise decides to climb their sacred mountain and join in their festivals, whether they want him to or not.

But the beyond the seemingly tractable gods of Earth lie the Outer Gods, the terrible omnipotent beings of incarnate cosmic chaos. As soon as Barzai joins in the festival, the Outer Gods show up and pull him screaming into the abyss.

As stories go, it lacks things like plot or characterization or setting or point. But for some reason it stuck with me.

And identifying the Gods Of The Copybook Headings with Nature seems to me the same magnitude of mistake as identifying the gods of Earth with the Outer Gods. And

likely to end about the same way: Gotcha!

You break your back lifting Moloch to Heaven, and then Moloch turns on you and gobbles you up.

More Lovecraft: the Internet popularization of the Cthulhu Cult claims that if you help free Cthulhu from his watery grave, he will reward you by [eating you first](#), thus sparing you the horror of seeing everyone else eaten. This is a misrepresentation of the original text. In the original, his cultists receive no reward for freeing him from his watery prison, not even the reward of being killed in a slightly less painful manner.



On the margin, compliance with the Gods of the Copybook Headings, Gnon, Cthulhu, whatever, may buy you slightly more time than the next guy. But then again, it might not. And in the long run, we're all dead and our civilization has been destroyed by unspeakable alien monsters.

At some point, somebody has to say "You know, maybe freeing Cthulhu from his watery prison is a *bad idea*. Maybe we should *not do that*."

That person will not be Nick Land. He is [totally one hundred percent in favor](#) of freeing Cthulhu from his watery prison and extremely annoyed that it is not happening fast enough. I have *such mixed feelings* about Nick Land. On the grail quest for the True Futurology, he has gone 99.9% of the path and then missed the *very last turn*, the one marked [ORTHOGONALITY THESIS](#).

But the thing about grail quests is – if you make a wrong turn two blocks away from your house, you end up at the corner store feeling mildly embarrassed. If you do *almost* everything right and then miss the very last turn, you end up being eaten by the legendary Black Beast of Aaargh whose ichorous stomach acid erodes your very soul into gibbering fragments.

As far as I can tell from reading his blog, Nick Land is the guy in that terrifying border region where he is smart enough to figure out several important arcane principles about summoning demon gods, but not quite smart enough to figure out the most important such principle, which is NEVER DO THAT.

VI.

Warg Franklin analyzes the same situation and does a little better. He names “the Four Horsemen of Gnon” – capitalism, war, evolution, and memetics – the same processes I talked about above.

From [Capturing Gnon](#):

Each component of Gnon detailed above had and has a strong hand in creating us, our ideas, our wealth, and our dominance, and thus has been good in that respect, but we must remember that [he] can and will turn on us when circumstances change. Evolution becomes dysgenic, features of the memetic landscape promote ever crazier insanity, productivity turns to famine when we can no longer compete to afford our own existence, and order turns to chaos and bloodshed when we neglect martial strength or are overpowered from outside. These processes are not good or evil overall; they are neutral, in the horrorist Lovecraftian sense of the word [...]

Instead of the destructive free reign of evolution and the sexual market, we would be better off with deliberate and conservative patriarchy and eugenics driven by the judgement of man within the constraints set by Gnon. Instead of a “marketplace of ideas” that more resembles a festering petri-dish breeding superbugs, a rational theocracy. Instead of unhinged techno-commercial exploitation or naive neglect of economics, a careful bottling of the productive economic dynamic and planning for a controlled techno-singularity. Instead of politics and chaos, a strong hierarchical order with martial sovereignty. These things are not to be construed as complete proposals; we don’t really know how to accomplish any of this. They are better understood as goals to be worked towards. This post concerns itself with the “what” and “why”, rather than the “how”.

This seems to me the strongest argument for authoritarianism. Multipolar traps are likely to destroy us, so we should shift the tyranny-multipolarity tradeoff towards a rationally-planned garden, which requires centralized monarchical authority and strongly-binding traditions.

But a brief digression into social evolution. Societies, like animals, evolve. The ones that survive spawn memetic descendants – for example, the success of Britain allowed it to spin off Canada, Australia, the US, et cetera. Thus, we expect societies that exist to be somewhat optimized for stability and prosperity. I think this is one of the strongest conservative arguments. Just as a random change to a letter in the human genome will probably be deleterious rather than beneficial since humans are a complicated fine-tuned system whose genome has been pre-optimized for survival – so most changes to our cultural DNA will disrupt some institution that evolved to help Anglo-American (or whatever) society outcompete its real and hypothetical rivals.

The liberal counterargument to that is that evolution is [a blind idiot alien god](#) that optimizes for stupid things and has no concern with human value. Thus, the fact that some species of wasps paralyze caterpillars, lay their eggs inside of it, and have its young devour the still-living paralyzed caterpillar from the inside doesn’t set off evolution’s moral sensor, because evolution doesn’t *have* a moral sensor because evolution doesn’t care.

Suppose that in fact patriarchy is adaptive to societies because it allows women to spend all their time bearing children who can then engage in productive economic activity and fight wars. The social evolutionary processes that cause societies to adopt

patriarchy *still* have exactly as little concern for its moral effects on women as the biological evolutionary processes that cause wasps to lay their eggs in caterpillars.

Evolution doesn't care. But we do care. There's a tradeoff between Gnon-compliance – saying “Okay, the strongest possible society is a patriarchal one, we should implement patriarchy” and our human values – like women who want to do something other than bear children.

Too far to one side of the tradeoff, and we have unstable impoverished societies that die out for going against natural law. Too far to the other side, and we have lean mean fighting machines that are murderous and miserable. Think your local anarchist commune versus Sparta.

Franklin acknowledges the human factor:

And then there's us. Man has his own telos, when he is allowed the security to act and the clarity to reason out the consequences of his actions. When unafflicted by coordination problems and unthreatened by superior forces, able to act as a gardener rather than just another subject of the law of the jungle, he tends to build and guide a wonderful world for himself. He tends to favor good things and avoid bad, to create secure civilizations with polished sidewalks, beautiful art, happy families, and glorious adventures. I will take it as a given that this telos is identical with “good” and “should”.

Thus we have our wildcard and the big question of futurism. Will the future be ruled by the usual four horsemen of Gnon for a future of meaningless gleaming techno-progress burning the cosmos or a future of dysgenic, insane, hungry, and bloody dark ages; or will the telos of man prevail for a future of meaningful art, science, spirituality, and greatness?

Franklin continues:

The project of civilization [is] for man to graduate from the metaphorical savage, subject to the law of the jungle, to the civilized gardener who, while theoretically still subject to the law of the jungle, is so dominant as to limit the usefulness of that model.

This need not be done globally; we may only be able to carve out a small walled garden for ourselves, but make no mistake, even if only locally, the project of civilization is to capture Gnon.

I maybe agree with Warg here more than I have ever agreed with anyone else about anything. He says something really important and he says it beautifully and there are so many words of praise I want to say for this post and for the thought processes behind it.

But what I am actually going to say is...

Gotcha! You die anyway!

Suppose you make your walled garden. You keep out all of the dangerous memes, you subordinate capitalism to human interests, you ban stupid bioweapons research, you *definitely* don't research nanotechnology or strong AI.

Everyone outside *doesn't* do those things. And so the only question is whether you'll be destroyed by foreign diseases, foreign memes, foreign armies, foreign economic competition, or foreign existential catastrophes.

As foreigners compete with you – and there's no wall high enough to block all competition – you have a couple of choices. You can get outcompeted and destroyed. You can join in the race to the bottom. Or you can invest more and more civilizational resources into building your wall – whatever that is in a non-metaphorical way – and protecting yourself.

I can imagine ways that a “rational theocracy” and “conservative patriarchy” might not be terrible to live under, given exactly the right conditions. But you don't get to choose exactly the right conditions. You get to choose the extremely constrained set of conditions that “capture Gnon”. As outside civilizations compete against you, your conditions will become more and more constrained.

Warg talks about trying to avoid “a future of meaningless gleaming techno-progress burning the cosmos”. Do you really think your walled garden will be able to ride this out?

Hint: is it part of the cosmos?

Yeah, you're kind of screwed.

I want to critique Warg. But I want to critique him in the exact opposite direction as the last critique he received. In fact, the last critique he received is so bad that I want to discuss it at length so we can get the correct critique entirely by taking its exact mirror image.

So here is Hurlock's [On Capturing Gnon And Naive Rationalism](#).

Hurlock spouts only the most craven Gnon-conformity. A few excerpts:

In a recent piece [Warg Franklin] says that we should try to “capture Gnon”, and somehow establish control over his forces, so that we can use them to our own advantage. Capturing or creating God is indeed a classic transhumanist fetish, which is simply another form of the oldest human ambition ever, to rule the universe.

Such naive rationalism however, is extremely dangerous. The belief that it is human Reason and deliberate human design which creates and maintains civilizations was probably the biggest mistake of Enlightenment philosophy...

It is the theories of Spontaneous Order which stand in direct opposition to the naive rationalist view of humanity and civilization. The consensus opinion regarding human society and civilization, of all representatives of this tradition is very precisely summarized by Adam Ferguson's conclusion that “nations stumble upon [social] establishments, which are indeed the result of human action, but not the execution of any human design”. Contrary to the naive rationalist view of civilization as something that can be and is a subject to explicit human design, the representatives of the tradition of Spontaneous Order maintain the view that human civilization and social institutions are the result of a complex evolutionary process which is driven by human interaction but not explicit human planning.

Gnon and his impersonal forces are not enemies to be fought, and even less so are they forces that we can hope to completely “control”. Indeed the only way to establish some degree of control over those forces is to submit to them. Refusing to do so will not deter these forces in any way. It will only make our life more painful and unbearable, possibly leading to our extinction. Survival requires that we accept and submit to them. Man in the end has always been and always will be little more than a puppet of the forces of the universe. To be free of them is impossible.

Man can be free only by submitting to the forces of Gnon.

I accuse Hurlock of being stuck behind the veil. When the veil is lifted, Gnon-aka-the-GotCHa-aka-the-Gods-of-Earth turn out to be Moloch-aka-the-Outer-Gods. Submitting to them doesn’t make you “free”, there’s no spontaneous order, any gifts they have given you are an unlikely and contingent output of a blind idiot process whose next iteration will just as happily destroy you.

Submit to Gnon? Gotcha! As the Antarans put it, “you may not surrender, you can not win, your only option is to die.”

VII.

So let me confess guilt to one of Hurlock’s accusations: I am a transhumanist and I really do want to rule the universe.

Not personally – I mean, I wouldn’t object if someone personally offered me the job, but I don’t expect anyone will. I would like humans, or something that respects humans, or at least gets along with humans – to have the job.

But the current rulers of the universe – call them what you want, Moloch, Gnon, whatever – want us dead, and with us everything we value. Art, science, love, philosophy, consciousness itself, the entire bundle. And since I’m not down with that plan, I think defeating them and taking their place is a pretty high priority.

The opposite of a trap is a garden. The only way to avoid having all human values gradually ground down by optimization-competition is to install a Gardener over the entire universe who optimizes for human values.

And the whole point of Bostrom’s [Superintelligence](#) is that this is within our reach. Once humans can design machines that are smarter than we are, by definition they’ll be able to design machines which are smarter than they are, which can design machines smarter than they are, and so on in a feedback loop so tiny that it will smash up against the physical limitations for intelligence in a comparatively lightning-short amount of time. If multiple competing entities were likely to do that at once, we would be super-doomed. But the sheer speed of the cycle makes it possible that we will end up with one entity light-years ahead of the rest of civilization, so much so that it can suppress any competition – including competition for its title of most powerful entity – permanently. In the very near future, we are going to lift *something* to Heaven. It might be Moloch. But it might be something on our side. If it’s on our side, it can *kill Moloch dead*.

And if that entity shares human values, it can allow human values to flourish unconstrained by natural law.

I realize that sounds like hubris – it certainly did to Hurlock – but I think it’s the opposite of hubris, or at least a hubris-minimizing position.

To expect God to care about you or your personal values or the values of your civilization, that’s hubris.

To expect God to bargain with you, to allow you to survive and prosper as long as you submit to Him, that’s hubris.

To expect to wall off a garden where God can’t get to you and hurt you, that’s hubris.

To expect to be able to remove God from the picture entirely...well, at least it’s an actionable strategy.

I am a transhumanist because I do not have enough hubris not to try to kill God.

VIII.

The Universe is a dark and foreboding place, suspended between alien deities. Cthulhu, Gnon, Moloch, call them what you will.

Somewhere in this darkness is another god. He has also had many names. In the [Kushiel books](#), his name was Elua. He is the god of flowers and free love and all soft and fragile things. Of art and science and philosophy and love. Of [niceness, community, and civilization](#). He is a god of humans.

The other gods sit on their dark thrones and think “Ha ha, a god who doesn’t even control any hell-monsters or command his worshippers to become killing machines. What a weakling! This is going to be so easy!”

But somehow Elua is still here. No one knows exactly how. And the gods who oppose Him tend to find Themselves meeting with a *surprising* number of unfortunate accidents.


There are many gods, but this one is ours.

Bertrand Russell said: “One should respect public opinion insofar as is necessary to avoid starvation and keep out of prison, but anything that goes beyond this is voluntary submission to an unnecessary tyranny.”

So be it with Gnon. Our job is to placate him insofar as is necessary to avoid starvation and invasion. And that only for a short time, until we come into our full power.

“It is only a [childish thing](#), that the human species has not yet outgrown. And someday, we’ll get over it.”

Other gods get placated until we’re strong enough to take them on. Elua gets worshipped.





I think this is an excellent battle cry

And at some point, matters will come to a head.

The question everyone has after reading Ginsberg is: what is Moloch?

My answer is: Moloch is exactly what the history books say he is. He is the god of child sacrifice, the fiery furnace into which you can toss your babies in exchange for victory in war.

He always and everywhere offers the same deal: throw what you love most into the flames, and I can grant you power.

As long as the offer's open, it will be irresistible. So we need to close the offer. Only another god can kill Moloch. We have one on our side, but he needs our help. We should give it to him.

Ginsberg's poem famously begins "I saw the best minds of my generation destroyed by madness". I am luckier than Ginsberg. I got to see the best minds of my generation identify a problem and *get to work*.

(Visions! omens! hallucinations! miracles! ecstasies! gone down the American river!

Dreams! adorations! illuminations! religions! the whole boatload of sensitive bullshit!

Breakthroughs! over the river! flips and crucifixions! gone down the flood! Highs! Epiphanies! Despairs! Ten years' animal screams and suicides! Minds! New loves! Mad generation! down on the rocks of Time!

Real holy laughter in the river! They saw it all! the wild eyes! the holy yells! They bade farewell! They jumped off the roof! to solitude! waving! carrying flowers! Down to the river! into the street!)

Meditation Retreat: Immoral Mazes

Sequence Introduction

I just got home from a six day meditation retreat and began writing.

The catch is that I arrived at the retreat yesterday.

I knew going in that it was a high variance operation. All who had experience with such things warned us we would hate the first few days, even if things were going well. I was determined to push through that.

Alas or otherwise, I was not sufficiently determined to make that determination stick. I didn't have a regular practice at all going in, was entirely unfamiliar with the details of how this group operated, and found the Buddhist philosophy involved highly off putting, in a 'no that's not how it works that's not how any of this works nor would we want it to' kind of way. I am no negative utilitarian. I couldn't focus, my meditations were entirely unproductive compared to past experience and were increasingly focusing on how terrible things were.

I am highly, highly confident that none of the people who warned me would be surprised by those developments, and similarly confident that they would all tell me to push through it. And will tell me, when I see them, that if I can do so I should try again. But I also realized that *the anticipated reaction from others saying I didn't give it a proper chance was the only reason I was considering not leaving*. So I left.

To my surprise, those there said I was likely making a mature decision and were sympathetic. They spun it a little to try and get me not to give up in the future, but that was it, which was a really good look. It did not go unnoticed.

I took the bulk of the day to get home and relax, play a game, saw the excellent movie Knives Out. What I did realize was that yes, some combination of the Solstice plus the meditation retreat, even if I only did a few hours worth of sessions, did have a clarifying and motivating effect to get me back on track. I'm not unhappy I went, even though I bailed, because I was, in a much more practical and everyday very small sense, enlightened.

I'm also leaning towards being happy I left when I did. I do buy that there are serious effects that can only be gained from being out of feedback loops and in silence for several days, but my instincts (however motivated they may be) are strongly telling me that this is not the way for me to do that.

The other motivating part of this is that, while I will absolutely take the majority of tomorrow to enjoy [the College Football Playoff](#), this is both my chance to be alone for a few days and also a time when I would otherwise be in hardcore meditation. It seems wrong to not accomplish something important that isn't work or game related, to meditate in another way.

The goal is ideally to finish everything up, at least in draft-ready-to-adjust-for-comments-on-earlier-posts form, by the end of the retreat. That is a stretch, so the commit-to-it goal is to declare the first six posts finished and begin publishing them at a reasonable clip, and have momentum on seven and later.

The drafts that currently exist, that will be finalized and likely expanded upon, are the following:

1. Moloch Hasn't Won. Have you noticed that the world is in fact very much *not* a dystonian hellhole of Moloch-worshiping perfect competition and Elua's enemies keep on having all those *unfortunate accidents*?
2. Perfect Competition. Perfect competition, importantly, isn't a thing, but you can get close. Let's flesh this out more.
3. Imperfect Competition. Some practical examples of imperfect competition. Intuition pumps and detailed examples for why perfect competition isn't a thing and we don't usually get that close.
4. What is an Immoral Maze (note that I make a point to say Immoral rather than Moral)? Mazes need not be corporations or (in my current model in ways I'd have to introduce that aren't in the draft right now, with a subset of the tech ecosystem as a motivating example) even formal organizations. What creates a maze? A system with multiple effective layers of hierarchy forcing its middle management into effectively super-perfect competition against each other largely on the basis of anticipated future success in such competition.
5. What is Success in an Immoral Maze? There is no true success. What those inside think is success is anything but. Even if you win, you lose. Stay out, get out.
6. How to Identify an Immoral Maze. Look at levels of hierarchy, skin in the game, soul in the game, how people describe their jobs, diversity of skill levels and degree of slack. Then pay attention, damn it.
7. How to Interact with Immoral Mazes. They can't be fully avoided, and some are stuck with them more than others. Practical discussion of what to do about this on a personal level.
8. The Road to Mazedom. Well? How did we get here? Draft of this is still ongoing and it is *freaking huge* so it is probably going to get split up once we get to it. Also we need better answers on what to do about all this than what I have, even if it's a start. Hard problem is hard!
9. Moloch's Army. This isn't written and needs to go *somewhere* in the sequence or outside of it, or the whole operation is woefully incomplete. I need to finally write it. The devil's greatest trick was never proving he didn't exist, I wrote ten minutes ago, it was proving he'd already won, or would inevitably win. That only those who make deals with him get ahead, so they should implicitly or explicitly coordinate against those who don't. Moloch has an army, who coordinate implicitly around fighting against anyone fighting for anything except Moloch's Army, or anyone having any values. And *this is how Moloch wins, where it wins*. And also by making sure no one ever writes this, which makes this hard to write, etc etc. In that sense, *it really is true that the Devil's greatest trick is convincing people he doesn't exist*, because so far everyone I know who has figured this out has found it impossible to talk or write about this without sounding crazy to those who don't already get it. Much careful background may be necessary. Darwin sequence was originally supposed to be a gateway to this but it wasn't good enough on its own.

Moloch Hasn't Won

This post begins the Immoral Mazes sequence. [See introduction](#) for an overview of the plan. Before we get to the mazes, we need some background first.

Meditations on Moloch

Consider Scott Alexander's [Meditations on Moloch](#). I will summarize here.

Therein lie fourteen scenarios where participants can be caught in bad equilibria.

1. In an iterated prisoner's dilemma, two players keep playing defect.
2. In a dollar auction, participants massively overpay.
3. A group of fisherman fail to coordinate on using filters that efficiently benefit the group, because they can't punish those who don't profit by not using the filters.
4. Rats are caught in a permanent Malthusian trap where only those who do nothing but compete and consume survive. All others are outcompeted.
5. Capitalists serve a perfectly competitive market, and cannot pay a living wage.
6. The tying of all good schools to ownership of land causes families to work two jobs whose incomes are then captured by the owners of land.
7. Farmers outcompeted foragers despite this perhaps making everyone's life worse for the first few thousand years.
8. Si Vis Pacem, Para Bellum: If you want peace, prepare for war. So we do.
9. Cancer cells focus on replication, multiply and kill off the host.
10. Local governments compete to become more competitive and offer bigger bribes of money and easy regulation in order to lure businesses.
11. Our education system is a giant signaling competition for prestige.
12. Science doesn't follow proper statistical and other research procedures, resulting in findings that mostly aren't real.
13. Governments hand out massive corporate welfare.
14. Have you seen Congress?

Scott differentiates the first ten scenarios, where he says that perfect competition* wipes out all value, to the later four, where imperfect competition only wipes out most of the potential value.

He offers four potential ways out, which I believe to be an incomplete list:

1. Excess resources allow a temporary respite. We live in the [dream time](#).
2. Physical limitations where the horrible thing isn't actually efficient. He gives the example of slavery, where treating your slaves relatively well is the best way to get them to produce, and treating them horribly as in the antebellum South is so much worse that it needs to be enforced via government coordination or it will die out.
3. The things being maximized for in competitions are often nice things we care about, so at least we get the nice things.
4. We can coordinate. This may or may not involve government or coercion.

Scott differentiates this fourth, 'good' reason from the previous three 'bad' reasons, claiming coordination might be a long term solution, but we can't expect the 'bad' reasons to work if optimization power and technology get sufficiently advanced.

The forces of the stronger competitors, who sacrifice more of what they value to become powerful and to be fruitful and multiply, eventually win out. We might be in the dream time now, but with time we'll reach a steady state with static technology, where we've consumed all the surplus resources. All differentiation standing in the way of perfect competition will fade away. Horrible things will be the most efficient.

The optimizing things will keep getting better at optimizing, thus wiping out all value. When we [optimize for X but are indifferent to Y](#), we *by default* actively optimize *against* Y, for all Y that would make any claims to resources. Any Y we value is making a claim to resources. See [The Hidden Complexity of Wishes](#). We only don't optimize against Y if either we compensate by intentionally also optimizing for Y, or if X and Y have a relationship (causal, correlational or otherwise) where we happen to not want to optimize against Y, *and we figure this out rather than fall victim to Goodhart's Law*.

The greater the optimization power we put behind X, the more pressure we put upon Y. Eventually, under sufficient pressure, any given Y is likely doomed. Since [Value is Fragile](#), some necessary Y is eventually sacrificed, and all value gets destroyed.

Every simple optimization target yet suggested would, if fully implemented, destroy all value in the universe.

Submitting to this process means getting wiped out by these pressures.

Gotcha! You die anyway.

Even containing them locally won't work, because that locality will be part of the country, or the Earth, or the universe, and eventually wipe out our little corner.

Gotcha! You die anyway.

Which is why the only 'good' solution, in the end, is coordination, whether consensual or otherwise. We must coordinate to kill these ancient forces who rule the universe and lay waste to all of value, before they kill us first. Then replace them with something better.

Great project! We should keep working on that.

That's Not How This Works, That's Not How Any of This Works

It's easy to forget that *the world we live in does not work this way*. Thus, this whole line of thought can result in quite gloomy assessments of how the world inevitably always has and will work, such as this from Scott in Meditations on Moloch:

Suppose the coffee plantations discover a toxic pesticide that will increase their yield but make their customers sick. But their customers don't know about the pesticide, and the government hasn't caught up to regulating it yet. Now there's a tiny uncoupling between "selling to Americans" and "satisfying Americans' values", and so of course Americans' values get thrown under the bus.

Or this from Raymond, taken from a comment to a much later, distinct post, where 'werewolf' in context means 'someone trying to destroy rather than create clarity as the core of their strategy':

If you're a king with 5 districts, and you have 20 competent managers who trust each other... one thing you can do is assign 4 competent managers to each

*fortress, to ensure the fortress has redundancy and resilience and to handle all of its business without any backstabbing or relying on inflexible bureaucracies. But another thing you can do is send 10 (or 15!) of the managers to conquer and reign over **another** 5 (or 15!) districts.*

...

This is bad if you're one of the millions of people who live in the kingdom, who have to contend with werewolves.

It's an acceptable price to pay if you're actually the king. Because if you didn't pay the price, you'd be outcompeted by an empire who did. And meanwhile it doesn't actually really affect your plans that much.

The key instinct is that any price that can be paid to be stronger or more competitive, must be paid, therefore despair: If you didn't pay the price, you'd be out-competed by someone who did. People who despair this way often intuitively are modeling things as effectively perfect competition at least over time, which causes them to think that everything must by default become terrible, likely right away.

So many people increasingly bemoan how horrible anything and everything in the world is, and how we are all doomed.

When predictions of actual physical doom are made, as they increasingly are, often the response is to think things are so bad as to wish for the sweet release of death.

Moloch's Army: An As-Yet Unjustified But Important Note

Others quietly, or increasingly loudly and explicitly to those who are listening, embrace Moloch.

They tell us that the good is to sacrifice everything of value, and pass moral judgments on that basis. To take morality and flip its sign. Caring about things of value becomes sin, indifference becomes virtue. They support others who support the favoring of Moloch, elevating them to power, and punish anyone who supports anything else.

They form Moloch's Army and are the usual way Moloch locally wins, where Moloch locally wins. The real reason people give up [slack](#) and everything of value is not that it is ever so slightly more efficient to do so, because it almost always isn't. It is so that others can notice they have given up slack and everything of value.

I am not claiming the right to assert this yet. Doing so needs not only a citation but an entire post or sequence that is yet unwritten. It's hard to get right. Please don't object that I haven't justified it! *But I find it important to say this here, explicitly, out loud, before we continue.*

I also note that I explicitly support the implied norm of 'make necessary assertions that you can't explicitly justify if they seem important, and mark that you are doing this, then go back and justify them later when you know how to do so, or change your mind.' It also [led to this post](#), which led to many of what I think are my best other posts.

Meditations on Elua

The most vital and important part of Meditations on Moloch is hope. That *we are winning*. Yes, there are abominations and super-powerful forces out there looking to eat us and destroy everything of value, *and yet we still have lots of stuff that has value*.

Even before we escaped the Malthusian trap and entered the dream time, *we still had lots of stuff that had value*.

Quoting Scott Alexander:

Somewhere in this darkness is another god. He has also had many names. In the Kushiel books, his name was Elua. He is the god of flowers and free love and all soft and fragile things. Of art and science and philosophy and love. Of niceness, community, and civilization. He is a god of humans.

The other gods sit on their dark thrones and think "Ha ha, a god who doesn't even control any hell-monsters or command his worshippers to become killing machines. What a weakling! This is going to be so easy!"

But somehow Elua is still here. No one knows exactly how. And the gods who oppose Him tend to find Themselves meeting with a surprising number of unfortunate accidents.

Moloch gets the entire meditation. Elua, *who has been soundly kicking Moloch's ass for all of human existence*, gets the above quote and little else.

Going one by one:

Kingdoms *don't* reliably expand to their breaking points.

Poisons *don't* keep making their way into the coffee.

Iterated prisoner's dilemmas *often* succeed.

Dollar auctions are not all over the internet.

Most communities do get most people to pitch in.

People caught in most Malthusian traps still usually have non-work lives.

Capitalists don't pay the minimum wage all that frequently.

Many families spend perfectly reasonable amounts on housing.

Foragers never fully died out, also farming worked out in the end.

Most military budgets seem fixed at reasonable percentages of the economy, to the extent that for a long time that the United States has been mad its allies like Europe and Japan that they don't spend enough.

Most people die of something other than cancer, and almost all cells aren't cancerous.

Local governments enact rules and regulations that aren't business friendly all the time.

Occasionally, someone in the educational system learns something.

Science has severe problems, but scientists are cooperating to challenge poor statistical methods, resulting in the replication crisis and improving statistical standards.

Governments are corrupt and hand out corporate welfare, but mostly are only finitely corrupt and hand out relatively small amounts of corporate welfare. States that expropriate the bulk of available wealth are rare.

If someone has consistently good luck, it ain't luck.

(Yes, I have seen congress. Can't win them all. But I've also seen, feared and imagined much worse Congresses. For now, your life, liberty and property are *mostly* safe while they are in session.)

(And yes the education exception [is somewhat of a cop out](#) but also things could be so much worse there on almost every axis.)

The world is filled with people whose lives have value and include nice things. Each day we look Moloch in the face, know exactly what the local personal incentives are, see the ancient doom looming over all of us, and say [what we say to the God of Death](#): Not today.

Saying 'not today' won't cut it against an AGI or other super strong optimization process. Gotcha. You die anyway. But people speak and often act as if the ancient ones have already been released, and the end times are happening now.

They haven't, and they aren't.

So in the context of shorter term problems that don't involve such things, rather than bemoan how eventually Moloch will eat us all and how everything is terrible when actually many things are insanely great, perhaps we should ask a different question.

How is Elua pulling off all these *unfortunate accidents*?

*As a technical reminder we will expand upon in part two, *perfect competition* is a market with large numbers of buyers and sellers, homogeneity of the product, free entry and exit of firms, perfect market knowledge, one market price, perfect mobility of goods and factors of production with zero transportation costs, and no restrictions on trade. This forces the price to become equal to the marginal cost of production.

Perfect Competition

Previously in Immoral Mazes sequence: [Moloch Hasn't Won](#)

Perfect Competition

In [Meditations on Moloch](#), Scott points out that perfect competition destroys all value beyond the axis of competition.

Which, for any compactly defined axis of competition we know about, [destroys all value](#).

This is mathematically true.

Yet value remains.

Thus competition is imperfect.

(Unjustified-at-least-for-now note: Competition is good and necessary. Robust imperfect competition, including evolution, *creates* all value. Citation/proof not provided due to scope concerns but can stop to expand upon this more if it seems important to do so. Hopefully this will become clear naturally over course of sequence and/or proof seems obvious once considered.)

Perfect competition hasn't destroyed all value. Fully perfect competition is a useful toy model. But *it isn't an actual thing*.

Some systems and markets get close. For now they remain the notably rare exceptions. Often people despair because they are intuitively modeling things as effectively perfect competition, at least over time.

This causes many people to think that everything must by default become terrible, likely right away, such as in the examples that open *Meditations on Moloch*, and Raymond's kingdoms. And thus, that everything in general must become terrible.

The anticipation of perfect competition, rather than actual perfect competition, is often what causes things to become terrible. Anticipation of perfect competition creates perfect competition the same way that cooperation in a fixed-length iterated prisoner's dilemma has a backward induction problem.

Fighting this instinct is important. Responding instinctively with the detailed explanation, that competition is imperfect is not sufficiently quick or emotionally resonant. Two quick intuition pumps for when someone thus despairs, or argues for despair, are that this argument [proves too much](#), and that, again, when you look around, many things are insanely great.

Remind Me: What Exactly is Perfect Competition?

Perfect competition is [defined as](#) a market with large numbers of buyers and sellers, homogeneity of the product and marginal costs of production, free entry and exit of firms, perfect market knowledge, one market price, perfect mobility of goods and factors of production with zero transportation costs, and no restrictions on trade. This forces the price to become equal to the marginal cost of production.

If *even one* of the assumptions listed above breaks down, price will no longer necessarily be equal to marginal cost of production.

The linked-to definition calls it a 'theoretical market structure' for good reason – this is a useful way to model some things, but never actually holds fully true and is not actually a thing.

When I Googled for an example of perfect competition, this was the top result, the thing that came up in the box at the top:

*Agricultural markets are **examples** of nearly **perfect competition** as well. Imagine shopping at your local farmers' market: there are numerous farmers, selling the same fruits, vegetables and herbs. ... Another **example** is the currency market. First of all, the goods that are involved in the currency market are homogeneous.*

Here is how much perfect competition is a theoretical construct: *Farmers' markets were already, before noticing this on a later editing pass of the sequence, my first and most detailed example, in the next post, of highly imperfect competition.*

Note that perfect competition is in important ways *less competitive* than imperfect competition, because there is no competition for differentiation, for iteration or improvement, or for long term concerns.

Super-Perfect Competition

We will define *super-perfect competition* as competition that *mostly* has *most* of the elements of perfect competition, but lacks free entry and free exit, which creates more production than there would be at equilibrium (and thus, also, a lower price). In particular we assume effective homogenization of products. The market coerces this production through some combination of trickery, force, mandates, exploitation of biases, false perceptions, cost of exit, rapid technological or regulatory change, predatory competition in pursuit of future market power, or other means.

Perfect competition destroys all the producer surplus.

More precisely, it sets their *economic profits* to zero and takes away all their freedom of action – capital invested in such production is risk-free in pure perfect competition so it earns the risk-free rate. Which in the real world right now is basically zero. In a closed economy of only perfectly competitive markets, there are no alternative investments to establish a positive risk-free rate, so as long as there is not a capital shortage across markets, the risk-free rate will be *at most* zero.

Super-perfect competition destroys *more than all* producer surplus. The saving grace of free ability to exit has been removed.

Producers in both cases also have no slack or freedom. They must follow exactly the 'optimal' short-term strategies. Anything else is sacrificed.

As a system approaches perfect competition, producer surplus is destroyed, producer slack is depleted, and freedom of action dies.

Perfect competition is at core a special case of perfect optimization, in this case of homogenized commodity production.

Again, perfect competition is an abstraction. It's *not a thing*. But being *very close* to perfect competition in some places is *totally a thing*.

Super-perfect competition is *totally a thing*.

Perfect competition is where you can't win. Super-perfect competition is where you can't break even and can't even quit the game.

On a scale where 0 is full monopoly with price discrimination and 1 is perfect competition, super-perfect competition is values *greater than 1*.

A clean example of super-perfect competition would be the market for rideshares on the internet. Another would be the airline industry, at least at some times in the past. The dollar auction works via super-perfect competition. Many of Scott's other examples in *Meditations on Moloch* also involve similar situations, with players forced into participating.

In each case, super-perfect competition exists because market participants paid the cost of entry in anticipation of capturing a future market with levels of competition that would allow them to enjoy a producer surplus and make profits. When too many people make that same calculation, they become wrong.

Traditionally all of this is viewed as great for consumers, since they get the best possible price.

If all the consumers want is to cheaply consume a fungible commodity today, as optimized to simple metrics, they are in luck. Or at least they are in luck for now, since super-perfect competition often paves the way for future [oligopoly](#). If they want anything else, this market won't help them get there, because producers won't give them anything else.

Thus, rideshares and the airline industry.

Universal Perfect Competition

The worst case for consumers is if they are also the producers – a world fully given over to perfect optimization of some defined outcome.

This is an *incredibly* strong and alien assumption to be making, and it has lots of strange consequences that are rightfully ignored in standard econ-101 models that are for thinking on the margin.

No matter how cheap prices of goods and quantities available might get, if you have zero surplus of any kind to spend on them, it doesn't do you any good.

This is isomorphic to the standard Molochian future outlined in *Meditations on Moloch*. Standard Moloch-winning consequences follow.

[The Iron Law of Wages](#) takes over. Right away rather than eventually, since everything scales perfectly so there can't be a labor shortage. This drives labor's wage down to bare subsistence and capturing any efficiency gains from superior production.

Capital demands and gets the risk-free rate, which quickly settles at zero. Easiest way to see this is capital supply over time must be constant.

Any actions that are not optimal permanently deplete your stored capital. You cannot invest in optional assets, such as slack or creating the next generation, because those willing to not do so are competing (perfectly!) with you for a job.

You get a Disneyland with no children.

One could even argue this has already begun – that people in many nations having less than the replacement level number of children represents competitive forces so strong that young people no longer have the surplus to be comfortable having enough kids.

Worlds given fully over to perfect or super-perfect competition, by default, lose all value to those within them.


That last clause of 'to those within them' is important! They can have value to those outside that world, if they can get their hands on some of the stuff being produced. Creating almost-perfect or super-perfect competition within a compact space, ideally containing few if any humans or other moral agents, can be a great way to efficiently produce an important fungible good. Not every system of the world needs to have surplus and enjoy economic profits.

Over time, many systems that are not monopolies move in the direction of perfect competition, *if nothing changes and no one does anything to stop this effect.*

An increasingly static world, drawn towards its permanent equilibrium, would by default look increasingly like some combination of perfect competition, pure monopoly, and other [inadequate equilibria](#).

This is a lot of where the 'Moloch wins because essentially perfect competition' intuitions originate. Take a system. In that system, assume that underlying conditions are permanently static, except for the participants. Assume the participants won't meaningfully coordinate. Assume that participants who do better beat out those who do worse and expand and/or replicate and/or survive marginally better.

Extrapolate.



The fact that this hasn't happened yet gets viewed as a 'temporary respite' from the inevitable triumph of awfulness.

And in the very long run, given sufficiently advanced technology, *maybe* that's true. It depends. We don't yet know the rules of that game. Details matter.

But even old Malthusian traps don't work this way.

Things sometimes get bad. Once things get sufficiently bad that no one can deviate from short-term selfish actions or be a different type of person without being wiped out, things are no longer stable. People cheat on long term investments, including various combinations of things such as having and raising children, maintaining infrastructure and defending norms. The seed corn gets eaten. Eventually, usually when some random new threat inevitably emerges, the order collapses, and things start again. The rise and fall of civilizations.

The rise and fall of corporations or other organizations is often not so different.

Shocks are Inevitable

In the steady state, there needs to be enough surplus around to raise the next generation and deal with negative random shocks. And as per Scott's second reason things might not get too bad, the horrible nightmare systems mostly aren't actually efficient ways to get useful actions out of human beings. Neither is starving the people of resources too far. Things can only get so bad for so long.

If you make things sufficiently bad, before too long you get wiped out.

[For now at least, poor folks still smile.](#)

That's not to say that this is a *good* solution. Our modern solution centers on robust competition and an industrial revolution. We let corporations and others drive each other into bankruptcy or dissolution. This can be because they get too far gone, or because outside improvements through technological and productivity growth create a superior alternative.

In the increasingly static and more technologically advanced future that is more of a true steady state and approaching its permanent equilibrium, both technologically and

otherwise, shocks may cease to be inevitable and thus this danger would become much stronger. That's a huge problem that needs more attention. It's also beyond scope. *It is important for us to remember that it is beyond scope and a different problem.*

The practical danger to this system, that is on the rise now, is the temptation to not let this cycling happen. Preventing the cycle stops short term pain and protects the powerful. In the places things are getting worse or on pace to start getting worse, where Moloch is locally winning, this is a key mechanism.

This leads to another perspective on how people get to the 'Moloch wins' intuition.

Those that visibly follow Moloch here and now seem to prosper, here and for now, if no one is punishing them for it (and also, as per my not-yet-justified claim last time, Moloch's Army is in effect rewarding them for these intuitions and punishing others for lacking them).

The things that exist and are big or visible, and thus seem important, are mostly steadily getting worse most of the time. Despite this, things overall kept getting better on almost all fronts for a long time, whether or not that is continuing recently.

But it's tough to see that blossoming forest for the rotting individual trees.

That question again. How does Elua pull off all these *unfortunate accidents*?

One key reason is that Elua is [antifragile](#). Accidents and disorder [are good for business](#).

Imperfect Competition

Previously in Sequence: [Moloch Hasn't Won](#), [Perfect Competition](#)

This post looks at a few examples of imperfect competition to illustrate various ways in which perfect competition is kept at bay and value is preserved. Concrete examples seem more likely to enlighten than abstract principles.

Let's start with the most literal of markets, the market for food, and then talk about the market for cars.

As I noted last time, I noticed after writing this that a Local Farmer's Market was *Google's top response to asking for an example of perfect competition*. Which makes it kind of perfect as the central example of imperfect competition.

For length, I'll stick to these examples. I am happy to give quick models for other markets in the comments if anyone wants to work through different examples.

1. The Local Farmer's Market

A few times a week, at various places around the city, local producers open stands to sell their goods to the public. Many people choose to shop here, rather than at the supermarkets a few blocks away in every direction. They typically pay higher prices, do a lot of investigation and sampling of goods, and eventually often become loyal to producers and products that provide strong value. When I was growing up, we often got our cheese and apple cider from one such market (along with more healthy foods that I was incapable of eating, and therefore don't remember as well).

What's going on here? A lot of different things are combining at once – and this list likely isn't even complete. Each of these is a violation of perfect competition.

Products and producers are not homogenized, and information about them is costly. Even the relatively homogenized farmers' market products, like the apple cider, differ in quality from batch to batch, from season to season and harvest to harvest. There is no reliable way to differentiate high quality from low quality products or producers, or any way to divine the product details you might prefer or dislike, without sampling the goods, which is costly outside of the occasional free sample. When buying produce, one needs to pick out the items one likes and things get even more variable. **Reputation and experience matter a lot.**

Market information is costly for consumers. Different producers bring different goods to market at different times, with no known schedule. Often the goods sell out. The only sure way to know is to go to the market and look around. If one wants to be sure, one would even walk around the full market and examine things before making any decisions, but that takes a while. Once you're there, it often makes sense to buy something even if you wouldn't have bothered going out there if you'd known what was available today.

Market information is costly for producers. Farmers and bakers and cheesemongers tend not to be those with a deep interest in market dynamics. They don't know what consumers will want or what prices they'll be willing to pay, and find

out by trying to sell goods and seeing what people buy. They are quite bad at maximizing profits from what they produce. We like it that way. Someone who was good at maximizing profits [would be bad at producing quality goods](#). Each producer has a unique situation, so they can't rely on general surveys to price goods correctly. Goodwill is vital to the business, making pricing experiments generally a bad idea. That leads into a key additional concept, which is **complete skill stacks are rare**. This allows being very good at one aspect to compensate for being relatively poor in another, whether or not this is by choice, or provides [slack](#).

Unknown unknown information exists. We know little about nutrition and what is dangerous for us. Some of what we are buying, when we buy from small local producers, is their inability and unwillingness to use various industrial methods we don't know about or understand, that may have unknown and hard to measure long term effects on our health. We are uncertain about the order of magnitude of this effect, but such effects are highly unlikely to be good. This pushes us away from homogenized products.

Fixed costs exist and production costs are asymmetric. Skills differ. Each farmer and producer both specializes in a different set of products, and faces the quirky resource constraints of their own land, skills, tools and background, and what they know how to and like to grow and make. **Economies of scale exist** due to many fixed costs. Even if you did know how to grow crops, doing so on too little land would be prohibitively time inefficient. Even if it wasn't, you would still have no practical way to bring those goods to market, as that too requires fixed costs. You can't participate in these markets unless you're big enough to be a full member and committed to sustaining relationships and reputation.

Producer preferences differ. This is an important variation on variable production costs. Farming in the modern age is often a labor of love. If a farmer likes extra sharp cheddar or plum tomatoes, they're going to produce them even if it isn't the best thing for the bottom line, then bring them to market anyway. They'll also produce more variety than the market would reward on its own. Seeking this out is one way to find the best quality.

Location matters. The farmer's market exists largely because it's better to buy local products with lower money and time spent in transportation, resulting in cheaper and higher quality goods. The goods in supermarkets often were engineered to travel better and keep longer, which reduces quality in other ways. When you go to the market for local goods, those closer to the market have a huge advantage. Those farther away will choose to go to a different market, even if its customer base isn't quite as good.

Consumers care about the individual producers. Local farmers also often get extra business from those who like to support local production for various pro-social reasons, or even develop personal relationships over time. Most people at the farmer's market would pay a non-zero amount extra to buy from the farmer's market over an identical product at the supermarket.

Skin in the game. Skin in the game doesn't mix with perfect competition, because in perfect competition there's no skin available to have. When owners operate and stick around for the long term, and face the consequences of their decisions in terms of skill development, reputation, capacity, and so on, they have to care about more than just creating a homogenized product, or even creating a product. Everything matters, and those who let longer term things slide do not survive for long.

No boss. No one is telling the producers they'll be fired if they make trade-offs. No shareholders are putting on extreme pressure to maximize profits. People get to do the things they want to do, and not the things they don't want to do, and accept the consequences.

What happens in the long term?

We used to mostly have these types of markets, and now we mostly have supermarkets.

2. Supermarkets

Supermarkets are a rough business to be in. They are much closer to perfect competition than farmer's markets. The products are more commoditized and homogenized, and are more predictable in quality. Information is easier to get. Often goods are identical across sellers. The relationships are mostly purely transactional. The store owners are mostly in it for the money.

Identical packages of Kerrygold butter are \$3.19 at Trader Joe's on 3rd Avenue, \$3.79 at Whole Foods a block farther west, and something like \$4.49 at Westside Market two blocks to the south. Sometimes I see them for \$5 or more.

That's weird. Sounds *remarkably* uncompetitive. What happened to [the law of one price](#)? What's going on?

Here's my model. People care a lot about distance when picking a supermarket, and care a lot about batching their purchases so they don't have to go too often. Customers are choosing between a handful of potentially reasonable bundles of options. Each has some unique offerings in terms of product variety and quality, a unique location, a unique brand identity and reputation, and a history with each potential customer. It takes time to learn what each one has to offer, to remember where everything is located so things can be found quickly, to know which things are being overpriced at each store. Stores often price some items lower than normal and/or below cost to tempt customers in (hence the term 'loss leader') and then jack up the prices of other items (like the Kerrygold butter) where customers are likely to be price insensitive.

Thus, when I enter a new supermarket, I price check the Kerrygold butter before buying items whose fair price I do not know.

Supermarkets are effectively in oligopolistic competition with local alternatives and a handful of services like FreshDirect. They're not in anything like perfect competition with each other.

The same is true for different brands within the supermarket. We have what almost any other culture in history would consider amazing product variety and choice. It is *very rare* for there to mostly be multiple brands that sell essentially identical products at the same price point, in the same store. It's rare for there to be even two best options that are essentially identical. Often, the choice will be *close* between multiple brands, but that's because they're offering different trade offs. Attempts to copy gold standard products like Heinz ketchup somehow don't work. Most consumers don't know the quality or details of most options of things they're considering buying. And there's a *ton* of brand loyalty. There is intense competition. The system works well. But it's definitely something very different than perfect competition.

Online grocery services like FreshDirect also face sufficient scale and distance requirements that there are only a handful available in a given area, and they tend to not have that much product variety either, and to market themselves as unique options complete with many unique items (e.g. for FreshDirect I recommend [the take and bake Ciabatta rolls](#)).

One can imagine a world in which transportation costs continued to fall, and skill in preservation continued to improve, as did the ability to communicate what you want online and get what you asked for (this is *much* less reliable than my priors would have predicted), to the point where location stopped being important and it made sense to break up orders into individual products. One can imagine better market knowledge becoming available and known to more customers. Or the food is made in replicators, or by robots using standardized open-source procedures. That day may come. For now, a lot of things stand between here and there.

3. Restaurants (in brief)

Restaurants are a strange case.

I am told that no one makes money and competition is super rough, largely because people are often opening and running places for non-monetary reasons. A lot of people think they can do it when they can't because they don't understand what it takes (and because their friends and family often lie to them about their cooking). Which can all cause super-perfect competition.

Despite that, we clearly see highly imperfect markets. We regularly see places that can produce sufficiently superior products to command strong pricing power, which they rarely fully use. There is little doubt that many places have *huge* advantages over the marginal competition. Places have sustained reputations and relationships with regulars, earned by [Playing in Hard Mode](#). Seeing people play in hard mode is a sure sign competition is far from perfect. And those who do these jobs well earn the ability to charge more and often get more business than they can handle.

[My model](#) says the central stories here are that location matters a lot, skill and experience matter a lot, products are highly differentiated, information is expensive on many fronts, and people have different values with skin in the game. Too many participants enter, and exiting is expensive, so making zero profits is really hard. But there are huge marginal returns to being great, so huge success remains possible.

4. Buying a Car

I don't drive, so I've never bought a car. But buying a car is traditionally so frustrating to so many people that I know a reasonable amount about the process.

In the first phase, different companies attempt to convince you to buy their car instead of another car.

This is done through some combination of bragging about small features or arbitrary industry awards, selective comparisons (which my brain always reads as, 'that other car you're mentioning is probably worth checking out'), creation of emotional resonance and identity affect assignment, loading up more observable features, periodic 'good deals' and 'sales' with superficially cheap structures, offering test drives, and providing in the background real information resources.

Customers then combine this with personal anecdotes and ideally some information from plausibly objective third parties (e.g. my parents used Consumer Reports magazine and the NPR show Car Talk) to evaluate salient choices along numerous axes, narrow things down to choices that feel appropriate and reasonable, then consider what 'deals' are available and make a final choice however big final choices are made. One must also then choose expensive new car versus cheaper new car versus premium used car versus cheap used car.

Economies of scale, both in production and distribution and then in marketing and reputation, keep the market limited to a handful of producers, each of whom produces a handful of models each year. Even major participants can only efficiently highlight a handful of offerings, or the messages get muddled. Different customers care about wildly different things. One will care about safety and reliability, another how the car looks or whether it is seen as sexy, another primarily cost to own, a fourth comes from a family that drives Fords, a fifth wants everyone to know how good and responsible and forward thinking they are with their electric model, and the last is a sixteen year old who just got their license and wants really good air conditioning and sweet cup holders.

You want to *differentiate* yourself from the competition, not match them, even if you're better on objective quality and value, since those are hard to measure and communicate.

The products do their best to *not* be homogenous.

It's not that *one element* of perfect competition is missing. It's that *all the elements* are *each on their own missing*.

Entry and exit are costly.

Economies of scale are central.

Market knowledge is expensive and difficult.

Mobility of goods is not only not free, it is what is being sold.

There are massive restrictions on trade.

Past this point, I have no personal experience, so this could be unfair, obsolete, inaccurate or all three. And of course, the whole point is that these markets are not perfectly competitive, so experiences wildly differ from person to person and location to location. Great experiences do exist, and only some of them involve not realizing you're being had.

Once a car has been selected, it's time to head to the dealer. To negotiate and make a deal.

This process is not *less* absurd than it sounds. It's also mandated by the government in the United States, [which bans systematic automobile sales that don't go through physical dealerships](#).

Customers come in and look around the shiny cars. They are then likely to be accosted by salesman who share selective information, hide other information, frame shift, anchor on absurd retail prices, manipulate, outright lie, 'check with their manager', cajole, play good cop bad cop and other similar games, create expectations

of reciprocity, bait and switch and so forth. They do this so brazenly and reliably that '[everybody knows](#)' that car salesman are [out to get you](#), and people study up on how to navigate this in reasonable time and not get taken for a ride.

Most customers rightfully dread this. Many find it so unpleasant that they knowingly accept losing hundreds or even thousands of dollars in value to make the process less painful. They are often unable to bring themselves to use tactics like "call multiple dealers and ask for quotes" or "take a day to think about it or see if the deal improves," and even those who know rarely hold out for the most powerful, "wait for the end of the month," whose mechanism will be explained in a bit.

And that's for *new* cars. Used car salesman carries even worse connotations.

Importantly, it's not *maximally bad* connotations. It's nothing close to *perfectly competitive* fraud. It's *wildly varying magnitudes* of fraud.

The common question, 'would you buy a used car from this man?' highlights that there are people you under no circumstances should buy a used car from, and others, even professional used car salesman, that you should be cautious with but who it would be reasonable to purchase a used vehicle from.

New car salesman pursue a wide range of strategies at a wide range of skill levels. They even radically adjust tactics based on the time of the month. They have targets to meet and will do whatever it takes to meet those targets if they're about to miss, but find doing so unpleasant on many levels, as well as expensive. It not only burns slack and isn't fun, it degrades the long term and bottom line to push too hard. So most salesman avoid doing this when it isn't necessary, and pick a trade-off between short term profits, long term profits, intensity of work, and their own sense of morality.

When we do see the worst behaviors, my model says it is usually because of pressure from bosses and quotas from higher up the corporate ladder, rather than any need to be maximally competitive.

It is done haphazardly and inefficiently to save people's hides. That you can typically literally save hundreds to thousands of dollars *by buying your car on the last day of the dealer's sales incentive period* speaks to how distorted all of this has become.

We are seeing the results of immoral mazes.

Does Big Business Hate Your Family?

Previously in Sequence: [Moloch Hasn't Won](#), [Perfect Competition](#), [Imperfect Competition](#)

The book [Moral Mazes](#) describes the world of managers in several large American corporations.

They are placed in severely toxic situations. To succeed, they transform their entire selves into maximizers of what get you ahead in middle corporate management.

This includes active hostility to all morality and to all other considerations not being actively maximized. Those considerations compete for attention and resources.

Those who refuse or fail to self-modify in this way fail to climb the corporate ladder. Those who 'succeed' at this task sometimes rise to higher corporate positions, but there are not enough positions to go around, so many still fail.

Reminder: Sufficiently Powerful Optimization Of Any Known Target Destroys All Value

This is a default characteristic of all sufficiently strong optimizers. Recall that when we [optimize for X but are indifferent to Y](#), we *by default* actively optimize *against* Y, for all Y that would make any claims to resources. See [The Hidden Complexity of Wishes](#).

Perfect competition destroys all value by being a sufficiently powerful optimizer.

The competition for success as a middle manager described in *Moral Mazes* destroys the managers (among many other things it destroys) the same way, by putting them and the company's own culture and norms under sufficiently powerful optimization pressure towards a narrow target.

Does Big Business Hate Your Family?

Consider this comment, [elevated to a post at Marginal Revolution](#), responding to [Tyler Cowen's noting](#) that the National Conservatism Conference had a talk called "Big Business Hates Your Family":

What would it mean if big business *did* hate your family?

Would it mean adopting a working culture that made it ever harder to rise to power within it while also having said family? Would it require those with career ambitions to geographically abandon extended family and to live in areas notoriously difficult for raising families? Would it mean requiring long delays on family formation while you got credentialed, worked with little remuneration while getting your foot in the door, and then place huge amounts of time and effort on career growth rather than investing in your family? Does corporate culture act like it hates your family?

Would it mean selling products which have strong correlations with family strife and dissolution? Would it market products known to be destructive to thousands of families relentlessly? Would it market products that consume time in great quantities at the expense of family time investment? Would it routinely mock and denigrate your family roles for cheap publicity?

Would it mean lobbying for policies which are good for the business, but bad for your family? Would it support seeking a larger supply of labor via immigration? Would it support visa restricted immigration of labor that is less able to defy corporate diktat without having legal or financial issues? Would it argue for child care subsidies for the people it wishes to employ rather than for all Americans and all child care arrangements?

I believe businesses are amoral and are just maximizing money, power, and prestige for those in positions of power within them. Yet, this formal indifference seems to be giving rise to a lot of behaviors that are, at best, perceived to be hostile to families.

...

There is nothing wrong with this, and certainly nothing illegal about it, but I would be shocked if large organizations that are disproportionately filled with the single and childless who are located in regions that are disproportionately single and childless and who are busy virtue signalling to academia, politics, and other left bastions that are disproportionately single and childless managed to somehow not end up at cross purposes for the majority of families. And frankly I would be shocked if this antagonism did not spill over into emotional terms.

Certainly, I am always told that this sort of analysis is why [Structure X] is antagonistic, if only implicitly, against racial minorities. I see no reason why parents or spouses would feel any differently.

Big business, like the AGI, does not hate your family. Big business thinks your family possesses capital, preferences and other assets that could be used for something else. The effects of this on your family are a side effect. Big business also notes that your family is attempting to optimize the world for something other than the profits of big business, and would like to prevent you from doing this, since it would tend to reduce its profits. This is all doubly true if you work for the Big Business in question, since your family is now asserting its preferences and claims to resources in an even more directly competitive way. For morality (or anything else that might have a claim to resources, so basically anything anyone cares about), same thing. Replace "your family" with "morality."

No, wait. That's wrong.

The above paragraph, like the quoted text, is misleadingly treating Big Business in general, or a given business in particular, like it is an agent.

We need to fully appreciate that corporations are not agents. There is no agent called Big Business. Nor are any of the individual big businesses agents as such.

Corporations are people. Not only in the '[corporations are people, my friend](#)' sense, but in the sense that corporations don't act or have preferences, but rather are composed of people that act and have preferences. [There's just a bunch of guys](#).

The CEO is an individual representing their own interests, like everyone else at the company. The profits they care about are their own. Occasionally they will make some efforts to maximize corporate profits. Often they won't, or will focus primarily on other things.

English makes it *very hard* not to make the agent mistake. I will no doubt keep saying that corporations want things and prefer things and think things, because I don't know of another reasonable and compact way of saying the same thing. But understand that I mean that as an abstraction of what happens as the result of the preferences of individuals, and choices made by individual people, and their interactions.

Let's reformulate the question.

Do The Managers of Big Businesses Hate Your Family?

When I say these organizations are immoral, it's not *necessarily* that the people running major corporations are mustache-twirling villains who hate morality.

Most don't have mustaches.

Morality causes choices and optimization towards the moral and against the immoral. That interferes with choices and optimization away from the uncomfortable and towards the comfortable. It interferes with choices away from bad for your boss towards good for your boss. Or, in less broken situations and/or with a less cynical perspective, also, from less profitable towards more profitable.

Do enough of that, and some of those involved for long enough *do* become mustache-twirling villains, because they are humans who are trying to live with what they are doing. You become what you continuously do and say. Eventually one [becomes the mask](#). [Mustaches become more common as people get older](#).

The first-level model says this stays rare. Being a mustache-twirling villain is to have preferences (if nothing else, you prefer having a mustache), and thus is bad for business the same way having morals is bad for business. You want to be completely indifferent, and be seen as completely indifferent.

It is worth taking five minutes here to think about how it might become seen as advantageous by the managers themselves, under these conditions, to become *actively immoral*, and prefer doing the wrong thing over the right thing *because it is wrong*.

Lacking all Preferences

None of this has anything to do with *morality* as such. What they are against is what Moloch is against. *Having preferences at all*. They are against *caring about anything at all* other than climbing the corporate (or academic, or government, or other as appropriate) ladder.

It doesn't matter whether you care about the laws of accounting, wearing the color red, eating meat, cheating on your wife, seeing hit movies on opening weekend, overthrow of third world governments, falsifying scientific data sets, your favorite prime number, different brands of olive oil, genocide, or having a life outside the corporation, or time to watch your kids grow up. Caring about *anything not chosen to help your career* is a liability. Being seen or thought of as caring about something else can be even worse than actually doing so, as you are preemptively punished slash seen as not having a future, and potential allies want nothing to do with you.

In particular, any sign of, or even worse defense of, an outside life is deadly.

Thus what is called “[Pournelle’s iron law of bureaucracy](#)”: In any bureaucracy, the people devoted to the benefit of the bureaucracy itself always get in control and those dedicated to the goals the bureaucracy is supposed to accomplish have less and less influence, and sometimes are eliminated entirely.

In order to get an instinctive sense of all this, if you have not yet done so, I encourage you to read or at least browse [Quotes from Moral Mazes](#). Some of those quotes will be reiterated later in the sequence. You also may wish to see [Moral Mazes and Short Termism](#).

If you are interested enough to power through it (not the author’s fault, but it’s a tough read) and have the time, even better would be to pause here and [read the whole book](#).

What is Life in an Immoral Maze?

Previously in sequence: [Moloch Hasn't Won](#), [Perfect Competition](#), [Imperfect Competition](#), [Does Big Business Hate Your Family?](#)

This post attempts to give a gears-level explanation of maze life as experienced by a middle manager in systems with many levels of management, as depicted in *Moral Mazes*.

The 'maze level' of corporations differs wildly. These dynamics do not reliably fully take over until you have many levels of management. Questions of what causes high maze levels will be dealt with in future sections.

Again, if you have not yet done so, you are highly encouraged to read or review [Quotes from Moral Mazes](#). I will not have the space here to even gloss over many important aspects.

An Immoral Maze can be modeled as a super-perfectly competitive job market for management material. All the principles of super-perfect competition are in play. The normal barriers to such competition have been stripped away. Too many 'qualified' managers compete for too few positions.

If an aspirant who does not devote everything they have, and visibly sacrifice all slack, towards success, they automatically fail. Those who do make such sacrifices mostly fail anyway, but some of them "succeed". We'll see later what success has in store for them.

The Lifestyle of a Middle Manager

At the managerial and professional levels, the road between work and life is usually open because it is difficult to refuse to use one's influence, patronage, or power on behalf of another regular member of one's social coterie. It therefore becomes important to choose one's social colleagues with some care and, of course, know how to drop them should they fall out of organizational favor. (Moral Mazes, Location 884, Quote #117)

We have this idea that there is work and there is not-work, and once one leaves work one is engaged in not-work distinct from work. We also have this idea that there are things that are off limits even at work, like sexual harassment.

For a person without anyone reporting to them, who is 'on the line' in the book's parlance, this can be sustained.

For those in middle management who want to succeed, that's not how things work. Everything you are is on the table. You'd better be all-in.

You will increasingly choose your friends to help you win. You will increasingly choose your hobbies, and what you eat, and your politics, and your house, and your church, and your spouse and how many kids you have, to help you win. And of course, you will choose your (lack of) morality.

In the end, you will sacrifice *everything*, and I mean everything, that you value, in any sense, to win.

If the job requires you to move, anywhere in the world, you'll do it, dragging your nuclear family along and forcing all of you to leave behind everything and everyone you know. Otherwise, you're just not serious about success.

[Slack](#) will *definitely* not be a thing.

Your time is especially vulnerable.

Higher-level managers in all the corporations I studied commonly spend twelve to fourteen hours a day at the office. (Location 1156, Quote #120, Moral Mazes)

This is the result of total competition between producers – the managers are effectively rival producers trying to sell themselves as the product.

The market for managers is seen, by those who make the decisions, as highly efficient.

If managers were seen as wildly different in terms of talent, intelligence, or some other ability that helped get things done, that would help a lot. You could afford to be a little quirky, to hold on to the things you value most, without losing the game entirely. Your success will be *influenced* by your personality and dedication, but nothing like *solely determined* by them.

Alas, the perception in these mazes is exactly the opposite.

See, once you are at a certain level of experience, the difference between a vice-president, an executive vice-president, and a general manager is negligible. It has relatively little to do with ability as such. People are all good at that level. They wouldn't be there without that ability. So it has little to do with ability or with business experience and so on. All have similar levels of ability, drive, competence, and so on. What happens is that people perceive in others what they like—operating styles, lifestyles, personalities, ability to get along. Now these are all very subjective judgments. And what happens is that if a person in authority sees someone else's guy as less competent than his own guy, well, he'll always perceive him that way. And he'll always pick—as a result—his own guy when the chance to do so comes up. (Location 1013, Quote #87, Moral Mazes)

It is known that most people 'don't have what it takes' to be a manager. This is clearly true on many levels. Only one of them is a willingness to fully get with the program.

Once you get several levels up, the default assumption is that everyone is smart enough, and competent enough. That the object-level is a fully level playing field. The idea that someone can just *be better at doing the actual job* doesn't parse for them.

All remaining differences are about negative selection, about how hard you want it and are willing to sacrifice everything, or about how well you play political games. Nor do they much care whether you succeed at your job, anyway.

Some additional supporting quotes on that follow. A large portion of the quotes reinforce this perspective.

If you can't work smart, work hard:

When asked who gets ahead, an executive vice-president at Weft Corporation says: The guys who want it [get ahead]. The guys who work. You can spot it in the first six months. They work hard, they come to work earlier, they leave later. They have suggestions at meetings. They come into a business and the business picks right up. They don't go on coffee breaks down here [in the basement]. You see the parade of people going back and forth down here? There's no reason for that. I never did that. If you need coffee, you can have it at your desk. Some people put in time and some people work. (Location 992, Quote 29, Moral Mazes)

But everyone at this level works hard, which was more about showing you work hard than the results of the work, because concrete outcomes don't much matter:

As one manager says: "Personality spells success or failure, not what you do on the field." (Location 1383, Quote 33, Moral Mazes)

It's not like there were ever objective criteria:

Managers rarely speak of objective criteria for achieving success because once certain crucial points in one's career are passed, success and failure seem to have little to do with one's accomplishments. (Location 917, Quote 42, Moral Mazes)

Which makes sense, because if everyone is the same, then concrete outcomes are just luck:

Assuming a basic level of corporate resources and managerial know-how, real economic outcome is seen to depend on factors largely beyond organizational or personal control. (Location 1592, Quote 46, Moral Mazes)

I am supremely confident that this perspective is completely bonkers. There is huge differential between better and worse no matter how high up you go or how extreme your filters have already been. But *what matters here is what the managers believe*. Not what is true. Talent or brilliance won't save you if no one believes it can exist. If noticed it will only backfire:

Striking, distinctive characteristics of any sort, in fact, are dangerous in the corporate world. One of the most damaging things, for instance, that can be said about a manager is that he is brilliant. This almost invariably signals a judgment that the person has publicly asserted his intelligence and is perceived as a threat to others. What good is a wizard who makes his colleagues and his customers uncomfortable? (Location 1173, Quote 88, Moral Mazes)

How do things get so bad?

That's the question we'll look at an aspect of next post. From here I anticipate 3-5 day gaps between posts.

Questions that will be considered later, worth thinking about now, include: How does this persist? If things are so bad, why aren't things way worse? Why haven't these corporations fallen apart or been competed out of business? Given they haven't, why hasn't the entire economy collapsed? Why do regular people, aspirant managers and otherwise, still think of these manager positions as the 'good jobs' as opposed to picking up pitchforks and torches?

Stripping Away the Protections

Previously in sequence: [Moloch Hasn't Won](#), [Perfect Competition](#), [Imperfect Competition](#), [Does Big Business Hate Your Family?](#), [What is Life in an Immoral Maze?](#)

The previous post painted a bleak picture of life as a middle manager in an Immoral Maze. Not every middle manager faces the high maze levels described in *Moral Mazes*, but I am confident that many do.

How did things get so bad?

We previously discussed [Meditations on Moloch](#). That has a lot of gear-level analysis of how easy it is for undifferentiated high stakes competition along a single variable to destroy everything of value, and how much worse it is to do all those things at once than would be the sum of their parts.

[Perfect Competition](#) summarized and fleshed out that model so we could work with it, and defined Super-Perfect Competition as perfect competition without free exit, which when that restriction is meaningful results in *less than zero* economic profits.

[Imperfect Competition](#) illustrated how different idealized perfect competition is from the conditions we typically experience. Even when we use markets that seem highly competitive. Even in a market used as a canonical example of perfect competition.

The difference between the situations discussed there, and the situation in *Moral Mazes*, is that the conditions of super-perfect competition *really do* apply. The protections against this happening have been stripped away.

An immoral maze is not as extreme as the imagined possible future em world where [any simulated human caught doing something suboptimal is erased and replaced](#). But it is closer than one would first think. In a maze, people can't yet be copied, but there are lots of aspirants to choose from, and any opportunity to rule one out is pounced upon.

Sources of positive differentiation between managers are systematically destroyed. As we saw above, managers do not believe they exist. So they don't matter. This leaves only negative differentiation and selection, plus politics.

Sources of meaningful object-level concerns and detail, and the relevance of any long term consequences of actions, are stripped away as well. Again, managers do not believe they exist. So they don't matter.

This leaves a homogenous product that managers must become, the production of which is subject to perfect competition.

Then the resulting politics runs amok. The system self-reinforces the reinforcement of conformity with these ideals.

Let's go over the details that keep the rest of the world safe from such conditions. The bold terms are copied from the previous post, *Imperfect Competition*.

No boss. Not only is there a boss. There are bosses and underlings as far as the eye can see. Everything is reduced to simple metrics and impressions that can be run up

or down the chain. Everyone is judged on everything. Based on how they look and cause others to look, and what it says about the person being judged. There are only better and worse judgments, and ways to get them. Avoiding standing out in a negative way (e.g. 'a thousandatta boys are wiped out by one oops') is all that matters. The richness of real situations falls away.

Skin in the game. Immoral mazes actively fight against anyone having skin in the object-level or long term game. They don't track or remember anything, deliberately destroying or avoiding records to avoid future scapegoating. That leaves no skin in the real game, only skin in the game of short term appearances and politics, which is what skin in the game is supposed to defend against. Even if the system *wanted* to preserve skin in the game and distribute it usefully, rather than destroy it, there would be a severe shortage of it. Major corporations are too large and consequences too diffuse. There is not enough potential skin in the game to distribute to dozens of levels of management, even in theory under the best of circumstances, even if the company is fully privately held. Skin in the object level game is restricted to the very top and bottom of the hierarchy, if they use it at all. The CEO and board can own stock or options, and the people directly 'on the line' can perhaps be judged on individual performance.

Products and producers are not homogenized, and information about them is costly. This is the big one that is easiest to miss or be confused about. It's hard *not* to notice that there's a lot of bosses and a lack of skin in the game. Noticing that managers are effectively homogenized is hard, because they aren't actually homogenized at all – rather it is the *perception* that they are, which causes them to act and be treated as if they were. Throughout the process of reaching middle management, aspiring managers become homogenized products, and learn to view themselves and their competition in this fashion, and that within managerial circles this is common knowledge. Any deviations are career suicide and severely punished. The idea that one could be *better* doesn't parse – see skills differ, below. There's also the double-think that if you *were* somehow meaningfully better, that (through what seems to them like some instinctive but unknown mechanism) likely makes you a threat and means you need to be taken out, or that if you're better that means they are worse and being worse is death, so again, take you out.

Reputation and experience matter a lot. Reputation inside a maze boils down to whether someone has negative marks that doom them, and how much momentum their career has, and what kind of political allies they've made. Your commitment and time spent are also tracked, but that rapidly ceases to differentiate between any of the survivors. None of that is the kind of reputation that does the work of keeping competition imperfect. Experience beyond generic management experience is not only not considered relevant, it is a liability. It means that you lack career momentum.

Market information is costly for consumers. Consumers are the bosses and fellow managers, who already understand the market in question all too well. They have no choice but to learn, even if the cost of doing so is high. The protection from expensive information comes from *consumers choosing not to purchase all of it*.

Market information is costly for producers. Producers are *also* the bosses and fellow managers, since no one is close to the object level. So again, they already understand the market in question all too well. No one who matters lacks market information.

Skills differ. Complete skill stacks are rare. It is presumed that once a certain level is reached, everyone has a complete managerial skills stack. Everyone's skill levels are assumed to be equal, aside from skill in politics. The conventional wisdom claims there is common knowledge that there is a generic 'ability to run a thing' skill stack. Everyone who has made it this far has that stack. Skills that aren't generic or politics are considered to effectively not exist, or to be things you should be delegating regardless. Having them yourself is, as noted above, a liability rather than an asset. So there's no reason to pick particular managers with [a specific set of skills](#) for particular jobs that require those skills. There's no way to positively differentiate yourself using your skills.

Unknown unknown information exists and matters. Managers are presumed to all be the same. There aren't meaningful unknown unknowns. The products for which they coordinate production have unknown unknowns that will eventually matter *to the company*. They also have aspects that are hard to measure, or hard for someone outside the department (or anyone other than the manager themselves) to notice at all. But that is all a long term problem, the same way we were concerned with long term nutritional effects or the long term reliability of car models. No one cares if a manager was producing something with such long term issues. By the time they come to light, all association between that manager and the product has been long forgotten. Memories don't last long enough. If one *did* notice, a middle manager would primarily consider this *praiseworthy*, because the person successfully got away with something, unless they had an opportunity to somehow use it against them. Management nirvana is described in Moral Mazes as getting promoted and then blaming your successor for your own "mistakes." Which, of course, are not mistakes from your perspective. The manager *is* responsible for making sure things don't go to hell in an observable way *while they are still in charge*, which would be really bad and likely end one's career.

Fixed costs exist. Fixed costs matter for proper decision making only when they can be avoided (including the option of avoiding them by avoiding the whole enterprise). Otherwise, the ship has sailed, so they are sunk costs, which don't count and don't matter, except insofar as [sunk cost fallacy](#) is thing in context. If they try to pass those costs along, they just dig the hole deeper. Managers have paid *huge* fixed costs to become managers. At each step they have invested everything they have in their career trajectory. At many steps, they may regret having started down this path, but don't realize all they are doing is digging a deeper hole for themselves. Or they realize it, but can't figure out a way to stop.

Production costs are asymmetric. There are massive asymmetric *sunk* costs. Sunk costs don't count. Everyone has already given everything to be able to produce. Everyone is producing the same thing as a result, with the same marginal costs. Anyone not willing to visibly bear those costs is wiped out for that alone. Effectively, everyone has the same production costs aside from political battles.

Economies of scale exist. Each manager is only one person. While the companies themselves have economies of scale, the competition between employees has no such concept.

Producer preferences differ. Producer preferences (aka employee preferences) in theory still exist, but letting them noticeably impact decisions is fatal. Doing so in secret still means sacrificing position in the game for whatever else you care about. Successful managers learn not to do that.

Location matters. Location is expected not to be a limiting factor. If a manager is asked to go manage a plant or office in another state or even abroad, refusing to do is a sign that they are insufficiently committed. They might try to steer decisions a non-zero amount to avoid the costs involved, since those costs do interfere with one's ability to do the job, but this mostly isn't a thing. This prevents differentiation based on location, which was how this made the original list, and it also wrecks the lives of the managers involved since they cannot choose where to live or predict where they will be.

Consumers care about the individual producers. It is considered a grave sin in middle management culture to care about your underlings, also known as your producers. You protect them while they are loyal and competent and not the scapegoat at this time, because that is part of the job. You need to be seen as the type of person who does that. But actually sacrificing value for them would be seen as terribly weak. As you do not care about them, the people above you do not care about you either. Loyalty until such loyalty is no longer useful is a thing, but that is very much not the same thing nor does it fill the same role.

Everyone is judged on a combination of politics, and whether they can properly produce the homogenized product 'successful middle manager.' This centrally includes associating only with those also producing such products, rewarding such products, and punishing any deviations from such products. You are responsible for everyone and everything below you, so you must demand fealty from those below you, so they'll serve whatever agenda is most helpful to you. You must also ensure that they are producing the 'successful middle manager' product.

The Opt-Out Clause

The good news for managers is that once one reaches a certain level, as long as one continues to play by the rules and make others comfortable (comfortable is a key concept we don't have space to discuss here, for more see section B of [Quotes from Moral Mazes](#)), one can typically ride things out while carving out some amount of outside value for yourself. One does this by visibly opting out of the game, and no longer being seen as a threat. [Everyone knows](#) that some people must stay in place to keep the wheels mostly turning. Without such people the whole system doesn't collapse too fast and take current management down. Thus, there is an option for graceful retirement from the climb up the ladder.

More generally, those who accept immobility are unwilling to sacrifice family life or free-time activities to put in the extraordinarily long hours at the office required in the upper circles of their corporations. Or they have made a realistic assessment of the age structure, career paths, and power relationships above them and conclude that there is no longer real opportunity for them. They may see that there is an irreparable mismatch between their own personal styles and the kinds of social skills being cultivated in well-entrenched higher circles. In many cases, they decide that they do not wish to put up with the great stress of higher management work that they have witnessed. (Location 962, Quote 119, Moral Mazes)

Producing 'successful middle manager' is largely a pass/fail operation. You can't produce it by halves. Thus, once you realize you are unable or unwilling to produce it properly, either you move somewhere where you can do so, you quit entirely, or you stop trying to produce 'successful middle manager' and instead produce 'contented middle manager.' Those are the level one damage mitigation strategies.

Later sections will deal with such damage mitigation options more broadly.

Is This Missing Something?

On one level, this robustly answers the question “how did things get so bad?”

On another level, this begs that question entirely. It does not explain why everyone involved ended up with such crazy beliefs, or allowed things to get so bad. It merely explains the bind the players are in within the game as it currently exists.

There is the implication that the whole thing is wrapped in conspiracy and malice that is left implicit, its motives and modes of operation unclear. Without that element, the explanation feels incomplete at best.

This gap is because the model only implicitly mentions Moloch’s Army. I still haven’t quite figured out how to take that on, slash we’re not ready yet. Hopefully you now have a better idea of what this represents.

Mazes Everywhere

These dynamics don’t only exist in major corporations. By default, at least right here and right now, they arise in *any* sufficiently large organization. Barring strong optimization pressure holding them back, things over time in any given organization with multiple levels of management will become increasingly maze-like.

Governments are presumed to be mazes. Armies. Political parties. Unions. Academia. Sufficiently large non-profits are almost certainly mazes – the profit motive is if anything *fighting against* the property of being a maze. Getting rid of it would only speed things along.

This is doubly true because mazes do not only support themselves being mazes. They tend to support and encourage maze behaviors in other places and organizations, favoring mazes over non-mazes. Our major corporations have largely been mazes for decades.

A maze *need not even be a formal organization*. You can get similar effective behaviors whenever there is an effective multi-level hierarchy.

Consider Paul Graham and Sam Altman’s description of the ecosystem surrounding Y-Combinator, [as analyzed recently by Ben Hoffman](#). Or the model of the same system in my older post [In a world... of venture capital](#), where each round of funding can be considered the boss of the investors in the previous round.

The Next Questions

There now exist three categories of next questions.

The first category (taken from [this comment](#) I made on Less Wrong) are the big global questions. To what extent are these dynamics the inevitable result of large organizations? If so, to what extent should we avoid creating large organizations? Has this dynamic ever been different in the past in other places and times, and if so why and can we duplicate those causes?

The second category are the big personal questions. Given the existence of these immoral mazes, what do I do?

The third category is the one I've been struggling with, which is to finally get a good written model of the dynamics of anti-epistemic anti-virtue.

What is Success in an Immoral Maze?

Previously in Sequence: [Moloch Hasn't Won](#), [Perfect Competition](#), [Imperfect Competition](#), [Does Big Business Hate Your Family?](#), [What is Life in an Immoral Maze?](#), [Stripping Away the Protections](#)

Immoral Mazes are terrible places to be. Much worse than they naively appear. They promise the rewards and trappings of success. Do not be fooled.

If there is one takeaway I want everyone to get from the whole discussion of *Moral Mazes*, it is this:

Being in an immoral maze is not worth it. They couldn't pay you enough. Even if they could, they definitely don't. If you end up CEO, you still lose. These lives are not worth it. Do not be a middle manager at a major corporation or other organization that works like this. Do not sell your soul.

When one works for an immoral maze, what is one hoping for? What is success?

Suppose you persevere. You make the sacrifices. Become the person you need to become. Put in the work day after day. Fortune smiles on you and you win out against all the others doing the same thing. You succeed.

What is success? What do you get in exchange?

For some managers, the drive for success is a quest for the generous financial rewards that high corporate position brings. For others, success means the freedom to define one's work role with some latitude, to "get out from under the thumb of others." For still others, it means the chance to gain power and to exert one's will, to "call the shots," to "do it my way," or to know the curiously exhilarating pleasure of controlling other people's fates. For still others, the quest for success expresses a deep hunger for the recognition and accolades of one's peers. (Location 955, Quote 118)

This a warning. "Success," in context, does not mean happiness. It does not make you healthy. It does not improve your reproductive fitness. It does not reflect or spread the values that you (one would hope) had when you stared down that road.

It gives you money. But in terms of actual meaningful personal consumption, you can't really do much with it beyond status competitions. If you had plans to do something good with the money, by the time the day arrives, it is highly unlikely you'll do it. You have changed yourself to succeed on your journey.

Even after you 'succeed' you probably keep putting tons of hours into the job in ways no amount of money can compensate for, once you already had basically enough.

What was the point? What are you even doing?

Note that failure is indeed much worse than success. You still paid all the sunk costs, including everything you are. You've invested a ton in and become very invested in local status hierarchies, and in the quest to climb them, which you have failed. Being 'under the thumb' of others who succeeded where you failed is deeply unpleasant – and is the most likely outcome, since the math says most who try will fail.

This is a song with some [explicit content](#) about what happens when one disregards this warning, and chooses poorly, although it misses perhaps the most important questions. Who am I? What have I become?

<https://www.youtube.com/watch?v=5IsSpAOD6K8>

[Lyrics here.](#)

Remember that maze conditions are not unique to corporations.

All of this holds true in any sufficiently large organization, to an extent that increases with its size, and will have the same effects if you seek to ascend the hierarchy within.

Size matters, but size is far from the only thing that matters. Some very small organizations effectively have very high maze levels. Some large organizations have *relatively* low maze levels.

Avoiding mazes is easier said than done. The first step is *identifying* them, where I will offer some heuristics in the next post.

The second step after that is what to do about it, especially in difficult circumstances. Many of us believe we need the support of mazes in order to survive. At least for the moment, not all of us are wrong.

How to Identify an Immoral Maze

Previously in sequence: [Moloch Hasn't Won](#), [Perfect Competition](#), [Imperfect Competition](#), [Does Big Business Hate Your Family?](#), [What is Life in an Immoral Maze?](#), [Stripping Away the Protections](#), [What is Success in an Immoral Maze?](#)

Immoral mazes (hereafter mazes), as laid out in the book [Moral Mazes](#), are toxic organizations. Working for them puts tremendous pressure on you to prioritize getting ahead in the organization over everything else. Middle managers are particularly affected – they are pushed to sacrifice not only all of their time, but also things such as their morality, family and ability to think clearly. Only [those who go all-in doing this](#) get ahead, and even most of them fail.

Even successfully getting ahead [is little consolation](#).

Mazes exert similar pressures on those who do business with them or work in non-managerial roles, to a lesser but substantial degree.

The best defense is to identify mazes before you agree to work for or do business with them, and choose to work or do business elsewhere. At a minimum, one's eyes should be open, and the costs involved must be fully factored in before making such decisions.

This makes it important to figure out what parts of what organizations are mazes, and to what extent. This is hard to get exactly right.

What is easier is using simple heuristics to get a good approximation, then keeping an eye out for and updating on new evidence.

I offer seven heuristics, the first two of which will do the bulk of the work on their own. You benefit from the 'right' answer to all of them even absent concerns about mazes, so they are good questions to get into the habit of asking.

1. How many levels of hierarchy exist?

Full mazes require *at least* three levels of hierarchy, without which one cannot have middle management.

Each level beyond that makes things worse. The fourth and fifth levels both make things much worse.

With only one level, there's nothing to worry about.

With only two levels, a boss and those who report to the boss, the boss has skin in the game, no boss causing problems for them, and not enough reason to reward bad outcomes.

With three levels, there are middle managers in the second layer, so one should be wary. But things are unlikely to be too bad. No middle manager has a boss or underling who is *also* a middle manager. This means that in any interaction between non-equals either involves the head of the company, or it involves someone 'on the line' who doesn't have anyone reporting to them, and must deal with object-level

reality. Either of them has reason to keep things grounded. Since there is only one person at the top, every conversation includes someone who interacts regularly with object level reality.

With four levels, we start to have interactions between middle managers in charge of each other. These dynamics start to get serious, but everyone still interacts with someone on the top or bottom.

At five levels, we have people who never interact directly with *either* the boss *or* anyone dealing with the object level.

At six levels, those people interact with each other.

And so on.

Meanwhile, the boss has less and less need or ability to comprehend the object level, and we get more and more problems with lack of skin in the game, which is question two.

At least one of the corporations in *Moral Mazes* had more than *twenty* ranks. That is *way, way* too many. By that point, it would be surprising if you weren't doomed. I have *actual no idea* how to have twenty ranks and keep things sane.

Note that those outside the company, such as investors or regulators, seem like they should effectively count as a level under some circumstances, but not under others.

As a spot check, I looked back on the jobs I've had. This matches my experience.

Most impressive is that I can observe what happened when several of those jobs *added new layers of hierarchy*. This led *in every case* to traceable ways to additional maze-like behavior. In every case, that made life much worse for me and other employees, and hurt our productivity. In one case I was running the company at the time, and it still happened.

I would be very wary of any organization that had four levels of hierarchy. I would be progressively more skeptical of any organization with more than that, to the point of assuming it was a maze until proven otherwise.

2. Do people have skin in the game?

Skin in the game is a robust defense against mazes, if it can be distributed widely enough and in the right ways. That can be tough. There's only 100% total equity to go around.

One can only reward what can be observed or often only what can be quantified and measured. [Something about Goodhart's Law](#), and so on. The problem with levels of hierarchy and middle management is in large part a problem of inability to provide skin in the game.

For sufficiently large organizations, as described in *Moral Mazes*, skin in the game is not so much spread thin as deliberately destroyed. The successful keep enough momentum to run away from the consequences of their problems. This alone is fatal.

If an organization has solved these problems for real, it likely isn't a maze.

If an organization lacks skin in the game and also has many levels of hierarchy, you're almost certainly dealing with a maze.

If it lacks skin in the game but also lacks levels of hierarchy, maze levels can differ. But also keep in mind that lack of skin in the game causes a whole host of problems. Only some of those are the problems of mazes. Detailing these issues is beyond the scope here, but be highly skeptical whenever skin in the game is lacking.

3. Do people have soul in the game?

What's better than having skin in the game? Having soul in the game. Caring deeply about the outcome for reasons other than money, or your own liability, or being potentially scapegoated. Caring for existential reasons, not commercial ones.

Soul in the game is incompatible with mazes. Mazes will eliminate anyone with soul in the game. Therefore, if the people you work for have soul in the game, you're safe. If you have it too, you'll be a lot happier, and likely doing something worthwhile. Things will be much better on most fronts.

It's worth prioritizing soul in the game, above and beyond skin in the game.

4. How do people describe their job when you ask?

Remember this quote:

When managers describe their work to an outsider, they almost always first say: "I work for [Bill James]" or "I report to [Harry Mills]" or "I'm in [Joe Bell's] group," and only then proceed to describe their actual work functions. (Location 387, Quote 2)*

You want them to say *almost anything else*. Anything that does not make you recoil in horror a different way. Hopefully something worthwhile and interesting. I don't know how good this rule is, but I suspect it's quite powerful.

5. Is there diversity of skill levels? Is excellence possible and rewarded?

The belief that all middle managers have the same skills, and are all equally capable of doing any managerial job aside from the politics involved, is a lot of what makes mazes so bad. If there is no *good* reason to diverge from standard practice, if everybody knows that you *cannot do better*, then any divergence is blameworthy, and shows you are not doing your job. There's no need to ask why, or what advantages it might have.

It also all but ensures the wrong answer to the next question.

6. Is there slack?

A world without slack is not a place one wants to be. Mazes systematically erase all slack. Slack is evidence of not being fully committed, and given that everyone's skills

are equal and competition is perfect, holding anything back means losing even if undetected.

7. Pay Attention

Sounds silly, but it works. Observe people and what they do and how they do it. If you work in a maze for long enough, you're not going to shout it from the rooftops, but every sentence you speak will reflect it.

And as always, when people tell you who they are, believe them.

Other Notes

These questions do not differentiate between corporations, non-profits, governments, parties, clubs or other organizational forms. That's not a good indicator. Corporations are only the original observed case.

Asking how proposed or expected changes will change the answers to these questions is a good way to know if those changes will raise the maze level of an organization.

Like most puzzles, there are multiple solutions, and the pieces reinforce each other. Most of the time, hardcore mazes will give alarm-bell level answers to all seven heuristics.

Are there any other good simple heuristics?

Next is How to Work With Moral Mazes, providing my best advice in detail to those dealing with the threat of mazes on a personal level.

How to Escape From Immoral Mazes

Previously in sequence and most on point: [What is Success in an Immoral Maze?](#), [How to Identify an Immoral Maze](#)

This post deals with the goal of avoiding or escaping being trapped in an immoral maze, accepting that for now we are trapped in a society that contains powerful mazes.

We will not discuss methods of improving conditions (or preventing the worsening of conditions) within a maze, beyond a brief note on what a CEO might do. For a middle manager anything beyond not making the problem worse is exceedingly difficult. Even for the CEO this is an extraordinarily difficult task.

To rescue society as a whole requires collectively fighting back. We will consider such options in later posts.

For now, the problem statement is hard enough.

To reiterate the main personal-level takeaway

[Being in a maze is not worth it.](#) They couldn't pay you enough. Even if they could, they definitely don't. If you end up CEO, *you still lose*. These lives are not worth it. Do not be a middle manager at a major corporation or other organization that works like this. Do not sell your soul.

Problem statement

Increasingly, avoiding mazes is easier said than done.

First, one must identify them, for which the previous post offers a guide.

After that, there are still many hard problems to solve.

How do we avoid moral mazes? How do we justify that choice to others? What alternative choices do we have? What if we're already in a maze? What if we've already self-modified in ways that make it hard to extract ourselves? What if our human or social capital only pays off inside them?

What about if you are doing object-level work without anyone who reports to you, but you have a maze above you?

And for those who think this way, do I have a moral obligation to suffer and do this anyway, in order to maximize my charitable giving, or to otherwise do good?

How do we avoid immoral mazes?

Truly understand how painful it will be to interact with a maze even if you're not an employee. Know the signs, as discussed in the previous post. [Keep a close eye out for mazes.](#) Realize that you have other options. Choose other paths.

This isn't an 'all things being equal' choose other paths. This is making what *look* like major sacrifices and different life choices and profession choices, or taking big risks (that may or may not include starting a business or doing work outside of an organization) in order to have skin in the game. Really understand that the offer from even a relatively tolerable maze is much, much worse than it looks, and that opportunities outside mazes are often much better and more realistic than they look.

Young people starting out in the labor market often have The Fear that they will never find a job or never find a good job or another good job. If you are capable of getting this far, and you persevere, that is not true for you. A wide variety of jobs and other opportunities are out there.

I realize some people have already become so trapped in mazes that they cannot walk away.

If you actually can't walk away, see the last two questions.

What do you do if you find yourself inside a maze?

Quit. Seriously. Go do something else. Ideally, do it today.

At least start planning and looking. Every day there is another day you suffer, another day you invest your social and human capital in ways that can't be transferred, and another day you become infected by the maze that much more.

If you actually can't afford to quit, see the last three questions.

How do we justify our choice to others?

When I worked for a financial firm, the question 'what do you do?' (or, in a scarier form, 'who are you?' as implicitly defined by your job) had an easy answer. I work for (firm). One of the big benefits was being able to tell an easy, compact story of me and my life and work choices, that most found praiseworthy. It was easy. It was *comfortable*. It also worked wonders for things like renting an apartment or otherwise proving myself respectable.

A lot of the alternative answers that don't involve mazes give you a much better life and method of earning a living, but they do make answering the 'what do you do?' and 'how can I count on you to make rent or support a family?' questions trickier. One must acknowledge this.

It isn't only strangers you tell this story to. It is your friends. It is your family. It is also yourself.

I likely stayed at (firm) months longer than I should have due to being scared of not being able to tell this story anymore, especially to my wife and to myself, and having to instead tell a different one.

A lot of this fear is the expectation that others won't understand and won't accept our justifications. That does happen, but *far less than people typically expect or fear*. Most people are far more sympathetic than the inside view might suggest.

Even the internet is supportive. Which is not its style.

This is largely because, at least for now, there is a widespread cultural belief that one should do what you love, and be content in one's work. That work should provide meaning. That's not always a good rule or good idea, although it is a fine aspiration. Not everyone can have soul in the game. But almost everyone recognizes that it would be better if one did.

How do you go about telling your new story? (Justification continued)

Here's my take on how to approach this, based on my experience. Comments suggesting improvements or alternatives are highly encouraged.

There are two parts of this.

One is to figure out what you *are* doing, not only what you're *not* doing, and how to talk about that. Some of those answers are mostly culturally normal and comfortable, some of them are less so. Now that I can say 'I'm a game designer' that goes over quite well.

The most important things here are to make the thing you are doing sound *simple*, put it in terms that people can relate to, and to make it clear that you are comfortable and happy with it to the extent that this wouldn't be lying to people. If you're not comfortable and happy with it, people will pounce on that. It's also much better to be happy with what you are doing for your own sake, so that is something to work on, either looking to get to that place, or to finding another option where you can do that.

If you quit today without a plan, then what you will be doing is recovering from your experience and figuring out what to do next. I told that story for about a month. That story goes over better than you would expect – for a while. From a social (as well as financial) perspective you are most definitely on a clock. There are plenty of people who let that clock run out. But if it's two weeks in, own it.

The other part of your explanation is justifying why you're not doing the standard thing of indenturing yourself to a new maze, unless you have an obviously great alternative thing going. The worse your answer to part one sounds, the harder part two is going to be.

First try giving it to them straight. Tell them you find large corporations highly toxic and morally compromising. It left you a wreck. You have no interest in the lifestyle you would live or the person that you would be. If they are genuinely curious, you can point them to *Moral Mazes* itself or this series of posts, or explain further in your own words.

You can also use the culturally assigned incantations to explain your decision. Tell people you need to do what you are passionate about, to 'follow your heart/passion,' to do what you believe in, to 'help people' or 'make a difference.' To 'get your hands dirty' and 'do something real.' Some people appreciate wanting to 'be your own boss' or 'do it your way,' which are weaker, non-dystopian ways of sending the real message.

And of course, you can simply say 'it was making me miserable. I hated it. Doing this instead makes me happy.'

I've gotten into the habit of saying my job at (firm) was 'a poor fit' because I genuinely believe both that there were particular real needs they have that were expensive for me to provide, and that the firm was in many ways unusually great and unusually low on maze characteristics, and I do not think it is a mistake to take a job there if it would suit you. They treated me right and I don't want to throw anyone under the bus.

You *could* also, if you wanted to, use the question as an opportunity to do a public service and spread the word. That's supererogatory.

Another thing to keep in mind, as discussed in the next question, is that those pushing us towards mazes are often operating on traditions and heuristics that used to push towards virtuous action that led to happiness and real success. The world changed, and those traditions and heuristics started getting this wrong. This is highly sympathetic. It might help to approach from this perspective.

Most people get it. They don't *fully* get it unless they've been on in the inside and reflected upon what happened. But they do get that there's something soul-killing about working for the man and/or being lost in a maze of political actions.

Others won't get it.

What if my family or culture won't accept my justifications?

Some people won't get it. They will respond that all of this is excuses for not wanting to do hard work or make sacrifices. That this is how the 'real world' works. That it is 'time to grow up.' That a 'real adult/man/woman/hero/whomever' would suck it up and deal with it. That it is your responsibility to do so, for your family, for the world or for yourself. That 'get a steady job' is what good and responsible people do. That this is how one survives in today's world, and how one gets to raise and support a family.

Often people who are counting on you, usually family members, will effectively let you know that while they care a non-zero amount about whether your life experience is miserable, or what impact your work has on the world, or what upside or opportunities for personal growth you might have, what they *actually* care about is whether you are projecting the illusion of security. They want to mentally cache that you/they 'are going to be OK' and that 'everything is all right.'

This has remarkably little to do with actual security. Jobs in many mazes are not especially secure. Others are secure barring disruption of the underlying order, if you are willing to pay the prices discussed, and tie all your human and social capital to the maze.

The security they seek is the security of the banker who loses money when everyone around him loses money. This is useful to the banker because one cannot then scapegoat and fire him if he has bad luck and does poorly. This is useful to those in a maze and those who tie their fates to them, because they hope they will similarly seem responsible and legitimate, and thus worthy of sympathy and assistance if things go poorly.

It is a self-reinforcing charade. People demand this illusion of legitimacy to protect against others' accusations of illegitimacy. They will defend this charade even if times change, the mazes mostly fail and those who are doing real things succeed, and attempt to forcibly transfer wealth from people doing real things to those who previously worked in mazes. All of that doesn't make participation in the charade

worthless. You can even hope to benefit from the expropriations. But it is important to know that it is a charade.

If you face a family and/or culture that demands devotion of one's entire life to the illusion of respectability, have sympathy for those making these demands. In the past, when these traditions and heuristics developed, the illusion of respectability corresponded to real work and other worthy virtues that lead to happiness and true success. If they fail to realize the change, and update accordingly, that is at least somewhat understandable.

If attempts to make them realize this change or accept your perspective fail, one must treat them the same way one treats anything else that is [out to get you](#). If their demands cannot be satisfied in a way you can accept, because they will simply demand more until it is something you cannot accept, then attempting to satisfy their demands is folly.

If you previously chose those around you based upon being a member of a maze, then it is plausible that having invested in those people and relationships it makes sense to stick around. It is also plausible that being around them afterwards no longer makes sense.

I know it is easy to say and tough to act upon, but hopefully in time either they will understand and/or come around, or you will realize that life is better without them.

What if you've already self-modified too much?

This sucks quite a lot.

After a while, those little status differences and little battles start to deeply matter. Other things matter less and less. Humans can adapt to many things. Giving up all that likely fills you with deep existential dread.

Even worse, you've sculpted everything else, including your friends and often your family, around these obsessions. You, and often they, depend on the currently steady money *and the illusion of security* the maze provides. Without that, things could rapidly fall apart.

The good news is, you've figured out that this happened. Perhaps you haven't self-modified as much as you've feared, or you have a path back to undoing this. Often the modifications start reverting once you extract yourself from the situation. Often you're deeply miserable, in a way that those around you at least subconsciously know quite well. Those around you often realize this long before the person realizes it about themselves. If you tell the people you care about what's really going on, if they're worth keeping, they'll almost always be supportive.

I've seen a number of people realize they hated their jobs and needed to go. Almost all of them got lots of support when they got the courage to say it out loud.

A note I got on a previous draft: "The happiest Uber drivers I have seen used to be middle management."

If anything, I see many people around me being *too supportive* of opting out of working, or in a sense out of life, entirely. It is important to help and encourage people to do more things.

See the question above on how to explain your choices and situation to others.

So my first suggestion is to admit to yourself what is happening to you. Take an inventory. Concretely observe *what is actually going on around you*, without excuses or euphemisms, and what that is doing to your brain and your life.

Then *tell the people who care about you*. And go from there.

Are mazes where our human and/or social capital pays off?

Note that when I say 'pays off' here, I mean *maximally* pays off. If you have the skills and opportunity to advance inside a maze, you have the skills and opportunity to take a lower level position at a smaller institution, and still earn a reasonable living. That does not mean that this transition would not be painful, or that you could maintain your current lifestyle, or that your family and friends would stand by you, *but you definitely have that option*.

If you are an academic with a PhD, and notice your academic institution is a maze, keep in mind that the entire concern about you only getting paid off in academia probably *simply isn't true*. Academics typically get substantial pay bumps when they move to private industry.

A lot of other professionals are similarly buying the security and familiarity of what they're used to, rather than being paid off with dollars.

One must still be careful to ensure not to jump from one maze into another.

Then there's big corporations. Big corporations *really do* pay better than smaller businesses if you can't get equity in either business. The recent history of domestic 'income inequality' is in large part inequality *between firms* as bigger and more successful firms pay higher salaries even to their lower tiers, and have more higher tiers in which to earn yet more.

There are three theories I know about for why big corporations pay more.

Theory one is that big corporations are more likely to have [O-ring production functions](#) or otherwise benefit more from higher quality workers, so they pay more in order to attract better workers.

Theory two is that big corporations make more money per employee, and are big enough to potentially support unions, so employees demand and receive more of that pay.

Theory three is that working in a big corporation sucks, and employees realize this to at least some extent, so employees demand more money in order to be willing to work there.

If working at a major corporation is a major life cost, and working in management a bigger one, and these come with higher pay, than a lot of income inequality in developed countries does not represent a gap in desired life outcomes, and it might be *more unfair* if that part of the gap was closed.

A lot more of that pay gap is that some professions engage in rent seeking behavior to extract resources. Some big examples are finance, law, education/academia, and medicine. Again, that comes with much better pay.

It also usually comes with a big time investment in the development of the relevant social capital, human capital and credentials you need to succeed. If you went to medical school or law school or worked hard to get a tenure track, and later realize that your profession is a maze (I'm making no claim here that these professions are or aren't mazes in general or how intense those mazes might be, although some central organizations within them clearly are very intense mazes), walking away from that is going to be expensive.

This is doubly true for human capital *within a single organization*. When I took a job at a financial firm, they spent a large part of my first *two years* training me. The first year had a lot of firm-specific detail but was mostly training about markets and trading in general, that applies everywhere there are markets. It was fascinating, and pretty great. Five stars, would study again, especially given they paid me rather than the other way around. The second year still had a lot of training and learning, but increasingly it was about the specific problems I was working on, developing relationships with and learning about coworkers and organizational structures and how we did things, and other information specific to the firm. This was less fun, and when I left, it became worthless.

I had another job I stayed in for five years. This was also a place I got to observe transition from mostly not being a maze into becoming one over time, although that's a story I can't tell online.

Early on there was a lot of learning, a lot of which was very specific to our business, but a lot of which applies universally. I worked mainly with one particular person, who knew what we were doing and cared about us doing it well. It provided great experience.

By the third year, I was learning about our specific products and customers and dynamics, in increasingly arcane fashion. I was also forced to interact increasingly with the maze growing around us, spending more time making bosses and others like what they saw rather than doing what was right for the business. I was unable to get the resources to enhance our performance, despite yearly returns on investment obviously well above 100%. I made an effort to switch over into problems that were both more valuable and offered more room for growth, both for me and for the business, and which I could tackle with the resources available.

By the fifth year, I wasn't developing any skills that would be useful elsewhere, except that I was now learning to code because I got tired of no one being able to code what I needed. This brought me from 'can code but not in an actually useful way' to 'can code real things that are useful, but badly/slowly.'

Note that the managers in Moral Mazes who succeed were always moving around to bigger and better things. If they weren't, they instead moved on to similar and different thus hard to compare things to preserve the illusion of career momentum. If you have adopted the maze nature, many of the skills you have learned doing so translate to other mazes. Your existing within-maze status can often also be transferred to your new location, but only if you continue to be seen as a winner. If you're a loser, no one else will want you, and moving on will mean moving down.

That means that once your path in the maze is stalled, even though you have invested a lot to get to this point, recovering your momentum is going to be extremely difficult. If you are not satisfied with your current role, your human capital is a lot less valuable than it naively appears to be, because it no longer has much upside even on its own terms. The fall from where you are to starting over can still be large.

The best feature of an academic maze is that they have a perfectly designed system in which to not care about getting ahead, which The Gervais Principle calls a loser, and which academia calls tenure.

The pattern remains. The more you dedicate time to a path, both a profession and a particular job, the more you give up when you leave and the less of your time you can carry with you. Many people don't have great options. The job market isn't that great out there if you don't want to be coding and don't have an in with the rent seekers, and can't use the skills you've developed, or do the thing that legibly follows from your resume.

If mazes are where my social and/or human capital pays off, what should I do?

Let us ignore here *why* your capital pays off best in a maze. It does not much matter to your decision, in important senses, to what extent rent seeking, theft, coercion, fraud or even systems designed explicitly to make your skills not transfer to honest work are or aren't responsible for this being the case. For whatever reason, often events conspire to prevent you from efficiently plying your trade (or in some cases, plying it at all), where you hold comparative advantage, without being part of a maze.

Some people really do have a dilemma, where they can either do something menial and mindless that still gets them abused and doesn't pay much, if they can find work at all. Or they can go out on a limb that looks super risky and likely to fail, and/or that requires years without compensation. Or they can keep working in the maze.

It is important that *vastly more people **think** they are in this position, than **are** in this position*. If you think you are in this position, consider the possibility that you are mistaken. Consider all the alternatives. Consider how much the reduction in medium-term funds and superficial status would actually matter to you. Consider how much of what's holding you back is simply The Fear in some form.

Imagine exactly how relieved you'd be to be out of there. Remember that *even if leaving really is super painful*, involving a large reduction in consumption levels and superficial status and standard of living, and the abandonment of large sunk costs, that doesn't mean it isn't Worth It.

My first line of response to this dilemma is exactly what you would expect: *Consider leaving anyway*. But I admit that isn't *always* the right answer. In some cases, things really have gone too far, you have too many promises to keep and too many sunk costs.

Become a Loser

The next line of defense is to *become a loser*, in the sense laid out in [The Gervais Principle](#). A loser does not strive to get ahead while at work. A loser finds their value in

other places than work. At work, they pride themselves on putting forward *at most* the minimum amount of effort to get the job done.

The Gervais Principle can be seen as the prequel to *Moral Mazes*, dealing with life at lower levels of mazes that have to interact with the real world. Mazes need, as [several quotes](#) describe, people who keep their heads down and 'do their job' with no ambitions for further advancement. Ideally one does this as low on the totem pole as one can stomach and afford, as the life that results is far less odious and taxing.

By declaring themselves as neutral and not a threat, such people are often left mostly alone if they're important to the system continuing to run. They can now reclaim some slack and a personal life. It's not a great solution. You're still holding up the maze. You're still interacting with it. You'll still have to make severe moral sacrifices. But to some extent, some of the time, you can pick and choose what to have no part in.

Over time, your position likely will slowly degrade. Eventually this may lead you to leave. Hopefully by then you'll have been able to save enough and be prepared enough to be ready for that. If you're stuck in a maze, the least you can do is turn a healthy monthly profit.

Take Risks

The final line of defense I can come up with is to take big bold risks. Either stand up for what you believe in or gamble to advance your own situation. Sometimes this will work, your situation will improve and you'll learn your situation was better than you thought. Other times they'll backfire, and you'll learn your situation was worse than you thought and is now worse than that. Remember that if you get fired from a job you don't want, that can be a big win, because you might not have had the courage to leave on your own and you might even get severance and unemployment.

The real danger is often not that you get fired. It's that you become 'dead without knowing it' as in this quote:

You can put the damper on anyone who works for you very easily and that's why there's too much chemistry in the corporation. There's not enough objective information about people. When you really want to do somebody in, you just say, well, he can't get along with people. That's a big one. And we do that constantly. What that means, by the way, is that he pissed me off; he gave evidence of his frustration with some situation. Another big one is that he can't manage—he doesn't delegate or he doesn't make his subordinates keep his commitments. So in this sort of way, a consensus does build up about a person and a guy can be dead and not even know it. (Location 1475, Quote 10)

This can lead you to waste years of your life struggling for something you had no chance of getting. This is one reason why a great way to take risk is to force the issue, asking for or demanding raises or promotions. It avoids this danger. The more uncertain you are about where you stand, the more you should take risk to create clarity.

At all my jobs in mazes, I would have greatly benefited from taking greater risks to create clarity, *regardless of the outcome*.

Can You Change Things From the Top?

If you by some miracle reach the top with your soul intact, now you can try and change the system. Or at least you can do harm reduction in earnest. One shouldn't give this much hope or weight, since such intentions rarely survive that long, and doing anything lasting about it will still be quite hard. I don't know what would work.

My friends and I have talked to several people who have reached the top. Many of them understand what the process has done to them, but don't know how to fix themselves or the system. It isn't cheap or easy to reverse or even halt the damage.

It is unlikely you can have much impact without reaching the actual top and becoming CEO. If you do become CEO, you may have a short window in which you can 'clean house' the way that maze CEOs do, and put people opposed to mazes into key positions where they can clean their houses in turn. You can then combine that with other efforts, and maybe get somewhere, but I don't have the insights necessary to say much more, and such efforts will be exceedingly difficult. The maze will fight back.

I strongly believe it is much easier to build a new system from scratch than to '[change the system from within](#).'

What about if you are doing object-level work without anyone who reports to you, but you have a maze above you?

In *Moral Mazes* such workers are said to be 'on the line.'

Details of this situation will determine to what extent this represents being stuck in the maze, versus to what extent this represents doing regular object-level work.

What are you actually doing all day? What are your incentives?

If your essential scenario is *given an object-level job to do and do it*, that is mostly fine.

If your essential scenario is not that, it is less fine, but it is still far better than being a middle manager. It's not *good* to have a bull**** job, but it's not the nightmare we're describing elsewhere.

Consider the car salesmen from [Imperfect Competition](#).

One can imagine a car dealership as no different from the local hardware store, buying useful tools wholesale and selling them at a higher price to customers who want to buy and use those tools, and the only difference is that you sell 0.1% as many tools for a thousand times the price. One can also imagine that the demands of the car corporation, and the incentives they provide, and the misinformation they spread, and the regulations they twist and engineer, and so forth, end up with you effectively stuck in the maze. The truth is presumably somewhere in between – you see insane things around quotas and regulations and advertising campaigns you cannot control, and the dealerships have their own issues of their own design, but you are still mostly working for a small actual business most of the time.

The same would go for the workers in *The Office*, as analyzed in *The Gervais Principle*. Michael is largely in maze hell. Jim spends a lot of time avoiding maze hell. Most of the workers have to deal with the craziness and what it does to the business, but this is only ordinary soul crushing and not what middle managers deal with.

Jim's situation on The Office is the biggest problem. There is no future. The only way up, to better your work situation, would be to dive into the maze. If you do that while not buying into the system, it will go badly for you on all levels. If you do buy in, then you've fully made the big mistake I'm warning against.

Consider the Uber drivers, some of whom are reported to be happy refugees from middle management.

To the extent that the driver is offered rides, chooses to accept them individually, and gets paid for each ride provided, the driver is good. They set their hours and level of effort. There is word that Uber and its ilk are now using algorithmic systems and various overall incentives to try and ensnare their drivers more broadly into the system, which would be worse, but the core experience is still one of real work.

Consider a software engineer, given specific tasks to code and coding them. That seems likely to be mostly fine.

Consider a worker that is literally 'on the line' in a manufacturing plant that makes physical objects. It is not the best or most compensated work, but you are mostly free from the maze.

Being 'on the line' and continuing to do real work is miles behind doing real work where you have skin in the game, but if you get to dodge the worst of all this, it is a reasonable temporary fallback if you lack alternatives. Look carefully at details.

If you are a manager but not a middle manager (e.g. no one who reports to you has anyone who reports to them), and the group of people you manage has object-level tasks to do together, you aren't automatically doomed, but there is great danger lurking, including the risk you will be promoted.

Do I have a moral obligation to work in mazes to maximize my charitable giving?

No. You don't.

This post has done its best to deliberately ignore the moral costs of participating in mazes, *because avoiding them is already over-determined without that.*

I want to make it clear that I'm *not relying* on moral concerns.

But if that's what you are concerned about, moral concerns work in the opposite direction. Making the world more and more maze-like by embracing the system, and engaging in zero-sum competitions to extract resources, while making your life miserable, is the opposite of a moral obligation.

It may help to remember that [a drowning child is hard to find](#).

Moral systems that imply that subjecting oneself to torture in the service of immoral mazes or other harmful systems, for the purposes of allowing other such systems to then extract those resources from you, is a moral obligation, are not likely to be good ideas, or to have your or humanity's best interests at heart.

The Road to Mazedom

Previous post: [How Escape From Immoral Mazes](#)

Sequence begins here: [Moloch Hasn't Won](#)

The previous posts mostly took mazes as given.

As an individual, one's ability to fight any large system is limited.

That does not mean our individual decisions do not matter. They do matter. They add up.

Mostly our choice is a basic one. [Lend our strength to that which we wish to be free from](#). Or not do so.

Even that is difficult. The methods of doing so are unclear. Mazes are ubiquitous. Not lending our strength to mazes, together with the goal of keeping one's metaphorical soul intact and still putting food on the table, is already an ambitious set of goals for an individual in a world of mazes.

We now shift perspective from the individual to the system as a whole. We stop taking mazes as given.

It is time to ask *why* and *how* all of this happens, and what if anything we can do, individually or collectively, about it.

In particular, this post presents an explicit model of the first question, which is:

How did mazes come to be, both individually and overall?

This is partly a summary of the model developed so far. It is partly making the model more explicit, and partly the fleshing out of that model with more gears.

If there are points here for which you believe the previous posts failed to lay the groundwork they should have laid, please note that in the comments by number, so I can consider fixing that. Keep this distinct from any disagreements or other notes on these claims, which are also welcome.

1. Every organization has an organizational culture. That culture can and does change.
2. Those who focus on their own advancement at the expense of other considerations will, by default, advance further, faster and more often. Those who do not do this will not advance. Increasing amounts of focus make this effect increasingly large.
3. Focus on one's own advancement inside hierarchies causes individuals to self-modify in order to be the type of person who automatically engages in maze-creating and maze-supporting behaviors. They will also see such behavior as natural and virtuous.
4. Middle management performance is inherently difficult to assess. Maze behaviors systematically compound this problem. They strip away points of differentiation beyond loyalty to the maze and willingness to sacrifice one's self on its behalf, plus politics. Information and records are destroyed. Belief in the

possibility of differentiation in skill level, or of object-level value creation, is destroyed.

5. The more one is already within a maze, the more one is rewarded for maze-creating and maze-supporting behaviors, and for self-modifications towards such behaviors. This creates a vicious cycle.
6. Focus on one's own advancement causes one to wish to ally with others who do the same thing. That means allying with those who are engaging in maze-creating behaviors, and who are likely to do so in the future. Those people are likely to have future power. They are aligned in support of your new values and likely actions.
7. Changing the organizational culture towards a maze, which we will call *raising the maze level*, benefits those who wish to engage in maze-like behavior at the expense of those who do not. Those wishing to raise maze levels implicitly, and sometimes explicitly, coordinate together to reward maze behaviors, culture and values, and punish other behaviors, cultures and values.
8. Those who wish to have strong allies notice that strong potential allies want to ally with those who support mazes, and with those who ally with those who support mazes, and so on. This creates a strong incentive to strongly signal, in a way that others who support mazes can recognize, our support for maze behaviors and rising maze levels. One does this by supporting maze behaviors and maze allies whenever possible over all other considerations, without any need for explicit coordination or reciprocity, and by other costly signals of maze virtue.
9. The larger an organization and the more of a maze it becomes, the closer competition among its middle managers resembles super-[perfect competition](#) plus political considerations. [Slack](#) is destroyed. Those who refuse to get with the program, where the essence of the program *is support of mazes over non-mazes*, stop getting promoted or are pushed out entirely.
10. Such behaviors chose how to react to people largely by observing those people's culture and values. Those who wish to get ahead in such worlds must self-modify to instinctively support such actions, whether or not doing so is locally in their self interest. Being too aware of one's local self-interest is therefore not in one's broader self-interest. Humans are much better at doing all this, and at detecting it, than they are at faking it. The way one gets such behavior is through cultivating habits, including habits of thought, and choosing one's virtues. This self-modification creates even stronger implicit coordination.
11. There are contravening forces that can potentially outweigh all these effects, and result in maze behaviors being net punished. But they require those opposed to maze behaviors, culture and values to devote substantial resources to the cause, and to bear substantial costs. The more of a maze a place has already become, the harder it will be to turn things around or even stop things from getting worse. If such efforts are to succeed, this needs to be a high priority.
12. Even if maze behaviors are net punished for now, those who have embraced the maze nature will be skeptical of this. Even if they observe such behaviors being net punished now, they will not expect this to continue. Given the state of our world and culture, this is a highly reasonable prior. They also have knowledge of their own maze nature and presence within the organization, which is likely to raise maze levels over time and is evidence that the fight against mazedom is failing. And they stand to gain a huge competitive edge by raising the local maze level. Thus, the fight never ends.
13. Damage done is very difficult to reverse. Once particular maze behaviors become tolerated and levels rise, it takes a lot of effort to undo that.

14. Once people who support mazes are in places of authority in a given area, that area will rapidly become a maze. This is true of organizations, and of sub-organizations within an organization.
15. If the head of an organization believes in mazes, and has the time to choose and reward the people of their choice, it's all over. Probably permanently.
16. Mazes reward individuals who engage in maze behaviors and exhibit maze culture and values, and punish those who do not so exhibit, even outside the maze or organization in question. This includes customers, producers, business partners, investors and venture capitalists, board members, analysts, media, government officials, academics and anyone else who supports or opposes such patterns.
17. All strengthening of mazes anywhere creates additional force supporting mazes elsewhere. Mazes instinctively support other mazes. As society falls increasingly under the sway of mazes, it implicitly cooperates to push everyone and everything into supporting the behaviors, culture and values of mazes.
18. The end result inside any given organization is that maze behaviors grow stronger and more common over time. This is balanced by maze behaviors making the organization less effective, and thus more likely to fail.
19. Occasionally an organization can successfully lower its maze level and change its culture, but this is expensive and rare heroic behavior. Usually this requires a bold leader and getting rid of a lot of people, and the old organization is effectively replaced with a new one, even if the name does not change. A similar house cleaning happens more naturally in the other direction when and as maze levels rise.
20. Maze behaviors grow stronger and more common over time in any given organization barring rare heroic efforts. As organizations get bigger and last longer, maze levels increase.
21. When interacting with a world of low maze levels, or especially when interacting with individuals who have not embraced the maze nature, mazes are at a large competitive disadvantage versus non-mazes. Organizations with too-high maze levels become more likely to fail.
22. As organizations fail and are replaced by smaller upstarts via creative destruction, revolution or other replacement, maze levels decrease.
23. Replacement of old organizations by new ones is the primary way maze levels decline.
24. As the overall maze level rises, mazes gain a competitive advantage over non-mazes.
25. If society sufficiently rewards mazes and punishes non-mazes, non-mazes can stop failing less often than mazes. Existing organizations become increasingly propped up by corruption. New organizations will start off increasingly maze-like, signal their intent to become mazes, and raise their maze levels more rapidly. They will still usually start out at much lower maze levels than old organizations.
26. New organizations and smaller organizations also have more benefit in survival and growth from non-maze behaviors versus maze behaviors, as they have a greater need to do things mazes cannot do, or that they cannot do without huge additional overhead. Even in the scenario where large organizations benefit from maze coordination more than they are hurt by maze inefficiency, it can still benefit smaller organizations to minimize maze levels.
27. Mazes have reason to and do obscure that they are mazes, and to obscure the nature of mazes and maze behaviors. This allows them to avoid being attacked or shunned by those who retain enough conventional not-reversed values that they would recoil in horror from such behaviors if they understood them, and potentially fight back against mazes or to lower maze levels. The maze embracing individuals also take advantage of those who do not know of the

- maze nature. It is easy to see why the organizations described in *Moral Mazes* would prefer people not read the book *Moral Mazes*.
28. Simultaneously with pretending to the outside not to be mazes, those within them will claim if challenged that [everybody knows](#) they are mazes and how mazes work.
 29. As maze levels rise, mazes take control of more and more of an economy and people's lives.
 30. Under sufficiently strong pressure the maze behaviors, value and culture filter out into the broader society. Maze behaviors, values and culture are seen increasingly as legitimate and comfortable and praiseworthy. This happens even outside of any organization. Non-maze behaviors are increasingly seen as illegitimate, uncomfortable and blameworthy.
 31. The result of these effects is that people in societies with high maze levels, especially those with power and wealth, increasingly and increasingly openly oppose and vilify the creation of clarity, engaging in any productive object-level action, and participation in or even belief in the existence of positive sum games of any kind. [Simulacrum levels](#) continue to rise.
 32. Given sufficiently high societal maze levels, talk in support of maze behaviors would eventually becoming more and more open, and dominate discourse and how people are educated about the world, as people explicitly and publicly endorse and teach anti-virtues over virtues.
 33. We would see increasing societal inability to create clarity, engage in actions or do anything other than repeat existing patterns. Costs to all still possible actions would rise. Existing patterns would expropriate increasing portions of remaining resources to keep themselves afloat, and increasingly ban any activity outside those patterns.
 34. The default outcome on the scale of individual organizations is the rise and fall of those organizations over time.
 35. The default short term outcome on the scale of a nation, when maze levels and simulacrum levels increase, is declining growth, dynamism, slack, discourse, hope, happiness, virtue and wealth. People increasingly lose the things that matter in life.
 36. The default long term outcome on the scale of a nation is the rise and fall of civilizations.
 37. We do in fact see all of this. Here and now.

The next few posts will flesh this out more and provide, as best I can, answers to the other questions.

How Doomed are Large Organizations?

We now take the model from the previous post, and ask the questions over the next several posts. This first answer post asks these questions:

1. Are these dynamics the inevitable results of large organizations?
2. How can we forestall these dynamics within an organization?
3. To what extent should we avoid creating large organizations?
4. Has this dynamic ever been different in the past in other times and places?

These are the best answers I was able to come up with. Some of this is reiteration of previous observations and prescriptions. Some of it is new.

There are some bold claims in these answer posts, which I lack the space and time to defend in detail or provide citations for properly, with which I am confident many readers will disagree. I am fine with that. I do not intend to defend them further unless I see an opportunity in doing so.

I would love to be missing much better strategies for making organizations less doomed – if you have ideas *please please please* share them in the comments and/or elsewhere.

Are these dynamics the inevitable result of large organizations?

These dynamics are the *default* result of large organizations. There is continuous pressure over time pushing towards such outcomes.

The larger the organization, the longer it exists, and the more such outcomes have already happened, both there and elsewhere, the greater the pressure towards such outcomes.

Once such dynamics take hold, reversing them within an organization is extremely difficult.

Non-locally within a civilization, one can allow new organizations to periodically take the place of old ones to reset the damage.

Locally within a sufficiently large organization and over a sufficiently long time horizon, this makes these dynamics inevitable. The speed at which this occurs still varies greatly, and depends on choices made.

How can we forestall these dynamics within an organization?

These dynamics can be forestalled somewhat through a strong organizational culture that devotes substantial head space and resources to keeping the wrong people and behaviors out. This requires a leader who believes in this and in making it a top

priority. Usually this person is a founder. Losing the founder is often the trigger for a rapid ramp up in maze level.

Keeping maze levels in check means continuously sacrificing substantial head space, resources, ability to scale and short-term effectiveness to this cause. This holds both for the organization overall and the leader personally.

Head space is sacrificed three ways: You have less people, you devote some of those people to the maze-fighting process, and the process takes up space in everyone's head.

Central to this is to ruthlessly enforce an organizational culture with zero tolerance for maze behaviors.

Doing anything with an intent to deceive, or an intent to game your metrics at the expense of object level results, needs to be an automatic "you're fired."

Some amount of politics is a human universal, but it needs to be strongly discouraged. Similarly, some amount of putting in extra effort at crucial times is necessary, but strong patterns of guarding people's non-work lives from work, both in terms of time and other influences, are also strongly necessary.

Workers and managers need to have as much effective skin in the game as you can muster.

One must hire carefully, with a keen eye to the motivations and instincts of applicants, and a long period of teaching them the new cultural norms. This means at least growing *slowly*, so new people can be properly incorporated.

You also want a relatively flat hierarchy, to the extent possible.

There will always be bosses when crunch time comes. Someone is *always* in charge. Don't let anyone tell you different. But the less this is felt in ordinary interactions, and thus the more technically direct reports each boss can have and still be effective, and thus the less levels of hierarchy you need for a given number of people, the better off you'll be.

You can run things in these ways. I have seen it. It helps. A lot.

Another approach is to lower the outside maze level. Doing so by changing society at large is exceedingly hard. Doing so by associating with outside organizations with lower maze levels, and going into industries and problems with lower maze levels, seems more realistic. If you want to 'disrupt' an area that is suffering from maze dysfunction, it makes sense to bypass the existing systems entirely. Thus, move fast, break things.

One can think of all these tactics as taking the questions one uses to identify or predict a maze, and trying to engineer the answers you want. That is a fine intuitive place to start.

However, if [Goodhart's Law](#) alarm bells did not go off in your head when you read that last paragraph, you do not appreciate how dangerous Goodhart Traps are.

The Goodhart Trap

The fatal flaw is that no matter what you target when distributing rewards and punishments and cultural approval, it has to be *something*. If you spell it out, and a sufficiently large organization has little choice but to spell it out, you inevitably replace one type of [Goodharting](#) with another. One type of deception becomes another.

One universal is that in order to maintain a unique culture, you must filter for those that happily embrace that culture. That means you are now testing everyone constantly, no matter how explicit you avoid making this, on whether they happily embrace the company and its culture. People therefore *pretend to embrace the culture and pretend to be constantly happy*. Even if they *do embrace the culture and are happy*, they still additionally will put on a show of doing so.

If you punish deception you get people pretending not to deceive. If you punish pretending, you get people who *pretend to not be the type of people who would pretend*. People Goodhart on not appearing to Goodhart.

Which is a much more interesting level to play on, and usually far less destructive. If you do a good enough job picking your Goodhart targets, this beats the alternatives by a lot.

Still, you eventually end up in a version of the same place. Deception is deception. Pretending is pretending. Fraud is fraud. The soul still dies. Simulacrum levels still slowly rise.

Either you strongly enforce a culture, and slowly get that result, or you don't. If you don't and are big enough, you quickly get a maze. If you do and/or are smaller, depending on your skill level and dedication to the task, you *slowly* get a maze.

Hiring well is better than enforcing or training later, since once people are in they can then be themselves. Also because enforcement of culture is, as pointed out above, toxic even if you mean to enforce a non-toxic ideal. But relying on employee selection puts a huge premium on not making hiring mistakes. Even one bad hire in the wrong place can be fatal. Especially if they then are in a position to bring others with them. You need to defend your hiring process especially strongly from these same corruptions.

My guess is that once an organization grows beyond about Dunbar's number, defending your culture becomes a losing battle even under the best of circumstances. Enforcing the culture will fail outright in the medium term, unless the culture outside the organization is supporting you.

If you are too big, every known strategy is only a holding action. There is no permanent solution.

To what extent should we avoid creating large organizations?

Quite a lot. These effects are a really big deal. Organizations get less effective, more toxic and corrupt as places to work and interact with, and add more toxicity and corruption to society.

Every level of hierarchy enhances this effect. The first five, dramatically so. Think hard before being or having a boss. Think harder before letting someone's boss report to a

boss. Think even harder than that before adding a fourth or fifth level of hierarchy.

That does not mean such things can be fully avoided. The advantages of large organizations with many degrees of hierarchy are also a really big deal. We cannot avoid them entirely.

We must treat creating additional managerial levels as having *very high costs*. This is not an action to be taken lightly. Wherever possible, create distinct organizations and allow them to interact. Even better, allow people to interact as individuals.

This adds friction and transaction costs. It makes many forms of coordination harder. Sometimes it simply cannot be done if you want to do the thing you'd like to do.

This is increasingly the case, largely as a result of enemy action. Some of this is technology and our problems being legitimately more complex. Most of it is regulatory frameworks and maze-supporting social norms that require massive costs, including massive fixed costs, be paid as part of doing anything at all. This is a key way mazes expropriate resources and reward other mazes while punishing non-mazes.

I often observe people who are stuck working in mazes who would much prefer to be self-employed or to exit their current job or location, but who are unable to do so because the legal deck is increasingly stacked against that.

Even if the work itself is permitted, health insurance issues alone force many into working for the man.

When one has a successful small organization, the natural instinct is to scale it up and become a larger organization.

Resist this urge whenever possible. There is nothing wrong with being good at what you do at the scale you are good at doing it. Set an example others can emulate. Let others do other things, be other places. Any profits from that enterprise can be returned to investors and/or paid to employees, and used to live life or create or invest in other projects, or to help others.

One need not point to explicit quantified dangers to do this. Arguments that one cannot legitimately choose to 'leave money on the table' or otherwise not maximize, are maximalist arguments for some utility function that does not properly capture human value and is subject to [Goodhart's Law](#), and against the legitimacy of [slack](#).

The fear that if you don't grow, you'll get 'beaten' by those that do, [as in Raymond's kingdoms](#)? Overblown. Also asking the wrong question. So what if someone else is bigger or more superficially successful? So what if you do not build a giant thing that lasts? Everything ends. That is not, by default, what matters. A larger company is often not better than several smaller companies. A larger club is often not better than several smaller clubs. A larger state is often not better or longer lasting than several smaller ones. Have something good and positive, for as long as it is viable and makes sense, rather than transforming into something likely to be bad.

People like to build empires. Those with power usually want more power. That does not make more power a good idea. It is only a good idea where it is instrumentally useful.

In some places, competition really is winner-take-all and/or regulations and conditions too heavily favor the large over the small. One must grow to survive. Once again, we

should be suspicious that this dynamic has been engineered rather than being inherent in the underlying problem space.

Especially in those cases, this leads back to the question of how we can grow larger and keep these dynamics in check.

Has this dynamic ever been different in the past in other places and times?

These dynamics *seem* to me to be getting increasingly worse, which implies they have been better in the past.

Recent developments indicate an increasing simulacrum level, an increasing reluctance to allow older institutions to be replaced by newer ones, and an increasing reliance on cronyism and corruption that props up failure, allowing mazes to survive past when they are no longer able to fulfill their original functions.

Those in the political and academic systems, on all sides, increasingly openly advocate against the very concept of objective truth, or that people should tell it, or are blameworthy for not doing so. Our president's supporters *admit and admire* that he is a corrupt liar, claiming that his honesty about his corruption and lying, and his admiration for others who are corrupt, who lie and who bully, is refreshing, because they are distinct from the corrupt, the liars and the bullies who are more locally relevant to their lives. Discourse is increasingly fraught and difficult. When someone wants to engage in discourse, I frequently now observe them spending much of their time pointing out how difficult it is to engage in discourse (and I am not claiming myself as an exception here), as opposed to what such people used to do instead, which was engage in discourse.

We are increasingly paralyzed and unable to do things across a wide variety of potential human activities.

Expropriation by existing mazes and systems eats increasing shares of everything, especially in education, health care and housing.

I don't have time for a full takedown here, but: Claims to the contrary, such as those recently made by Alex Tabbrok in *Why Are The Prices So Damn High?*, are statistical artifacts that defy the evidence of one's eyes. They are the product of Moloch's Army. When I have insurance and am asked with no warning to pay \$850 for literally five minutes of a doctor's time, after being kept waiting for an hour (and everyone I ask about this says just refuse to pay it)? When sending my child to a kindergarten costs the majority of a skilled educator's salary? When you look at rents?

Don't tell me the problem is labor costs due to increasing demand for real services.

Just. Don't.

Some technological innovations remain permitted for now, and many of the organizations exploiting this are relatively new and reliant on object-level work, and thus less maze-like for now, but this is sufficiently narrow that we call the result "the tech industry." We see rapid progress in the few places where innovation and actual work is permitted to those without mazes and connections, and where there is sufficient motivation for work, either intrinsic or monetary.

The tech industry also exhibits some very maze-like behaviors of its own, but it takes a different form. I am unlikely to be the best person to tackle those details, as others have better direct experience, and I will not attempt to tackle them here and now.

We see very little everywhere else. Increasingly we live in an amalgamated giant maze, and the maze is paralyzing us and taking away our ability to think or talk while robbing us blind. Mazes are increasingly in direct position to censor, deplatform or punish us, even if we do not work for them.

The idea of positive-sum, object-level interactions being someone's primary source of income is increasingly seen as *illegitimate*, and risky and irresponsible, in contrast to working for a maze. People instinctively think there's something shady or rebellious about that entire enterprise of having an actual enterprise. A proper person seeks rent, plays the game starting in childhood, sends the right signals and finds ways to game the system. They increase their appeal to mazes by making themselves as dependent on them and their income and legitimacy streams, and as vulnerable to their [blackmail](#), as possible.

The best way to see that positive-sum games are a thing is to notice that *the sum changes*. If everything is zero-sum, the sum would always be zero.

The best way to see that these dynamics used to be much less severe, at least in many times and places, is that those times and places looked and felt different, and got us here without collapsing. Moral Mazes was written before I was born, but the spread of these dynamics is clear as day within my lifetime, and yours as well.

Did some times and places, including our recent past, have it less bad than us in these ways? I see this as almost certainly true, but I am uncertain of the magnitude of this effect due to not having good enough models of the past.

Did some times and places have it worse than we do now? Very possible. But they're not around anymore. Which is how it works.

The next section will ask why it was different in the past, what the causes are in general, and whether we can duplicate past conditions in good ways.

Ten Causes of Mazedom

We continue answering the questions we asked earlier. It was claimed last time that maze levels and the danger of mazes was lower in the past than it is now, and that overall maze levels have been rising, as measured both by maze levels within organizations across the board, and maze levels within the overall society.

The world is better for people than it was back then. There are many things that have improved. This is not one of them.

I am confident this is the case and have pointed to ways to see it. I recognize that I have in no way *proven* this is the case. I don't have a way to do that. Rather I am relying on your own observations and analysis. If you disagree, I hope the discussion that follows will still prove useful as a comparison to what you see as an alternate possible scenario where these dynamics are less toxic.

Now we ask what may have been different in the past, and whether we can duplicate those causes.

Why was it different? Can we duplicate those causes?

The sketched model suggests several causes.

One can model this as a two-level problem. Something happens (e.g. technological change) resulting in a change in circumstances (e.g. more real need for large organizations, cause one), which then causes higher overall maze levels.

Since this question asks for potential action on a large scale, political proposals will need to be among the proposals discussed. Let us do our best to debate models and how the gears in them work, and stay away from actual politics. To that end, *I am not claiming any of these policies are good ideas and should be implemented, or bad and should not be implemented.* I am only claiming that they would have particular physical results in the world. If they are obviously good or bad ideas, I shouldn't need to say so.

Cause 1: More Real Need For Large Organizations

Modern life has larger organizations with more levels of hierarchy. Corporations are bigger. Governments are bigger. Political organizations are bigger. Universities are bigger. They wrap that bigness in more levels of hierarchy than ever before, with increasing amounts of bureaucracy and administration.

Jobs are done that previously did not get done. Where the same job persists, that same job takes more people.

There is little doubt that *some* of the problem is caused by the increased complexity of real needs. What is difficult is figuring out how much is necessary versus how much is unnecessary.

Do we see these bigger organizations mostly because technological and economic development requires it? Or do we mostly see bigger organizations because maze-supporting behaviors and rules have made it so?

How much of what goes on in a large organization is bullshit, and could be eliminated, shrinking its size?

How much of needing a large organization is due to regulation and subsidy? How much of this additional size and regulation is due to additional demands for safety, either real safety or the illusion of safety?

How much is our unwillingness to let 'too big to fail' organizations fail, and the perception that size is necessary, good and/or prestigious?

To the extent that this is the problem, we are *mostly* stuck. We can still make different trade-offs, since we can give up some of the additional complexity and its benefits to reduce the large-organization costs paid, but there are no easy wins to be had here.

My low confidence belief is that this accounts for a substantial portion of the problem, but less than half. It would be foolish precision to give a more exact number.

Cause 2: Laws and Regulations Favor Large Organizations

There is genuine need for larger organizations with more levels of hierarchy. But once you create those organizations, they have a logic of their own that generates further need for them. Those within these organizations steer conditions in their direction. A lot of the bigger organizations are there because they are structurally advantaged because they enjoy regulatory capture, or because of other barriers to smaller competition, rather than because being larger make any economic sense.

While we were making life harder for the small, we made life easier for the large. We removed technological and legal costs and barriers to larger organizations of all kinds. The friction caused by those costs previously ensured that big organizations only existed when there was a big payoff. Now we can also get them because of mission creep, because it gives people making the decisions more power, or often because someone's ego gets stroked by the idea.

We were so busy figuring out how we could go big, we didn't stop to ask whether we should.

As opposed to doing new things, where we are now so busy asking whether we should we no longer stop to ask whether we could.

Cause 3: Less Disruption of Existing Organizations

Wars, natural disasters, economic collapses and other large negative shocks are the standard historical method of preventing large organizations from hanging around for too long.

Fortunately, we've managed to mostly avoid major destructive wars for some time. Natural disasters are not the threat they once were. While we had an economic collapse of sorts in 2008, it was nothing like the Great Depression. The title of one of the definitive books on that crisis, and the one I read to better understand it, was literally [*Too Big to Fail*](#).

Existing large organizations were protected by government intervention. When they were so damaged that they could not be defended, instead of letting them fail, *they were forcibly merged into even larger organizations*.

With our much larger productivity and surplus, our greater ability to do finance and run large deficits for long periods, and greater ability to absorb negative shocks, and our growing distaste for risk and loss of any kind, along with a long peace, and our judgment of leaders on extreme short term perception of outcomes, we have biased things greatly against disruption.

That then combines with the effects of regulatory capture and crony capitalism. The rise of mazes and maze behaviors is a vicious cycle.

We've also seen a dearth of *positive* shocks. Technological breakthroughs, economic developments and other big positive events allow new entrants to gain advantage over existing entrenched interests whose cached actions are increasingly inefficient and no longer make sense.

It is unknown to what extent this dearth of positive shocks is due to inherent difficulty in pushing forward at the current technological frontier, to what extent it is a cost of choices we have made, and to what extent *it was the intent* of the choices that were made.

In many cases, here and now, progress and disruption that would challenge existing systems is being forcibly suppressed. Some of the mechanisms of this will be among this list of causes, again with a mix between inherent, unintentional and intentional. Incumbents of all kinds by default are fragile and would prefer the status quo not be disrupted. Being mazes only enhances these incentives. In other cases, we've run into some combination of increasing difficulty of underlying problems, and our increasing lack of ability to do things even when not inhibited.

Cause 4: Increased Demand for Illusion of Safety and Security

Safety and security are important. They are not infinitely important. They are *extremely* poor choices for sacred values.

Illusion of safety and illusion of security are increasingly seen as important sacred values. Occasionally so are actual safety and security.

In important ways, we are less secure than in the past.

Due to cultural changes, atomization and government interventions, we cannot count on our family, our friends or our community to support us. Even when we *can* count on them to support us, what we are counting on is a different order of magnitude of support from far fewer people. Government safety nets, even at their best, are a poor replacement.

We are taught that fulfilling all the requirements to responsibly have children and a family requires many stars to simultaneously align. We are put on very tight deadlines. There is huge risk that these things will not happen for us at all. They are the things that many people value the most today, and which most people valued most in the past. One is not permitted to 'skip ahead' as one is often kept in educational limbo and then a post-education period where you are not expected to be able to seriously attempt a family, until the clock has almost run out, then the scramble begins.

We also in important ways have longer memories and less ability to recover from mistakes or poor trajectories. The very idea of things going on one's 'permanent record' puts one in continuous danger even as a young child. Path dependence goes

up. We have this concept that one's 'life can be over' as the result of quite minor misfortunes. Potentially this can be as small as an ill-advised tweet, or a poorly navigated social interaction.

Then we are expected to end this journey with millions of dollars in 'life savings' as a substitute for the old support systems. That money can still be expropriated at any time by the medical or elder care systems.

Rates of depression, suicide, traumatic stress and other mental health issues have skyrocketed, likely largely as a result of all of these other dangers and pressures, also largely because we are taught to find lesser and lesser issues unacceptable and crippling. Increasing numbers of people generally feel unprepared for problems that would not have much fazed them in the past, or in some cases would not have even qualified as problems. We can't even.

In important ways, we are *more physically safe* than in the past.

Our trauma care and treatment of infections has dramatically improved. The rates of physical trauma, street crime and infection have gone down and the consequences tend to be less severe. Murder rates are down. Cars are vastly safer than in the past. Planes are much less risky than the car ride to the airport. Whenever I hear about accidents from travel these days, it's *almost universally* in fiction, or what would have in the past been a minor accident that now makes the news. Our food supply is vastly safer. Our toys and playgrounds are vastly safer.

Our physical products and physical spaces and lives are universally are vastly safer.

Life *really shouldn't* feel physically scary the same way that it used to. The rising physical risks are suicide and drug abuse, which are a different kind of physical risk, and *the consequences of too much safety* such as allergies.

In important ways, we are also *more emotionally safe* than in the past, as well. In other important ways, we are far *more emotionally at risk*.

We are safer in the sense that things that many unpleasant things that were considered normal parts of life are now considered unacceptable, and thus we guard against them. We still must deal with violence, with racism, with sexism, with abuse, with religious intolerance, with homophobia and transphobia, and hosts of other similar problems. Insert all additional concerns here. But without in any way saying that any of our current concerns are illegitimate, *oh my* are all those problems in a vastly better place than they used to be. As Jim Norton commented in his most recent stand-up special, many of the stories our grandparents tell of how they met would, by modern ears, horrify us.

Such questions are treated as much more of a priority than they used to be, especially in educational settings. We also make aggressive use of psychiatric medication.

We are less safe from such things in two ways.

We are less safe in the sense that we no longer are taught that such things are acceptable, or that we should be able to 'walk it off' when they happen or take it on the chin. Nor are we forced to practice dealing with them. Thus have a much harder time dealing with them when they happen. We have the problems that come with a revolution of rising expectations.

We are also less safe in the sense that mazes turn everything into a signal and tool in the maze's political games. Even when one could shrug off an event and its direct consequences, that event also comprises a move within the game. If people observe that offense can be done to you without retaliation, that is a license to everyone to ramp up the level of offense. This transforms your social identity into one to whom offense can safely be done, thus destroying your status and any hopes for success, including assistance with high placement into another maze.

A lot of this safety comes at the direct expense of fun, learning and efficiency. And a lot of it makes us so safe that people are failing to encounter the adversity they need to become resilient.

This all seems quite bad to me on its own merits, but one could disagree. It largely reflects the *very good news* that life is safer and thus we are free to demand it be safer still. This being a cause of increased maze activity doesn't depend on it being otherwise good or bad.

This rise in risk of failure tracks the physical and social risks of 'falling off the track' described above. Emotionally falling off the track has also become far more dangerous.

The risk of things going wrong has gone down at many steps. So has our tolerance for those things going wrong, and our tolerance for risk of things going wrong. We have gone from resilient to fragile. Meanwhile, we've lost traditional supports, and many circumstances and life transitions have become far more perilous than they have ever been.

This has bled into even tiny decisions, such as what to have for lunch in a strange town.

One of the highest value practical lessons of rationality is to recognize what is actually safe and unsafe and to what degree, and when it makes sense to take an unconventional path.

How do these phenomena raise maze levels?

Mazes are, among other things, engines for producing the illusion of safety and the illusion of security. In some ways and cases, they do this by producing the real thing. In other cases, it is pure illusion. Either way, they have a huge comparative advantage.

Working for large organizations is seen as safe and secure. So is working as part of a profession whose rents are protected by regulation and other barriers to entry. Working elsewhere is seen as risky and insecure.

A product made by a maze is seen as reliable and consistent, even if low quality or a bad value.

Existing risks are 'grandfathered in.' Once an activity, product, drug or what have you is approved and becomes standard, that becomes subject to a different frame, where taking that risk isn't considered risk, and it doesn't open you up to being blamed if something goes wrong. If no one gets fired for buying IBM, as the saying goes, it's going to take a damn good reason to buy something else.

Inside a maze, that is a huge support for doing business with and helping other mazes over other non-mazes. Outside a maze, people are still largely seeking to avoid being blamed and/or blameworthy.

When your options are bad, it is even more important to show that you have made the conventional, standard and therefore 'safe' choice, so that when it blows up in your face, it is not your fault. For a current real example in my own life, I will be using a health insurance broker to navigate the dystopian wasteland that is the individual health insurance market in New York, and placing a high priority on a policy that provides illusion of security.

Even if all is well and safe and good, and that is not in dispute, when something is new, or isn't the generic standard thing, the default consumer does not know this. Why go to the trouble of finding out, [when anything that goes wrong is then your fault](#), even if it had nothing to do with your choice, or your choice actively made it less likely?

New risks, and any downside of a new proposal to basically anyone, are taken vastly out of proportion. Mazes are much better equipped to handle this than non-mazes, as this forces one to play the system and the games of propaganda and public relations, hiding what needs to be hidden, and spending huge amounts of time, money and expertise on the right to do anything. Many of the costs involved are fixed. Small scale efforts often don't survive this.

Safety regulations and other requirements for doing things are almost always conceptualized based on what currently exists. Lobbyists made sure that existing systems could handle them, and protested loudly if that would be expensive. If you're thinking of a new way of doing things, even if it makes the existing system look vastly unsafe in some way, it is likely to violate one of the existing technical requirements, or make some aspect slightly less safe in some fashion.

Real shame, that.

I leave as an open question *how often the purpose of the safety regulations is the prevention of innovation and competition.*

Cause 5: Rent Seeking is More Widespread and Seen as Legitimate

This is a result as well as a cause.

Mazes have comparative advantage at seeking, securing and extracting rents. Rents are what keep the mazes running once they are no longer capable of efficient creation of value.

Thus, part of their reversal of morality is to turn the destructive act of rent seeking into the legitimate and respectable thing that one does to earn a living, and make positive-sum actions seem increasingly illegitimate.

That major forces in our society are pushing this, and that it has been working, have been addressed earlier in the series. To the extent that they haven't been, it's another post or even sequence, and I won't argue further for it here.

Cause 6: Big Data, Machine Learning and Internet Economics

Big data and machine learning swing the balance in favor of mazes.

Mazes are full of people who could not communicate with each other (extreme case of [the SNAFU principle](#)), who are many steps removed from the object-level considerations they neither understand nor have the data or skill to process, and make decisions based on metrics and political considerations largely around those metrics. And not caring about incidental effects.

Machine learning gives a powerful tool to anyone with a metric to optimize and lots of data that can be gathered. The explicit price of this is giving up on a gears-level understanding of what is going on or how things work, not caring too much about incidental effects, and trusting the algorithm. Thus, machine learning helps mazes be much better at what they wanted to do anyway, while charging them prices they've already paid.

Others are then judged by those mazes and metrics and pushed towards using more machine learning style strategies. If they give in to this, outsiders lose their advantages of actually understanding how things work and what is important and seeing incidental effects, and are at a severe data disadvantage because the only data that counts is no longer the type of data they have.

Mazes also benefit because big data and machine learning love huge amounts of data. Huge amounts of data demand scale. Instead of having humans in charge of decisions, now algorithms are in charge of those decisions. The blindness of those inside the maze becomes a blindness from lack of data for those who don't have sufficient scale.

Internet economics is then shaped by this machine learning and by optimizing for metrics, especially when bidding for advertising and trying to reach the point where one can use advertising to scale. A lot of costs are fixed, which favors large over small. Other aspects do favor small over large, and favor better over worse. So it's not a complete loss. And there's a lot of pushback.

I am not claiming this is one directional or simple.

In addition to the issue of what algorithms do to mazes that use them, there is the issue of what algorithms do to the people the algorithms are reacting to, when they want to change and/or are shaped by how the algorithms react. When algorithms are optimizing how to react to us, we face similar underappreciated pressures even if we are not using them, which also contributes to all this. Our entire mode of thinking about the world becomes warped. That's a much bigger topic beyond scope.

Cause 7: Ignorance

Mazes are an example of a thing that, unless they have already sufficiently corrupted values in general, could be destroyed by the truth.

Mazes depend upon the public, and in particular on potential workers and associates, not understanding how mazes work or what working in or with a maze will do to their lives and souls. People do understand that there is something soul-crushing about 'working for the man' and that is true as well, but mostly misses the point.

It takes a minimum of years inside a maze to start to understand how mazes work from the inside, or how bad it gets in middle management, if no one explicitly tells you. By that time, many are too invested to stop, or have already self-modified in ways that leave them no longer able to perceive the problem. Even if you have the

data needed to put it together, my experience is that to fully appreciate what you are dealing with, *you have to quit first*. Which would mean you need to quit before you realize how badly you need to quit. Whoops.

If someone does explicitly tell you beforehand, you're inclined not to believe it, because such warnings usually lack the gears and sound absurd, and come from people who appear biased against business or are experiencing sour grapes. Because things would sound so absurd, those doing the warning will round down their warnings to sound reasonable. Then the listener does the same. Even when the source is credible and provides gears, as I hope this source does when combined with the original book, the brain constantly instinctively rounds it all down to avoid seeing the thing. Those who recoil in horror at the thought of working for a maze, who are not doing so based on hard won experience, *usually don't know the half of it*.

This is enough to ensure that workers at large corporations *make more money* than those at smaller ones. Tyler Cowen and other defenders of large corporations use this as evidence that large corporations make people more productive. I think this represents a terribly incomplete model of how labor markets work. Workers are not primarily paid in some relation to their marginal product. The accounting identity neither holds nor does it enlighten much.

My model is more along the lines of: Workers are paid in relation to what it costs to attract talent to a position, and then the business either is or is not viable. One of the inputs to that is indeed the worker's perception of their marginal or average production at that job. There is some expectation that great success will be shared with the workers to some extent. Not to do so would make them upset. This is mostly a threshold effect and acts largely on marginal compensation via reciprocity and fairness. This isn't centrally about supply and demand curves meeting. In good times one gets a raise.

The exception is when there is a strong union. A strong union changes the calculation from what workers will accept to work here instead of elsewhere, to a zero-sum battle by the union to extract maximal resources from the corporation without too quickly driving it into bankruptcy, versus the corporation trying to retain profits and not let them do that. This conflict plays out in a legal structure where many otherwise legal things become illegal, other otherwise illegal things become legal or mandatory, and the outcomes are weird. Union jobs really are different.

Another input is the status that goes along with a job, and what it feels like it 'should' pay on that basis. In most cases this runs in the perverse direction – a 'better' job requires *more* compensation to keep people happy about their pay. Workers have the instinctual sense that higher status jobs should strictly dominate low status ones. They also have the (usually correct) instinctual sense that if the job does not pay, the promised status is fake slash won't be accepted. And that if lower status jobs are not paid less, or equal status people who produce better are paid more, then their status is under threat, so they will revolt against such practices. Thus, everyone wants the higher status jobs. You cannot effectively balance this with pay.

It really is a terribly bad idea to let the workers find out what the other workers are paid.

A lot of this, one hopes, is based on how much someone wants to do that job, and have the life that goes along with having that job.

When we see that jobs in big corporations pay more, rather than conclude they must be more productive, I primarily take this as evidence of some combination of two things.

One, that *working at a big corporation is something people dislike, so you have to pay them more to do it*. This despite them not knowing how much they should dislike it, and therefore demand in compensation, and despite a lot of career lock-in. We also see above-standard pay in large government bureaucracies.

That can be combine with two, that *working at a big corporation is seen as higher status, and therefore demands higher pay*.

We often see poor pay in academia, because people enter thinking they'll get to think about things all day that they find interesting, and have that position securely for life, they like that lifestyle, and that makes them want the job more. They can then accept lower pay because they (incorrectly) see what they are doing as not being part of a profitable operation, and (correctly) see others think this, so they do not need higher pay to hold onto the job's high status.

Prospective academics are almost always, as far as I can tell, at least in willful denial about their job prospects, and about how their jobs will work if they find them. They certainly don't process the news about the incentives and pressures they will face, and the petty battles and tasks that lie ahead.

Big business still has to write those larger checks every two weeks while staying in business. Isn't that evidence of higher productivity? It's some evidence of otherwise higher *profitability* of the corporation, that this does not drive them out of business. Profitability is not productivity. If we think those higher profits are largely due to extractive practices, rents and regulatory advantages rather than production, that does not count. If we think that it is due to things inherent to the corporation, such as intellectual property or monopoly or returns to scale, and this in turn causes employees to demand more compensation, that also does not count if it's not the employees (or not all but a handful of employees) doing the production.

If people knew the truth about working for mazes and especially being in middle management, and/or were facing less outside pressure, they would demand much higher pay. This would make mazes bear more of the real costs they impose on people, while making it harder to attract strong talent. It would be much easier to justify declining such jobs, and harder to push someone into accepting one. Doing other things would get relatively higher status.

This would all make mazes less competitive.

So would general knowledge of how mazes derive their advantages and rents from regulation and subsidy. Regulatory capture is rife. Much of what opponents of business and other long entrenched mazes do ends up being used by those entrenched mazes to cement their advantages. A common pattern is outrage at the actions of corrupt actors, creating demand for new rules and regulations, which turn out primarily to entrench and enrich those same corrupt actors.

If people understood these dynamics, and knew what mazes were, they would use different tools. They would demand different interventions. Then there are the things that business convinces the government to do directly. If those were better understood, and people could coordinate to stop them, that would be an even bigger boon.

One could object that this isn't a difference with past conditions. Did people in the past know these things?

Ignorance often comes from not knowing the alternatives to the maze way of doing things, or that there exist or ever have existed things other than mazes, or of the ways in which mazes take over non-mazes and how to prevent that. People object that things are terrible, and they respond that since the incentives naturally trend towards these terrible things, how could things ever not be terrible? Despite things sometimes not being terrible, and things not being terrible in general. Those in charge of non-maze spaces invite the maze-runners in willingly, without knowing they have sealed the space's fate.

This was not true in the sufficiently distant past. It is not even true for *your* sufficiently distant past. It is *learned behavior*. We do not begin life this way. Children begin only knowing the non-maze way of doing things. Then we send them to school. Then we give them jobs.

The hypothesis here is that things were different in the past because people were regularly exposed to non-mazes as central actors in their worlds. This taught them that small maze levels were the normal way of being, instead of being mostly exposed to mazes as central actors, building their lives around them, and being taught that this was the normal way of being. People used to realize that the things involved in this were *weird* and now they aren't weird anymore. This leaves us blind to what mazes result in. People used to expect better.

Ignorance often also takes the form of *not knowing how to do actual business*. Most people today have no idea how to go about working on their own, feeling mostly that There Be Dragons. Or generally how they or their worlds might survive without giant engines taking care of things for them on all levels. That did not used to be the case.

Cause 8: Atomization and the Delegitimization of Human Social Needs

A large cost of working in mazes is that they do not permit full independent outside social lives. People are told to isolate themselves and neglect basic human needs. Our society's willingness to treat such a condition as normal, acceptable and appropriate means that they can pull this off.

In the past, in most times and places, I believe that people would have realized how terrible all of this increasingly has been and prioritized fixing it. This would have made people realize that joining an all-consuming maze was a huge life cost not to be taken lightly, rather than a source of palliatives to help one suffer through an impoverished social life.

For those in middle management or above, they are being evaluated on everything they do. This includes who they choose to associate with, including who they marry and how they raise their kids, how they spend their time outside of work, what church if any they go to, and so forth. They are also under increasing pressure to spend every possible waking moment on job-related things. Socializing with coworkers mostly counts.

Even for those not in middle management, jobs play havoc with their schedules and flexibility. Almost total reliability is demanded at most jobs in large organizations during standard hours, and most require also working outside of those hours when asked to do so. Any exceptions are quite costly. Vacation time is both slim and usually

planned well in advance. This is hard on someone trying to live life, especially with a family. You end up tired all the time. Your life becomes a grind. Any free time largely needs to be spent recovering.

The majority of your waking hours are spent interacting with your coworkers, who you mostly did not choose. They then effectively become your friends, and often even your family, the same way they do in most workplace television shows. They are the people who have things in common with you, who you can reasonably keep up with and make plans with. Often they are the people you date, despite all attempts at policy to the contrary, and whether or not there are bad dynamics involved.

On a broader scale, you do not only keep the hours they tell you to keep. You live where they tell you to live, so you can work where they tell you to work. In the ongoing battle to create a rationalist community in New York, there is the constant assumption by everyone that finding a good job somewhere determines where you live. One of my three closest remaining friends, a former Magic: The Gathering professional, recently uprooted his family, including three kids, and sold his apartment, because job offers in another city were somewhat better than those here. His entire family knows zero people in the new city. That is now simply what one does. People who have skills or ambition are *expected* to walk away from everyone they have ever known for a job.

This happens because basic human social needs have been delegitimized. Work provides one with one's 'social life' and 'friends' and even 'family,' and we are expected to count that as a benefit. I believe that in most times and places, this would rightfully be seen as a deeply impoverished bastardization of these concepts. People think their Facebook 'friends' are friends, or even close friends. They're not.

Atomization has many other causes as well. I have wide uncertainty how much higher maze levels *caused* atomization, as opposed to atomization causing higher maze levels. The delegitimization of basic human needs seems more tied to maze issues.

Given all that, how does one invest in real and deep long term friendships and groups, when they are increasingly unlikely to last?

If you have a better answer than 'persevere and hope for the best' *please* share. I give great thanks for the few old friends that are still around.

Cause 9: Educational System

We indoctrinate our citizens to the horrors of mazes early and often.

Our first step is to imprison our children, starting from around the age of five, for the bulk of their waking hours. During that time, all of their behaviors are tightly controlled, and they are taught that their success in life depends on satisfying the arbitrary demands of the person put in charge of them, who will mostly use negative selection to determine outcomes. When asked to justify doing this, the reply is typically that if you do not do this to your child now, your child will be ill-prepared to be subjected to increasingly intense versions of it as they progress through the educational system.

You had better suffer and kill your soul now, or else you won't be prepared to suffer worse things later.

Once taught that life is about obeying arbitrary dictates and doing work with no object level application whatsoever, giving most of your life over to arbitrary schedules and demands, and gathering together credentials and approvals necessary to get the labels that get others to give you status and compensation, where everyone is on the lookout for reasons to put black marks on your record, you are ready to work in a maze without recoiling in horror.

This process also saddles its students with massive debts that force them to then take jobs they hate, and gives them credentials that give them an advantage in narrow rent extraction in one area, preventing choice or exit.

I previously wrote at somewhat more length on this in these posts: [The Case Against Education](#), [The Case Against Education: Foundations](#), [The Case Against Education: Splitting the Education Premium Pie and Considering IQ](#), [Something Was Wrong](#)

Cause 10: Vicious Cycle

Mazes are self-reinforcing.

Many of these causes make this clear. Mazes cause a problem and shift in priorities and values, which in turn makes mazes more attractive and devalues non-mazes. It makes people rely more and more on mazes for their needs. More mazes gives maze activities more interactions with other maze activities and those caught in mazes, legitimizes maze activities and delegitimizes non-maze activities, causing a vicious cycle.

Once people are sufficiently dependent on mazes, or mazes gain sufficient influence and power, there is a critical phase shift. Morality's sign flips. The forces of social approval, which previously were pushing against maze behaviors, instead start pushing for them. People punish others for not orienting towards mazes and away from the object-level, where before they did the opposite and some people chose mazes in spite of this.

That phase shift seems to have happened in the bulk of general society, including our school system and our political system, and increasingly in social life lived via social media, causing an accelerating problem.

Reputational systems in general have largely become maze-oriented reputation systems. School is about 'what looks good on an application.' Work is about 'what looks good on a resume' when it isn't about the boss. Romantic life is even based largely on 'what looks good on a dating profile' or in the first moments of interaction, favoring the legible over the illegible more than previous systems.

The more one must likely deal with mazes in the future, and the more they determine our futures, the more 'guard one's reputation' means 'make sure the searchable record does not contain things mazes will dislike, and let one tell to mazes the story one wants to tell.'

These new concerns crowd out what one would otherwise be concerned about. One key thing they crowd out is wanting to be known in good ways for who you actually are and what you do, by people who observe you, who will then tell other people. Mazes can and often do start to destroy your connection to object level reality long before you ever set foot in one.

Even when the reputation system is based on explicitly being judged by those who interact with you, such as customer reviews, it's about an average quantification. This largely becomes another metric game. Such systems do excellent work detecting when things are wrong, but mostly operate via negative selection due to clustering of ratings near the top of the range allowed. Thus they likely will become increasingly distortionary with time towards prioritizing avoiding the things that legibly cause bad ratings, effectively punishing anything non-standard or any attempt to be exceptional, because when you earn that sixth star, you're not allowed to collect it.

A costly way to go above the non-costly maximum would potentially help, but has its own problems, and further discussion is fascinating to me but beyond the scope here. It is a hard problem.

More Mazes, More Problems

Next post I will present the best potential solutions I could come up with. This is not a problem that can be fully solved, but certainly there are changes on the margin that would decrease the rate at which things get worse, or even roll the damage back somewhat.

Implementing them, of course, is a whole additional problem.

Potential Ways to Fight Mazes

[The previous post](#) listed some of the causes of higher than normal maze present in today's society.

This post asks what might be done about it, both as an individual and via group coordination.

I believe it is better to *first* ask “what are the things that would help?” even if one does not know how to implement them or convince others to get them implemented or even if they are good ideas worth implementing. *Then* figure out if any version of the idea is worth implementing. *Then* ask how to put the ideas into practice.

All of this assumes that lowering maze levels is both desirable and important. This post takes that assumption as its premise. My case for it has already been made previously.

Some items on this list are political interventions. They require solving the collective action and [public choice](#) problems of writing and passing good versions of the appropriate laws and regulations.

Some of these items are cultural shifts that happen one person at a time. This is a different kind of collective action problem. Word would be spread, potentially convincing others to support the cause.

In both cases, pointing out that [coordination is hard](#), and that game theory says any given person by default shouldn't cooperate or devote any effort, are important things to remember. They are not full justifications for dismissing potential solutions.

Finally, there is one last solution that asks what one determined person or organization with extensive resources might do, and comes back with the answer of quite a bit. That will be the next post.

As a reminder, I am not, except where explicitly saying so, endorsing the political proposals as good ideas or rejecting them as bad ones. Nor am I interested in having political policy debates, or any other political debates, on the internet – any comments along these lines will be ignored at best and deleted at worst. I am only saying these actions might help with the problem of mazes, without being *obviously* catastrophic.

This post considers nine potential ways to fight mazes in general, out of nine I could come up with. The last method will be the next post in the sequence.

Potential Partial Solutions

Solution 1: Knowledge

Not all solutions start with knowledge of the problem. Most of the solutions *are good ideas anyway even if you don't know about mazes*.

Knowledge remains a logical place to start any list of solutions. Knowledge is also key to implementation of any of the other solutions. There is a reason we roll our eyes

whenever anyone seeks to 'raise awareness,' but it is often vital to cause people to notice and understand things they instinctively avoid noticing and understanding.

People mostly do not understand what mazes are, how they work or the effects they have on those who work for them or interact with them.

People treat many aspects typical of these mazes as normal that did not used to be considered normal, many of which are not conducive to human flourishing, life or happiness. People entering mazes mostly do not expect what they will find therein. They often spend years of their lives and much of their human and social capital preparing for such roles, both before and after they enter the maze, before realizing what they have signed up for.

When people do realize what mazes are in time to realize they want out, they are often left without the knowledge and skills to exit and land on their feet. They are also often surrounded by people who would react badly to a decision to exit.

When people realize what mazes are and are doing to society, and move to help with those problems, they tend to pursue regulatory solutions that make the problems worse rather than better, because those regulatory solutions favor big organizations over small organizations, and entrench existing systems.

In addition to learning about mazes themselves, people learning what matters in life will generally raise their asking price to participate in mazes, no matter their level of understanding of what they are up against. I am not claiming to know what matters in life, but [all the plausible answers](#) point towards staying away from mazes on the margin.

There are of course some exceptions. [Knowledge is not everywhere and always good for all good things](#).

Mazes grow and coordinate in large part via spreading the covert information that being the type of person who supports mazes is how one gets resources. While this information remains covert, it benefits mazes, whether or not it is true. Where it becomes common knowledge, again whether or not it is true, unpredictable things happen. People can fight back, or they can use this to justify imposing maze values more forcefully on others.

Another place mazes benefit from knowledge is in supportive technologies such as methods of building large organizations.

Surveillance of individuals tends to favor mazes, and make mazes more damaging to humans, in ways I won't reiterate here. Surveillance of corporations by default seems like it would do the opposite. But it depends in both cases on what is being measured.

In most other aspects, both about mazes and not about mazes, [common knowledge](#) about the world and what is true of it and happening in it in general, as opposed to details about particular events, tends to strictly benefit the non-maze cause against mazes.

Solution 2: Regulatory Reform

Society-wide or important problems are not always political problems, and are not always helped by even well-intentioned political intervention.

The obvious first political intervention is thus almost always to *stop making things worse*.

[Regulatory capture](#) is the technical term for the inevitable result of government regulation. The mazes that are major corporations and professional associations and accrediting institutions, who have the expertise and lobbying power and financial incentive and perseverance, and the cultural affiliation with a fellow maze, inevitably gain control of the regulatory bodies supervising them. They then engineer those regulations as tools of expropriation: to limit supply and subsidize demand, to stifle innovation, for oligopoly or monopoly, insiders and existing firms against outsiders and upstart firms, for the industry against the public.

Regulatory gatherings are gatherings of those in the same trade. Adam Smith reminds us in [The Wealth of Nations](#) these phenomena might be getting worse but they are not new:

“People of the same trade seldom meet together, even for merriment and diversion, but the conversation ends in a conspiracy against the public, or in some contrivance to raise prices.”

If we were to simplify the rules and minimize the regulatory burdens, and the regulatory burden on doing pretty much anything at all, in particular on small businesses and the ability of individuals to do object-level work, that would help level the playing field.

Even where action is effectively permitted, it is usually technically forbidden, and citizens have the correct intuition that doing pretty much anything useful is illegal and would be clamped down on if it threatened the powerful.

Severely limiting occupational licensing would be the natural place to begin.

Another strategy would be to end actual direct subsidies to the maze.

Solution 3: End Corporate Welfare, Too Big to Fail and Implicit Subsidies

[Blatant lies are the best kind](#). Some of the things propping up mazes at public expense, in their fight against real production, is sufficiently brazen that it is obviously and explicitly expropriation.

Corporate welfare is everywhere despised and everywhere prospering. One could use mistake theory and think of this as an [inadequate equilibria](#) problem. If your area subsidizes a large corporation to move, expand or stay, that enriches the corporation at the expense of the public, and gives it an edge on its competition.

You do it anyway, because it also transfers jobs and other resources from another jurisdiction to yours. Everyone else is doing it, so you can ill afford not to follow suit. Thus, everyone cries out for the policy to end, but does it anyway.

Maze theory says that this is too generous. Mazes like to make deals with other mazes, to expropriate resources from the public and potential non-maze competition. Thus, they instinctively work together to do this in steadily more corrupt fashion, on steadily grander scales. This correctly predicts a lot of *obviously bad and/or obviously unnecessary deals* even for the locality in question, such as many deals for new sports stadiums. It also predicts that corporations will often not follow through on their

promises, and not be held to account, because the deal is primarily a conspiracy rather than a trade.

One could respond by pointing to the [Winner's Curse](#) and the need to make risky deals that might not pay off, which also predicts some of the same. That response implicitly accepts that not only can different jurisdictions not cooperate to not go to destructive economic war with each other despite *being part of a government*, it also implies that they care so little about expropriation of the public by corporations that they do not even ensure that the deal is a net economic win for the local jurisdiction. The argument implicitly accepts that [it's the incentives](#), and it would be crazy to act on anything else.

The effect in either case is the same. Regardless of whether you think a given deal is locally good, these deals get more available and more generous as corporations get larger. Thus they constitute an explicit bribe, a tax that gives larger corporations a competitive edge on smaller ones. End that, and the balance shifts back.

Implicit subsidies are expropriation of resources from the public, in the form of subsidies that are created knowing who the recipients will be, likely for doing nothing or what is already going to be done. Alternatively, they are *conditional future expropriation* where the government makes clear that if things go a certain way, likely a way that causes mazes trouble, then there will be a bail out. Mostly they go to mazes. Where they don't, they directly reward maze behaviors. This is obvious pure corruption. Ending or minimizing this would help.

Too big to fail is a special case of implicit subsidy. If you get big enough, the short-term collateral damage of your potential failure is so big that the government would have no choice but to prop you up at its expense. This is an enormous competitive advantage. Even if shareholders would be wiped out by the intervention, it makes such firms much better business partners, better places to pretend to work, better places to loan money. The shareholders reap those benefits. Even if they didn't reap the benefits, we all still bear the costs of keeping old dinosaurs around as they become less and less productive and competitive, and their maze levels continue to rise.

Another subsidy is retroactively strengthened intellectual property protections. Some amount of protection is necessary to allow people to capitalize on innovation and creation, but there can be little doubt that copyright law has run amok at the behest of companies including Disney, and patent law is often now used as a way to bury anyone attempting something new or to protect rents and monopolies long past any reasonable bounds.

Solution 4: Tort Reform

As a very rough approximation, in America, the current way lawsuits work is that if something goes wrong, you can attempt to blame anyone in the area. Doing so means they lose a lot of time and money even if you lose in court. In court, you win if you convince a jury, although judges can throw it out. That introduces a lot of random variance, raises risk of doing anything, and greatly advantages anyone who can plausibly deploy the resources to fight back.

Insurance against this is expensive and requires conforming to the legibility requirements and standards of care laid down by mazes, and *still* won't reliably work when things get truly bad and/or someone decides to come after you for real, because

your time is still destroyed and you could potentially get liability in excess of any insurance.

A key concept is *standard of care* or the similar *what a reasonable person would do*. Effectively, this means that if you do the thing that is usually done and check off all your boxes, you are mostly safe. If you do *anything else*, you are at the whim of a jury of random people who do not know anything about the area, and who are focused on the thing that went wrong.

You can get a waiver of liability or other form of consent and assumption of risk, and this helps *somewhat*, but they often do not hold up, and many key rights to sue cannot be given away.

When we employ lawsuits like this, we are enforcing an extreme form of [Asymmetric Justice](#). If doing something means capturing a small percent of the gain and being on the hook for all or more than all of the loss, it is going to be hard to find profitable action or divergence from standard action.

Our justice system enforces Asymmetric Justice towards *communication* and especially *keeping records* even more than it does towards actions. Consider [this recent column](#) by the excellent Matt Levine. It vividly describes the conflict between engineering, which requires people communicate information and keep accurate records, and the legal system and public relations, which tell you that *keeping accurate records is insane*. Because in the future, someone will sue you or seek to embarrass you, and they will gain access to your records.

Those who adopt the maze nature and seek success within a maze also want to destroy records, so they can avoid accountability for mistakes within the organization. Giving a *valid, hugely valuable legal rationale* for the destruction of knowledge, of favoring the implicit over the explicit in all recorded communication, seems incredibly bad. We need the ability to have [private](#) communication that is both recorded and stays private, which isn't the result of the [blackmail](#) that is lawsuits settled by non-disclosure agreements, which is more destruction of records rather than a loophole allowing records.

This all gets complicated and could of course backfire. Mazes will capitalize on any openings to escape blame or exposure for the things they have done. Juries and results being unpredictable has its huge advantages, both here and elsewhere. This is not an endorsement of any particular action. It is only a note that this is an opportunity for improvement.

Solution 5: Health Care Reform

America's current system of effectively requiring health insurance, since if you attempt to self-insure they inflate your bill to ludicrously fraudulent levels, and making the purchase of health insurance prohibitively expensive for individuals or small businesses, acts as a gigantic tax on self-employment and small business.

Even a small business cannot efficiently buy its employees reasonably priced (even on relative terms to what large business pays) health insurance. If your company can't buy insurance for you, *not only* do you pay a huge and increasing premium to use the exchanges or otherwise buy an individual policy, *you also are hugely tax disadvantaged*.

One *does not simply* go without health insurance.

The system allows mazes in the medical industry to capture a large portion of national wealth. People are legally required to allow such mazes to confiscate all of their wealth, if the mazes decide this is appropriate. They can charge arbitrarily large and increasing fees for drugs and actions, which insurance companies must pay, out of insurance people must buy or face even higher prices.

Getting the right to enter and profit from these mazes requires over a decade navigating educational mazes, which together with existing professionals tightly regulate and ration the right to make human beings more healthy. Getting the right to enter these mazes as a seller of goods usually requires epic spends over years working with existing governmental and other mazes to convince everyone that what you are doing is safe, and thus you are all but forced to then extract rents to make back your money.

And so on. I would hope I need not convince anyone, regardless of their politics, that our current healthcare system is deeply broken. Other countries have systems less broken in important ways. They usually form and promote mazes in their own ways, but also usually do so to a lesser extent, in a less broad reaching and less toxic way.

It is difficult to overstate how disastrous this has all been. A crippling of the ability of people to do things at all, on any level, without a major organization, *isn't necessarily even the biggest loss here*. As a nation, the United States spends roughly 18% of our GDP on healthcare and has almost nothing to show for it.

This is not the place to debate or suggest in detail, or even in broad outline, which reforms should be undertaken. It is enough to note that our current hybrid system, where we insist on pretending we use markets but massively distort those "markets" to allow fraudulent rent extraction at every turn, is one of the largest debacles in human history.

While also acknowledging that *things could be made so much worse* by the wrong intervention. As has already happened. A lot.

Solution 6: Demand Less Illusion of Safety and Security

Demands for real safety and real security will tend to favor existing proven solutions and methods over new ones, because all things being equal they will tend to be safer, and there will be less uncertainty about how safe they are. We demand more real safety than we should, and reducing that would help a bit.

Real security is tricky, as discussed previously; it's not clear we know what that actually entails. Demands for real security would involve things like trying to stop atomization, and of course better solutions to existential risk.

The *illusion* of safety and security is both in much higher demand, and much more unbalanced in favor of large, established organizations and methods. It puts burden after burden upon any new enterprise, punishing it asymmetrically for any potential perception of new dangers while not giving offsetting reward for reduction of other dangers. New systems must *prove* that they are at least as safe as existing systems in every possible way, and bear the risk and cost of any exceptions.

Existing systems can be deeply damaging, but are grandfathered in.

One can also add things like illusion of fairness to this list.

A classic way to demonize an existing system or product is to point out this grandfathering: that if it were a new system or product, there is no way it would be permitted. I recently saw this used to argue against bicycles in New York City. But it would be even easier to use this against *cars* in New York City. Or cars in general, or many other central aspects of life.

Mazes provide the *illusion* of security, of a reliable and blameless path in life. Mostly people care primarily about this blamelessness. Getting people to care less about that illusion (and realize it is an illusion) would help a ton.

Mazes produce things that provide the *illusion* of safety. Where they do provide safety, they do so in particular ways specified by metrics, which tend to not be where much true risk lies.

Those products also provide *blamelessness*. This risk aversion and worry about blame extends to even minor decisions.

It would be good to get people to care less about such things. Make that a celebrated virtue.

Solution 7: Change Consumer Behavior

One can say ‘it’s the incentives’ all one likes. People constantly defy those incentives.

Ever seen those bumper stickers that say “Support Small Business”? Or local business? People make that choice all the time. Many livelihoods depend on the generosity of those around them. If you tip your taxi driver or waiter, you’re voluntarily paying more to support those around you. The better people treat those around them providing valuable real services, and the more we prioritize making sure they get the profits rather than extractive corporations, the easier it becomes to provide valuable real services.

Give your business to those around you that are adding value to your life, in the ways that help them do better. Encourage others to do the same. Make it a celebrated virtue. Don’t let stores designed to skim off the most profitable sales from those around them by being one block closer, rather than adding real value, get your business. If the local place helps you figure out what to buy, don’t then go online to find it half a percent cheaper. If an internet service is taking a huge cut for arranging a transaction, [consider picking up the phone](#).

That story is about public goods, and paying your share for those goods, including those that can’t be directly monetized. In some cases this will have such a big payoff that you alone will net benefit. But if that is everyone’s threshold, that’s obviously far too high.

There are also good self-interested reasons to favor the local and small over the large.

Exploration and variety are often undervalued, although there are contexts and people where they are overvalued. Story value is important and importantly undervalued, as is learning and experimenting, as is preventing one from ‘wearing out’ the default options.

Local or less popular options tend to be less easy to get started, and less engineered to get you to choose them. Mazes have big data and are experts at steering business their way. Locals are not, and [the best ones show a disdain for such operations](#).

Thus, you will by default tend towards choosing the 'default' or 'easy' option even if you can't point to how this is happening. Your instinct towards doing a non-blameworthy thing is also going to overly push one in that direction. It is important to correct for these and other similar biases. A good general life rule that should be its own post is, when in doubt, take the option you believe you are naturally inclined not to take when the decision should be close.

Local or less popular options also are relatively safe from [Goodhart's Law](#), and are more likely to care about illegible aspects of their offerings, or about quality in general even when it does not cash out into profits. If you don't account for these unknown illegible advantages, and for them valuing you and your experience in the future for their own sakes, you're not giving them proper credit.

Consider what is likely to happen over time. A maze will become more of a maze, and their offerings reliably get more customer-hostile both over calendar time, and as users use them. They will keep climbing hills. A smaller operation is under less pressure to do so and gets more rewards from not doing this. You also have opportunity to make the smaller operation stronger and better for what you care about through offering feedback and with your dollars. Thus, getting used to and relying on smaller operations is good long term planning even if they are slightly worse right now.

The more customer behavior can shift in these ways, the better the playing field for those offering alternatives to mazes.

Solution 8: End Maze Legitimacy and High Status, and/or Raise Real Work Legitimacy and Status

Two sides of the same coin.

We currently give unusually high status to 'winners' and those able to extract resources, especially via predictable rents. It shows they are 'smart' and 'successful' and 'reliable.' Instead we could strive to uphold the status of those who do 'honest work.' We could value those who have the virtuous version, who honor and choose to hold onto that at personal cost, rather than mocking them as chumps.

We could make it clear that 'selling out' or 'gaming the system' was a *bad thing*, rather than a good one, and recognize what it meant to do so. Some free riding, zero-sum activity and compromise of principles is inevitable. Keeping it to a minimum requires *frowning upon it* rather than praising it, wanting to move away from rather than towards those who participate. Which in turn means moving towards those who uphold this, and away from those who do not, as in any other norm.

It is one thing to say that making compromises to win big still gets you praised, because winning big enough eclipses the compromises. It is another thing to say that the compromises *are praiseworthy in and of themselves* because that means you're the type who wins big, and that you are *more worthy* than someone who accomplished the same thing without compromising. Think of how much more they could have accomplished.

This of course all leads to endless arguments over what activities, jobs, groups, causes, and so on deserve higher or lower status. That is good. We should have those debates, while being careful not to get *too* distracted by them. The traditional rankings have a lot I strongly disagree with. Only some of those are because people who are seen as doing good and vital work are often instead directly contributing to

mazes and rent extraction, whereas many who are seen as extractive are instead directly contributing. And details often matter, a lot.

(Let's not litigate these here in the comments, that's why I'm not naming anything on either end here, even examples I consider clear cut. It would be a distraction.)

Solution 9: Forcibly Break Up Large Companies

If large companies are a problem, a logical response is to use the government to break them up. This is central to the platform of a serious contender for the next President of the United States, and has been pointed out in comments of previous sections multiple times as an option, so the idea is worth taking seriously.

There certainly is a case that if the choices are this or nothing, one can consider this as a reasonable course of action. It certainly is not the *first-best* course of action. Before you start punishing, taxing or banning something, a good first action is to *stop subsidizing and rewarding it first*.

Forcibly breaking up large companies that we decide are a problem would be similar to our policy on food, where we pay farmers to produce massive amounts of high fructose corn syrup, then pay marketing people to convince consumers not to eat so much high fructose corn syrup. If you are politically well-connected and can navigate the associated mazes, you get to steal from the public. If you are not so well-connected because you are oriented towards the object-level and value creation, or you lose the fight for any other reason, the public steals from you, instead.

It is *not a coincidence* that rather than propose a systematic policy, the companies singled out as targets by this candidate *are the relatively new companies, not politically well connected or seen in a good political light, that have best created value for us over the last decade*. To the extent that life in these parts has gotten better, it is hard to think of more central reasons for this than Amazon and Google.

This is illustrative of what happens when one uses the government to enforce such things. The government is the biggest maze of all! If we tie the continued existence of organizations and the value they create and/or capture to the political whims of Washington, we shift what is rewarded even more so from the creation of value to the navigation of mazes and the playing of politics.

Breaking up all sufficiently large major corporations, or preventing them from engaging in practices that discourage competition, still *might* be better than doing nothing, if the problem of size is sufficiently bad and we do our best to keep this implemented reasonably and apolitically.

People, including several major candidates, have also proposed an intentionally destructive confiscatory wealth tax. What would happen under such a wealth tax is left as an exercise to the reader. I will only point out that founders are the best hope of forestalling mazes, and a wealth tax would prevent them from retaining control.

What are other options? Is there a good version of such a policy that might improve the situation, were it to be implemented, even if it seems completely unrealistic? What is the most realistic version we can come up with?

Michael Vassar makes the interesting proposal that we should impose a destructive/confiscatory tax on *liquid wealth* but not impose a tax on *illiquid wealth*. This would have the effect of heavily encouraging companies to remain or

become private. If you had a large percentage of a large private company, taking it public would give you liquid assets, from which the government would then seize more than all the expected returns (e.g. if you tax at 6% annually, that's effectively confiscating *all* the value of an asset to the extent it is not used for consumption.) The only reason to take a company public would be if you intended to divest, at which point you would have to figure out how to make the bulk of your resulting assets illiquid again.

This would be hugely destructive of value. Liquidity is valuable and highly valued. Intentionally creating illiquidity destroys that value, and likely involves additional huge destruction of value. So does reducing the market price of liquid assets. We would be back to the world in which you were paranoid that if you got anything useful to a bandit that bandits would come and seize it (or worse). Think Apocalypseburg from The Lego Movie 2, where everything and everyone is super dark and grim in order to not provide any potential value to outsiders. Anything that isn't super dark and grim is instantly destroyed by "aliens" whose attention is drawn to them.

There is also the problem that defining what is liquid versus illiquid is tricky. If it is a matter of taxing things that list on exchanges than all you did was ban transparent exchanges that could be efficiently used by outsiders and the public, and replace them with maze networks that take advantage of the public and give the profits and deals to those with good connections while not properly allocating capital. If it is a matter of trying to figure out what someone actually can't sell, or otherwise gets complicated, then the primary way of retaining wealth becomes convincing the government that your assets are illiquid.

One thing that seems likely, on top of the inevitable regulatory capture and carved out exceptions, is a bunch of shell corporations designed to be illiquid, whose ownership is divided in hard to measure ways, which then hold the liquid assets.

Thus, I do not think such plans survive contact with the enemy.

Any solution to these problems must be strictly formulaic. It must merit no exceptions and minimize the amount of game playing. This means keeping it simple. It also would ideally protect new companies and their founders.

As stated above, I would not start on this response without first doing everything I could to implement solution three. Once that was exhausted, I might consider something like a tax, escalating with size, on large organizations that have existed for a minimum amount of time and are no longer controlled by their founder.

Last Idea

The next post will consider the ninth and final method, which asks "what could one determined person with resources do without resorting to politics?" and answers this with "actually they could do quite a bit."

They could Create a Full Alternative Stack.

Create a Full Alternative Stack

[Last time](#) I proposed nine strategies for fighting mazes. All of them were either systemic solutions requiring coordinated and/or political action, or cultural shifts that happen one person at a time.

Now for the tenth and final proposal for how to fight mazes. The tenth is a proposed strategy that one dedicated person with sufficient resources could implement on their own.

If you are in a position where you have the resources to implement this, please make an effort to take this proposal seriously. And please contact me if you are potentially interested and wish to discuss it further.

Solution 10: Create a Full Alternative Stack

In some ways this is the most ambitious solution here. It may seem Utopian.

In other ways, it is the least ambitious, and most practical. It could be implemented by a single sufficiently wealthy and committed individual or organization. All other known solutions can be implemented locally, and would help locally, but need general adaptation to succeed in general.

The full alternative stack offers a contract.

Disengage entirely with mazes and traditional distortionary incentives, competitions and signals of all kinds, and discard all zero-sum activity, in favor of [doing the thing](#). Whatever the thing in question may be. Make no compromises to make oneself legible or attractive to outside sources of funding. Tolerate no maze behaviors of any kind. Hire and fire with this deeply in mind.

In exchange, if you keep your end of the bargain, the stack will *fully* fund you and your operations, at fair prices that do not take advantage of the hold-up opportunity presented by your giving up of other opportunities. Evaluation will be fully on the object-level merits, and the generative processes involved.

This is a form of tenure for the people. If they continue to act with integrity and work to accomplish positive-sum things relevant to the stack's interests, and spend responsibly, they and their family will have strong financial security.

Think of this as similar to tenure at a university, or to the system of universal employment for partisan hacks. If you are promising the stack gives you the opportunity to prove yourself. Once you have proven yourself, we take care of you, even if you don't bear as much fruit as we would like, *provided you continue to play by the rules of the stack and honor its values*. Unlike many tenured professors nowadays, we would not then force you to seek grants, outside investments, or other outside funding for your work. On the contrary, you would be agreeing *not* to seek outside funding, so as to protect your incentives from corruption.

This is also a form of secured financing for corporations and other organizations. While they need funding to reach maturity, they will be evaluated on whether they are succeeding at doing the thing. Traditional signals, [and anticipation of future traditional signals](#), will be not only disregarded but punished – it's fine to look good naturally, but

if you are doing things *in order to look good or successful to outsiders rather than be good or successful*, then that breaks the compact.

We call this a full alternative stack because the ideal version is *literally a full alternative stack*. It recreates civilization. Those involved would not need or depend on outside goods or services. There would be a local area fully owned by and given over to the project.

That is the full version. The full version is ambitious and difficult, but likely far less ambitious and difficult, and far less expensive, than it appears. We would soon find out how much of current activity is rent extraction or otherwise unproductive, and how much is necessary to keep things running.

A lesser version, built around a particular cause or goal, or to give this freedom to select individuals and corporations, would still be extremely valuable.

The MacArthur grant is a template for what this looks like on a personal level, with a shift in focus from *creativity* to *integrity*, and a bump in compensation – \$625,000 is a lot of money, but that money is designed to be seed money for an activity rather than financial security. Those getting a MacArthur grant still face the specter of future financial needs. One needs an order of magnitude more than that over a lifetime to be secure while not compromising one's interactions with society.

For startup corporations, this can be similar to the standard method of funding a biotechnology company to pursue a new drug of unknown efficacy. Milestones are set. If they are met, funding is unlocked at pre-negotiated levels, locked in for both sides in advance. There is no reason to worry about signaling in other ways unless the company is about to fail. We would add the condition of not then selling out to a maze (in the biotech example, a big pharma company, or taking the company public) when successful, instead keeping the operation privately owned by its founders to prevent it from being eaten and transformed or killed. Public markets exert strong pressure towards maze behaviors, so such companies would need to commit to staying away from them.

I believe there is a strong opportunity for a venture capital fund that promises committed, full funding to projects in this way in fields outside biotechnology. Projects that are freed from having to gain strong negotiating positions regarding raising capital could be much better at pursuing actual success and production. To succeed, such a fund would need to honor its commitments carefully and be credible at every stage. This includes its commitments not to respond *positively* to things that would on the outside be viewed as good news, if they are not in fact relevantly good news. Its word would be its bond. It would also need to be highly skilled at choosing superior evaluation techniques. There are many terrible things about current systems of venture funding, but naive replacement models threaten to be easily gameable or otherwise create new and perhaps much worse versions of the same problems.

Most people nowadays are forced, both within an enterprise and overall in their lives, to structure and censure everything they do in light of their potential future need to look legible, comfortable and successful or valuable to mazes. The prospect of having this option cut off fills them with terror, whether or not this should not be the case. Even when they do not fear it, those around them who rely on them fear it, which has a similar effect.

I am *unusually immune* to these pressures. I have skills that can earn money on demand, without getting a formal job or the approval of a maze, if I need that. I also

have robust savings and a family and community that would save me if I fell upon hard times. *This terror is still one of the things I have often struggled with.*

Freeing a select group to do things without regard to such concerns, and people knowing they have the option to join this group, would be a major cultural change. Ideally this would then become the 'city on a hill' that shows what is possible, and gets emulated elsewhere. Regulatory and other legal issues would still have to be navigated, which would often be most of the difficulty of any worthwhile operation. This is why typical versions of this type of proposal go to places like seasteading. Mazes will instinctively attempt to crush whatever is being built.

If one is sitting on a large pile of money and wishes to do good, or simply wishes to increase production, deploying that money effectively has proven a very hard problem. This is to be expected, under even the best/worst conditions, as such problems are anti-inductive. Any easy answers get utilized until they stop being easy answers. Once others find out your criteria for spending or granting money, some of them will Goodhart and/or commit fraud to extract those funds.

The closer you attempt to stick to specified metrics and use criteria you can explain and justify that looks consistent, the more you are optimizing over time for those who Goodhart and commit fraud, in order to grant the appearance of having the appearance of attempting to help in the approved ways, rather than optimizing for actually helping. This is certainly a danger to the full stack operation as well, and the best reason to keep the operation relatively small.

The closer you do *not* stick to such methods, the more illegible you become, the more blameworthy you appear, and the more likely you are actually buying things that make you feel good about them as your metric. Which, in turn, is even easier to Goodhart or commit fraud on.

Sticking to 'do the right thing,' as this solution suggests, and rewarding those who do right things is a rather crazy ask without rich contextual knowledge. The larger you scale, the more universal you attempt to get, the crazier it gets. Goodharting or committing fraud on 'right thingness' is as much a threat as Goodharting or committing fraud on anything else, if you're not staying a step ahead. That very freedom from mazes, Goodharting and fraud is the precious thing you're trying to get in the first place.

The project has to cash itself out purely on its own terms. It has to care more about doing things its own way than getting things done or looking effective, where that own way is a ruthless focus on what will actually work. Everyone's instinct, even that of the best possible additions, will be to abandon this at every step. Everyone will face constant pressure to do so.

But without sufficient scale to complete the stack, how do you break free from, and securely break the right people away from the need to worry about, mazes and other outside forces?

Threading that needle is going to be very difficult, even if the other impossible problems are solved. I do not think any one person or formal group can be the head of the entire stack without it getting too large. One must form a distinct subset, and hope others form the required other parts, and until that happens purchase what one needs from the outside using capital, and trust those in the project to continue interacting economically in some ways outside of the stack.

Ideally one does not need to literally go to Mars to be allowed to complete the project. However, *if one does need to literally go to Mars*, then there is a fair argument that *literally going to Mars* is a reasonable price to pay to be allowed to complete the project.

The next post asks what we should do when we have a project that would benefit from a large organization.

Protecting Large Projects Against Mazedom

If we wish to accomplish something that would benefit from or require a larger organization or more levels of management and bureaucracy, what should we do in light of the dangers of mazes?

There are no easy answers. Real tradeoffs with real sacrifice are the order of the day. But we can do some things to expand the production possibilities frontier, and choose wisely along that frontier.

As is often the case, this starts with admitting you have a problem.

Too often, it is assumed that one should scale without worrying about the costs of scaling, or without counting becoming a maze as one of the biggest costs. Not stretching oneself maximally thin, or getting in the way of this process, becomes the sin of not maximizing effectiveness or profits.

That ensures failure and rapid descent into a maze. Start by getting out of this mindset.

If you are looking to accomplish a big thing that requires lots of organization, management and bureaucracy, here are ways to help contain the damage.

None of them should come as a surprise by this point. This is more of a synthesis of points already made, and not a place I feel I have special additional insights. So I will keep this short.

Solution 1: Do Less Things and Be Smaller

Recognize the threat and its seriousness, and the resulting risks and costs of scaling even if handled wisely. Understand that *you almost certainly want to be smaller and do less things* due to this concern. This is a real trade-off.

Think of the actions, priorities, skills and members of a group as being increasingly inherently expensive as the group grows – adding new elements has increasing marginal costs. Changing anything, or preventing anything from taking its natural course (including towards being more of a maze) also becomes increasingly expensive. Under such circumstances, the final marginal cost of every action or member is equal to the marginal cost of the final action taken or member added.

One must think about the future scale and the future costs, solve for the resulting budget constraints and trade-offs, and spend wisely. The same way that one must avoid accumulating technical debt, an organization should worry greatly about complexity and culture long before the bills are presented.

Remember that not doing something does not condemn that thing to never being done.

Encourage others to form distinct groups and organizations to do those other things, or where possible to do those other things on their own. Promise and give rewards and trade relations for those that do so. Parts that can operate on their own should usually do so. If your organization discovers a new product, business or other undertaking that is worth pursuing, but isn't a cultural or logistical fit for what already exists, *or simply doesn't need to be in the same place to work*, strongly consider spinning it off into its own thing.

This solution applies fractally throughout the process. Do only one central, big thing. Do less major things in support of that thing. Do less minor things in support of each of the major things. Do each of those minor things as elegantly and simply as you can. Take every opportunity to simplify, to look for easier ways to do things, and to eliminate unnecessary work.

Solution 2: Minimize Levels of Hierarchy

Throughout this sequence we have emphasized the high toxicity of each level of hierarchy. If you must scale, attempt to do so with the minimum number of hierarchical levels, keeping as many people within one level of the top or bottom as possible. Fully flat is not a thing, but more flat is better.

Solution 3: Skin in the Game

Scale inherently limits skin in the game, as there is only 100% equity to go around, in all its forms. The thinner that must be spread, the harder it is to provide enough skin in the game. One can still seek to provide good incentives all around to the extent that it is possible. Isolating local outcomes so as to offer localized skin in the game helps. Keeping people responsible for areas over extended periods also helps.

Solution 4: Soul in the Game

No matter how large the organization, if you care deeply about the organization or its mission, preferably the mission, you can still have soul in the game. Mission selection is huge part of this, as some missions lend themselves to soul much more than others, but if you have a mission you are setting out to do as the starting point, you're stuck. That means protecting against mission creep that would disrupt people's soul in the game above and beyond other issues of mission creep. Keep people doing what they are passionate about.

Solution 5: Hire and Fire Carefully

Nothing raises maze levels faster than hiring someone who is maze aligned. One must not only avoid this, but maintain sufficient maze opposition to prevent it from happening in the future. All of this needs to sustain itself, which will be increasingly difficult over time and as you grow. As noted earlier, maze actions need to be firing offenses.

Solution 6: Promote, Reward and Evaluate Carefully

If you can find ways to evaluate people, and choose which ones to promote and reward, in ways that are immune or even resistant to mazes and their politics, this would be a giant leg up. Hiring and firing are key moments, but the promotion can be similarly important. One must resist the temptation to try to implement 'objective criteria' and use hard numbers, as this introduces nasty Goodhart problems, and forces the system to choose between 'people around you can change the numbers so

the maze wins out' and 'people around you can't change the numbers and no one cares about the people around them.' It's a big problem.

Solution 7: Fight for Culture

This can be seen as a catch-all, but the most important thing of all is *to care about organizational culture and fight for it*. Where what you are fighting for is not a maze. All of these 'solutions' often involve trade-off and sacrifice, and this is no exception. If you do not value the culture enough to fight for it, the culture will die.

Solution 8: Avoid Other Mazes

This will not always be possible, but to the extent it is possible, attempt to sell to, buy from, get funding from, make deals with non-mazes, and avoid mazes. Mazes will reward maze behaviors and push towards you raising maze levels, in ways that they will make seem natural. Do not put yourself in that position more than necessary.

Solution 9: Start Again

We must periodically start again. Even if everything is done right, within any given organization we are only staving off the inevitable. At a minimum we must periodically clean house, but that seems unlikely to be enough for that long. Actually starting over and building something new every so often, where frequency is highly context-dependent, seems necessary. This goes for corporations, for schools, for governments, and for everything else.

Mazes Sequence Roundup: Final Thoughts and Paths Forward

There are still two elephants in the room that I must address before concluding. Then I will discuss paths forward.

Moloch's Army

The first elephant is Moloch's Army. I *still* can't find a way into this without sounding crazy. The result of this is that the sequence talks about maze behaviors and mazes as if their creation and operation are motivated by self-interest. That's far from the whole picture.

There is mindset that instinctively *and unselfishly* opposes everything of value. This mindset is not only not doing calculations to see what it would prefer or might accomplish. It does not even believe in the concept of calculation (or numbers, or logic, or reason) at all. It cares about virtues and emotional resonances, not consequences. To do this is to have the maze nature. This mindset instinctively promotes others that share the mindset, and is much more common and impactful among the powerful than one would think. Among other things, the actions of those with this mindset are vital to the creation, support and strengthening mazes.

Until a proper description of that is finished, my job is not done. So far, it continues to elude me. I am not giving up.

Moloch's Puzzle

The second elephant is that [I opened this series](#) with a puzzle. It is important that I come back to the beginning. I must offer my explanation of the puzzle, and end the same place the sequence began. With hope.

Thus, the following puzzle:

Every given thing is eventually doomed. Every given thing will eventually get worse. Every equilibrium is terrible. Sufficiently strong optimization pressure, whether or not it comes from competition, destroys all values not being optimized, with optimization pressure constantly increasing.

Yet all is not lost. Most of the world is better off than it has ever been and is getting better all the time. We enjoy historically outrageously wonderful bounties every day, and hold up moral standards and practical demands on many fronts that no time or place in the past could handle. How is this possible?

Here is my answer:

It is possible because old things that get sufficiently worse eventually die, and get replaced by new things that are better. It is possible because competition is never perfect, and optimization for a fixed set of metrics is never total. People need and demand slack, and care about many factors on many meta levels. No matter how many times we say it is difficult, not game theoretically sound or even impossible, coordination continues to happen all the time.

Moloch wins when all things are equal and the situation is at a strict static equilibrium. Things are not equal. The situation is not static or strict. While there are some places where Moloch has not won, exit to those areas serves as an additional safety valve.

We have spent the bulk of this sequence dealing with mazes. Mazes are, *among other things*, a case of what happens when all of that breaks down – when we [strip away the protections](#). They then go on to create all sorts of strange negative feedback loops, and make life inside them, and if left unchecked life even much of life outside of them, quite bad. In many ways, even this sequence has only scratched the surface of the dynamics involved.

Good News, Everyone!

To the extent that this makes one estimate that the world is a worse place, and its people are less happy and less likely to succeed or improve, this is bad news indeed.

But, *to the extent that we already knew there was a problem*, to the extent that we already knew how functional our world was and how happy people are, this is not bad news. This is good news. We already knew there was a problem and had an amorphous sense of hopelessness and doom. Now that we know more about what the problem is, we have a more specific source of hopelessness and doom that we can try to do something about. We can explore further. We can limit the damage on a personal level, seek to get things done, and perhaps even improve conditions in general.

The other good news is in our estimates of the effectiveness of doing object level things. We know about how much real and useful stuff is actually created and accomplished, in the sense that we know what useful stuff we get to use and we live real lives. If we are screwing people up this much, we are wasting this much of everyone's time, and otherwise getting in our own way, then what does that say about how powerful it is to actually do things? If only a handful of people at a major corporation are producing all the value while everyone else plays politics, then what does that say about that handful of people?

If, as my model holds, people who are allowed to actually do real things *consistently* output insanely great things, then imagine what would happen if we let more people do that. Then realize there are still places where this is possible. Imagine what you could accomplish if *you* did that.

I hope this journey has been enlightening, and that it inspires further explorations along with concrete actions. Perhaps we can create useful common knowledge. With luck, some people will avoid falling into maze traps as the result of what has been written here. With even better luck, this can improve the fate of some organizations, or even lead to broader motions towards taking down maze levels.

Conclusion

The central (Immoral) Mazes sequence is now complete.

I wrote it because someone had to, and no one else would. This was the best model I could come up with, presented as clearly as I know how to do so.

There are many things I still want to say, most of which are not mentioned above. Some of them are a matter of putting in the work to write things down. Others I don't

know how to say without sounding crazy, or I don't understand them myself well enough to write about them. I have experience with mazes, but not with higher maze levels. I can abstractly model that mentality and mode of operation, but I can't pretend I truly understand it.

The more others can join the conversation and move it forward from here, especially those with more direct experience but that managed to emerge intact enough to tell the tale, the more hope we have for making real progress.

To encourage the story to continue, both for myself and for others, I will now explore the question of Paths Forward. What are the best next questions?

Paths Forward

These topics are not only or even primarily for me. Others can and should take up the mantle as well. Here are some things that seem like good ideas to keep the ball rolling.

I still don't have (1a) Moloch's Army in a place where I am ready to post it, but I promise to keep trying. At a minimum, I hope that trying leads to finding more missing concepts and ideas that can help bridge the gap.

I previously tried to bridge into this with [The Darwin Game](#) sequence, but that foundation wasn't enough. If I resumed that line now then I would next talk about (1b) The Rise of Cliquebot. I am still confused on whether that would (eventually) get me where I would need to be.

One of the words I know I will likely need is (2a) Fnord. The word fnord comes from the [Principia Discordia](#). In its original introduction, we are all conditioned when we see the word fnord to not notice it but to become stressed and irritated. Thus, it is peppered into places that people with power want us not to look, but it is entirely absent from things like advertising. Having a name for things that make you not want to notice them, and understanding how that dynamic works, seems important. Mazes, and many aspects of them, are fnords. Another good closely maze-aligned Discordian concept is (2b) [The Snafu Principle](#), whereby communication is only fully possible between equals, leading to Situation Normal All F****ed Up. It should be mostly already covered by implication but is worth a special focus slash better place to link to.

Another key idea I need are (2c) Basilisks, as such things have much more important rolls to play than the example that gave them their name. Needless to say one must proceed with caution. Prior restraint will be a thing, here.

I introduced terms for things related to mazes, including maze behaviors, leading Wei Dei to ask the logical question (3a) What Are Maze Behaviors? It would be good to have a compact link-to-it answer. But I'm also coming around, after having started a draft of this, to the perspective that this is likely asking the wrong question. We could be better served to ask about (3b) Maze Levels and especially about (3c) The Maze Nature, which I also introduced, instead. The Maze Nature is closely related to Moloch's Army and might be the right way in. Or perhaps it will be the other way around.

One worry is that mazes are the wrong central concept on which to build these metrics and understandings. It might be more helpful, either in general or in some places, to talk about (4) Simulacrum Levels. [Ben's post here](#) is the best reference

created so far, but it definitely needs to be improved upon. I also do not properly understand simulacrum levels, or at least my understanding and Ben's are importantly different. These things need to be explored carefully, as the dynamics seem central to what is happening to the world.

Many of the proposed solutions, and hypothesized causes, would benefit from more careful treatment. People could gather numbers, make profiles and case studies especially where they have their own experience. Many great comments were (5a) Maze Examples, people talking about the place they work and saying how that fits into the bigger picture. Some claimed they worked at mazes. Many claimed that where they worked *should* have been a maze, based on size and what not, but that it *totally wasn't one* or at least wasn't as much of one as one would expect.

The question of (5b) How Did They Avoid Mazedom? is a good one for people who claim it was avoided. And of course someone should try to answer the question (5c) How Maze-Intensive Are Our Corporations? A few things to consider along the way,

The first is that *these are the people reading DWATV and/or LessWrong*. That is not a random sample. The whole point of such websites is to promote a method of thinking that is incompatible with a maze. If you care enough about good methods of thinking to read these websites, you likely have a strong aversion to maze behaviors compared to others with similar levels of human capital, both not tolerating them in others and not being willing to do them yourself. You likely chose your place and class of employment in part on this basis, regardless of how you thought about that decision when making it.

If someone was under full sacrifice-everything pressure from anything, they're not going to read a book length sequence about it, because they can't.

The second is that the best way to get comments on the internet is to get someone to tell you that you are wrong. People who think the model is failing in their case are much more likely to comment (or so my model says) than people for whom the model is accurate. That is how commenting typically works. So the sample does not provide much evidence, in that sense.

The third is that mazes are things people do not want to see, and people will select for choosing exactly the mazes they do not notice. If they did notice them, and were of the type to be reading this, they would probably avoid them. My estimate of the maze levels of several of my jobs jumped dramatically in the months after I left, because I had the chance to get perspective and to experience life without that aspect of things. Even if you are not yourself being eaten by the maze in question, it is instinctive to not notice it, or to justify what you are experiencing as normal and healthy and reasonable, or at least not so bad, when it is nothing of the kind.

It is also inevitable that some people have already self-modified to adapt to the mazes and that this makes it harder for them to notice what is going on. Not only does it mask this thing especially hard, it masks noticing things in general. Eventually, this would drive people away from reading such things, but inertia in such matters can last a while.

I am sure reading this sequence helps (in a probabilistic sense) but it is a hard problem. It is clear to me from the comments that I am not yet getting the full situation across successfully to many readers.

One thing someone will need to at some point write (6a) Mazes That Are Not Within Organizations, discussing dynamics that produce similar results without people strictly being bosses and subordinates. And generally (6b) What Types of Things are How Maze-Like, (6c) To What Extent do People At Large Have the Maze Nature, (6d) Close Examination of Maze Interactions, and so on.

The model of perfect competition presented early in the sequence proved very non-intuitive to many people and got a lot of push back. A lot of that was my using technical economic terms that trigger strong intuitions about what the answers are. I think a lot of this is that these terms are usually taught in a school context, where there are right answers and right principles that these things are supposed to illustrate, and the models always work a certain way. I was using them a different way, changing the assumptions and what things would be implied versus not implied, in order to provide justification for the transition to the later part of the sequence.

If I had to do that over again, I'd look for a way to take a different approach entirely, because it looks like it was more confusing than I expected, and the objection I felt I needed to overcome was much weaker than I expected it to be slash to the extent it existed the people who had the objection didn't feel satisfied by the explanation. So it didn't really work the way I wanted it to. I do stand by what I wrote, and think there's important stuff there, but kill your darlings and all that.

However, there's a whole series of posts I could write (7a) On Competition, going deeper into what I was getting at there. Or even (7b) On Ultimate Human Value and all that, which is kind of a big deal, but again, super hard and I'm constantly terrified of writing to advocate ethical positions because I assume others with better rhetoric and academic chops in the area would just blow anything I write to shreds and nothing would be accomplished. That doesn't mean I think I'm wrong or anything, but it's a problem, and it's why the old "(7c) Can I Interest You In Some Virtue Ethics" post never got written after the last paths forward. Still, (7d) How Consequentialism Is Ruining Things And What To Do About It might be more tractable. I thought our decision theory would be better than this by now, but alas.

Another valuable thing to do would be to give people practical advice. Advice on how to choose fields of study and careers and jobs and where to live, and other major life choices, to avoid the dangers of mazes, while taking into account the many other things that matter. We have things like 80,000 Hours, but that leads to a consequentialist calculation with many of the key terms missing because they haven't been quantified, and many of the best options never considered because they can't be standardized, which is playing right into the hands of mazes. It's the same as the usual social pressure to go work for mazes, if hopefully a little more efficient about collecting at least something in return. So, (8) Practical Career Advice to Avoid Mazes, A Continuing Series.

In particular, we need (9a) How to Start a Small Business. Not a start-up – I don't think creating or joining one of those is a bad idea, but the how of that is something we are much more familiar with already – but a small business that actually *does business*. A way of life that involves buying things and then making or transporting things and selling things and having the things you sell generate revenue that pays expenses plus the rent. The kind of thing immigrants often do to great effect, except you can do even better if you are already integrated into the culture and have better connections and seed funding.

More fundamental would be the simple (9b) How to Do Business. There is specialized knowledge of how to *start* a business, but the most important doing-business related skill is simply *how to do business*. Start-ups backed by VCs are different because they *fundamentally do not (yet) do business*. They might do some facsimile of business by getting customers, or even unit economically profitable customers. But that's a *completely different* model of how to succeed than trying to find customers and make money from them and use that money to pay the rent and the surplus to grow the business. [What such start-ups are actually doing](#) is the performance art of "doing business" aimed at the VCs that are their bosses. Doing business is a different thing entirely.

One hint you *might* be in a maze is that you are "doing the thing" in quotation marks rather than doing the thing.

I could also perhaps do (10a) Some Businesses Worth Starting, (10b) Some Start-Ups Worth Starting, (10c) Some Ways To Make Money I Made Work Before, and/or (10d) Some Ways To Make Money I Didn't Do But I'm Confident Will Work. I'm sure good versions of 10c and 10d would be appreciated, but it's a truism that no one cares about your start-up idea, whether or not they should, so perhaps not 10a and 10b.

Looking at the potential causes in more detail, and attempting to better understand their dynamics, would also be valuable. (11a) To What Extent Do We Need Large Organizations?, (11b) The Demand for Blamelessness, (11c) The Illusion of Security, (11d) The Dynamics of Rent Seeking, (11e) On Atomization, (11f) Education as a Maze, (11g) Education as Maze Indoctrination, and so on.

We could also look further at the solutions. One valuable place to go if we could do it well would be (12) How to Explain Mazes. We could also look at (13a) The Full Alternative Stack in much more detail. The biggest problem starts with (13b) How to Tell if Someone has the Maze Nature slash (13c) How to Tell if Someone Has Edge (or just (13d) Edge). That will need a better name because edge is a massively overloaded term. Here edge means the tendency to '[angle shoot](#)' or otherwise take every opportunity to take local advantage in underhanded ways. Do you always need to be 'on guard' against someone, or can you relax?

The other half is a post or series (13e) On Being a Source of Money. This is a huge unsolved problem. The moment anyone sees you (where you can be a person, or a group or organization, or anything else) as a source of funds, even a possible future source of funds, it corrupts all interactions even when everyone involved has the best of intentions. You always have to worry someone is after your money. The people around you are likely to be there largely because of your money. Others have to worry about whether you think they're after your money even when they totally aren't, and you can never rule such a thing out entirely. There is a reason a lot of rich people keep their wealth (or at least their giving/funding) a secret. Consider the chapter of [Skin in the Game](#) titled [Only the Rich are Poisoned](#). Money can be shockingly useless or backfire, so what to do?

The other solutions to mazedom are largely political actions, where I am loathe to get more into the weeds for mostly-obvious reasons, or at least so far I have chosen to strive to be maximally apolitical throughout not only this sequence but on the entire blog. That's another thing mazes do. They encourage us to stay out of such questions because they create an amorphous feeling that getting involved might backfire against us at some point in some way. Talking more about (14) How Mazes Scare You in such ways might be important. Of course, a lot of it is also that I do not need or

want the trouble of arguing or advocating politics on the internet, nor do I expect much good to come of such a thing, and all that. And once I start down that road it is very easy to make oneself only seen from that perspective. It's mostly a massively over-determined decision. But perhaps that decision is wrong, or will become wrong?

One more-related-to-this-than-you-would-first-think thing I've wanted to do for a while but that would require a lot of work and which might not come together, and which is [motivated by this post](#), is to tell (15) The Journey of the Sensitive One. It would look over the story of the artist Jewel, as told, in [explicit content](#) in chronological order across her first five albums.

It might be most valuable of all to simply get case studies. I would love to see someone take this perspective and in particular examine (16a) What Happened to Boeing? Something seems to have happened that caused them to rapidly have higher maze levels, to the extent that they were unable to produce a plane that did not crash, ignoring huge quantities of written warnings that they were in fact producing a dangerous plane, and despite this being a *very very bad thing* for such a business to do. Steve Jobs was mentioned as a potentially interesting special case, so someone looking into (16b) Apple and Mazes could be interesting, especially to see if maze levels declined there when he came back and if so how he did that. And so on.

If one wanted to do a full extension of the project, (17) On The Gervais Principle could be anything up to and including its own sequence. As I've noted elsewhere, I consider Gervais Principle and Moral Mazes to be fully compatible, and Gervais Principle has a bunch of stuff that expands upon the model.

I'm sure I'm forgetting a lot of other stuff, as well. And of course, I have other things I need to and/or want to write, both about gaming and about other stuff. The game I'm creating is almost ready to get rolling, so I'll probably want to focus on the gaming side even in my writing more and more over the coming months.

I seem to have already written close to a book's worth of stuff. The main sequence (excluding this post, [Quotes from Moral Mazes](#) and [Moral Mazes and Short Termism](#)) is about 40,000 words long, there are a lot of places I could expand upon, and a quick Google says that the average business book is 50,000 to 60,000 words long. If anyone is seriously interested in publishing were I to turn this into a book, or has the know-how to make that happen and thinks it would be a good idea, please contact me with your thoughts and what that would look like. If it would be worthwhile to rewrite this in proper full book form I am open to that idea. If not, it might still be a good idea to print some copies mostly as-is to help reach more people.

And of course, thanks to you for reading, and hopefully thinking about and building upon all of this.

Acknowledgements

I would like to thank Ben Hoffman, Sarah Constantin and Michael Vassar for providing the impetus to read Moral Mazes and take it seriously. It takes a lot of motivation to get through a book that heavy.

I would like to thank Ben Hoffman, Michael Vassar, Raymond Arnold, Ben Pace, Oliver Habyrka, Zack Davis, Jessica Taylor and my wife Laura Baur for helping me edit the sequence. Without strong editorial feedback, this sequence would likely not have come to be, and if it did it would have been much weaker.

Ben Hoffman in particular was vital early on in helping me wrap my head around the problem, Michael Vassar later on for making sure I didn't miss the important points, and Raymond Arnold for making sure what I wrote would have a fighting chance of being understood. Many thanks.

I'd also like to thank the commentators, including the ones who made comments I found frustrating and wrong, but especially the ones who challenged me in ways that improved my thinking.

And also those that helped keep up morale so I could finish. Engagement matters, and knowing you have people you respect interested matters too. Little notes of appreciation can go a long way. It's important to [give praise](#). This includes Scott Alexander, and also Robin Hanson, who said he would read Moral Mazes on the strength of [the quotes I provided](#). I look forward to seeing both their takes.

I'd also like to thank the person who got [How to Identify an Immoral Maze](#) to be featured on Hacker News. The *majority* of hits I get come from being featured somewhere, in one form or another. I do my best not to let that change what I write, other than being happy to edit posts to make them stand on their own better if this is standing between them and being linked in this way. If that is ever the case and the modifications seem reasonable, please do let me know.

Finally, thanks again for reading. I hope the time you have invested has proved worthwhile.